CS 162: Operating Systems and Systems Programming

Lecture 2: Fundamental Concepts of Operating Systems

Sept 3, 2019
David Culler
https://cs162.eecs.berkeley.edu

Reading: A&D CH2.1-7, 2.10

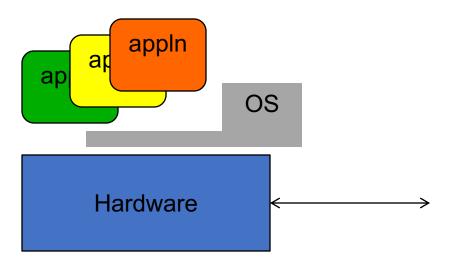
8.1-2 (skim)

HW:0 due 9/6

Autograder Registration: TODAY

Recall: What is an operating system?

- Special layer of software that provides application software access to hardware resources
 - Convenient abstraction of complex hardware devices
 - Protected access to shared resources
 - Security, Authentication, Isolation, Fault tolerance
 - Communication amongst logical entities

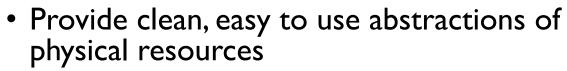


Recall: What is an Operating System?



- Referee
 - Manage sharing of resources, Protection, Isolation
 - Resource allocation, isolation, communication





- Infinite memory, dedicated machine
- Higher level objects: files, users, messages
- Masking limitations, virtualization



Glue

- Common services
 - Storage, Window system, Networking
 - Sharing, Authorization
 - Look and feel

Recall: Protecting Processes

- Use two features offered by hardware:
 - I. Dual Mode Operation
 - Processes run in user or kernel mode
 - Some operations prohibited in user mode
 - 2. Address Translation
 - Each process has a distinct and isolated address space
 - Hardware translates from virtual to physical addresses
- Policy: No program can read or write memory of another program or of the OS

Recall: Virtual Machines

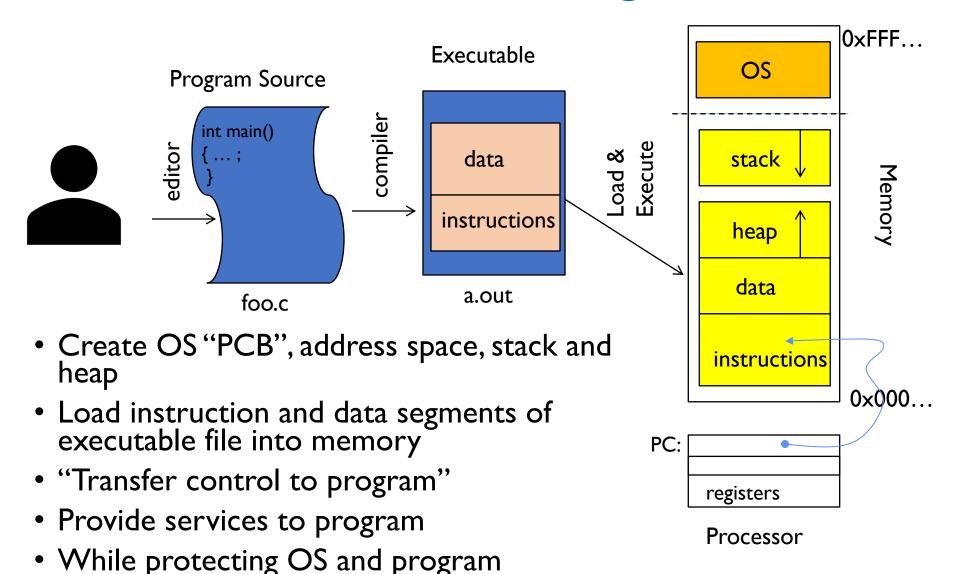
- Definition: Software emulation of physical machine
 - Programs believe they own the machine
 - Simulated "hardware" can have features we want
 - So perfect can run OS on it (guest OS)
- Versus Processes and Containers

Today: Four Fundamental OS Concepts

Thread: Execution Context

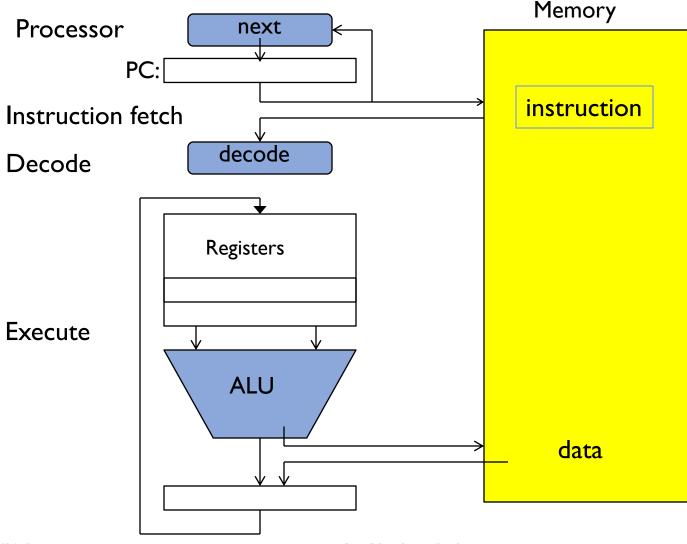
- Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads + ...
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

OS Bottom Line: Run Programs



Review (6 I C): Instruction Fetch/Decode/Execute

The instruction cycle



Today: Four Fundamental OS Concepts

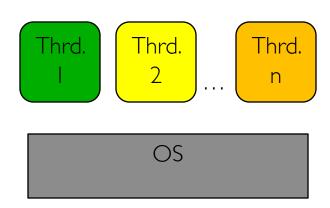
Thread: Execution Context

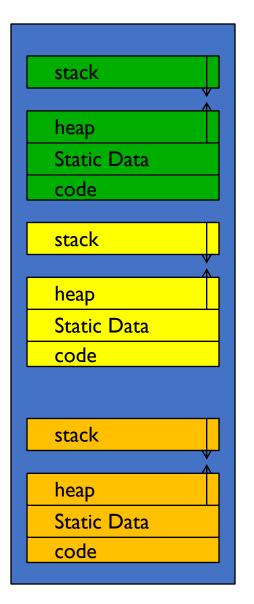
- Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

Thread of Control

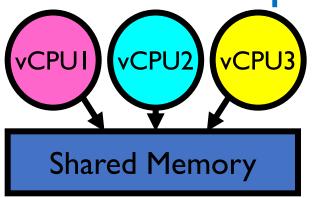
- Definition: A single, unique execution context
 - Program counter, registers, stack
- A thread is executing on a processor (core) when it is resident in that processor's registers
- Registers hold the root state of the thread:
 - The rest is "in memory"
 - Including program counter & currently executing instruction
- Registers point to thread state in memory:
 - Stack pointer to the top of the thread's (own) stack

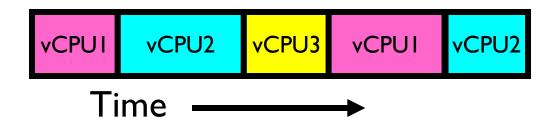
Multiprogramming - Multiple Threads of Control





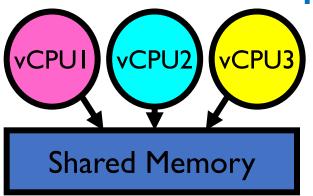
Illusion of Multiple Processors

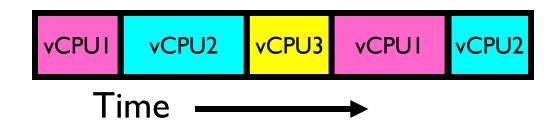




- Threads are virtual cores
- Multiple threads: Multiplex hardware in time

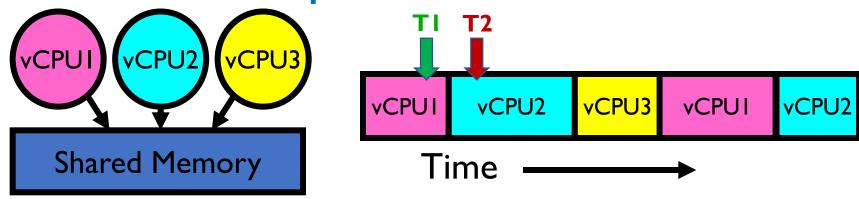
Illusion of Multiple Processors





- Contents of virtual core (thread):
 - Program counter, stack pointer
 - Registers
- Where is it?
 - On the real (physical) core, or
 - Saved in memory called the *Thread Control Block (TCB)*

Illusion of Multiple Processors



- At TI: vCPUI on real core, vCPU2 in memory
- At t2: vCPU2 on real core, vCPU1 in memory
- What happened?
 - OS Ran [how?]
 - Saved PC, SP, ... in vCPUI's thread control block (memory)
 - Loaded PC, SP, ... from vCPU2's TCB, jumped to PC

OS object representing a thread?

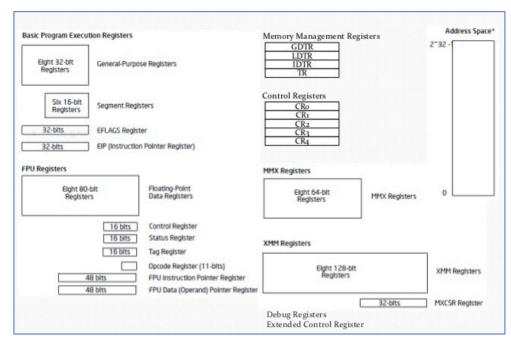
- Traditional term: Thread Control Block (TCB)
- Holds contents of registers when thread is not running
- What other information?

PINTOS? — read thread.h and thread.c

Registers: RISC-V => x86

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	-
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	-
x4	tp	Thread pointer	-
x5	to	Temporary/alternate link register	Caller
x6-7	t1-2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	si	Saved register	Callee
x10-11	a0-1	Function arguments/return values	Caller
x12-17	a2-7	Function arguments	Caller
x18-27	s2-11	Saved registers	Callee
x28-31	t3-6	Temporaries	Caller

Load/Store Arch with software conventions



Complex mem-mem arch with specialized registers and "segments"

- cs61C does RISC-V. Will need to learn x86...
- Section this week

Very Simple Multiprogramming

- All vCPU's share non-CPU resources
 - Memory, I/O Devices
- Each thread can read/write memory
 - Perhaps data of others
 - can overwrite OS?
- Unusable?
- This approach is used in
 - Very early days of computing
 - Embedded applications
 - MacOS I-9/Windows 3.I (switch only with voluntary yield)
 - Windows 95-ME (switch with yield or timer)

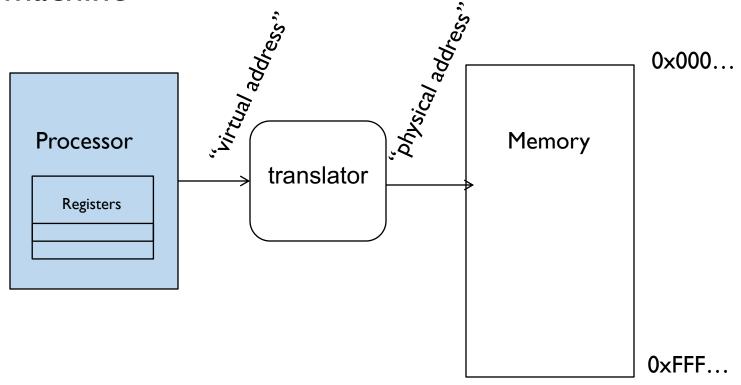
Four Fundamental OS Concepts

- Thread: Execution Context
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

18

Key OS Concept: Address Space

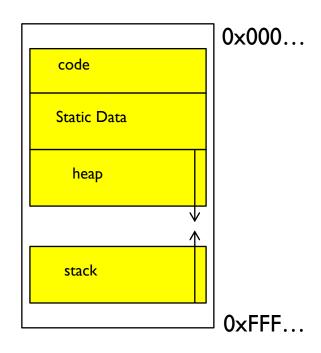
 Program operates in an address space that is distinct from the physical memory space of the machine



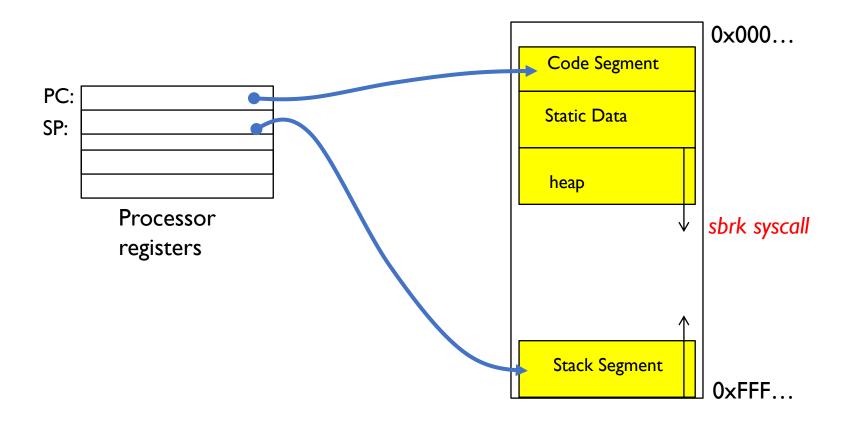
Address Space

- Definition: Set of accessible addresses and the state associated with them
 - $2^{32} = \sim 4$ billion on a 32-bit machine
- What happens when you read or write to an address?
 - Perhaps acts like regular memory
 - Perhaps causes I/O operation
 - (Memory-mapped I/O)
 - Causes program to abort (segfault)?
 - Communicate with another program

•

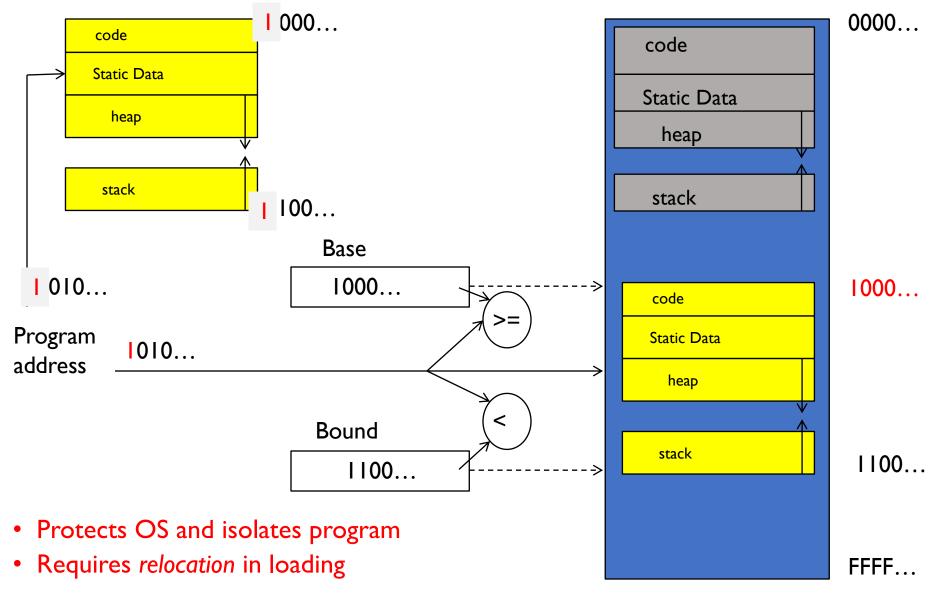


Address Space: typical structure

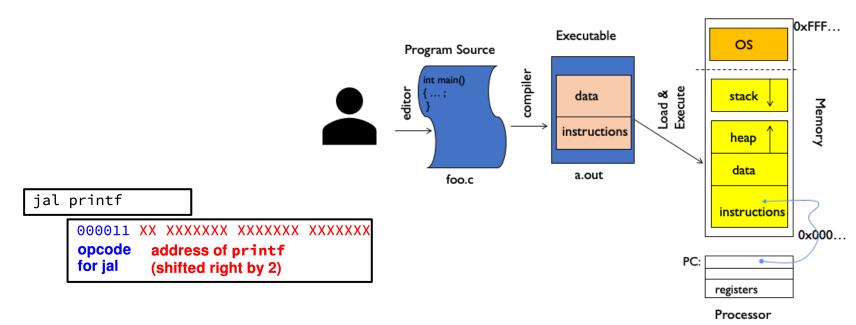


What can the hardware do to help the operating system protect itself from programs? Programs from others?

Strawman: Base and Bound (no translation)

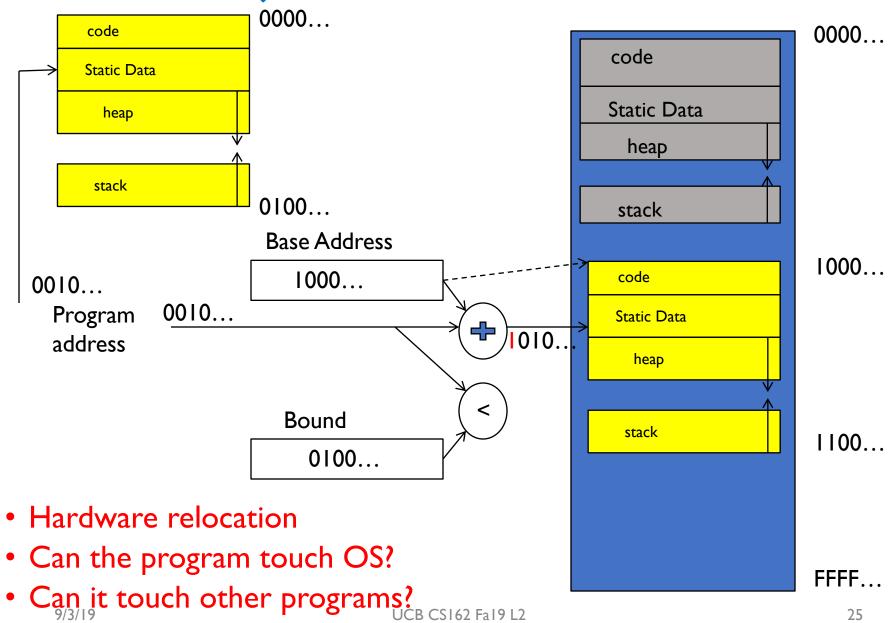


6 I C Review: Relocation



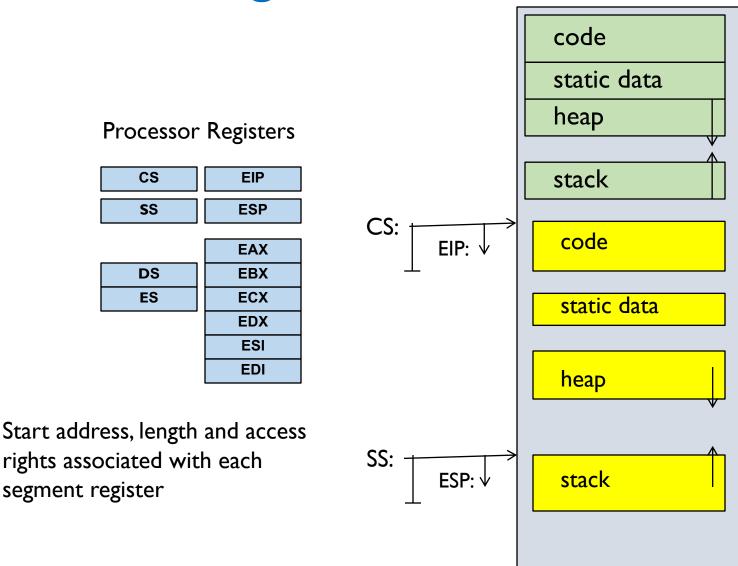
- Compiled .obj file linked together in an .exe
- All address in the .exe are as if it were loaded at memory address 00000000
- File contains a list of all the addresses that need to be adjusted when it is "relocated" to somewhere else.

What if we just added in the base?



25

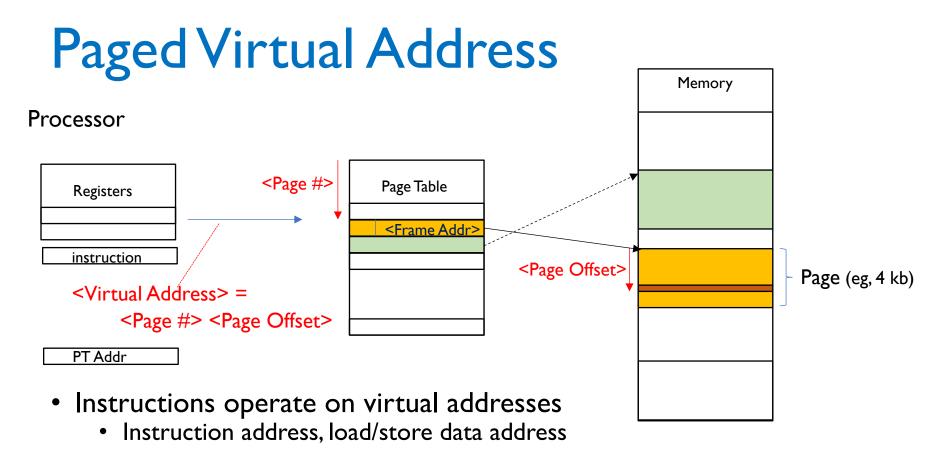
x86 - segments and stacks



Paged Virtual Address Space

- What if we break the entire virtual address space into equal size chunks (i.e., pages) have a base for each?
- Treat memory as page size frames and put any page into any frame ...

Another cs6IC review



- Translated to a physical address (or Page Fault) through a Page Table by the hardware
- Any Page of address space can be in any (page sized) frame in memory
 - Or not-present (access generates a page fault)
- Special register holds page table base address (of the process)

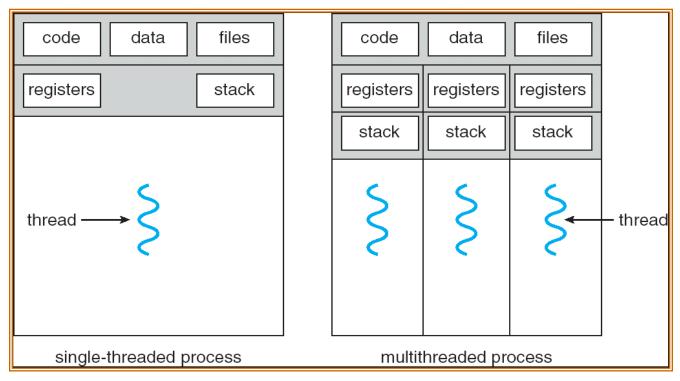
Four Fundamental OS Concepts

- Thread: Execution Context
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

The Process

- Definition: execution environment with restricted rights
 - Address Space with One or More Threads
 - Owns memory (mapped pages)
 - Owns file descriptors, file system context, ...
 - Encapsulates one or more threads sharing process resources
- Application program executes as a process
 - Complex applications can fork/exec child processes [later]
- Why processes?
 - Protected from each other. OS Protected from them.
 - Execute concurrently [trade-offs with threads? later]
 - Basic unit OS deals with

Single and Multithreaded Processes



- Threads encapsulate concurrency: "Active" component
- Address spaces encapsulate protection: "Passive" part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

Kernel code/data in process VAS?

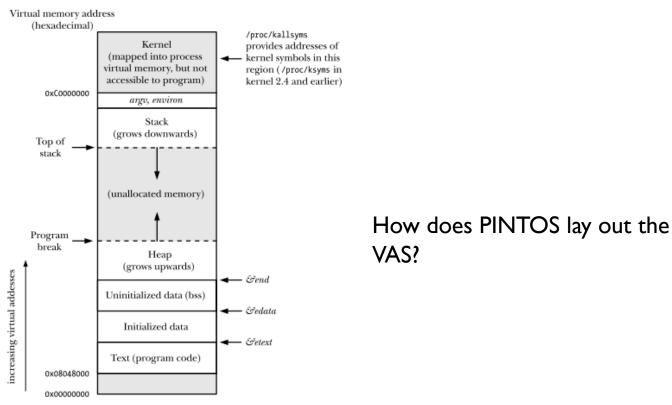


Figure 6-1: Typical memory layout of a process on Linux/x86-32

• Unix: Kernel space is mapped in high - but inaccessible to user processes

Peer discussion: How does kernel access physical resources?

- The actual RAM
- (memory mapped) I/O devices

More cs61c review

Protection (Isolation)

- Why?
 - Reliability: buggy programs only hurt themselves
 - Security and privacy: trust programs less
 - Fairness: enforce shares of disk, CPU
- Mechanisms:
 - Address translation: address space only contains its own data
 - Syscall processing (e.g., enforce file access rights)
 - Thread "traps" to the OS
 - Privileged instructions, registers (???)

Typical system (Unix) structure – what happens where?

User Mode		Applications	(the users)		
		Standard Libs	shells and commands mpilers and interpreters system libraries		
		system-call interface to the kernel			
Kernel Mode	Kernel	signals terminal handling character I/O system terminal drivers	file system swapping block I/O system disk and tape drivers	CPU scheduling page replacement demand paging virtual memory	
		kerne	are		
Hardware		terminal controllers terminals	device controllers disks and tapes	memory controllers physical memory	
9/3/19		UGB-65462	Fa 1 9 L2	35	

Four Fundamental OS Concepts

- Thread: Execution Context
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

Dual Mode Operation: HW Support

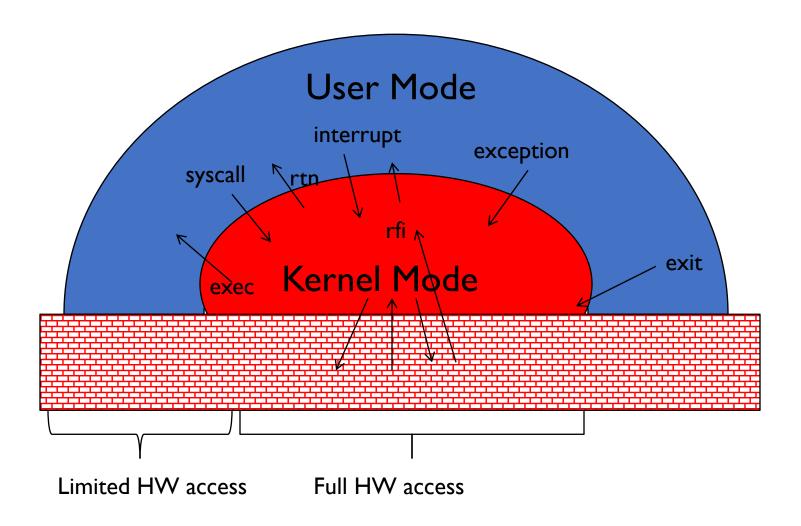
- bit of state (user/kernel mode)
- Certain actions only permitted in kernel mode
 - Which page table entries can be accessed
- User => Kernel mode sets kernel mode, saves user PC
 - Only transition to OS-designated addresses
 - OS can save the rest of the user state if necessary
- Kernel => User mode sets user mode, restores user PC, ...

Logistics

- Course expanded to 487 students
 - Waitlist order
- HW0 due on Friday 9/6, 9:00 PM
- C & 6 I C Review tonight at 7
- Project I Pre-Released
- HWI release 9/7
- Group formation enabled on autograder
 - Opens 9/7, deadline 9/11
- Requested and in process
 - Alternate Midterm: Thurs 10/10 4-6 pm
 - Alternate Final: Mon 12/16 3-6pm

Break

User/Kernel (Privileged) Mode



3 types of U→K Mode Transfer

Syscall

- Process requests a system service, e.g., exit
- Like a function call, but "outside" the process
- Does not have the address of the system function to call
- Like a Remote Procedure Call (RPC) for later
- · Marshall the syscall id and args in registers and exec syscall

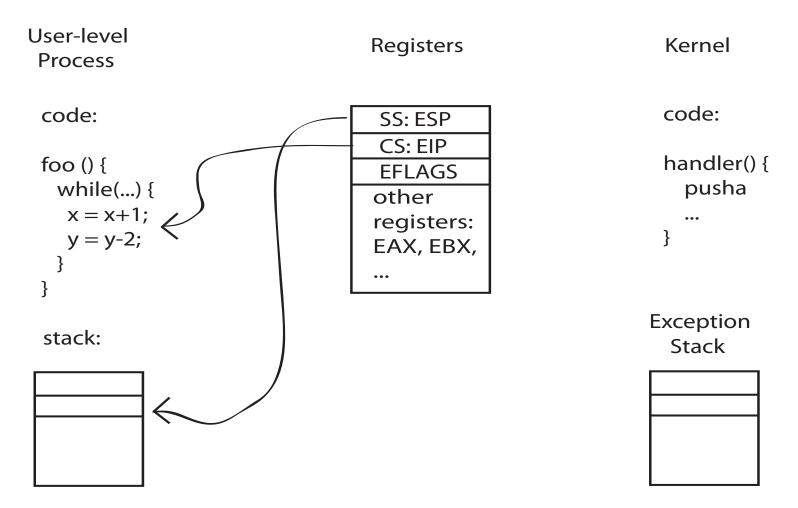
Interrupt

- External asynchronous event triggers context switch
- eg.Timer, I/O device
- Independent of user process

Trap

- Internal synchronous event in process triggers context switch
- e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 exceptions are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go?

Before Exception (Interrupt / Trap)

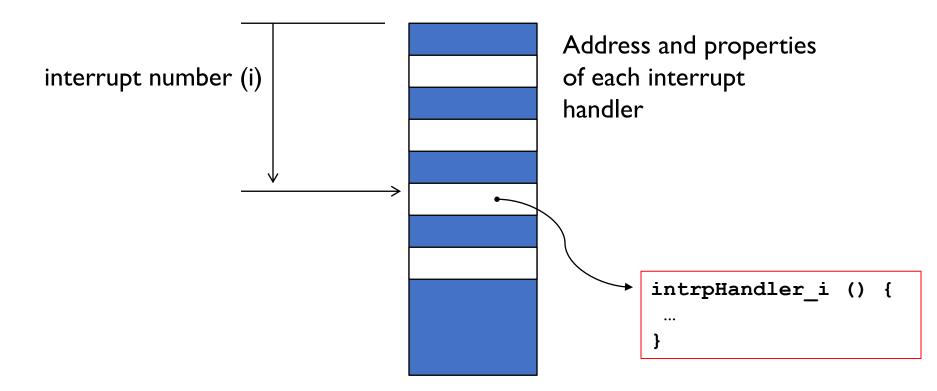


During Exception

```
User-level
                                                               Kernel
                                  Registers
 Process
 code:
                                                               code:
                                   SS: ESP
                                   CS: EIP
                                                               handler() {
 foo () {
                                   EFLAGS
                                                                 pusha
  while(...) {
                                  other
   x = x+1;
                                  registers:
   y = y-2;
                                  EAX, EBX,
                                                             Exception
 stack:
                                                               Stack
                                                                 SS
                                                                 ESP
                                                               EFLAGS
                                                                 CS
                                                                 EIP
                                                                error
```

Where do U→K mode transfers go?

- Cannot let user program specify (why?)
- Solution: Interrupt Vector



Implementing Safe Kernel Mode Transfers

- Carefully constructed kernel code packs up the user process state and sets it aside
- Must handle weird/buggy/malicious user state
 - Syscalls with null pointers
 - Return instruction out of bounds
 - User stack pointer out of bounds
- Should be impossible for buggy or malicious user program to cause the kernel to corrupt itself
- User program should not know interrupt has occurred (transparency)

How do K→U transfers go back?

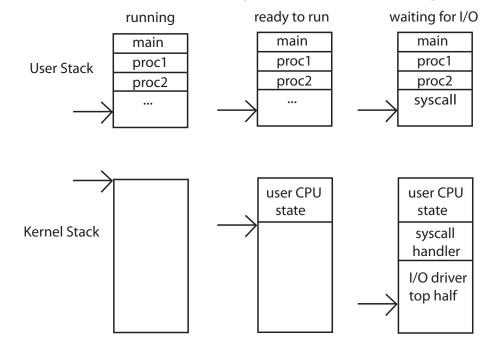
- "Return from interrupt" instruction
- Drops mode from kernel to user priviledge
- Restores user PC (possibly stack)

The Kernel Stack

- Interrupt handlers need a stack
- System call handlers need a stack
- Can't just use the user stack [why?]

The Kernel Stack

- Solution: two-stack model
 - Each OS thread has kernel stack (located in kernel memory) plus user stack (located in user memory)
- Place to save user registers during interrupt



Kernel System Call Handler

- Vector through well-defined syscall entry points!
 - Table mapping system call number to handler
- Locate arguments
 - In registers or on user (!) stack
- Copy arguments
 - From user memory into kernel memory carefully checking locations!
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - Into user memory carefully checking locations!

Hardware support: Interrupt Control

- Interrupt processing not visible to the user process:
 - Occurs between instructions, restarted transparently
 - No change to process state
 - What can be observed even with perfect interrupt processing?
- Interrupt Handler invoked with interrupts 'disabled'
 - Re-enabled upon completion
 - Non-blocking (run to completion, no waits)
 - Pack up in a queue and pass off to an OS thread for hard work
 - wake up an existing OS thread

Hardware support: Interrupt Control

- OS kernel may enable/disable interrupts
 - On x86: CLI (disable interrupts), STI (enable)
 - Atomic section when select next process/thread to run
 - Atomic return from interrupt or syscall
- HW may have multiple levels of interrupts
 - Mask off (disable) certain interrupts, eg., lower priority
 - Certain Non-Maskable-Interrupts (NMI)
 - e.g., kernel segmentation fault
 - Also: Power about to fail!

How do we take interrupts safely?

- Interrupt vector
 - Limited number of entry points into kernel
- Kernel interrupt stack
 - Handler works regardless of state of user code
- Interrupt masking
 - Handler is non-blocking
- Atomic transfer of control
 - "Single instruction"-like to change:
 - Program counter
 - Stack pointer
 - Memory protection
 - Kernel/user mode
- Transparent restartable execution
 - User program does not know interrupt occurred

Four Fundamental OS Concepts

- Thread: Execution Context
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

Process Control Block

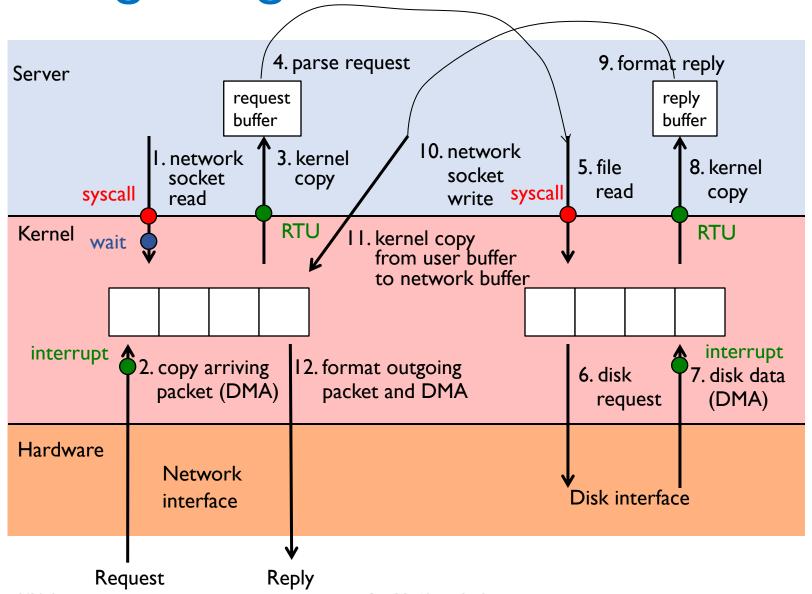
- Kernel representation of each process
 - Status (running, ready, blocked)
 - Register state (if not running)
 - Thread control block(s)
 - Process ID
 - Execution time
 - Memory translations
- Scheduler maintains a data structure of PCBs
 - Or threads
- Scheduling algorithm: Which thread/process should the OS run next?

Scheduler

```
if ( readyProcesses(PCBs) ) {
   nextPCB = selectProcess(PCBs);
   run( nextPCB );
} else {
   run_idle_process();
}
```

- Scheduling: Mechanism for deciding which processes/threads receive the CPU
- Lots of different scheduling policies provide ...
 - Fairness or
 - Realtime guarantees or
 - Latency optimization or ..

Putting it together: web server



Conclusion: Four OS Concepts

- Thread: Execution Context
 - Program Counter, Registers, Execution Flags, Stack
- Address space (with translation)
 - Program's view of memory is distinct from physical machine
- Process: an instance of a running program
 - Address Space + One or more Threads
- Dual mode operation / Protection
 - Only the "system" can access certain resources
 - Combined with translation, isolates programs from each other

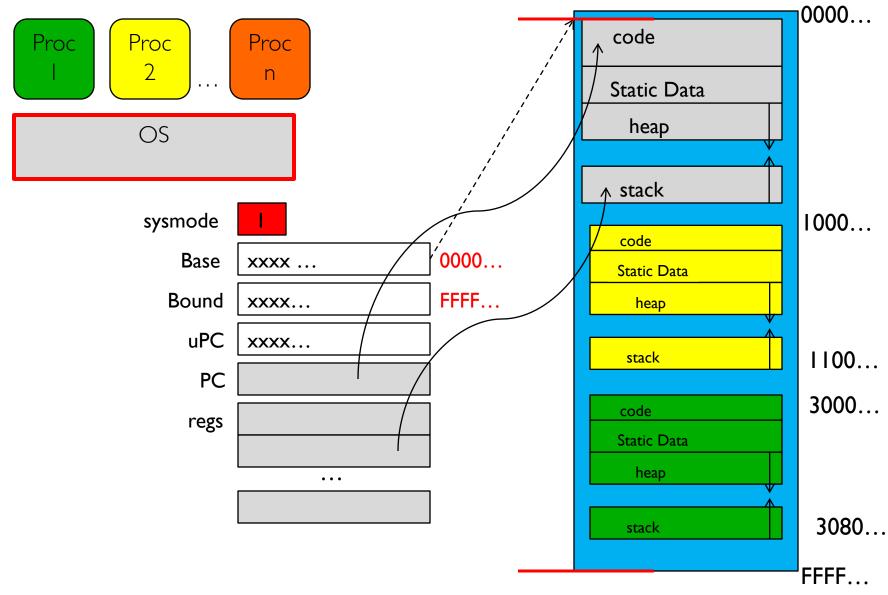
Additional Material

Explore on your own

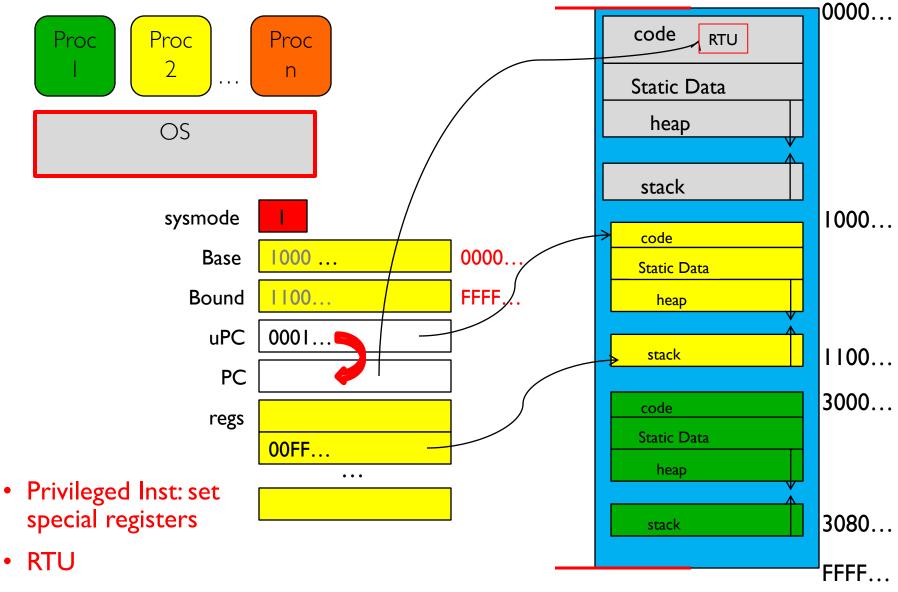
Putting It Together: Mode Transfer & Translation

- Mode transfer should change address translation mapping
- Examples:
 - Ignore base and bound in kernel mode
 - Page tables with "kernel mode only" bits

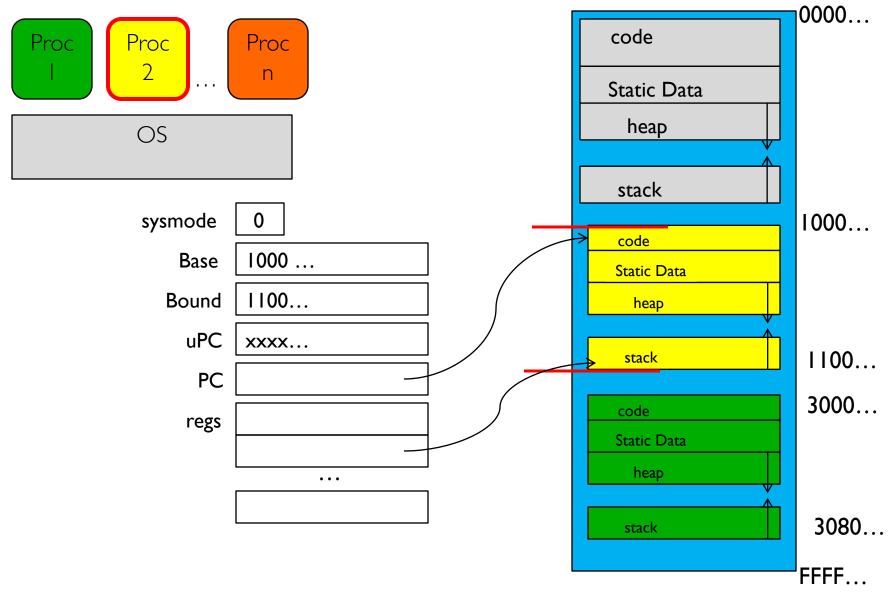
Base & Bound: OS Loads Process



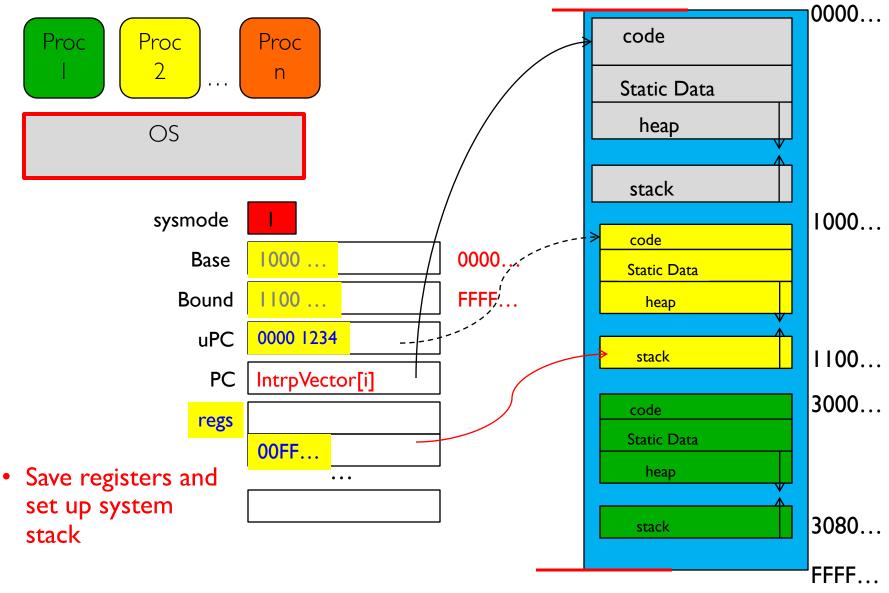
Base & Bound: About to Switch



Simple B&B: User Code Running



Base & Bound: Handling Interrupt

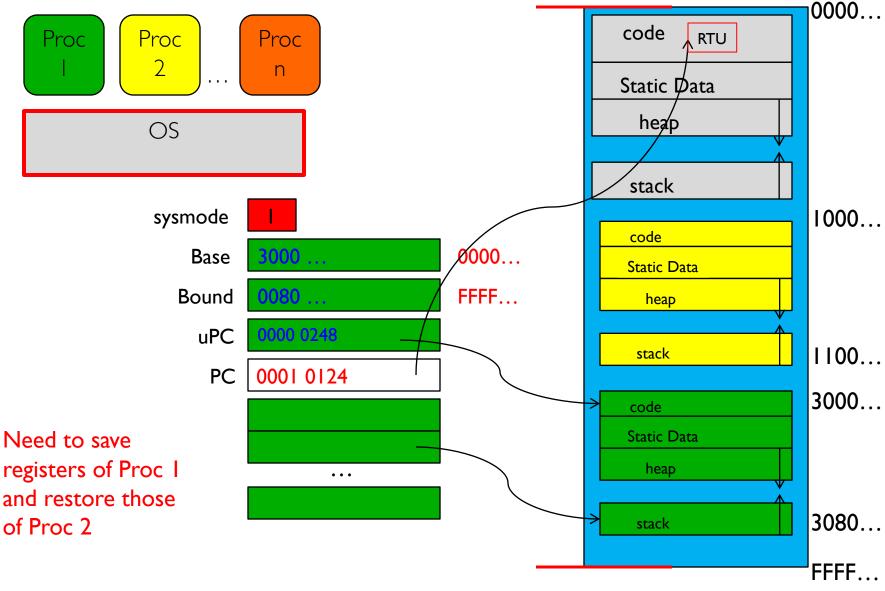


How do we switch between processes?

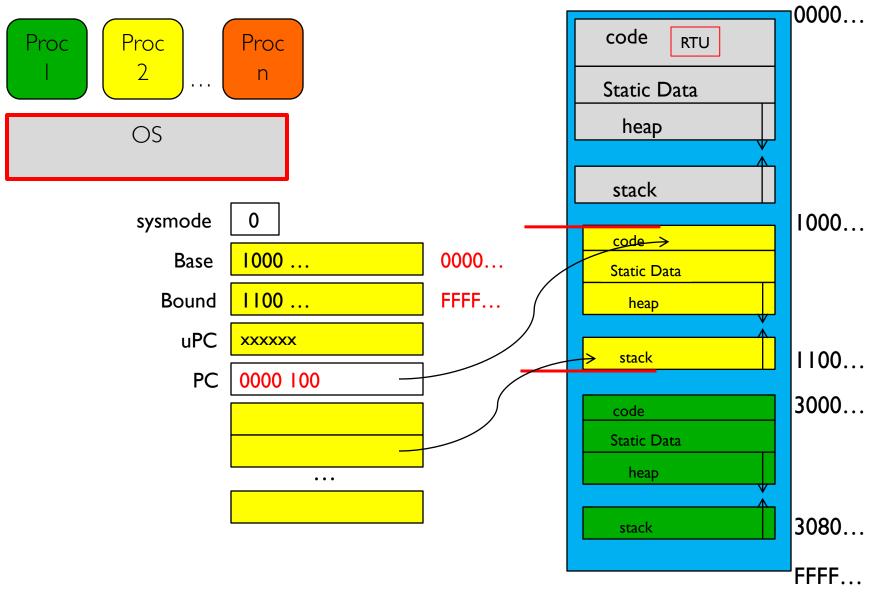
 We already have all the necessary machinery!

• Just requires two mode transfers

Base & Bound: Switch User Process



Base & Bound: Switch User Process



Representing Processes: PCB

