



---

# Processes - Representation in the Operating System & Syscalls

**David E. Culler**

**CS162 – Operating Systems and Systems Programming**

<http://cs162.eecs.berkeley.edu/>

**Lecture 3**

September 5, 2019

Read A&D Ch3, 4.4-6  
HW0 Due 9/6  
Early drop: 9/6  
HW 1 out 9/7, Due 9/18

# Recall: Four Fundamental OS Concepts

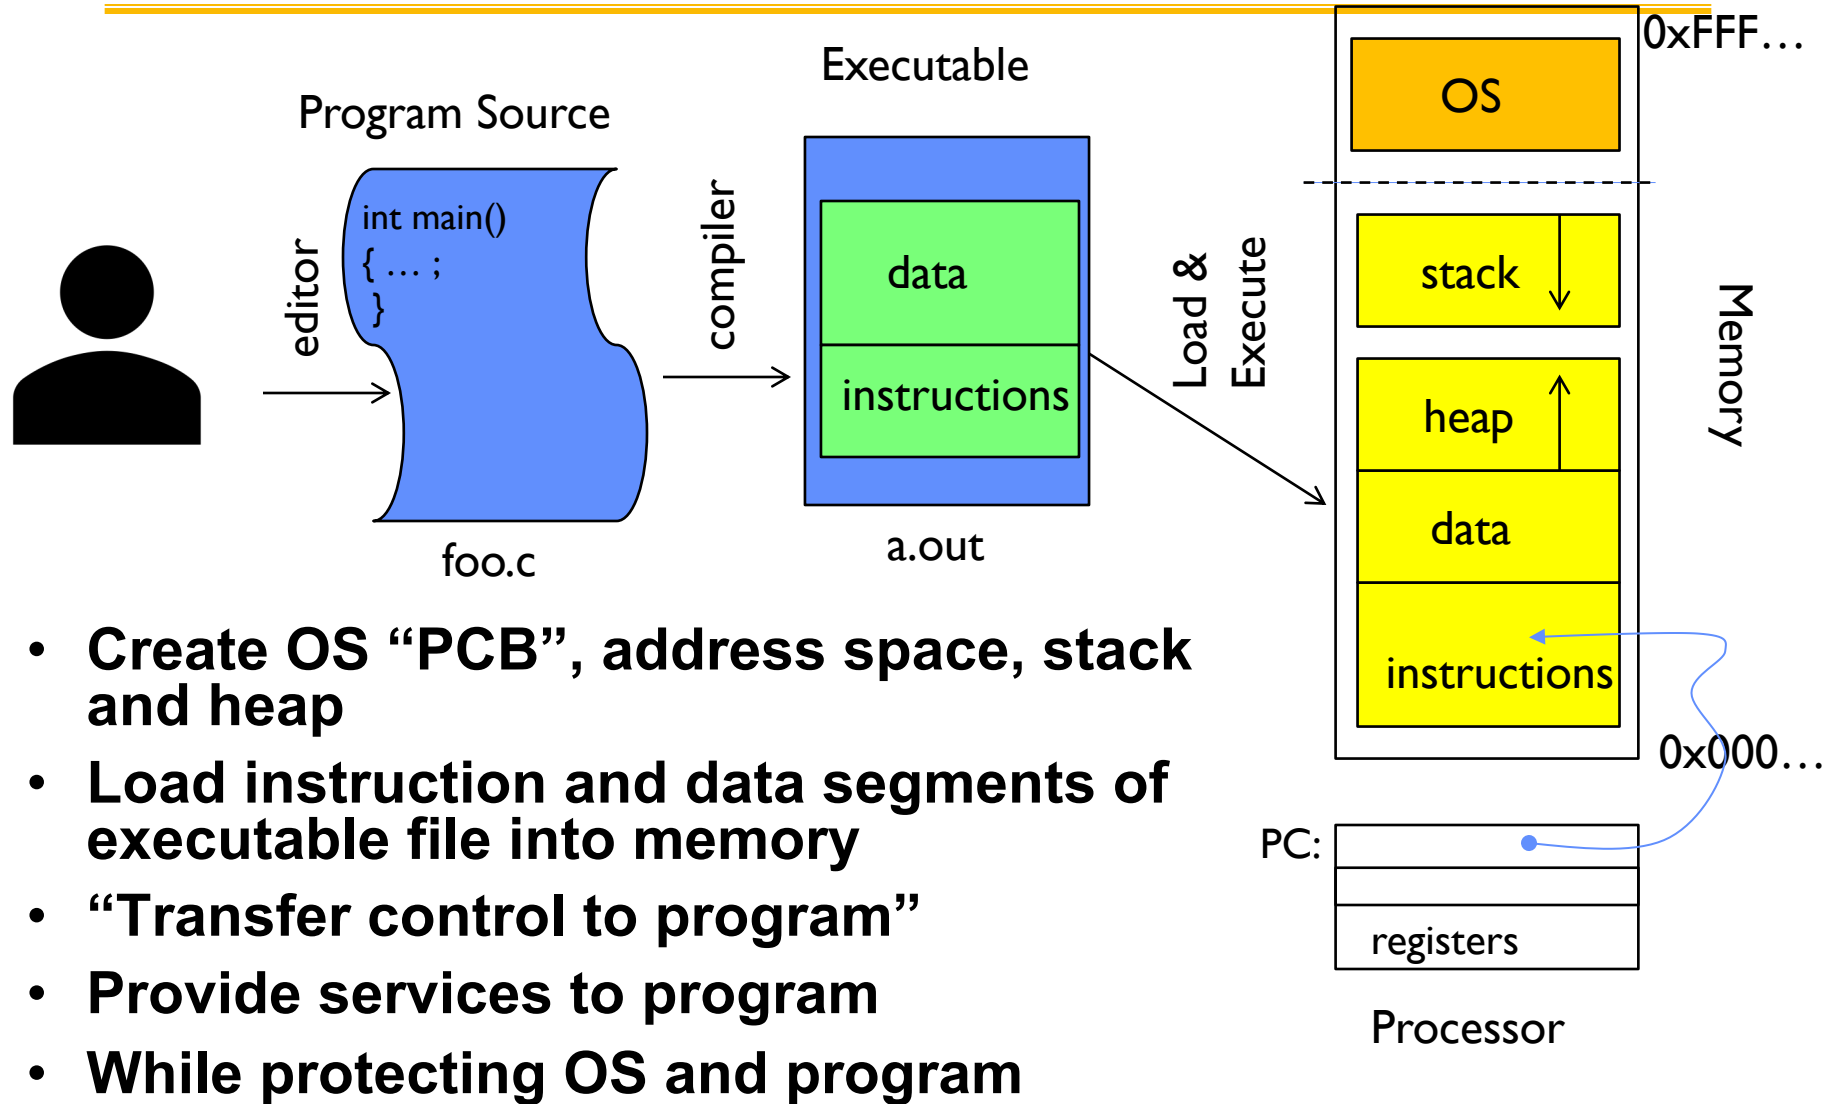
---



- **Thread: Execution Context**
  - Program Counter, Registers, Execution Flags, Stack
- **Virtual Address space**
  - Program's view of memory is distinct from physical machine
- **Process: an instance of a running program**
  - Address Space + One or more Threads + ...
- **Dual mode operation / Protection**
  - Only the “system” can access certain resources
  - Combined with translation, isolates programs from each other



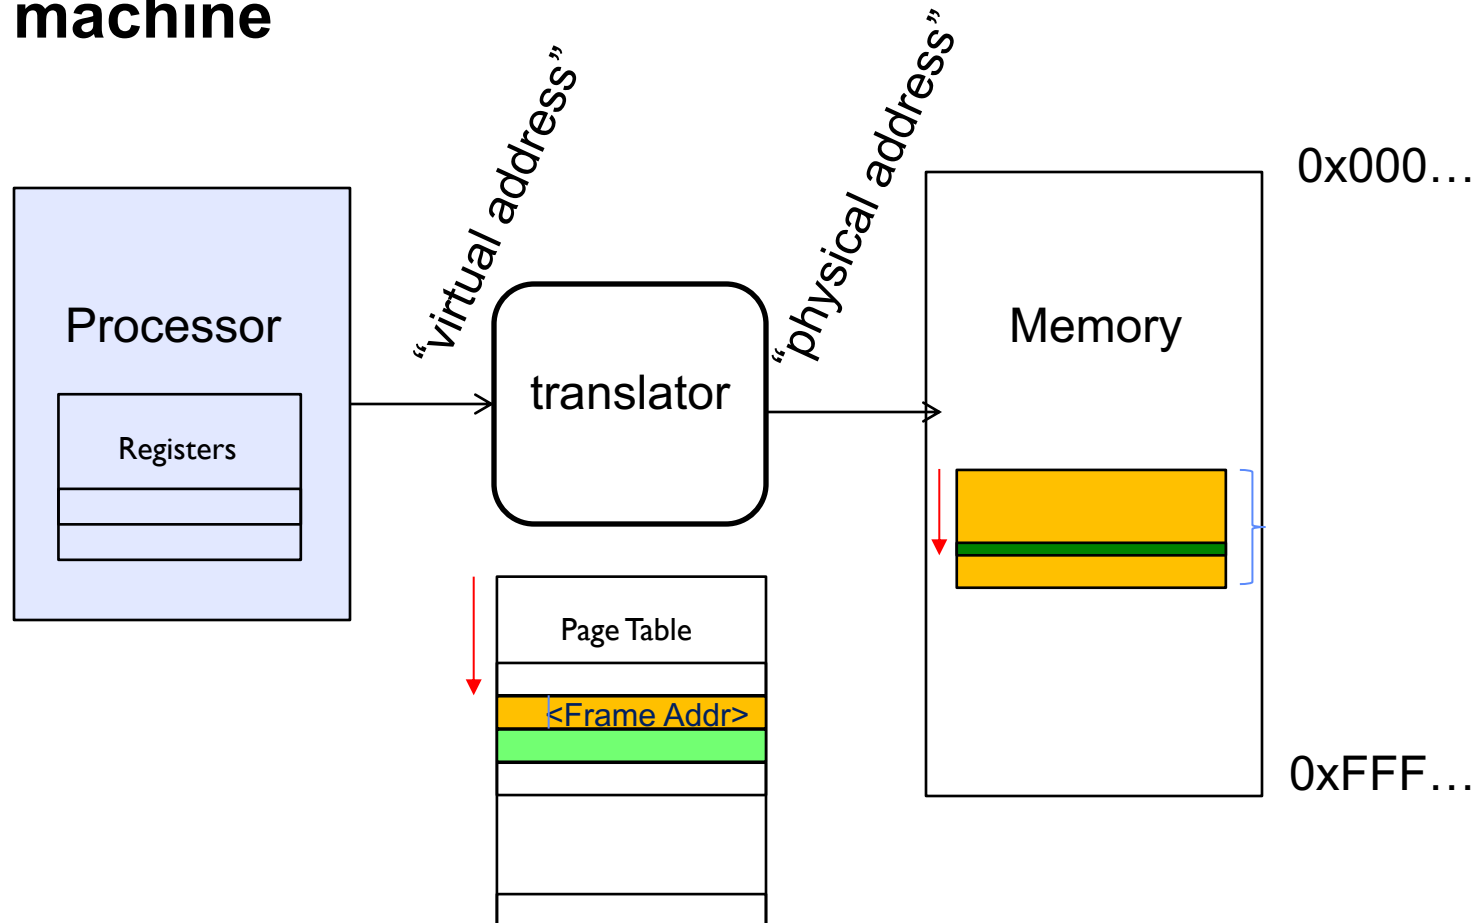
# Recall: OS Bottom Line: Run Programs





# Key OS Concept: Address Space

- Program operates in an address space that is distinct from the physical memory space of the machine





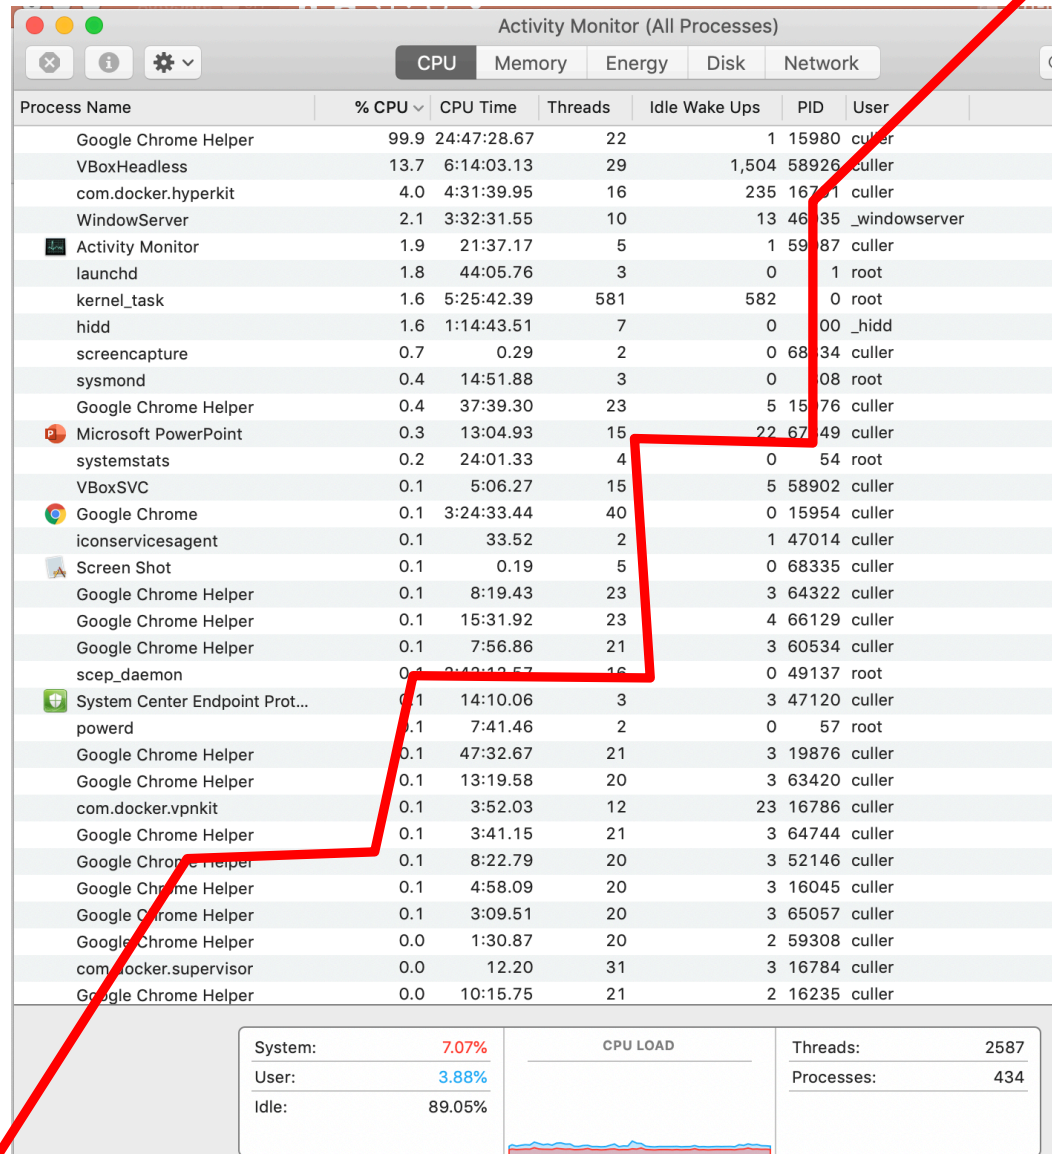
# Recall: The Process

---

- **Definition: execution environment with restricted rights**
  - **Address Space with One or More Threads**
    - » *Page table per process!*
  - Owns memory (mapped pages)
  - Owns file descriptors, file system context, ...
  - Encapsulates one or more threads sharing process resources
- **Application program executes as a process**
  - Complex applications can fork/exec child processes [later]
- **Why processes?**
  - Protected from each other. OS Protected from them.
  - Execute concurrently [ trade-offs with threads? later ]
  - Basic unit OS deals with



# What's beneath the Illusion?





---

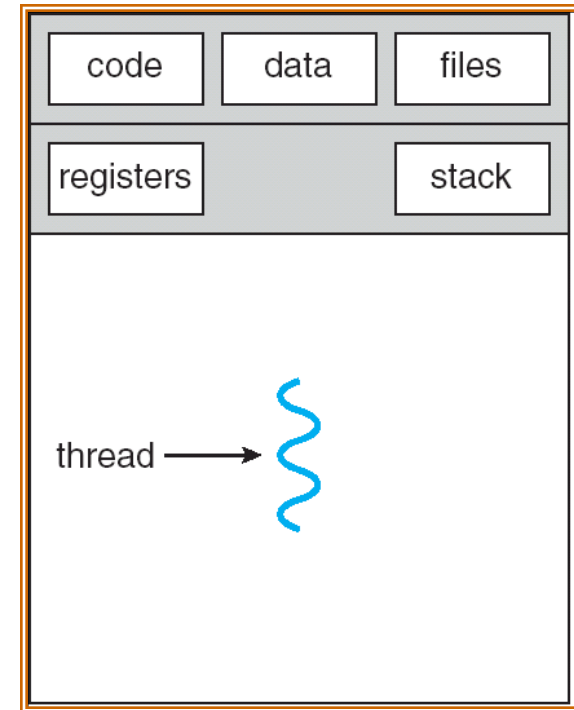
# Today: How does the Operating System create the Process Abstraction

- **What data structures are used?**
- **What machine structures are employed?**
  - Focus on x86, since will use in projects (and everywhere)



# Starting Point: Single Threaded Process

- **Process: OS abstraction of what is needed to run a single program**
  1. **Sequential program execution stream**
    - » Sequential stream of execution (thread)
    - » State of CPU registers
  2. **Protected resources**
    - » Contents of Address Space
    - » I/O state (more on this later)

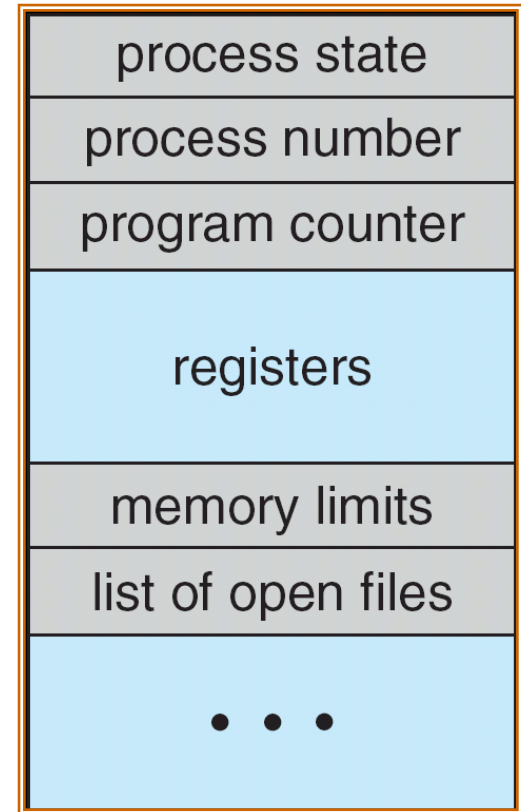






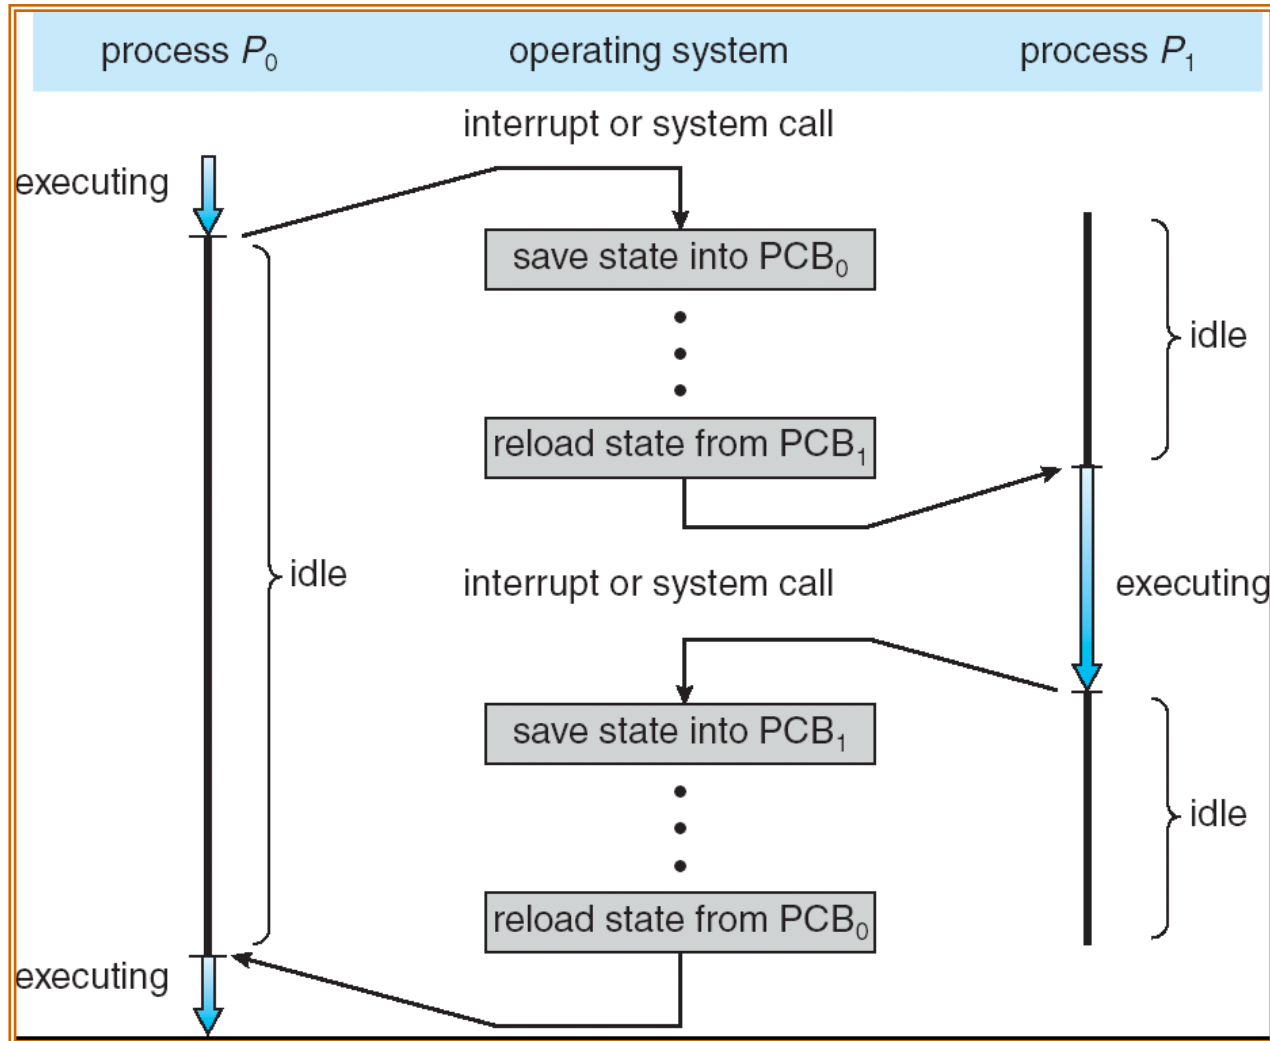
# Multiplexing Processes

- **Snapshot of each process in its PCB**
  - Only one active at a time (per core...)
- **Give out CPU to different processes**
  - Scheduling
  - Policy Decision
- **Give out non-CPU resources**
  - Memory/IO
  - Another policy decision



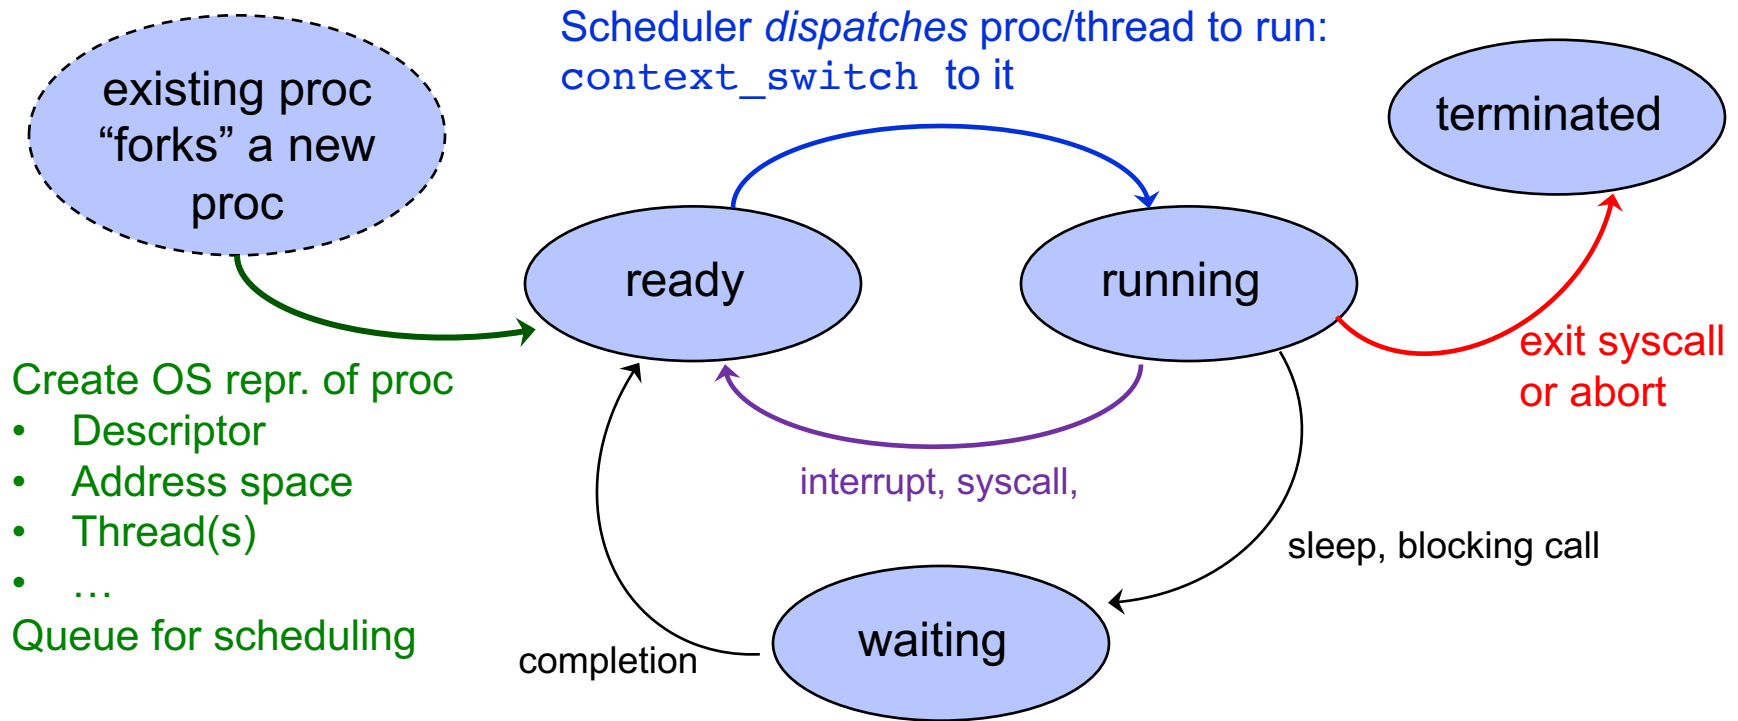
Process  
Control  
Block

# Context Switch





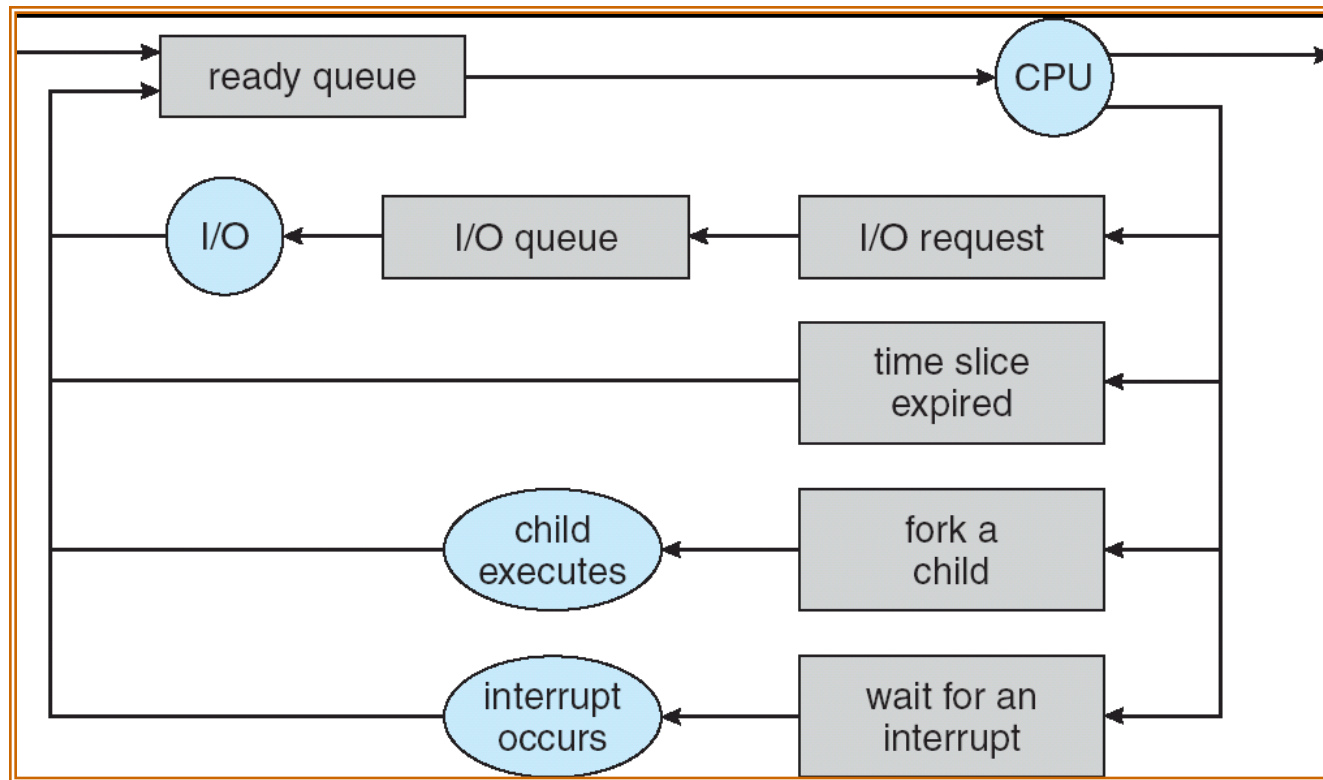
# Lifecycle of a process / thread



- OS juggles many process/threads using kernel data structures
- Proc's may create other process (fork/exec)
  - All starts with init process at boot

Pintos: process.c

# Scheduling: All About Queues

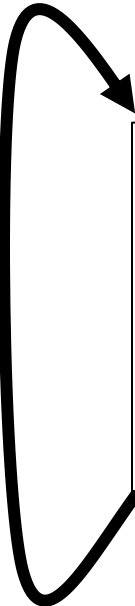


- **PCBs move from queue to queue**
- **Scheduling: which order to remove from queue**
  - Much more on this soon



# Recall: Scheduler

---

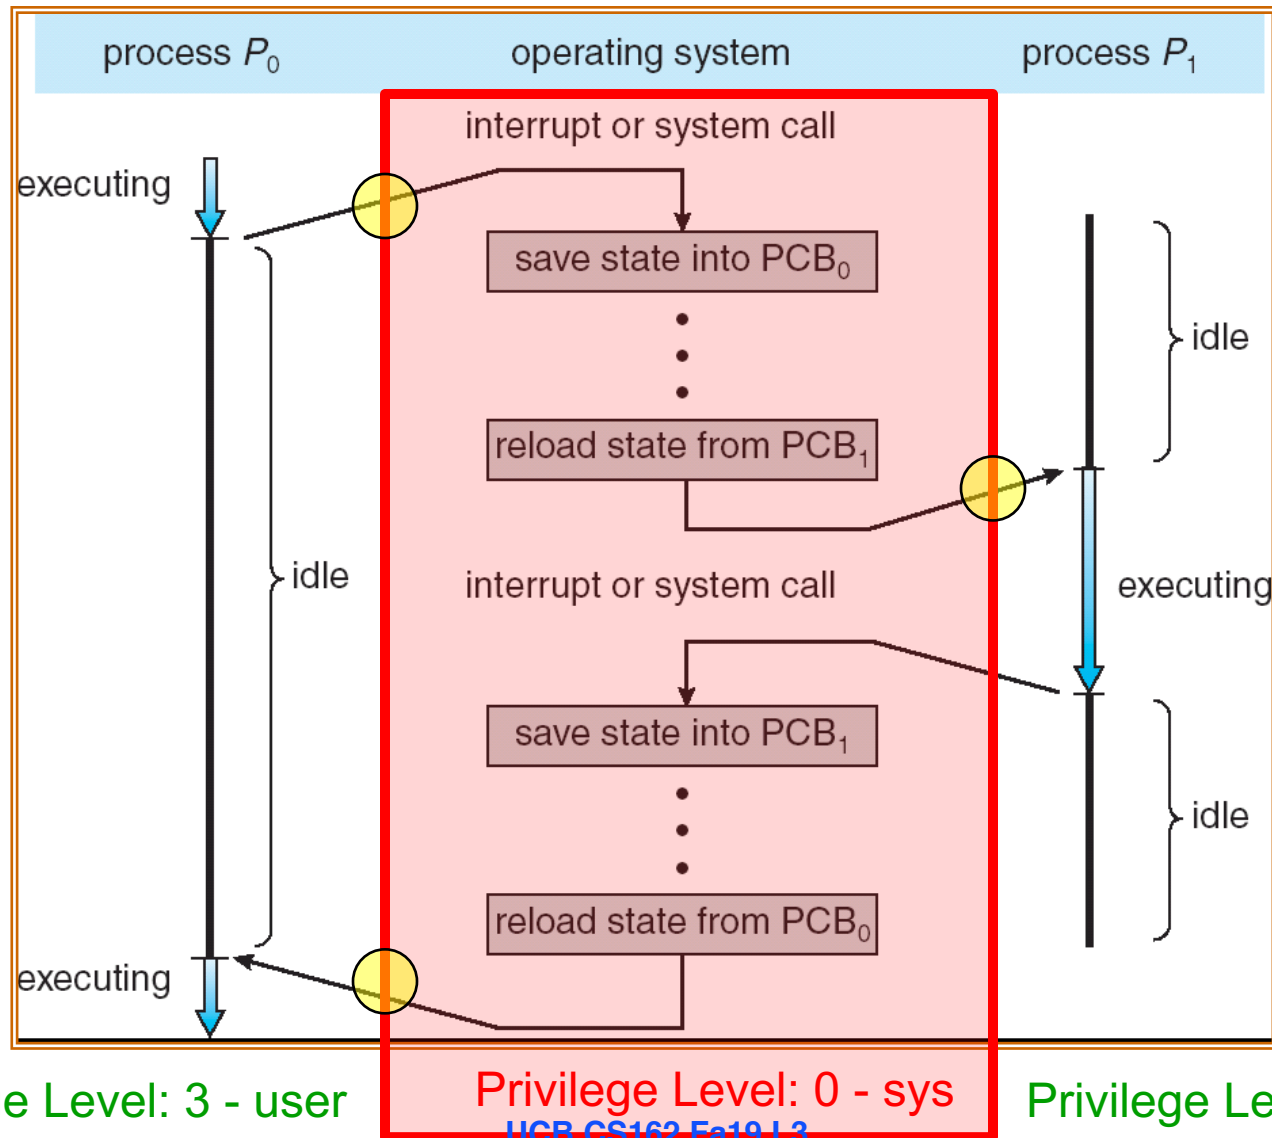


```
if ( readyProcesses(PCBs) ) {  
    nextPCB = selectProcess(PCBs);  
    run( nextPCB );  
} else {  
    run_idle_process();  
}
```

- **Scheduling: Mechanism for deciding which processes/threads receive the CPU**
- **Lots of different scheduling policies provide ...**
  - Fairness or
  - Realtime guarantees or
  - Latency optimization or ..



# Context Switch





# Process Control Block

---

- **Kernel representation of each process**
  - Status (running, ready, blocked)
  - Register state (if not running)
  - Thread control block(s)
  - Process ID
  - Execution time
  - Address space
  - Open files, etc
- **Scheduler maintains a data structure of PCBs**
- ***Scheduling algorithm*: Which process should the OS run next?**


How is this represented?



# Process Creation

---

- **Allocate and initialize Process object**
- **Allocate and initialize kernel thread mini-stack and associated Thread object**
- **Allocate and initialize page table for process**
- **Load code and static data into user pages**
- **Build initial User Stack**
  - Initial register contents
- **Schedule (post) process thread for execution**
- ...
- **Eventually switch to user thread ...**
- **Several lists of various types**



Pintos: process.c, thread.c





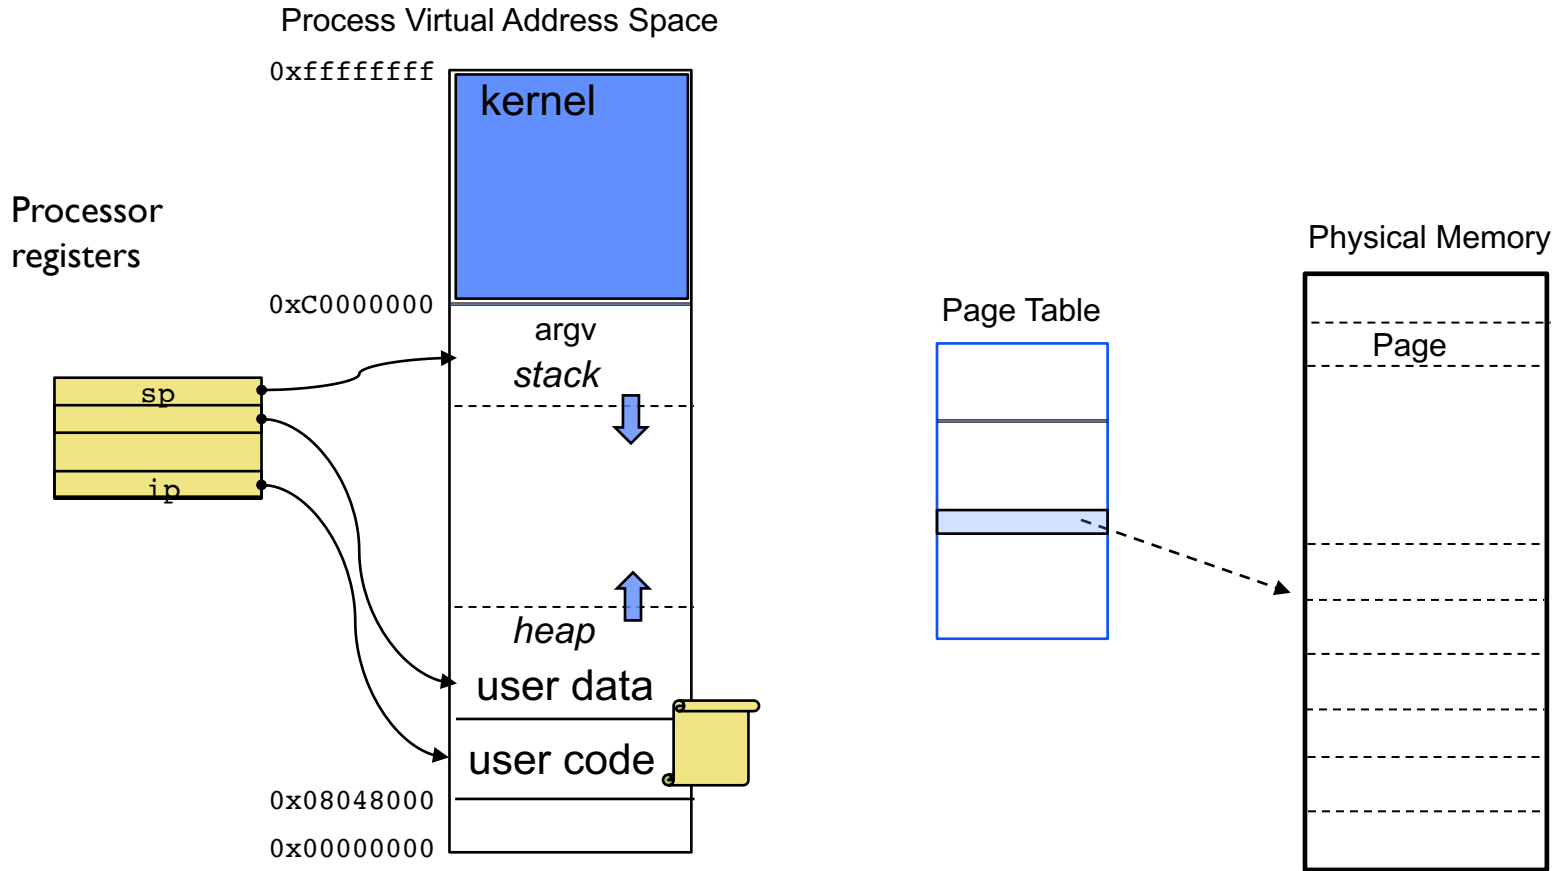
# Understanding "Address Space"

---

- **Page table is the primary mechanism**
  - **Privilege Level determines which regions can be accessed**
    - Which entries can be used
  - **System (PL=0) can access all, User (PL=3) only part**
  - **Each process has its own address space**
  - **The "System" part of all of them is the same**
- => All system threads share the same system address space and same memory**

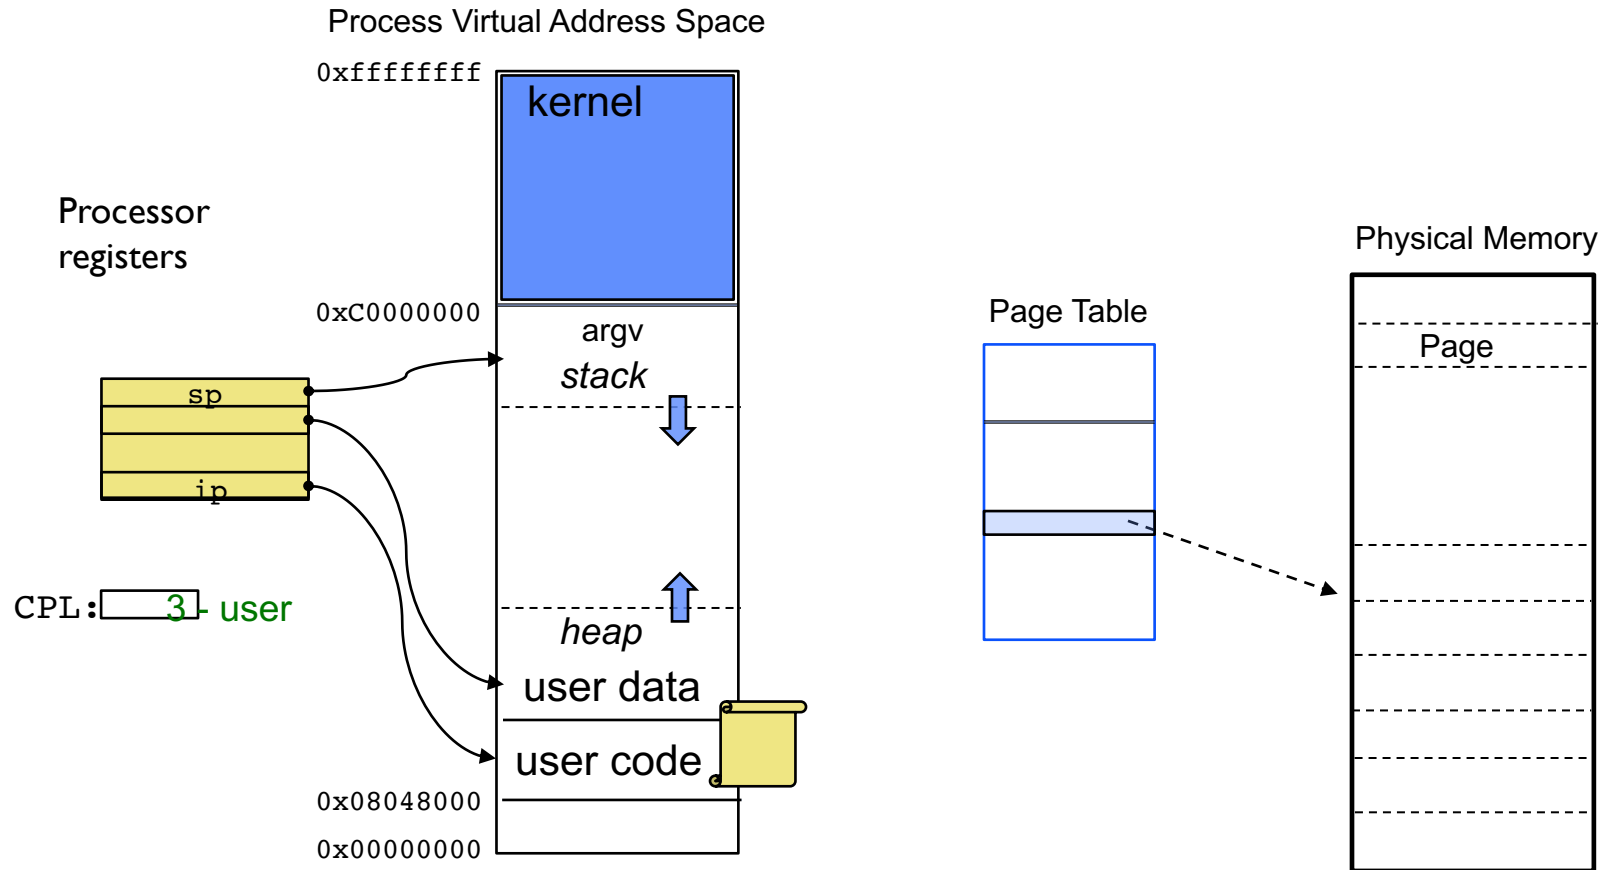


# User Process View



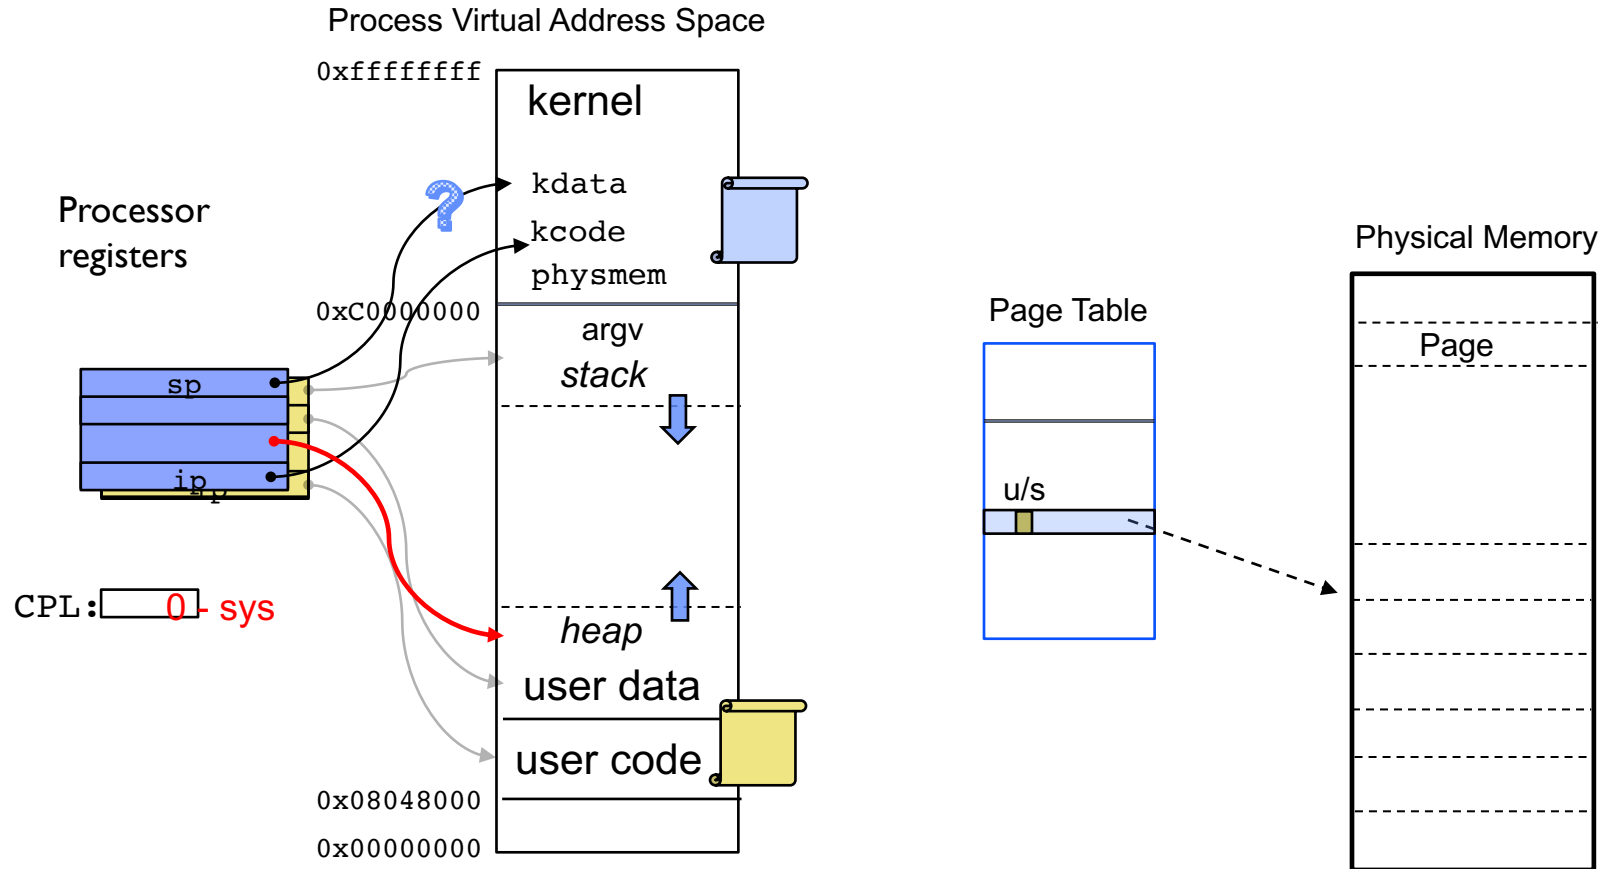


# Processor Mode (Privilege Level)



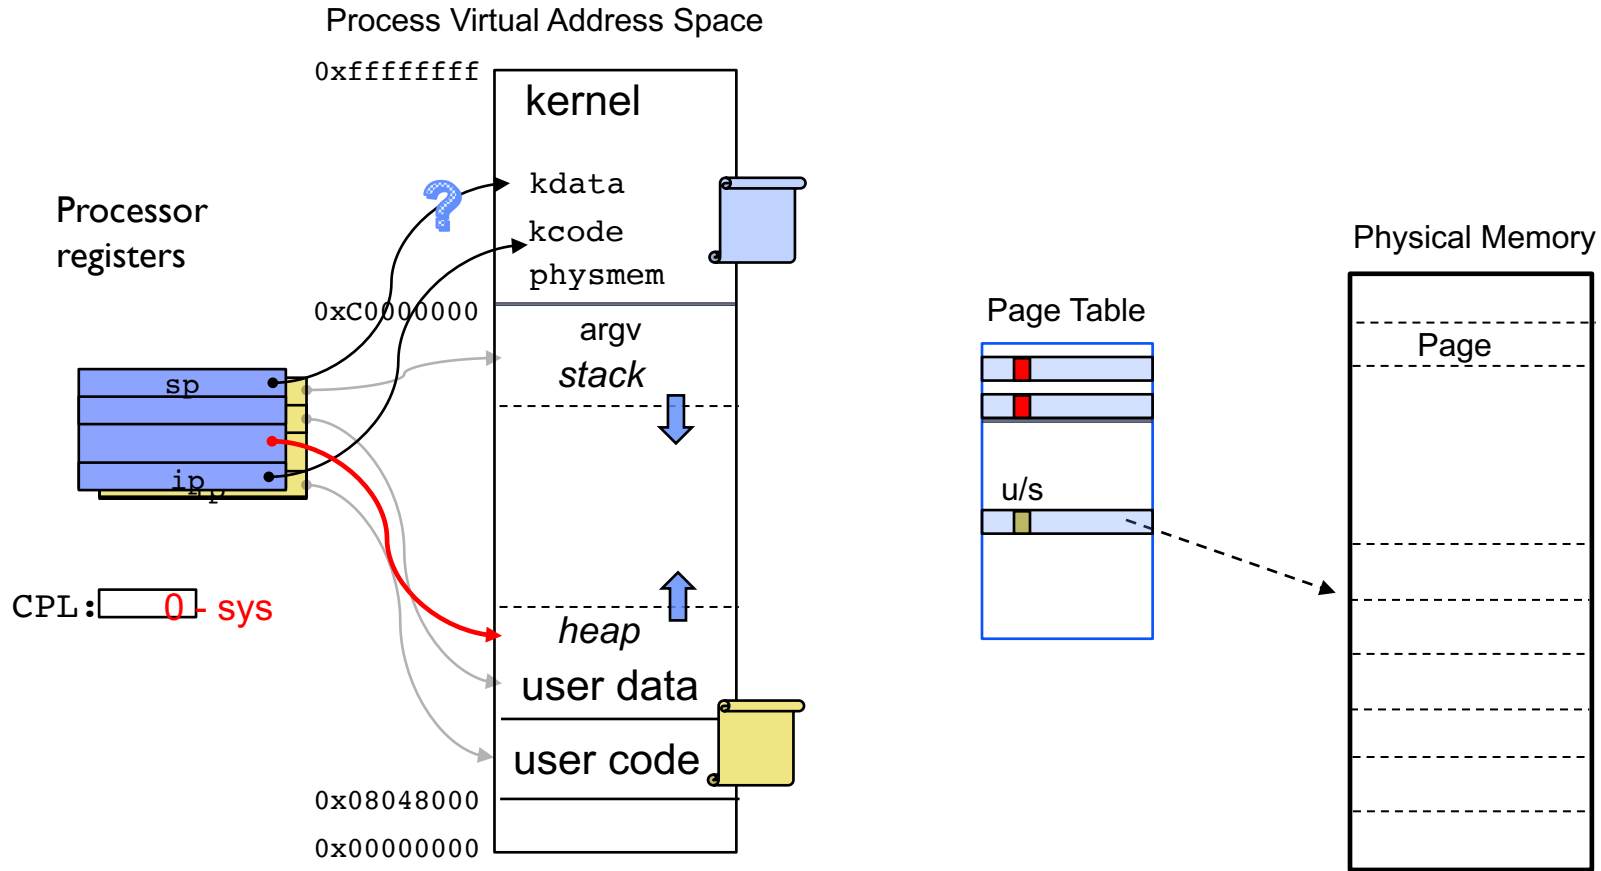


# User → Kernel: PL = 0





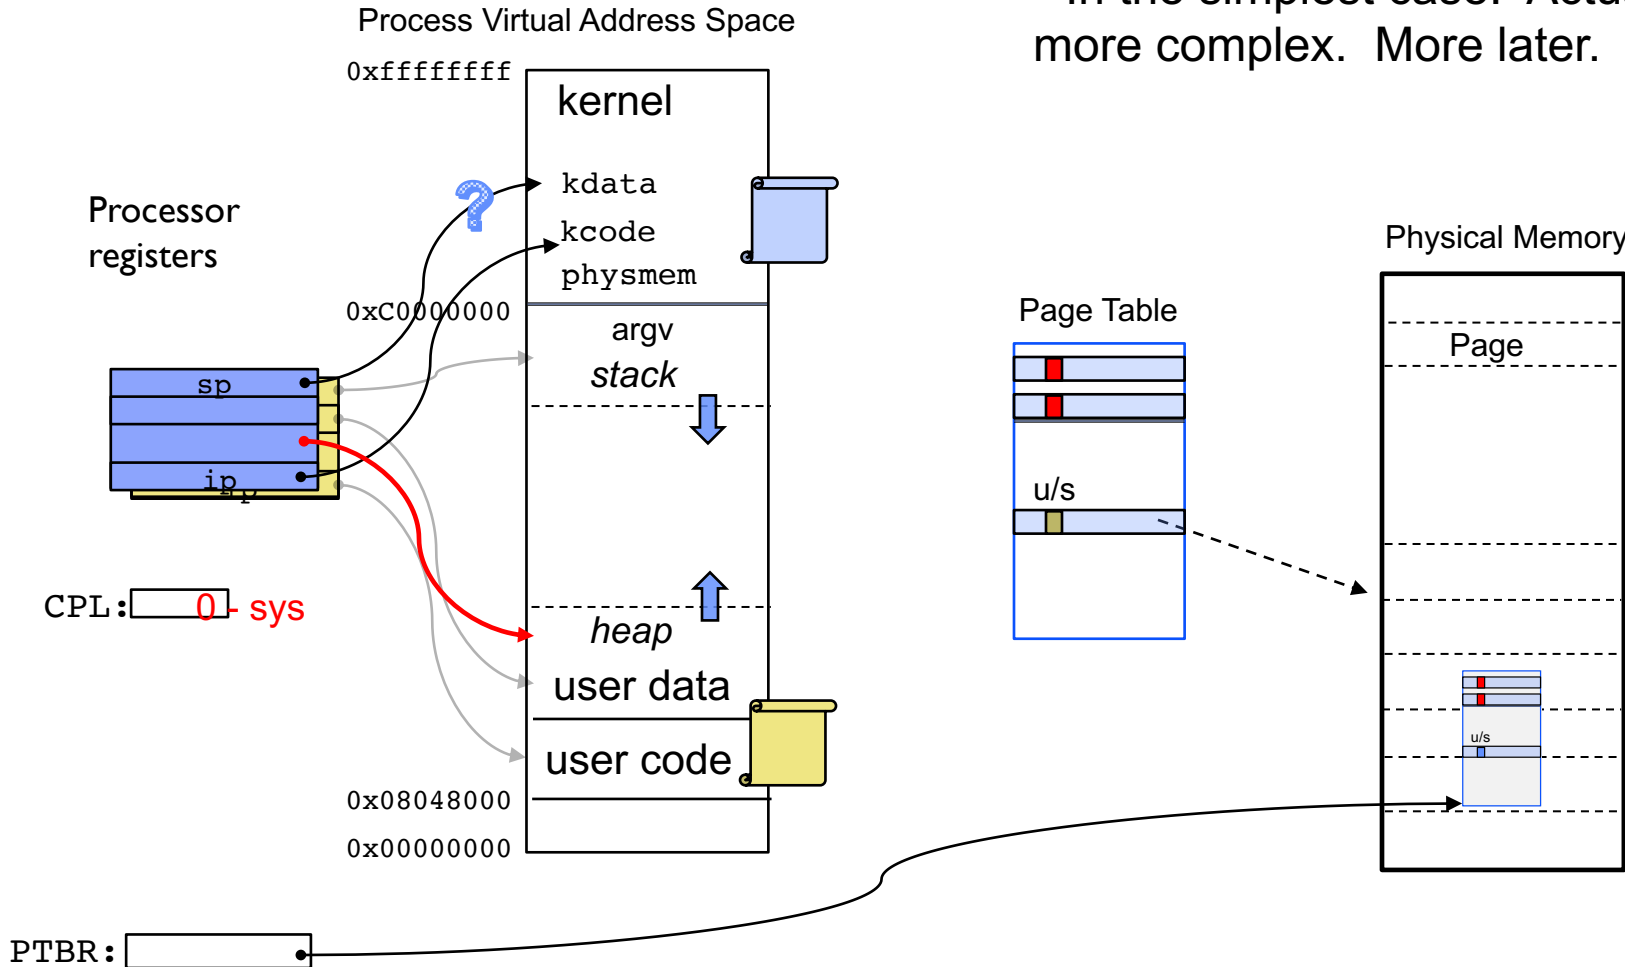
# Page Table enforces PL





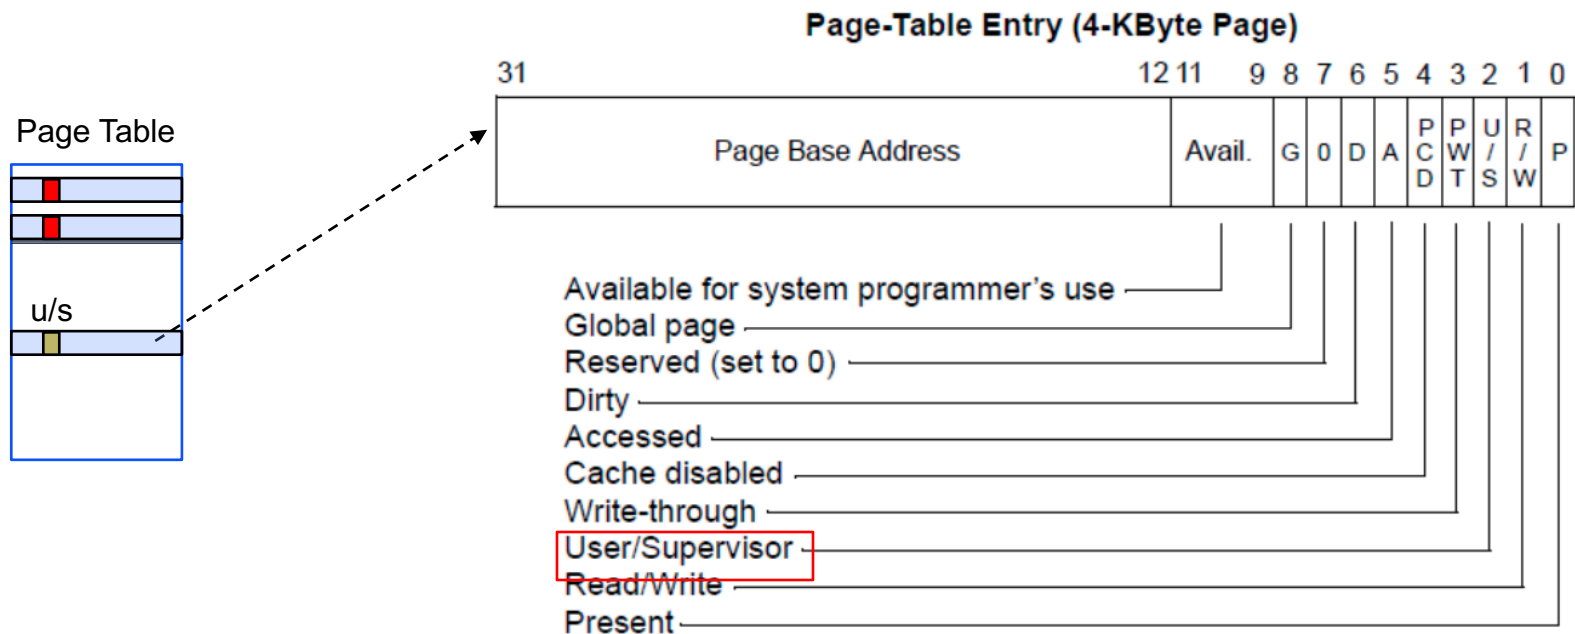
# Page Table resides in memory\*

\* In the simplest case. Actually more complex. More later.





# x86 (32-bit) Page Table Entry



- Controls many aspects of access
- Later – discuss page table organization
  - For 32 (64?) bit VAS, how large? vs size of memory?
  - Use sparsely. Very very fast HW access

Pintos: page\_dir.c



# Kernel Portion of Address Space

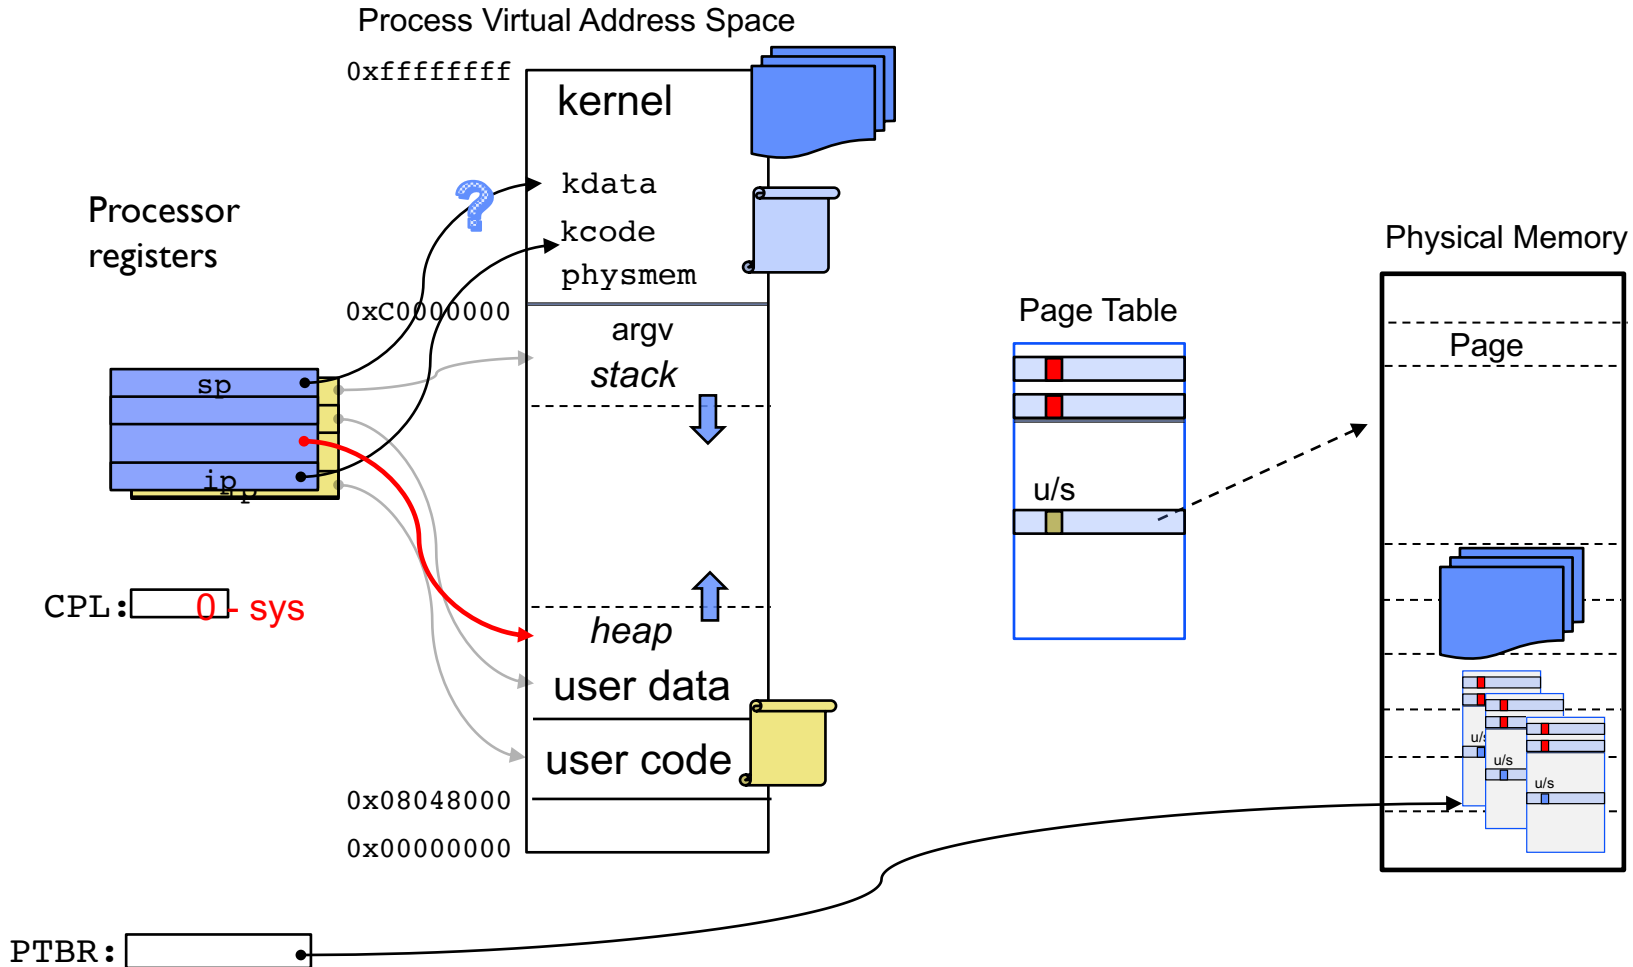
---

- **Contains the kernel code**
  - Loaded when the machine booted
- **Explicitly mapped to physical memory**
  - OS creates the page table
- **Used to contain all kernel data structures**
  - List of all the processes and threads
  - The page tables for those processes
  - Other system resources (files, sockets, ttys, ...)
- **Also contains (little) stacks for “kernel threads”**
  - Early OS design serviced all processes on a single execution thread
    - » Event driven programming
  - Today: Each Process Thread supported by (little) Kernel Thread





# 1 Kernel Code, many Kernel “stacks”



# From Machine Structure to OS Data Structures

---

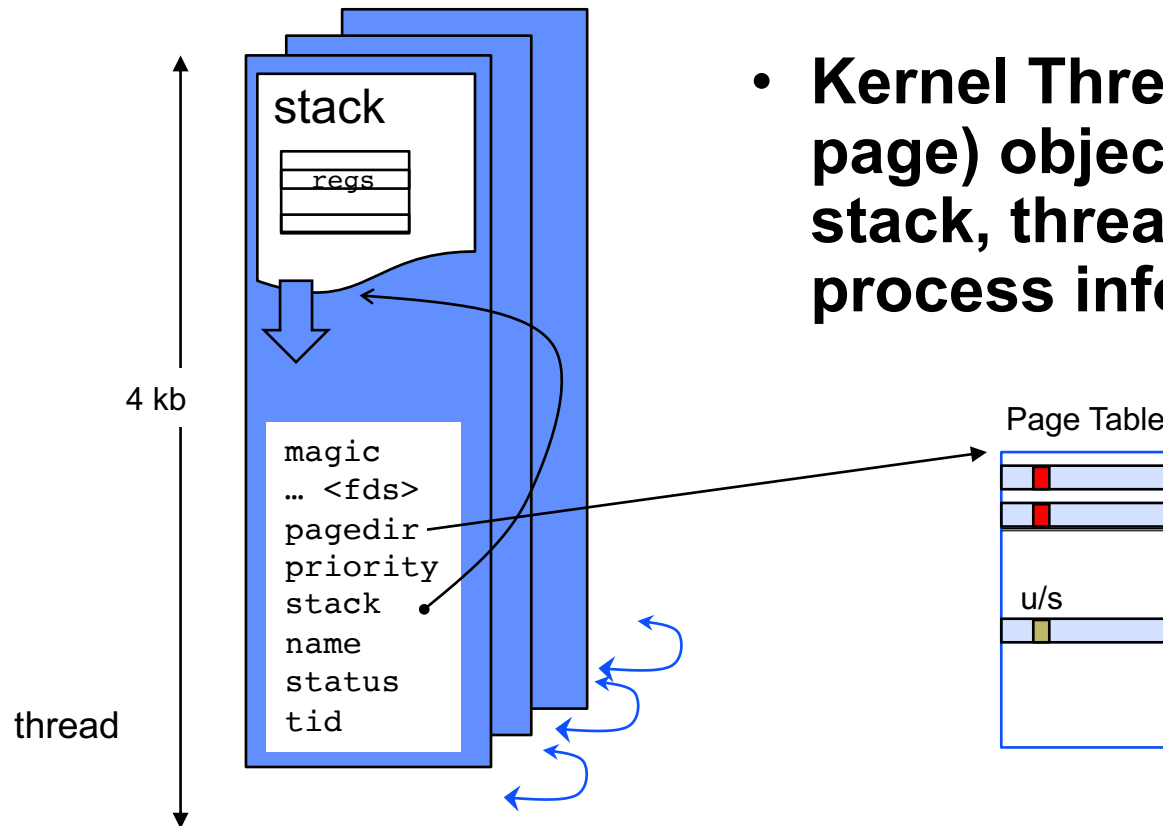


- Traditional Unix, etc. design maintains a **Process Control Block (PCB)** per process
- Each with a **Thread Control Block (TCB)** per thread of that process
- Today, assume single thread per process
  - PINTOS model
  - Linux also organized around threads with “groups of threads” associated with a process



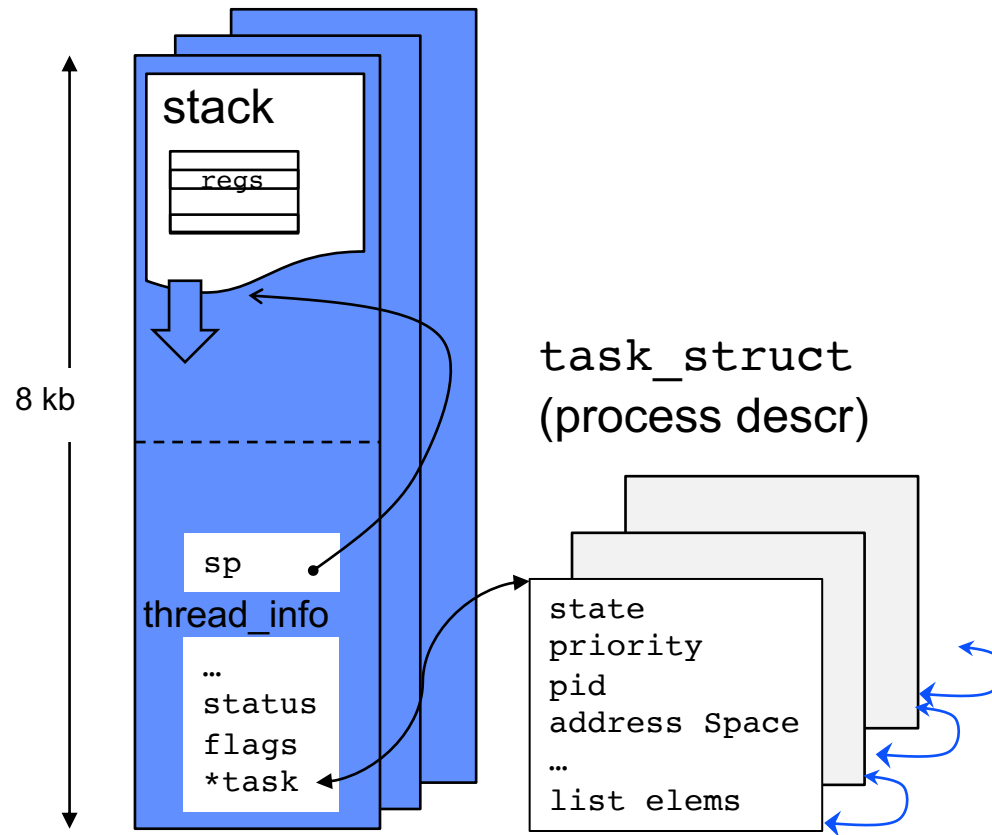
# PINTOS Thread

- Kernel Thread as 4 kb (2 page) object containing stack, thread info, process info



Pintos: thread.c

# Linux “Task”



- **Kernel Thread as 8 kb (2 page) object containing stack and thread information + process descriptor**



# Process Creation

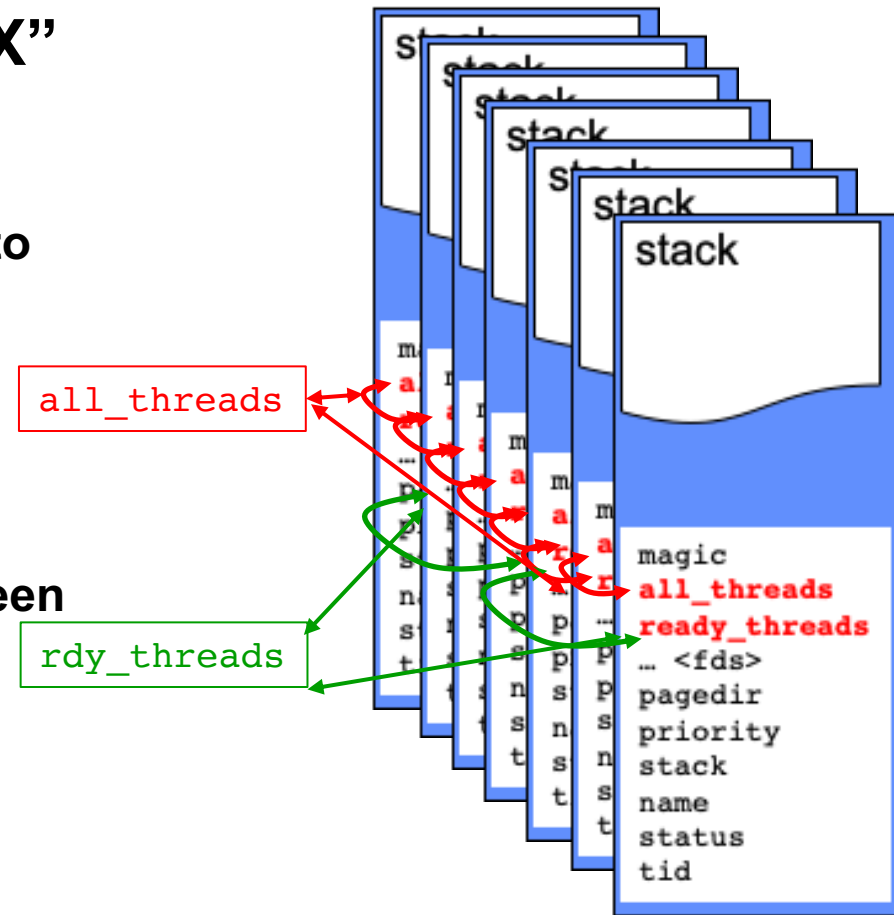
---

- **Allocate and initialize Process object**
- **Allocate and initialize kernel thread mini-stack and associated Thread object**
- **Allocate and initialize page table for process**
  - Referenced by process object
- **Load code and static data into user pages**
- **Build initial User Stack**
  - Initial register contents, argv, ...
- **Schedule (post) process thread for execution**
- ...
- **Eventually *switch* to user thread ...**
  
- **Several lists of various types**



# Aside: Polymorphic lists in C


- Many places in the kernel need to maintain a “list of X”
  - This is tricky in C, which has no polymorphism
  - Essentially adding an *interface* to a package (ala Go)
- In Linux and Pintos this is done by embedding a `list_elem` in the struct
  - Macros allow shift of view between object and list
  - You’ll practice in HW1 – before getting into PINTOS



Pintos: list.c

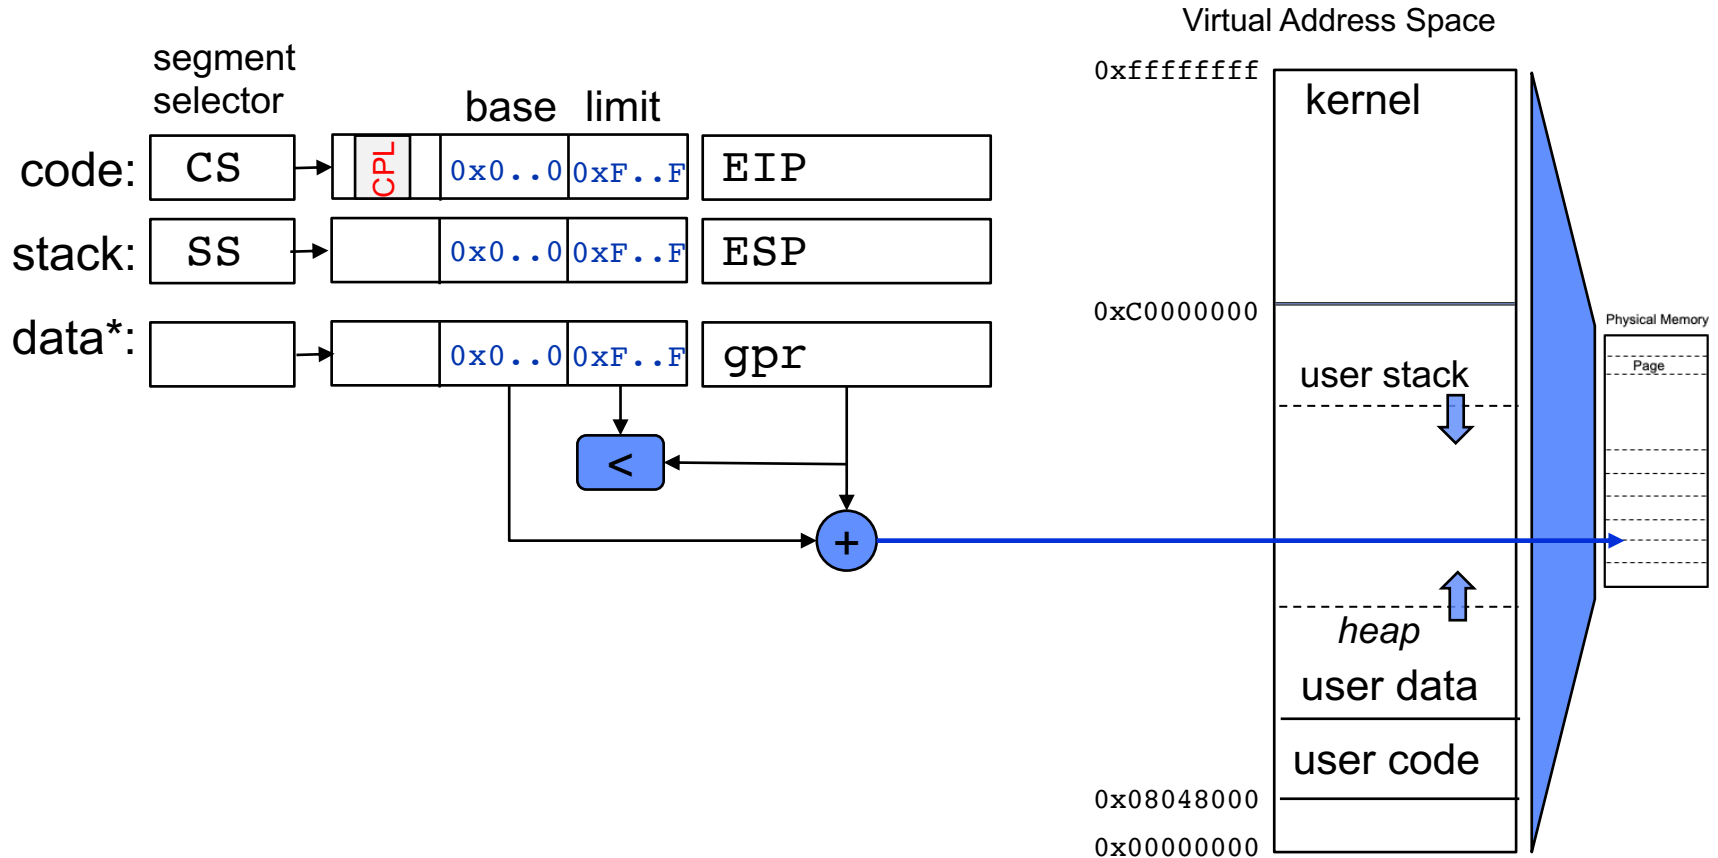


# Bit of x86 thread/process/VAS management





# Bit of x86 thread/process/VAS management



Pintos: loader.h






# Recall: 3 types of U→K Mode Transfer

---

- **Syscall**
  - Process requests a system service, e.g., exit
  - Like a function call, but “outside” the process
  - Does not have the address of the system function to call
  - Like a Remote Procedure Call (RPC) – for later
  - Marshall the syscall id and args in registers and exec syscall
- **Interrupt**
  - External asynchronous event triggers context switch
  - eg. Timer, I/O device
  - Independent of user process
- **Trap**
  - Internal synchronous event in process triggers context switch
  - e.g., Protection violation (segmentation fault), Divide by zero, ...
- **All 3 *exceptions* are an UNPROGRAMMED CONTROL TRANSFER**
  - Where does it go? (To handler specified in interrupt vector)
  - Are interrupts enabled or disabled when get there?



# Stack for Thread Transition



**Cyan** = User Stack; **Red** = Kernel Stack

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping(); /* Do any cleanup */  
}
```

**Scheduling: Policy Decision**



# Hardware context switch support

- **Syscall/Intr (U  $\rightarrow$  K)**
  - PL 3  $\rightarrow$  0;
  - TSS  $\leftarrow$  EFLAGS, CS:EIP;
  - SS:SP  $\leftarrow$  k-thread stack (TSS PL 0);
  - push (old) SS:ESP onto (new) k-stack
  - push (old) eflags, cs:eip, <err>
  - CS:EIP  $\leftarrow$  <k target handler>
- **Then**
  - *Handler then saves other regs, etc*
  - *Does all its works, possibly choosing other threads, changing PTBR (CR3)*
  - kernel thread has set up user GPRs
- **iret (K  $\rightarrow$  U)**
  - PL 0  $\rightarrow$  3;
  - Eflags, CS:EIP  $\leftarrow$  popped off k-stack
  - SS:SP  $\leftarrow$  user thread stack (TSS PL 3);

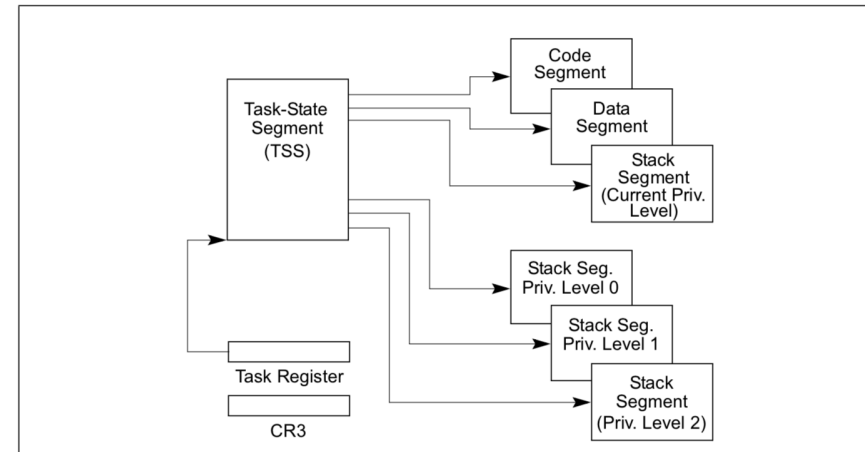
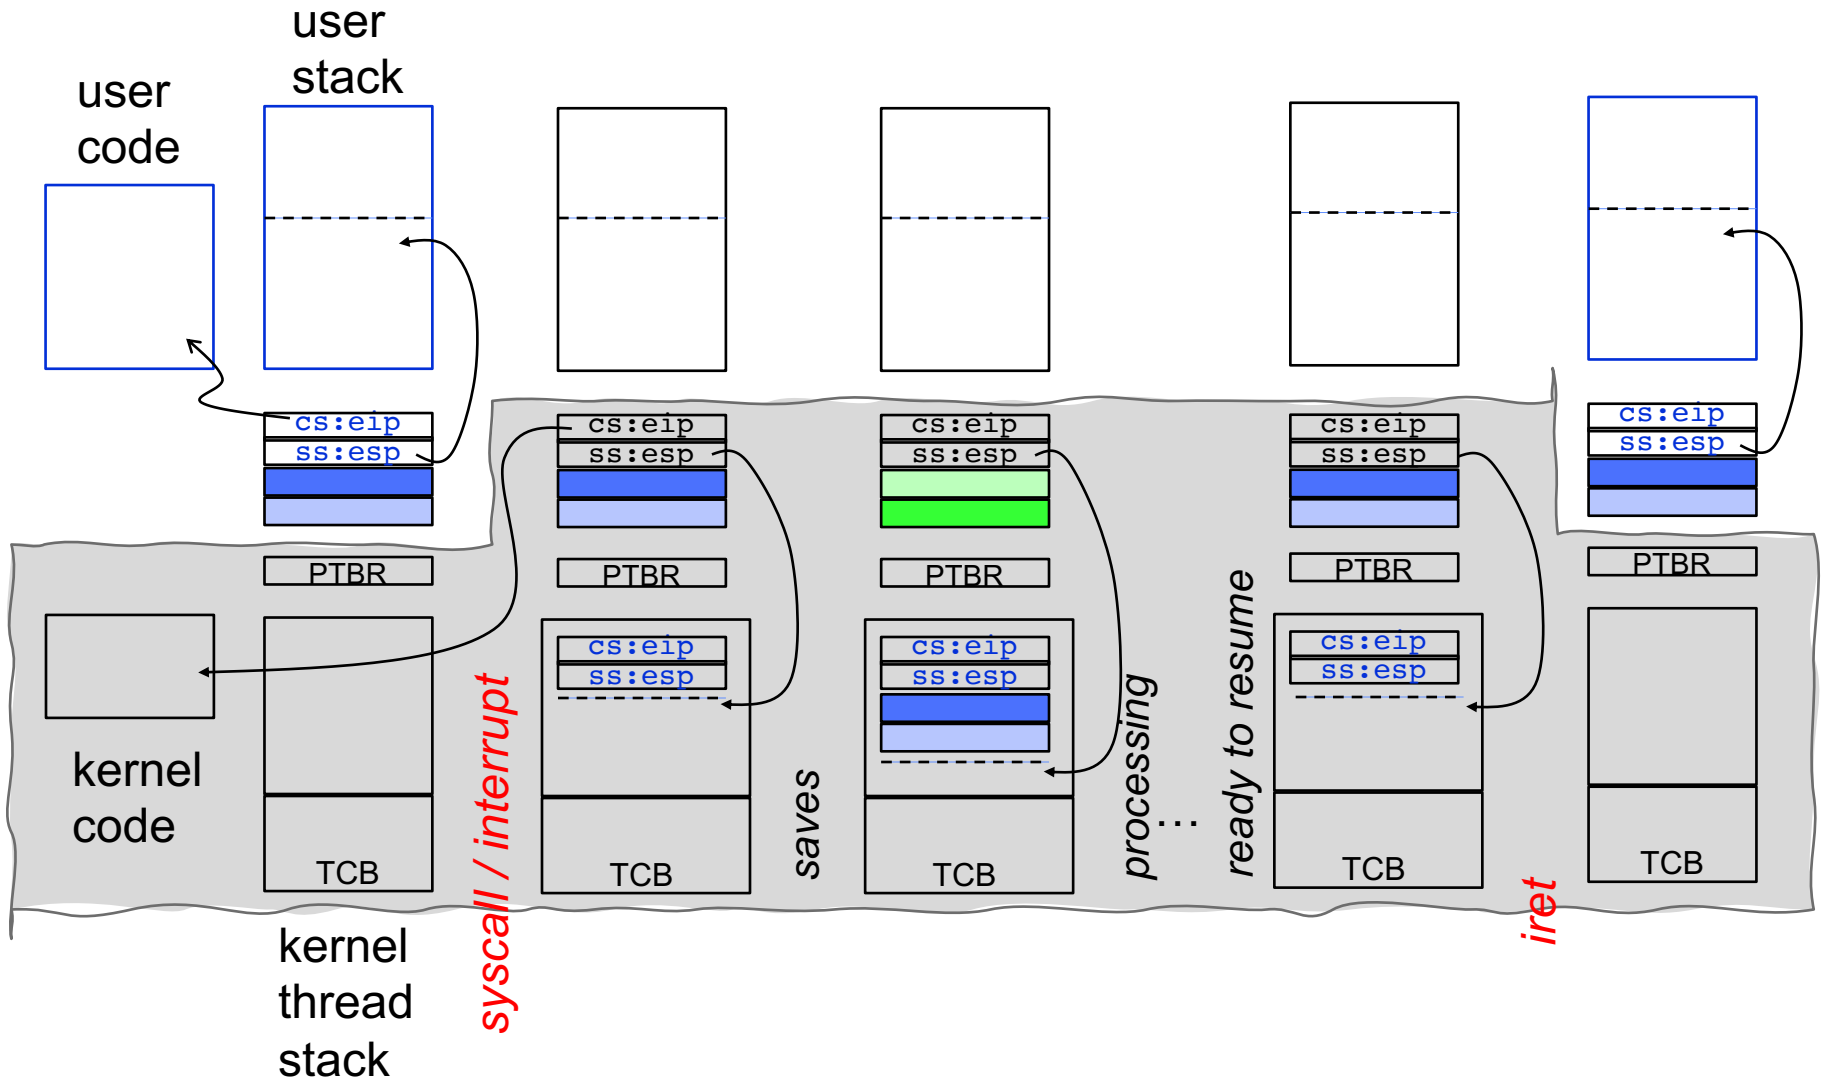


Figure 7-1. Structure of a Task

Pintos: tss.c, intr-stubs.S

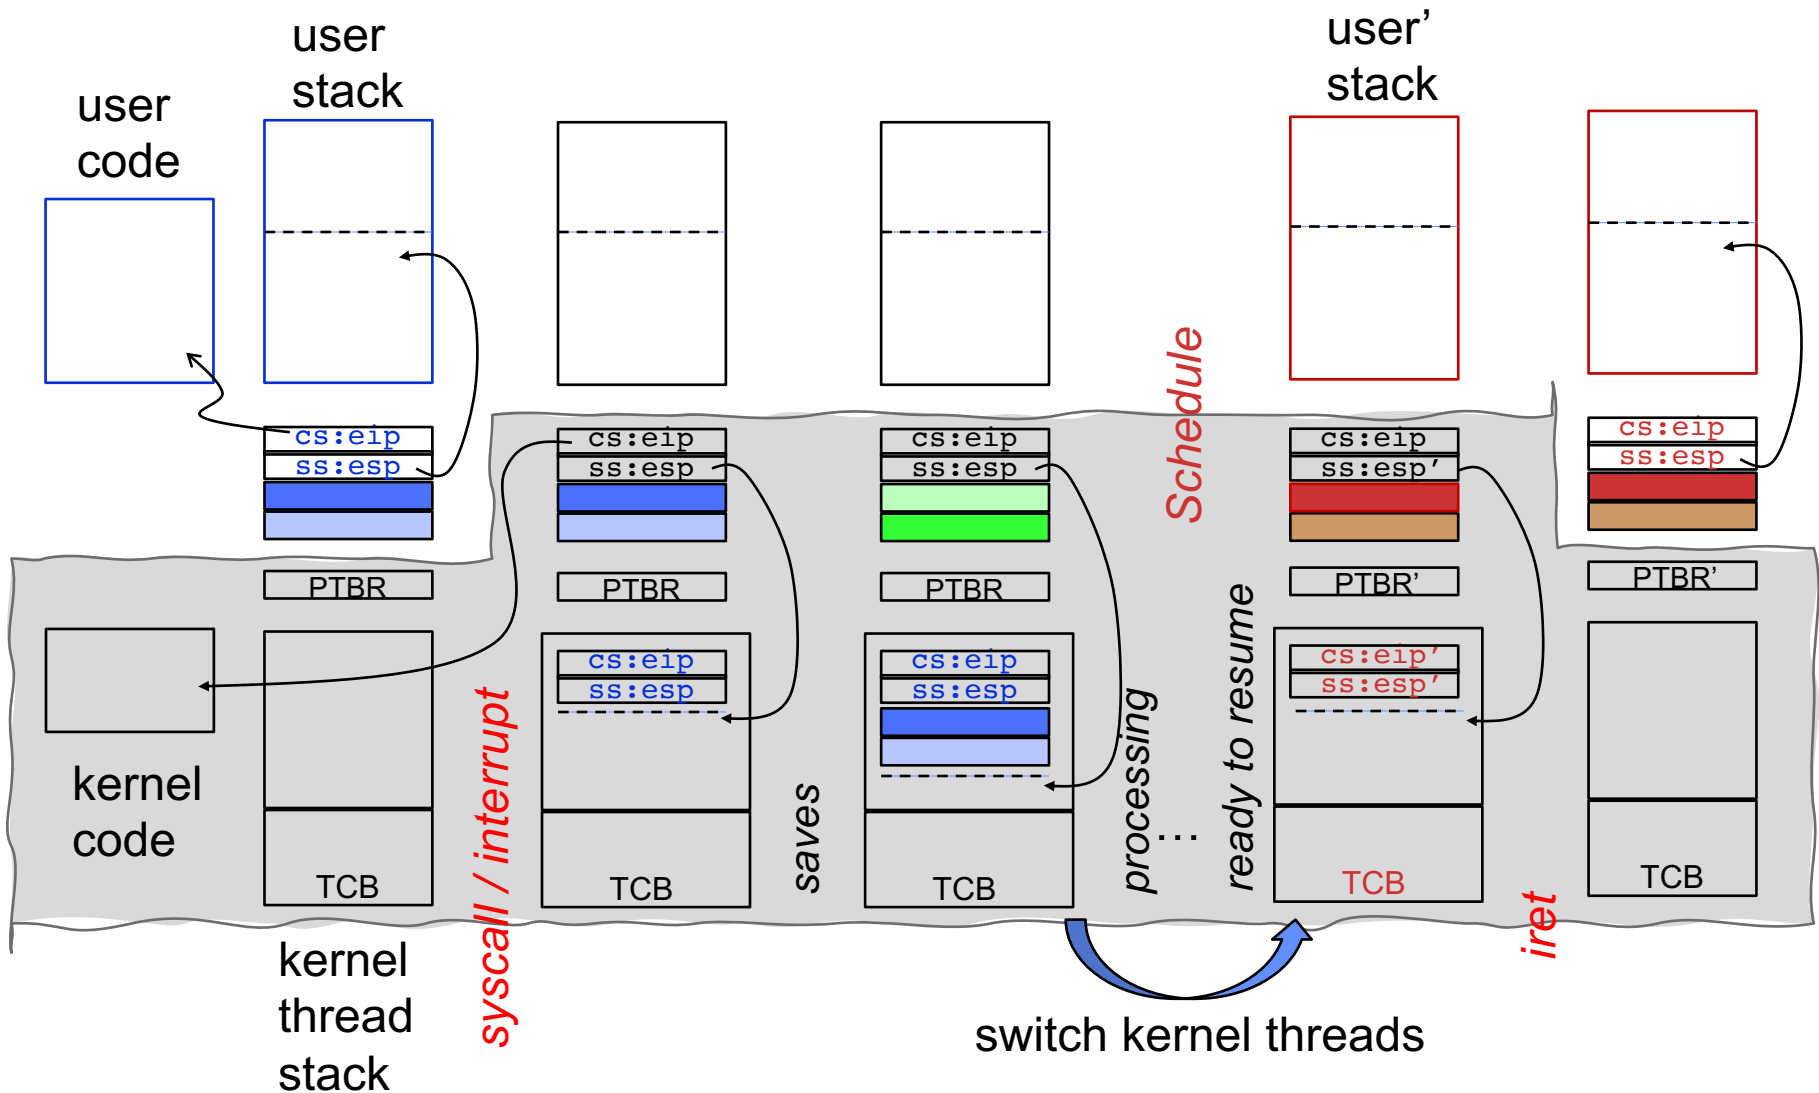


# Context Switch – in pictures





# Context Switch – Scheduling



Pintos: `switch.S`

# Context Switch between K-threads



```
Switch(tCur,tNew) {  
    /* Unload old thread */  
    TCB[tCur].regs.r7 = CPU.r7;  
    ...  
    TCB[tCur].regs.r0 = CPU.r0;  
    TCB[tCur].regs.sp = CPU.sp;  
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/  
  
    /* Load and execute new thread */  
    CPU.r7 = TCB[tNew].regs.r7;  
    ...  
    CPU.r0 = TCB[tNew].regs.r0;  
    CPU.sp = TCB[tNew].regs.sp;  
    CPU.retpc = TCB[tNew].regs.retpc;  
    return; /* Return to CPU.retpc */  
}
```



# Concurrency

---

- **But, ... ???**
- **With all these threads in the kernel, won't they step on each other?**
  - For example, while one is loading a program, other threads should run ...
  - Processes are isolated from each other, but all the threads in the kernel share the kernel address space, memory, data structures
- **We will study synchronization soon**
- **The kernel controls whether hardware interrupts are enabled or not**
  - Disabled on entry, selectively enable
  - Atomic operations, ...



# Dispatch Loop

---

```
Loop {  
    RunThread();  
    ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

- ***Conceptually*** all the OS executes
- **Infinite Loop**
  - When would we ever "exit?"
  - Can we assume some thread is always ready?





# Dispatch Loop

---

```
Loop {  
    RunThread();  
    ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

## How to run a new thread?

- Load thread's registers into **CPU**
- Load its environment (address space, if in different process)
- Jump to thread's **PC**

## How does dispatch loop get control again?

- Thread returns control voluntarily – **yield**, **I/O**
- External events: thread is *preempted*



# Thread Operations in Pintos


---

- `thread_create(name, priority, func, args)`
  - Create a new thread to run `func(args)`
- `thread_yield()`
  - Relinquish processor voluntarily

*More later,  
incl. synch ops*
- `thread_join(thread)`
  - Wait (put in queue) until thread exits, then return
- `thread_exit`
  - Quit thread and clean up, wake up joiner if any

# Peer question

- Which kind of thread is performing these operations?
  - “user level thread” with its full stack and user address space?
  - “system proxy thread” for a “user level thread”





# Tout Question

---

- **Process is an instance of a program executing.**
  - The fundamental OS responsibility
- **Processes do their work by processing and calling file system operations**
- **Are there any operations on processes themselves?**
- **exit ?**



# pid.c

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[])
{
    pid_t pid = getpid();    /* get current processes PID */

    printf("My pid: %d\n", pid);

    exit(0);
}
```

ps anyone?



# Can a process create a process ?

---

- Yes
- Fork creates a copy of process
- What about the program you want to run?



# Break

---





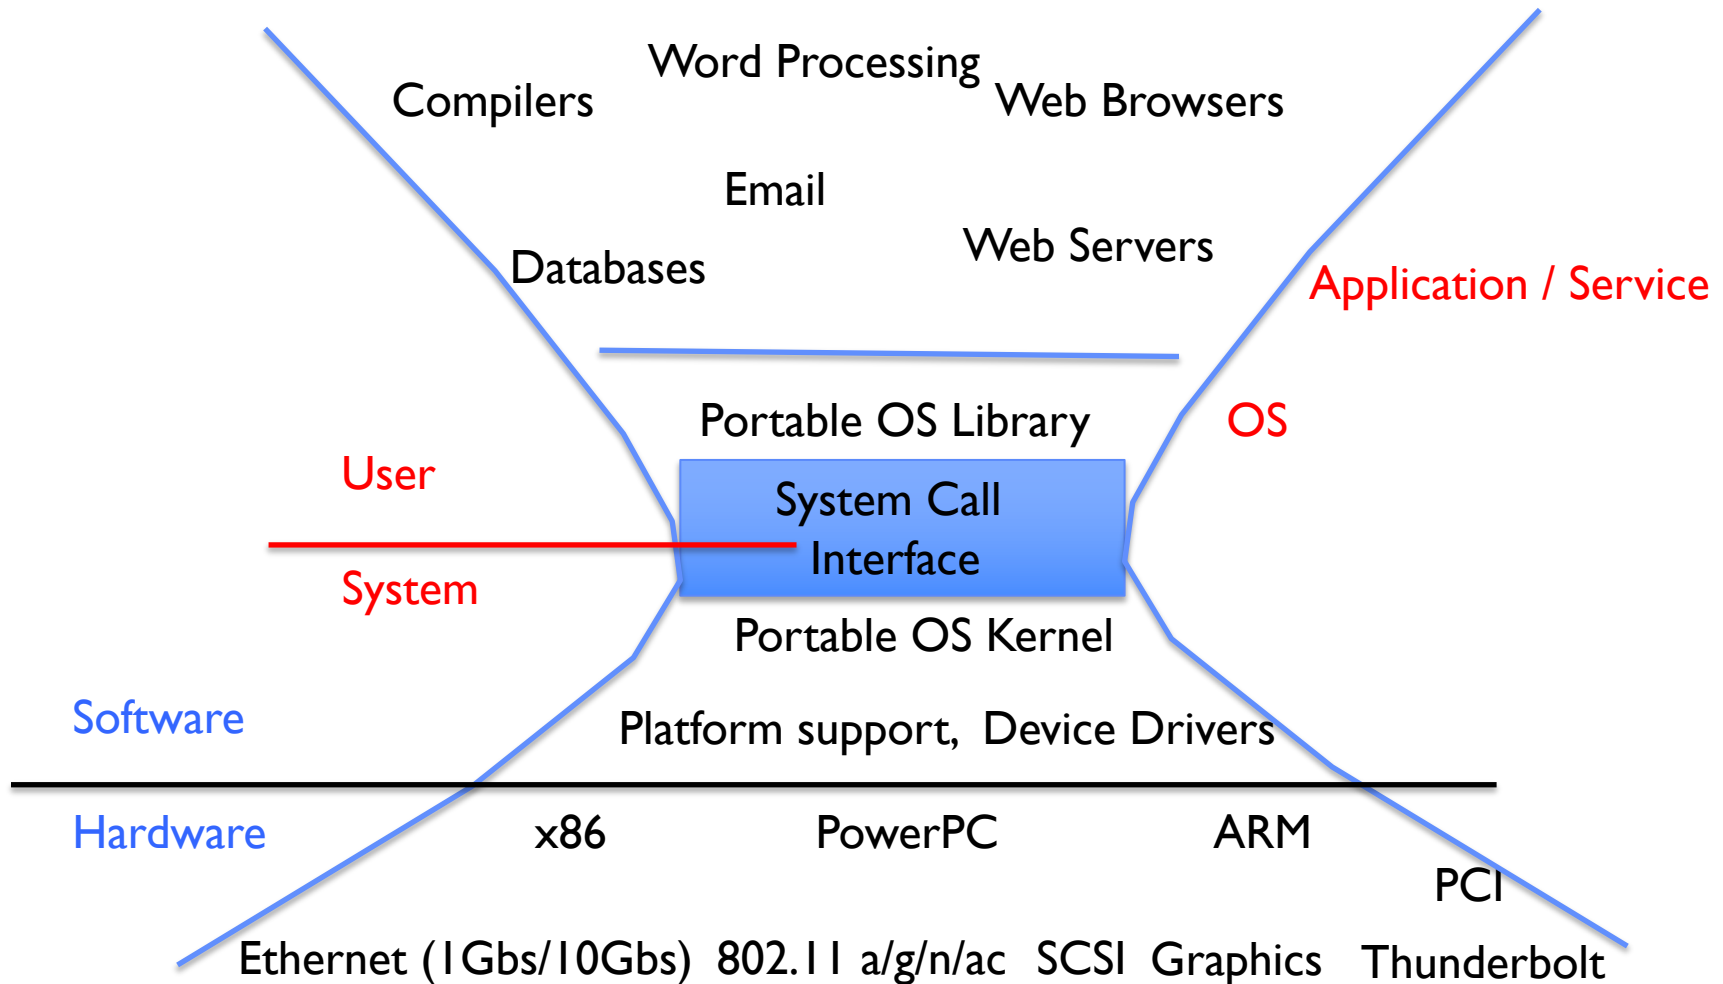
# OS Run-Time Library







# A Narrow Waist





# POSIX/Unix

---

- **Portable Operating System Interface [X?]**
- **Defines “Unix”, derived from AT&T Unix**
  - Created to bring order to many Unix-derived OSs
- **Interface for application programmers (mostly)**



# System Calls

Application:

```
fd = open(pathname);
```

Library:

```
File *open(pathname) {  
    asm code ... syscall # into ax  
    put args into registers bx, ...  
    special trap instruction
```

Operating System:

```
    get args from regs  
    dispatch to system func  
    process, schedule, ...  
    complete, resume process
```

```
    get results from regs  
};
```

Continue with results

Pintos: userprog/syscall.c, lib/user/syscall.c



# SYSCALLs (of over 300)

%eax	Name	Source	%ebx	%ecx	%edx	%esi	%edi
1	sys_exit	<a href="#">kernel/exit.c</a>	int	-	-	-	-
2	sys_fork	<a href="#">arch/i386/kernel/process.c</a>	<a href="#">struct pt_regs</a>	-	-	-	-
3	sys_read	<a href="#">fs/read_write.c</a>	unsigned int	char *	<a href="#">size_t</a>	-	-
4	sys_write	<a href="#">fs/read_write.c</a>	unsigned int	const char *	<a href="#">size_t</a>	-	-
5	sys_open	<a href="#">fs/open.c</a>	const char *	int	int	-	-
6	sys_close	<a href="#">fs/open.c</a>	unsigned int	-	-	-	-
7	sys_waitpid	<a href="#">kernel/exit.c</a>	pid_t	unsigned int *	int	-	-
8	sys_creat	<a href="#">fs/open.c</a>	const char *	int	-	-	-
9	sys_link	<a href="#">fs/namei.c</a>	const char *	const char *	-	-	-
10	sys_unlink	<a href="#">fs/namei.c</a>	const char *	-	-	-	-
11	sys_execve	<a href="#">arch/i386/kernel/process.c</a>	<a href="#">struct pt_regs</a>	-	-	-	-
12	sys_chdir	<a href="#">fs/open.c</a>	const char *	-	-	-	-
13	sys_time	<a href="#">kernel/time.c</a>	int *	-	-	-	-
14	sys_mknod	<a href="#">fs/namei.c</a>	const char *	int	<a href="#">dev_t</a>	-	-
15	sys_chmod	<a href="#">fs/open.c</a>	const char *	<a href="#">mode_t</a>	-	-	-
16	sys_lchown	<a href="#">fs/open.c</a>	const char *	<a href="#">uid_t</a>	<a href="#">gid_t</a>	-	-
18	sys_stat	<a href="#">fs/stat.c</a>	char *	<a href="#">struct old_kernel_stat *</a>	-	-	-
19	sys_lseek	<a href="#">fs/read_write.c</a>	unsigned int	<a href="#">off_t</a>	unsigned int	-	-
20	sys_getpid	<a href="#">kernel/sched.c</a>	-	-	-	-	-
21	sys_mount	<a href="#">fs/super.c</a>	char *	char *	char *	-	-
22	sys_oldumount	<a href="#">fs/super.c</a>	char *	-	-	-	-
23	sys_setuid	<a href="#">kernel/sys.c</a>	<a href="#">uid_t</a>	-	-	-	-
24	sys_getuid	<a href="#">kernel/sched.c</a>	-	-	-	-	-
25	sys_stime	<a href="#">kernel/time.c</a>	int *	-	-	-	-
26	sys_ptrace	<a href="#">arch/i386/kernel/ptrace.c</a>	long	long	long	long	-
27	sys_alarm	<a href="#">kernel/sched.c</a>	unsigned int	-	-	-	-
28	sys_fstat	<a href="#">fs/stat.c</a>	unsigned int	<a href="#">struct old_kernel_stat *</a>	-	-	-
29	sys_pause	<a href="#">arch/i386/kernel/sys_i386.c</a>	-	-	-	-	-
30	sys_utime	<a href="#">fs/open.c</a>	char *	<a href="#">struct utimbuf *</a>	-	-	-

Pintos: syscall-nr.h



# Recall: Kernel System Call Handler

---

- **Locate arguments**
  - In registers or on user(!) stack
- **Copy arguments**
  - From user memory into kernel memory
  - Protect kernel from malicious code evading checks
- **Validate arguments**
  - Protect kernel from errors in user code
- **Copy results back**
  - into user memory



# Process Management

---

- `exit` – terminate a process
- `fork` – copy the current process
- `exec` – change the *program* being run by the current process
- `wait` – wait for a process to finish
- `kill` – send a *signal* (interrupt-like notification) to another process
- `sigaction` – set handlers for signals



# Creating Processes

---

- `pid_t fork();` -- copy the current process
  - New process has different pid
- Return value from `fork()`: pid (like an integer)
  - When  $> 0$ :
    - » Running in (original) **Parent** process
    - » return value is **pid** of new child
  - When  $= 0$ :
    - » Running in new **Child** process
  - When  $< 0$ :
    - » Error! Must handle somehow
    - » Running in original process
- **State of original process duplicated in *both* Parent and Child!**
  - Address Space (Memory), File Descriptors (covered later), etc...



# fork1.c

---

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>


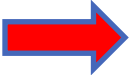
int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```





# fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>


int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
      cpid = fork();
    if (cpid > 0) {                 /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {         /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```



# fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                 /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```





# fork\_race.c

---


```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```

- What does this print?
- Would adding the calls to `sleep` matter?



# Fork “race”

```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```





# Process Management

---

- **fork** – copy the current process
- **exec** – change the *program* being run by the current process
- **wait** – wait for a process to finish
- **kill** – send a *signal* (interrupt-like notification) to another process
- **sigaction** – set handlers for signals



# fork2.c – parent waits for child to finish

```
int status;
pid_t tcpid;
...
cpid = fork();
if (cpid > 0) {                                /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    tcpid = wait(&status);
    printf("[%d] bye %d(%d)\n", mypid, tcpid, status);
} else if (cpid == 0) {                        /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
}
...
```



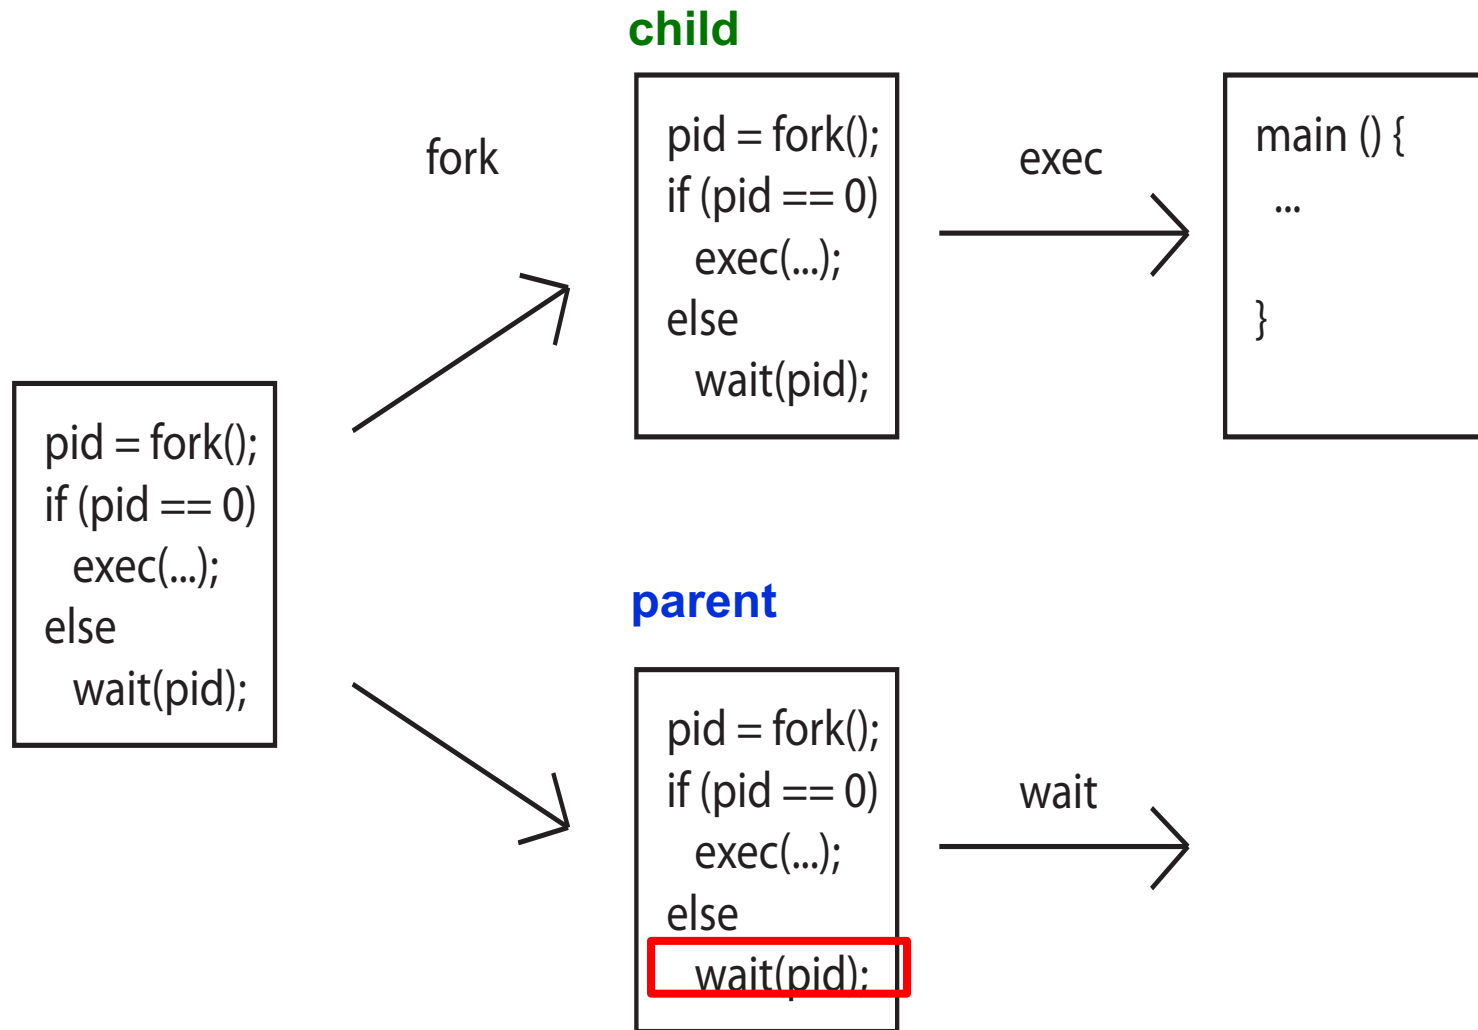
# Process Management

---

- **fork** – copy the current process
- **exec** – change the *program* being run by the current process
- **wait** – wait for a process to finish
- **kill** – send a *signal* (interrupt-like notification) to another process
- **sigaction** – set handlers for signals



# Process Management







# fork3.c

---

```
...
cpid = fork();
if (cpid > 0) {                                /* Parent Process */
    tcpid = wait(&status);
} else if (cpid == 0) {                        /* Child Process */
    char *args[] = {"ls", "-l", NULL};
    execv("/bin/ls", args);
    /* execv doesn't return when it works.
       So, if we got here, it failed! */
    perror("execv");
    exit(1);
}
...
```



# Shell

---

- **A shell is a job control system**
  - Allows programmer to create and manage a set of programs to do some task
  - Windows, MacOS, Linux all have shells

- **Example: to compile a C program**

```
cc -c sourcefile1.c
```

```
cc -c sourcefile2.c
```

```
ln -o program sourcefile1.o sourcefile2.o
```

```
./program
```

A blue starburst graphic with multiple points, containing the text "HW3" in white.

HW3



# Process Management

---

- **fork** – copy the current process
- **exec** – change the *program* being run by the current process
- **wait** – wait for a process to finish
- **kill** – send a *signal* (interrupt-like notification) to another process
- **sigaction** – set handlers for signals



# inf\_loop.c

---

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <signal.h>

void signal_callback_handler(int signum) {
    printf("Caught signal!\n");
    exit(1);
}

int main() {
    struct sigaction sa;
    sa.sa_flags = 0;
    sigemptyset(&sa.sa_mask);
    sa.sa_handler = signal_callback_handler;
    sigaction(SIGINT, &sa, NULL);
    while (1) {}
}
```



# Common POSIX Signals

---

- **SIGINT – control-C**
- **SIGTERM – default for `kill` shell command**
- **SIGSTP – control-Z (default action: stop process)**
  
- **SIGKILL, SIGSTOP – terminate/stop process**
  - Can't be changed with `sigaction`
  - Why?



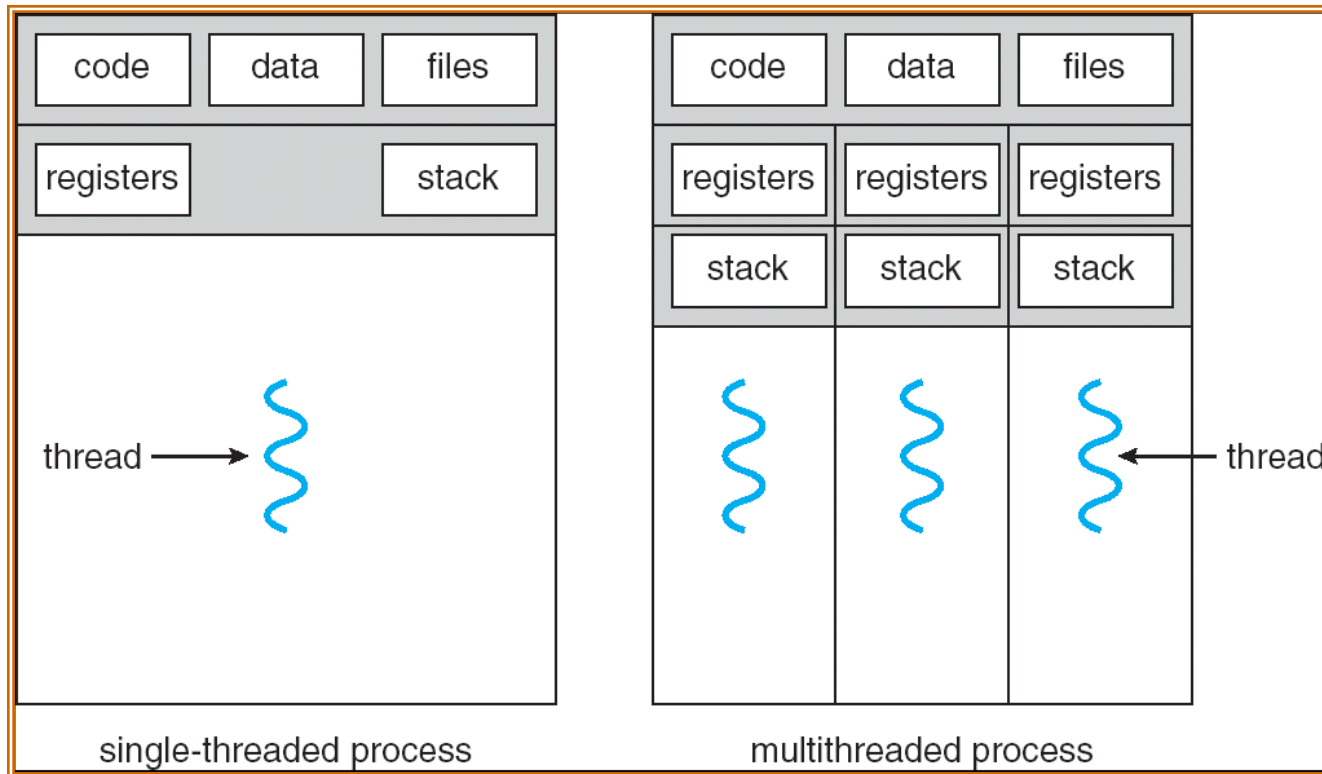
# Modern Processes: Multiple Threads

---

- **Thread: execution stream within a process**
  - Used to be called "lightweight processes"
  - Shares address space with other threads belonging to the same process
- **Why separate concepts of threads and processes?**
  - Threads: Concurrency
  - Processes: Protection



# Single vs. Multithreaded Processes





# Summary

---

- **Process consists of two pieces**
  1. **Address Space (Memory & Protection)**
  2. **One or more threads (Concurrency)**
- **Represented in kernel as**
  - Process object (resources associated with process)
  - Thread object + (mini) stack
  - Hardware support critical in  $U \rightarrow K \rightarrow U$  context switch
  - Different privileges in different modes (CPL, Page Table)
- **Variety of process management syscalls**
  - fork, exec, wait, kill, sigaction
- **Scheduling: Threads move between queues**
- **Threads: multiple stacks per address space**
  - Context switch: Save/Restore registers, "return" from new thread's switch routine
  - So far, we've only seen kernel threads