

CS 162: Operating Systems and Systems Programming

Lecture 4: Threads and Concurrency

Sept 10, 2019

Instructor: David E. Culler

<https://cs162.eecs.berkeley.edu>

Read: A&D 5.1-3, 5.7.1

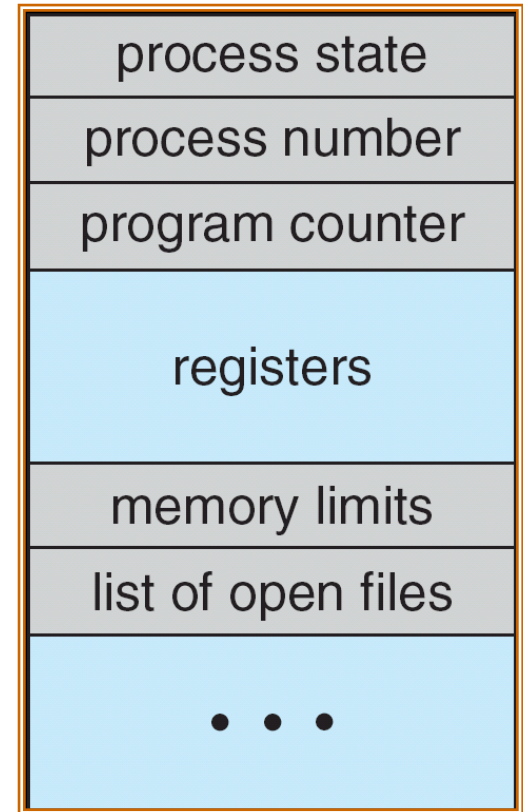
HW 1 due 9/18

Groups and Section Pref **Wed**

Proj 1 released, Design doc **Tues**

Recall: Multiplexing Processes


- Snapshot of each process in its PCB
 - Only one **thread** active at a time per core...
- Give out CPU to different processes
 - **Scheduling**
 - **Policy Decision**
- Give out non-CPU resources
 - Memory/IO
 - Another **policy decision**



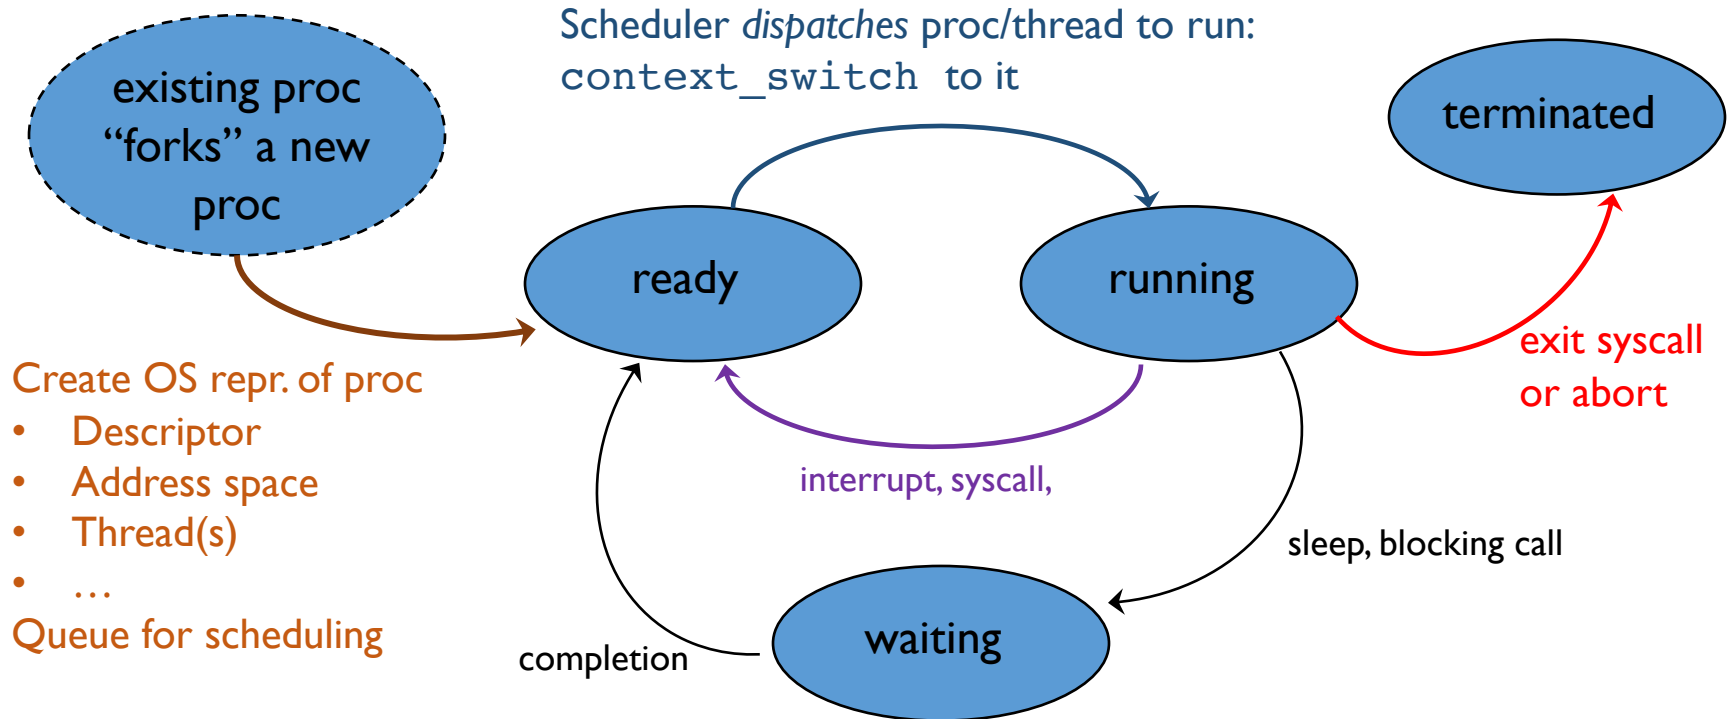
Process
Control
Block

Recall: Context Switch

TCB, Stacks and
Register Mgmt



Recall: Lifecycle of a process / thread



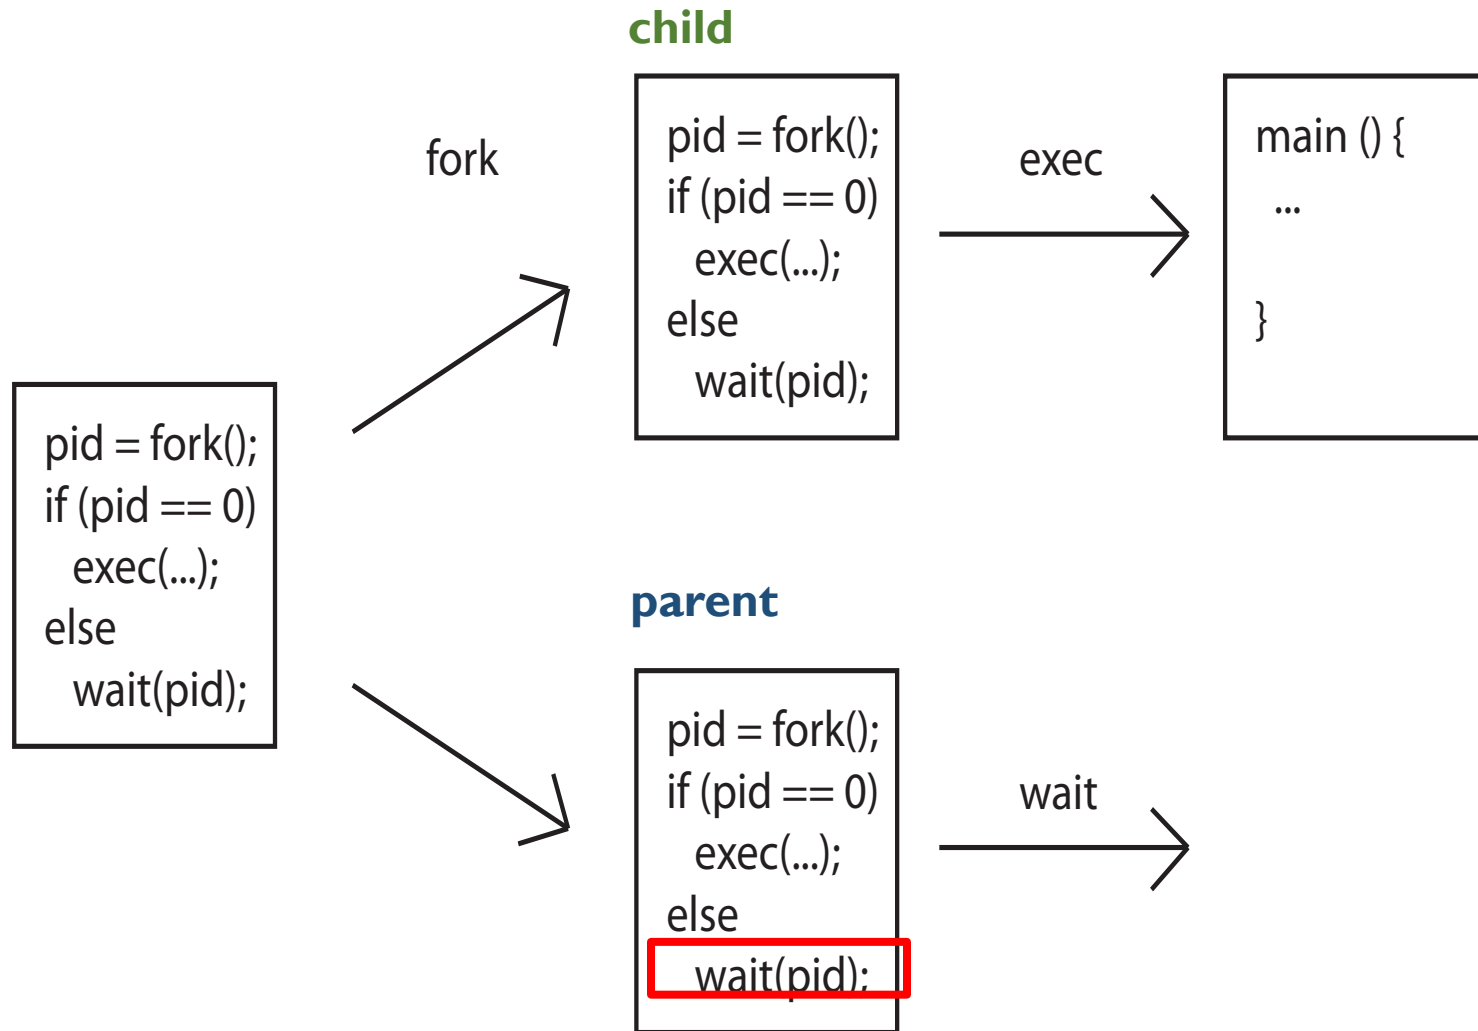
- OS juggles many process/threads using kernel data structures
- Proc's may create other process (fork/exec)
 - All starts with init process at boot

Pintos: process.c

Recall: Process Management

- `exit` – terminate a process
- `fork` – copy the current process
- `exec` – change the *program* being run by the current process
- `wait` – wait for a process to finish
- `kill` – send a *signal* (interrupt-like notification) to another process
- `sigaction` – set handlers for signals



Recall: Process Management




User/OS Threading Models

Simple One-to-One Threading Model

Almost all current implementations

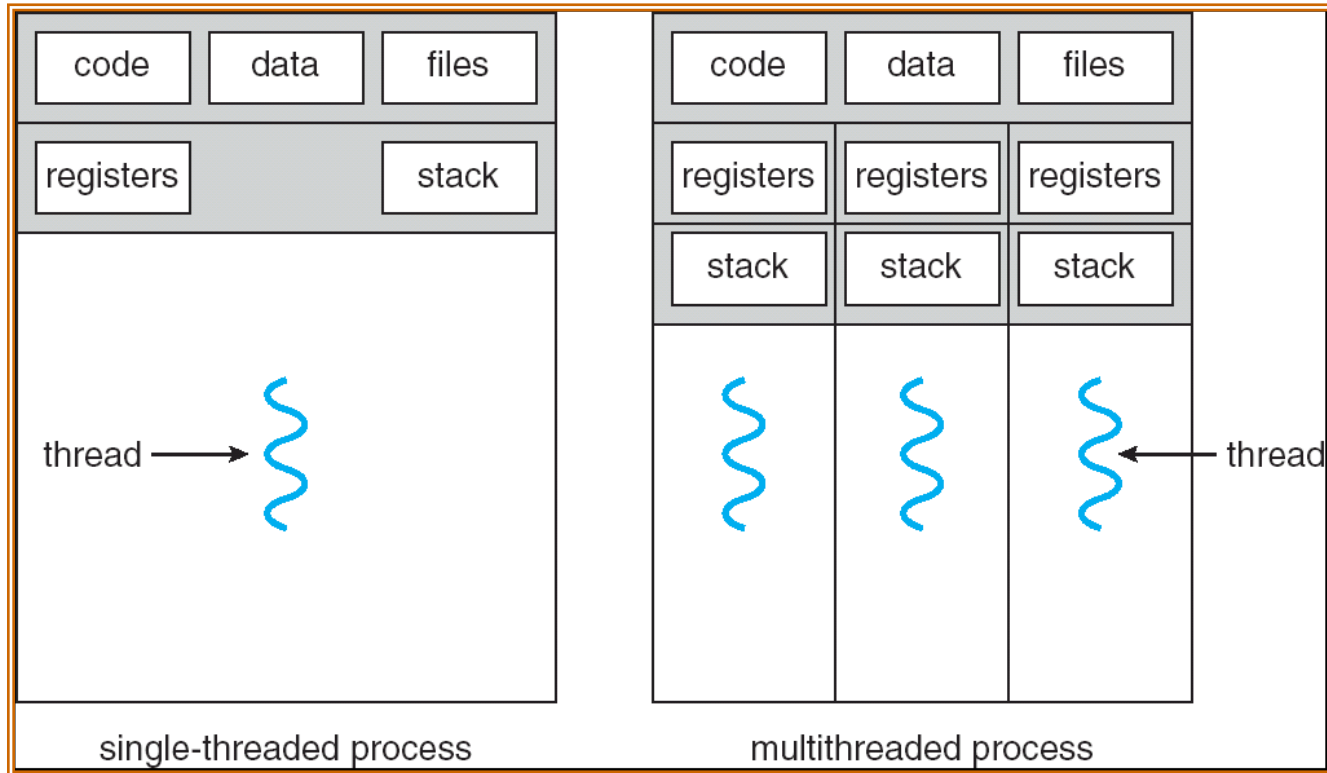


Many-to-One



Many-to-Many

Single vs. Multithreaded Processes



Today

- What, Why, and How of Threads
- Kernel-Supported User Threads
- Coordination among Threads
 - Synchronization
- Implementing Synchronization
- User-level Threads

Definitions

- A *thread* is a single execution sequence that represents a separately schedulable task
- Protection is an orthogonal concept
 - Can have one or many threads per protection domain
 - Single threaded user program: one thread, one protection domain
 - Multi-threaded user program: multiple threads, sharing same data structures, isolated from other user programs
 - Multi-threaded kernel: multiple threads, sharing kernel data structures, capable of using privileged instructions

Threads Motivation

- Operating systems need to be able to handle *multiple things at once* (MTAO)
 - processes, interrupts, background system maintenance
- Servers need to handle MTAO
 - Multiple connections handled simultaneously
- Parallel programs need to handle MTAO
 - To achieve better performance
- Programs with user interfaces often need to handle MTAO
 - To achieve user responsiveness while doing computation
- Network and disk bound programs need to handle MTAO
 - To hide network/disk latency
 - Sequence steps in access or communication

Silly Example for Threads

Imagine the following program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("classlist.txt");  
}
```

- What is the behavior here?
 - Program would never print out class list
 - Why? ComputePI would never finish

Adding Threads

- Version of program with Threads (loose syntax):

```
main() {  
    thread_fork(ComputePI, "pi.txt" );  
    thread_fork(PrintClassList, "classlist.txt");  
}
```

- `thread_fork`: Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs



Time 

More Practical Motivation

Back to Jeff Dean's "Numbers everyone should know"

Handle I/O in
separate thread,
avoid blocking
other progress

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Little Better Example for Threads?

Imagine the following program:

```
main() {  
    ...  
    ReadLargeFile("pi.txt");  
    RenderUserInterface();  
}
```

- What is the behavior here?
 - Still respond to user input
 - While reading file in the background

Voluntarily Giving Up Control


- I/O – e.g. keypress
- Waiting for a signal from another thread
 - Thread makes system call to *wait*
- Thread executes `thread_yield()`
 - Relinquishes CPU but puts calling thread back on ready queue

Adding Threads

- Version of program with Threads (loose syntax):

```
main() {  
    thread_fork(ReadLargeFile, "pi.txt" );  
    thread_fork(RenderUserInterface, "classlist.txt");  
}
```

- `thread_fork`: Start independent thread running given procedure
- What is the behavior here?
 - Now, you would actually see the class list
 - This *should* behave as if there are two separate CPUs

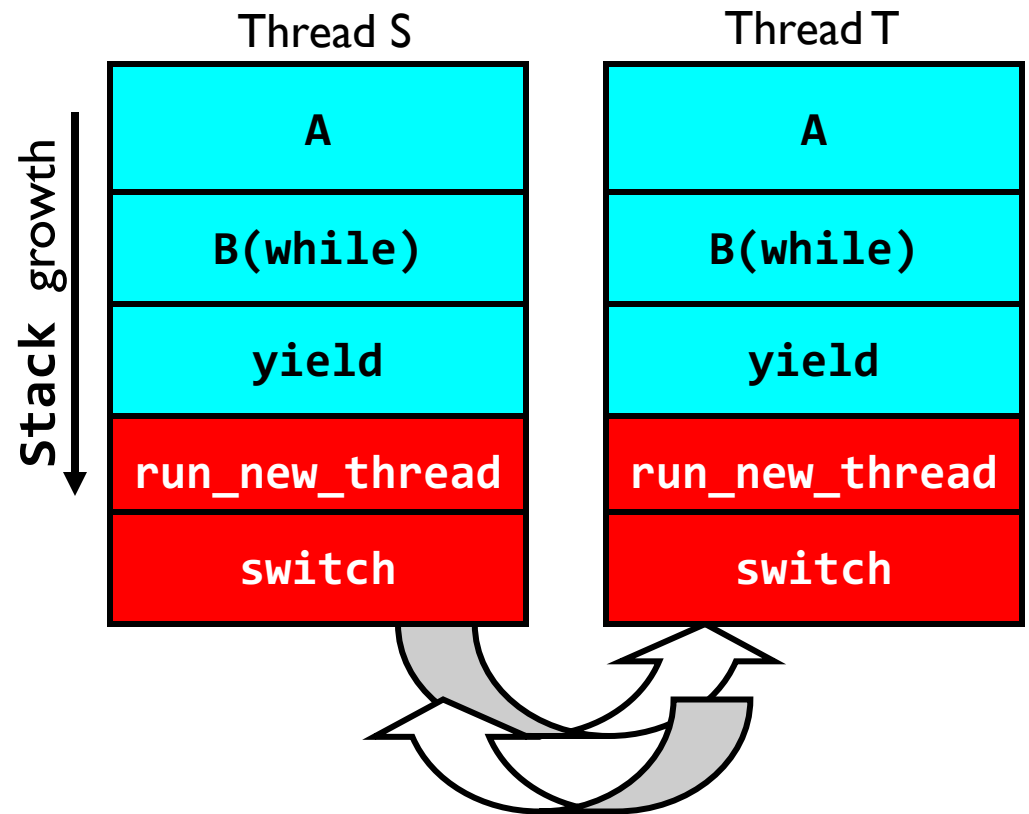


Switching Threads

- Consider the following code blocks:

```
func A() {  
    B();  
}  
func B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

- Two threads, S and T, each run A




Thread S's switch returns to
Thread T's (and vice versa)

Aren't we still switching contexts?


- Yes, but **much cheaper** than switching processes
 - No need to change address space
- Some numbers from Linux:
 - Frequency of context switch: 10-100ms
 - Switching between processes: 3-4 μ sec.
 - Switching between threads: 100 ns

Processes vs. Threads



- Switch overhead:
 - Same process: **low**
 - Different proc.: **high**
- Protection
 - Same proc: **low**
 - Different proc: **high**
- Sharing overhead
 - Same proc: **low**
 - Different proc: **high**

Processes vs. Threads



- Switch overhead:
 - Same process: **low**
 - Different proc.: **high**
- Protection
 - Same proc: **low**
 - Different proc: **high**
- Sharing overhead
 - Same proc: **low**
 - Different proc: **high**


Example: Multithreaded Server

```
serverLoop() {  
    connection = AcceptNewConnection();  
    (thread_)fork(ServiceWebPage, connection);  
}
```

- One process/thread per connection, many concurrent connections
- Process (isolation) vs Thread (performance)
- How fast is creating threads?
 - Better than `fork()`, but still overhead
- Problem: What if we get a lot of requests?
 - Might run out of memory (thread stacks)
 - Schedulers usually have trouble with too many threads


Web Server: Thread Pools

- **Bounded** pool of worker threads
 - Allocated in **advance**: no thread creation overhead
 - **Queue** of pending requests
 - **Limited number** of requests in progress



Multiprocessing vs Multiprogramming

- Multiprocessing: Multiple cores
- Multiprogramming: Multiple Jobs/Processes
- Multithreading: Multiple threads/processes
- What does it mean to run two threads concurrently?
 - Scheduler is free to run threads in any order and interleaving



Thread vs. Process State

- Process-wide state:
 - Memory contents (global variables, heap)
 - I/O bookkeeping
- Thread-"local" state:
 - CPU registers including program counter
 - Execution stack
 - Kept in **Thread Control Block**

Shared vs. Per-Thread State


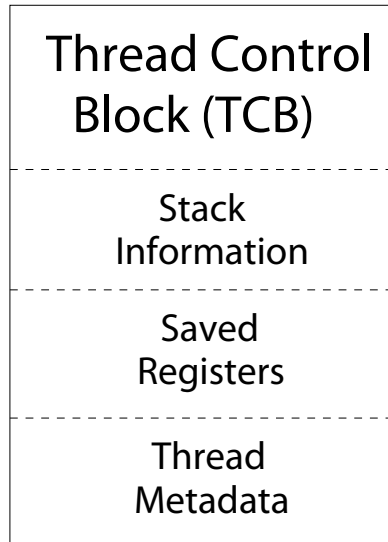
Shared State

Heap


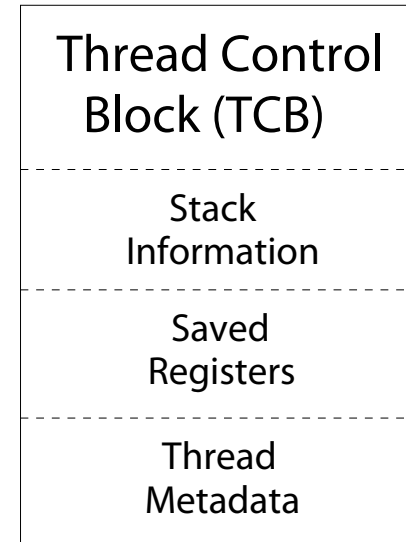
Global Variables

Code

Per-Thread State




Per-Thread State




Memory Footprint: Two Threads

- Two sets of CPU registers
- Two sets of Stacks
- Issues:
 - How do we position stacks relative to each other?
 - What maximum size should we choose for the stacks?
 - What happens if threads violate this?
 - How might you catch violations?
- User threads need ‘proper’ stacks
 - System threads may be very constrained



Yield is covered, what about I/O?

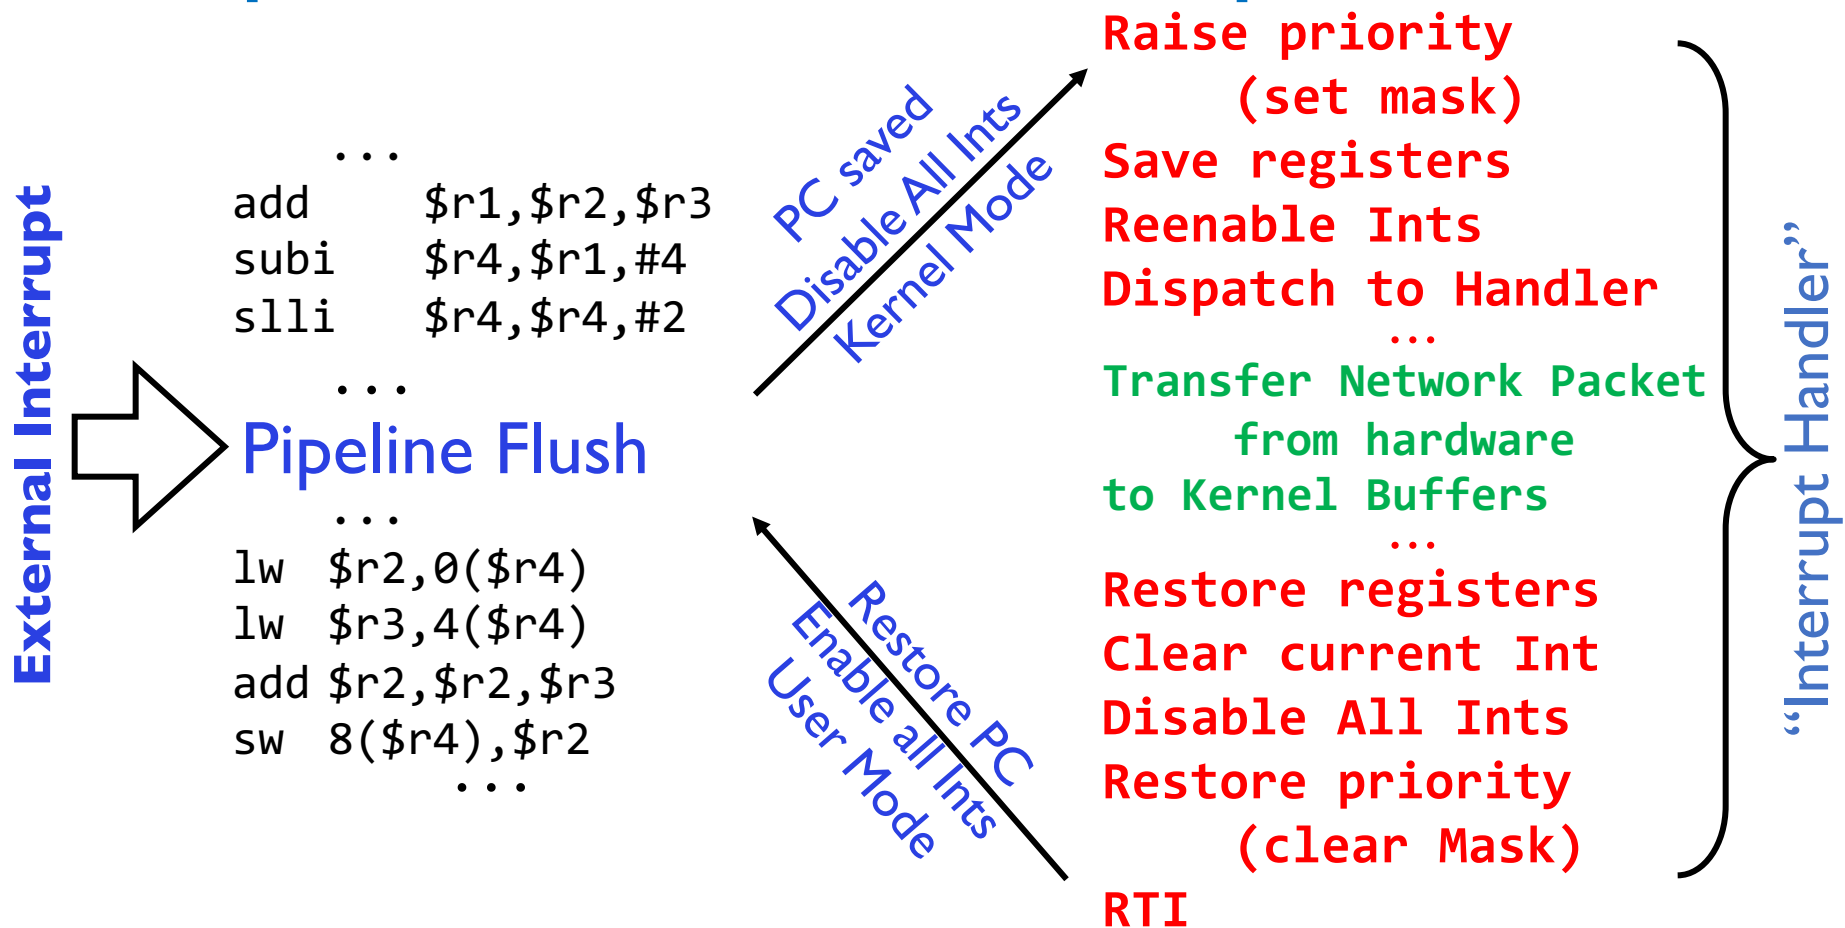


- User code invokes syscall
- IO operation initiated (more later)
- Run new thread, switch
- Really, same thing as before
 - Just put the thread on a different queue

Preempting a Thread

- What happens if thread never does any I/O, never waits, and never yields control?
 - Must find way that dispatcher can regain control!
- **Interrupts:** signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some milliseconds
- Interrupt is a hardware-invoked mode switch
 - Handled immediately, no scheduling required


Example: Network Interrupt



- An interrupt is a hardware-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately

Switching Threads from Interrupts

- Prevent thread from running forever with **timer interrupt**



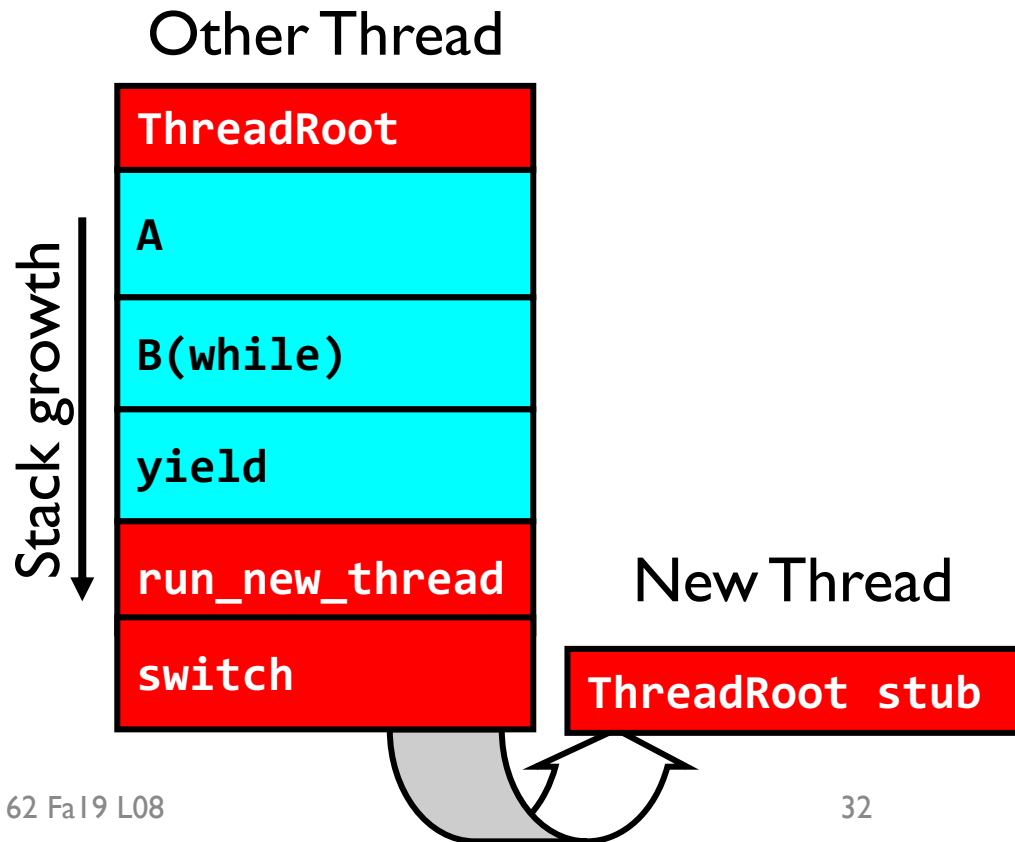
```
TimerInterrupt() {  
    DoPeriodicHouseKeeping();  
    run_new_thread();  
}
```

- Same thing from IO interrupts
 - Example: immediately start process waiting for keypress

How does a thread get started?

- Can't call `switch()` without starting a thread
- How do we make a **new** thread?

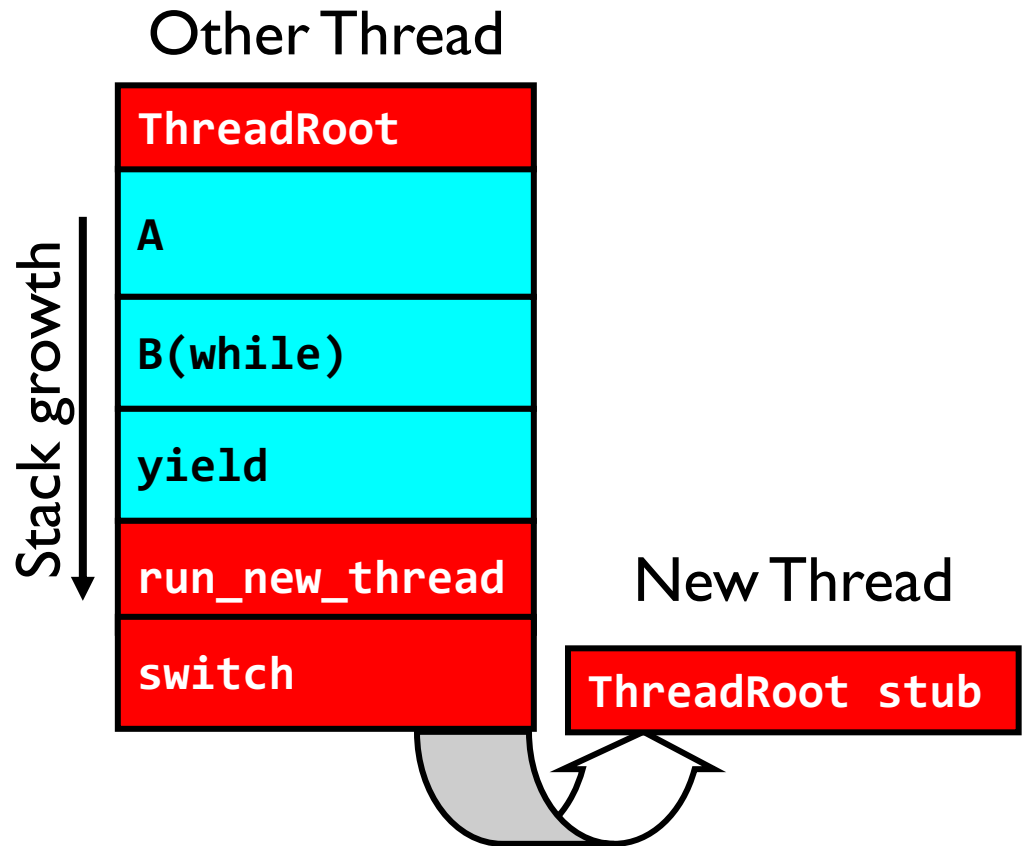
```
SetupNewThread(tNew) {  
    ...  
    TCB[tNew].regs.sp =  
        newStack;  
    TCB[tNew].regs.retpc =  
        &ThreadRoot;  
}
```



How does a thread get started?


- So when does the new thread really start executing?

`run_new_thread`
selects this thread's TCB,
"returns" into beginning
of `ThreadRoot`



Bootstrapping Threads: ThreadRoot

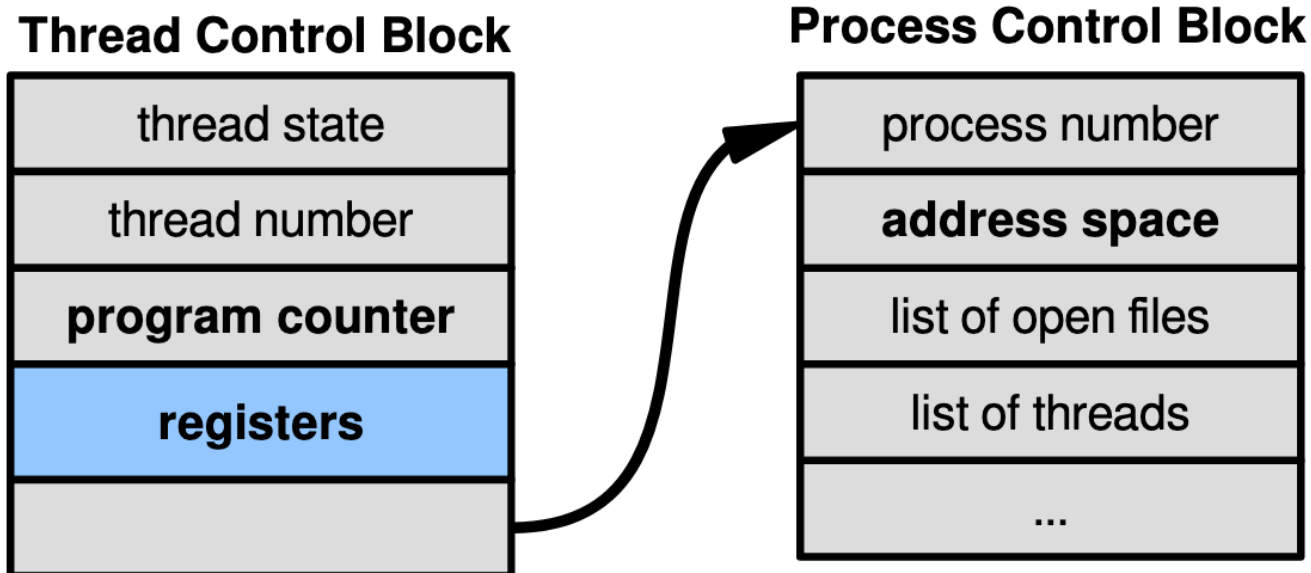
```
ThreadRoot() {  
    DoStartupHousekeeping();  
    UserModeSwitch(); /* enter user mode */  
    call fcnPtr(fcnArgPtr);  
    ThreadFinish();  
}
```



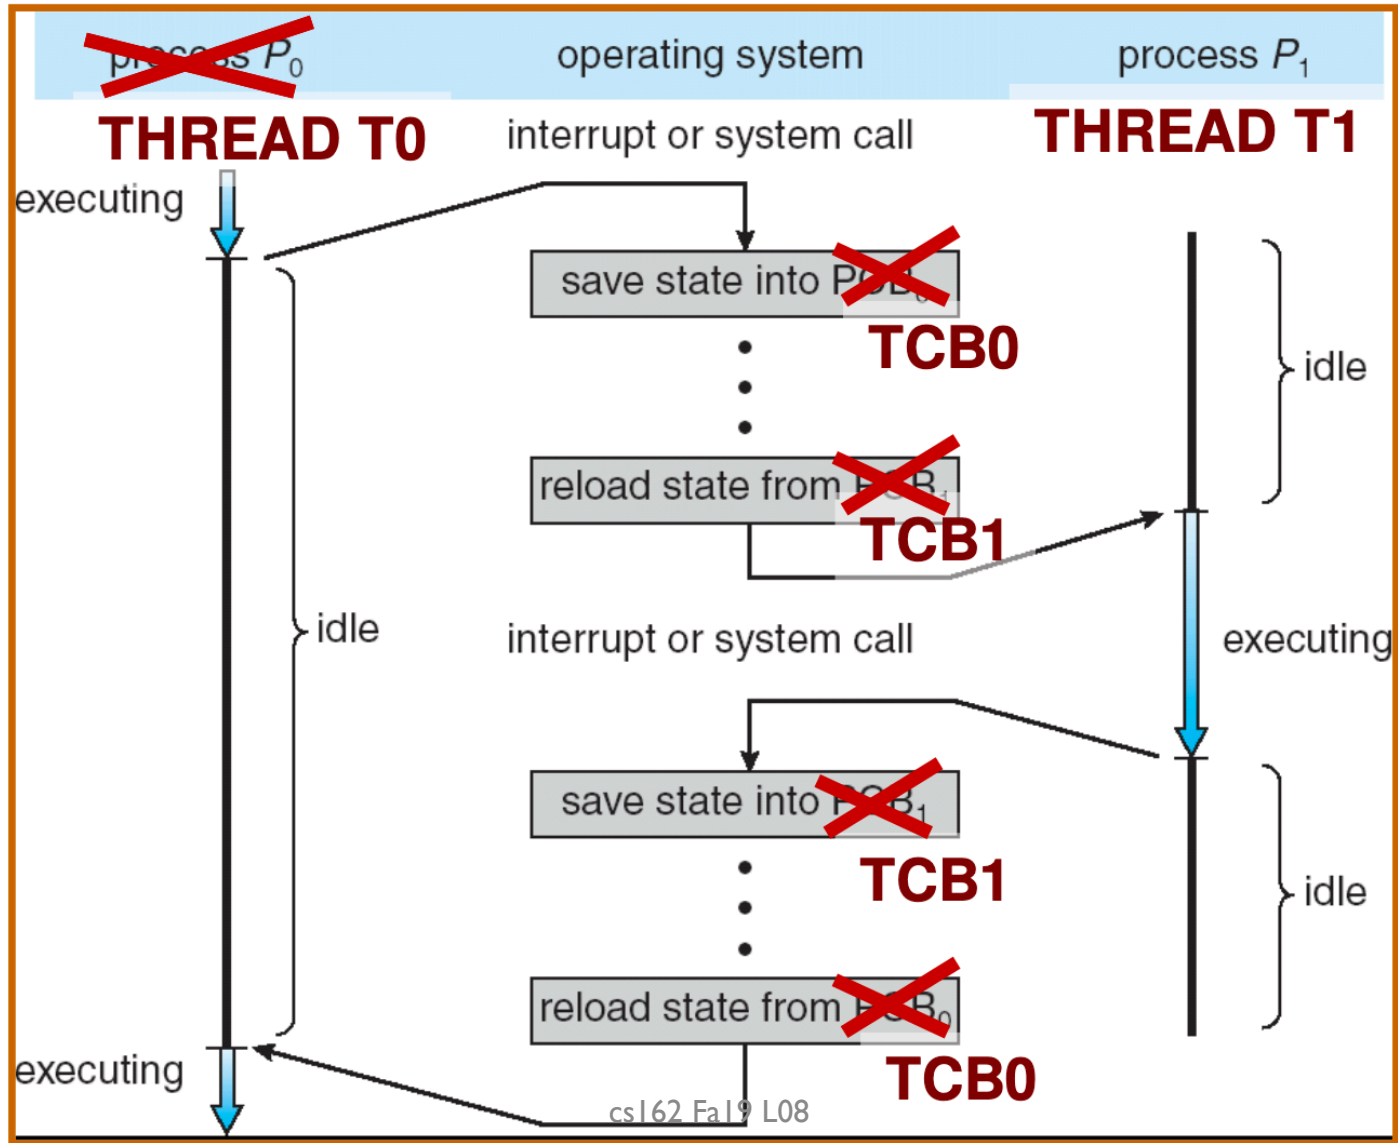
- Stack will grow and shrink with execution of thread
- `ThreadRoot()` never returns
 - `ThreadFinish()` destroys thread, invokes scheduler

Kernel-Supported Threads

- Each thread has a **thread control block**
 - CPU registers, including PC, pointer to stack
 - Scheduling info: priority, etc.
 - Pointer to **Process control block**
- OS scheduler uses TCBs, not PCBs



Kernel-Supported User Threads



User-level Multithreading: *pthread*s

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void*), void *arg);`
 - thread is created executing *start_routine* with *arg* as its sole argument. (return is implicit call to `pthread_exit`)
- `void pthread_exit(void *value_ptr);`
 - terminates and makes *value_ptr* available to any successful join
- `int pthread_join(pthread_t thread, void **value_ptr);`
 - suspends execution of the calling thread until the target *thread* terminates.
 - On return with a non-NULL *value_ptr* the value passed to [*pthread_exit\(\)*](#) by the terminating thread is made available in the location referenced by *value_ptr*.

man pthread

<https://pubs.opengroup.org/onlinepubs/7908799/xsh/pthread.h.html>

Little Example

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <string.h>
```

```
int common = 162;
```

```
void *threadfun(void *threadid)
{
    long tid = (long)threadid;
    printf("Thread #%lx stack: %lx common: %lx (%d)\n", tid,
           (unsigned long) &tid, (unsigned long) &common, common++);
    pthread_exit(NULL);
}
```


```
int main (int argc, char *argv[])
{
    long t;
    int nthreads = 2;
    if (argc > 1) {
        nthreads = atoi(argv[1]);
    }
    pthread_t *threads = malloc(nthreads*sizeof(pthread_t));
    printf("Main stack: %lx, common: %lx (%d)\n",
           (unsigned long) &t, (unsigned long) &common, common);
    for(t=0; t<nthreads; t++){
        int rc = pthread_create(&threads[t], NULL, threadfun, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    for(t=0; t<nthreads; t++){
        pthread_join(threads[t], NULL);
    }
    pthread_exit(NULL);
}
```

```
[(base) CullerMac19:code04 culler$ ./pthread 4
Main stack: 7ffee2c6b6b8, common: 10cf95048 (162)
Thread #1 stack: 70000d83bef8 common: 10cf95048 (162)
Thread #3 stack: 70000d941ef8 common: 10cf95048 (164)
Thread #2 stack: 70000d8beef8 common: 10cf95048 (165)
Thread #0 stack: 70000d7b8ef8 common: 10cf95048 (163)
```

How to tell if something is done?
Really done?
OK to reclaim its resources?

Fork-Join Pattern




- Main thread *creates* (forks) collection of sub-threads passing them args to work on, *joins* with them, collecting results.

Interleaving & Nondeterminism


Programmer vs. Processor View

Programmer's View	Possible Execution #1
.	.
.	.
.	.
$x = x + 1;$	$x = x + 1;$
$y = y + x;$	$y = y + x;$
$z = x + 5y;$	$z = x + 5y;$
.	.
.	.
.	.


Possible Executions



a) One execution



b) Another execution




c) Another execution

Correctness with Concurrent Threads

- Non-determinism:
 - Scheduler can run threads in **any order**
 - Scheduler can switch threads **at any time**
 - This can make testing very difficult
- *Independent Threads*
 - No state shared with other threads
 - Deterministic, reproducible conditions
- *Cooperating Threads*
 - Shared state between multiple threads
- **Goal: Correctness by Design**

Remember: Multiprogramming



- Scheduler can run threads in any order
- And with multiple cores:
 - Even more interleaving
 - **Could truly be running at the same time**

Race Conditions

- What are the possible values of x below?
- Initially $x = y = 0$;

Thread A

$x = 1$;

Thread B

$y = 2$;

- Must be 1. Thread B cannot interfere.

Race Conditions

- What are the possible values of x below?
- Initially $x = y = 0$;

Thread A

$x = y + 1$;

Thread B

$y = 2$;

$y = y * 2$;

- 1 or 3 or 5 (non-deterministic)
- Race Condition: Thread A races against Thread B

Atomic Operations

- Definition: **An operation that runs to completion or not at all**
 - Need some to allow threads to work together
- `counter++;` **// atomic?**
 - x86 has memory-to-memory instructions, but that still doesn't make them atomic
- Some store instructions are not atomic
 - Ex: double-precision floating point store

Real-Life Analogy: Too Much Milk

Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

See “Additional Materials” and text...

Break

Relevant Definitions

- **Mutual Exclusion:** Ensuring only one thread does a particular thing at a time (one thread *excludes* the others)
- **Critical Section:** Code exactly one thread can execute at once
 - Result of mutual exclusion

Relevant Definitions

- **Lock:** An object only one thread can hold at a time
 - **Provides** mutual exclusion
- Offers two **atomic** operations:
 - `Lock.Acquire()` – wait until lock is free; then grab
 - `Lock.Release()` – Unlock, wake up waiters

Using Locks

```
MilkLock.Acquire()  
if (noMilk) {  
    buy milk  
}  
MilkLock.Release()
```

But how do we implement this?

First, how do we use it?

Pthreads - mutex

```
#include <pthread.h>
```

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

Our example

Critical section



```
int common = 162;
pthread_mutex_t common_lock = PTHREAD_MUTEX_INITIALIZER;

void *threadfun(void *threadid)
{
    long tid = (long)threadid;
    pthread_mutex_lock(&common_lock);
    int my_common = common++;
    pthread_mutex_unlock(&common_lock);

    printf("Thread #%lx stack: %lx common: %lx (%d)\n", tid,
        (unsigned long) &tid,
        (unsigned long) &common, my_common);
    pthread_exit(NULL);
}
```

Semaphores



- Semaphores are a kind of generalized lock
 - First defined by Dijkstra in late 60s
 - Main synchronization primitive used in original UNIX (& Pintos)
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
 - **P() or down()**: atomic operation that waits for semaphore to become positive, then decrements it by 1
 - **V() or up()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any

P() stands for “*proberen*” (to test) and **V()** stands for “*verhogen*” (to increment) in Dutch

Two Important Semaphore Patterns

- **Mutual Exclusion:** (Like lock)
 - Called a "binary semaphore"
 - `initial value of semaphore = 1;`
 - `semaphore.down();`
 - `// Critical section goes here`
 - `semaphore.up();`
- **Signaling** other threads, e.g. `ThreadJoin`

Initial value of semaphore = 0


```
ThreadJoin {  
    semaphore.down();  
}
```

```
ThreadFinish {  
    semaphore.up();  
}
```




Think of *down* as *wait()* operation


What can we conclude... over “All Possible Executions” ?



a) One execution



b) Another execution



c) Another execution

Implementing Locks: Single Core

- Idea: A context switch can only happen (assuming threads don't yield) if there's an **interrupt**
- “Solution”: **Disable interrupts** while holding lock
- x86 has `cli` and `sti` instructions that only operate in system mode (PL=0)
 - Interrupts enabled bit in FLAGS register

Naïve Interrupt Enable/Disable

```
Acquire() {  
    disable interrupts;  
}
```

```
Release() {  
    enable interrupts;  
}
```

- Problem: can stall the entire system

```
Lock.Acquire()  
While (1) {}
```

- Problem: What if we want to do I/O?

```
Lock.Acquire()  
Read from disk  
/* OS waits for (disabled) interrupt! */
```

Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

Discussion

- Why do we need to disable interrupts at all?
 - Avoid interruption between checking and setting lock value
 - Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

} Critical
Section

- Note: unlike previous solution, the critical section (inside `Acquire()`) is very short
 - User of lock can take as long as they like in their own critical section: doesn't impact global machine behavior
 - Critical interrupts taken in time!

Implementing Locks: Single Core

- Idea: Disable interrupts for **mutual exclusion** on accesses to `value` indicating lock status

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        run_new_thread()  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone waiting) {  
        take a thread off queue;  
    } else {  
        Value = FREE;  
    }  
    enable interrupts;  
}
```

Reenabling Interrupts When Waiting

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        run_new_thread()  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

enable interrupts →

enable interrupts →

- Before on the queue?
 - Release might not wake up this thread!
- After putting the thread on the queue?
 - Gets woken up, but immediately switches away

Reenabling Interrupts When Waiting

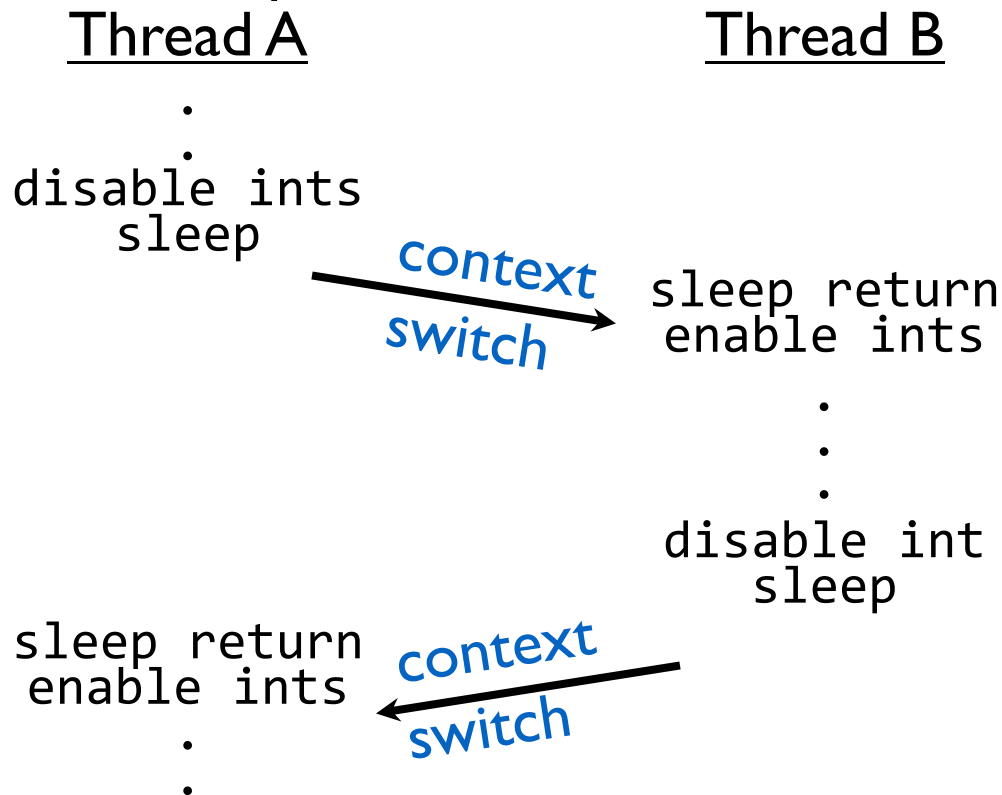
```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        run_new_thread()  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

enable interrupts →

- Best solution: after the current thread suspends
- How?
 - run_new_thread() should do it!
 - Part of returning from switch()

How to Re-enable After Sleep()?

- In scheduler, since interrupts are disabled when you call sleep:
 - Responsibility of the next thread to re-enable ints
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts




Recall: 61c

- Hardware provides certain atomic operations
 - Swap, Compare&Swap, Test&Set, Fetch&Add, LoadLocked/StoreConditional
 - More on optimized synchronization ops later
- System threads need more than the atomic operation
 - May need to manipulate scheduling queues too
 - Requires combination of HW and SW to do it right
- Pintos implements “semaphores”
 - Builds locks and CVs on top of them

Multithreaded Processes

- PCB may be associated with multiple TCBs:



- Switching threads within a process is a simple thread switch
- Switching threads across blocks requires changes to memory and I/O address tables.

So does the OS schedule processes or threads?

- We've been talking about processes assuming the "old model" -> one thread per process
 - And many textbooks say this as well
- Usually it's really: **threads** (e.g., in Linux)
- More on some of these issues later
- One point to notice: switching threads vs. switching processes incurs different costs:
 - Switch threads: Save/restore registers
 - Switch processes: Change active address space too!
 - Expensive
 - Disrupts caching

User-level threads?


- Can multiple threads be implemented entirely at user level?
- Most other aspects of system virtualize.

Kernel-Supported Threads

- Threads run and block (e.g., on I/O) independently
- One process may have multiple threads waiting on different things
- Two mode switches for every context switch (expensive)
- Create threads with syscalls
- Alternative: multiplex several streams of execution (at user level) on top of a single OS thread
 - E.g., Java, Go, ... (and many many user-level threads libraries before it)

User-Mode Threads

- User program contains its own scheduler
- Several user threads per kernel thd.
- User threads may be scheduled **non-preemptively**
 - Only switch on yield
- Context switches cheaper
 - Copy registers and jump (switch in userspace)



User-Mode Threads: Problems

- One user-level thread blocks on I/O: they all do
 - Kernel cannot adjust scheduling among threads it doesn't know about
- Multiple Cores?
- Can't completely avoid blocking (syscalls, page fault)
- One Solution: *Scheduler Activations*
 - Have kernel inform user-level scheduler when a thread blocks
- Evolving the contract between OS and application.

Classification

# threads Per AS:	# of addr spaces:	One	Many
One		MS/DOS, early Macintosh	Traditional UNIX
Many		Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, HP-UX, Win NT to 8, Solaris, OS X, Android, iOS

- Real operating systems have either
 - One or many address spaces
 - One or many threads per address space

Summary

- Process consists of two components
 1. Address Space (Protection)
 2. One or more threads (Concurrency)
- Threads: unit of concurrent execution
 - Useful for parallelism, overlapping computation and IO, organizing sequences of interactions (protocols)
 - Require: multiple stacks per address space
 - Thread switch:
 - Save/Restore registers, "return" from new thread's switch routine
 - Challenging to write correct concurrent code:
 - **Arbitrary interleavings**
 - Could access shared resources while in bad state
 - Kernel threads, Kernel-supported User Threads, User-mode Threads
- Synchronization
 - Building block: atomic operations
 - Mutual exclusion (locks) & Signaling (exit->join, semaphore)
- Scheduling: Threads move between queues
 - Synchronization and scheduler deeply interrelated

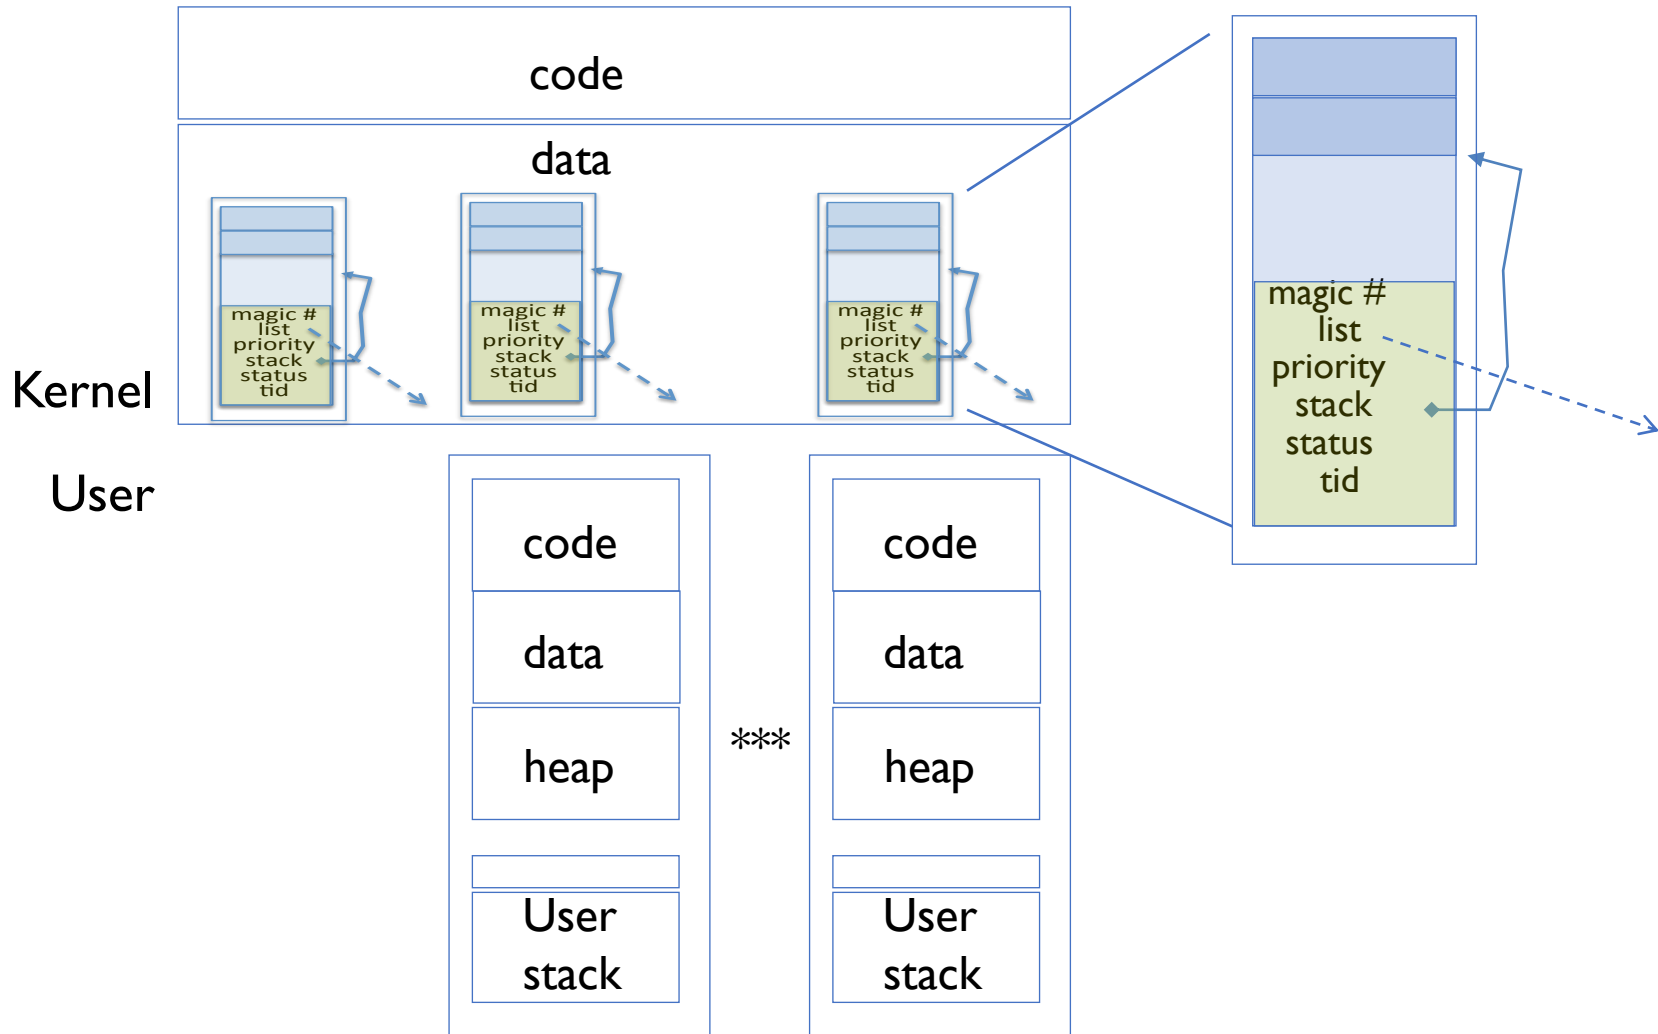
Additional Materials

Deeper Review: User/Kernel Threads in Pintos



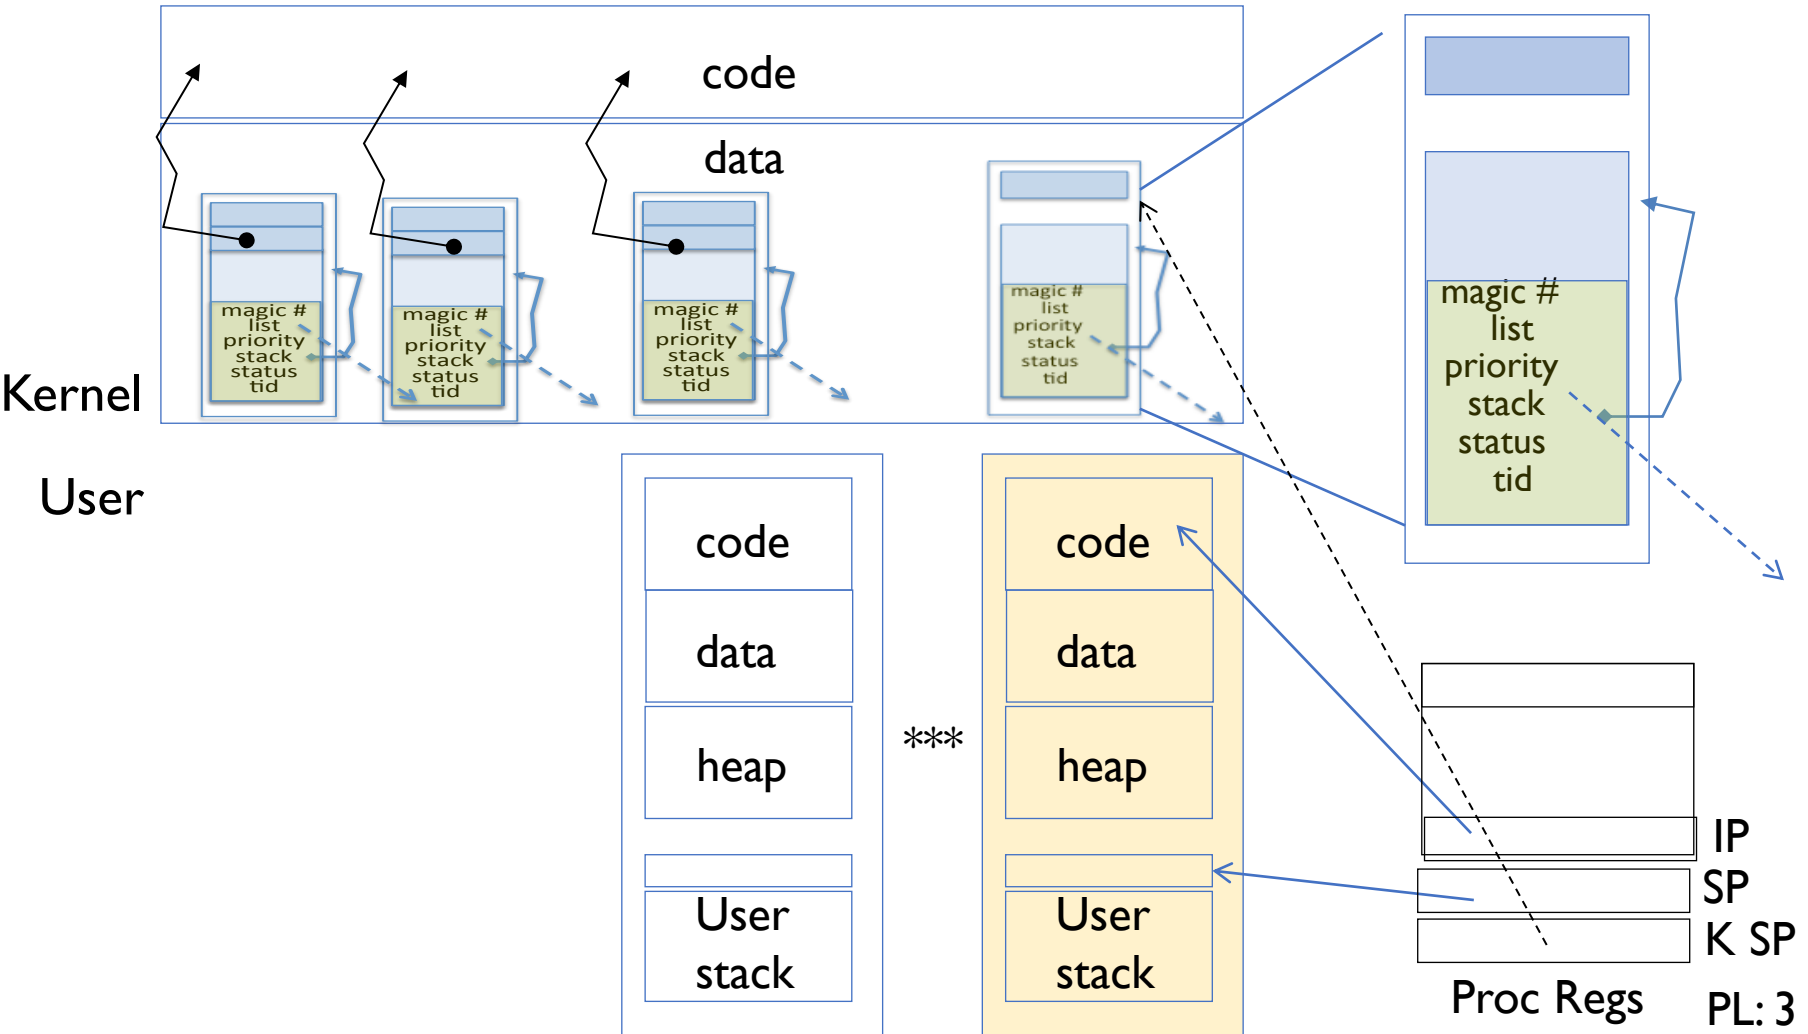
- Now that you're reading the code, let's do a quick picture of what's going on

MT Kernel IT Process ala Pintos/x86



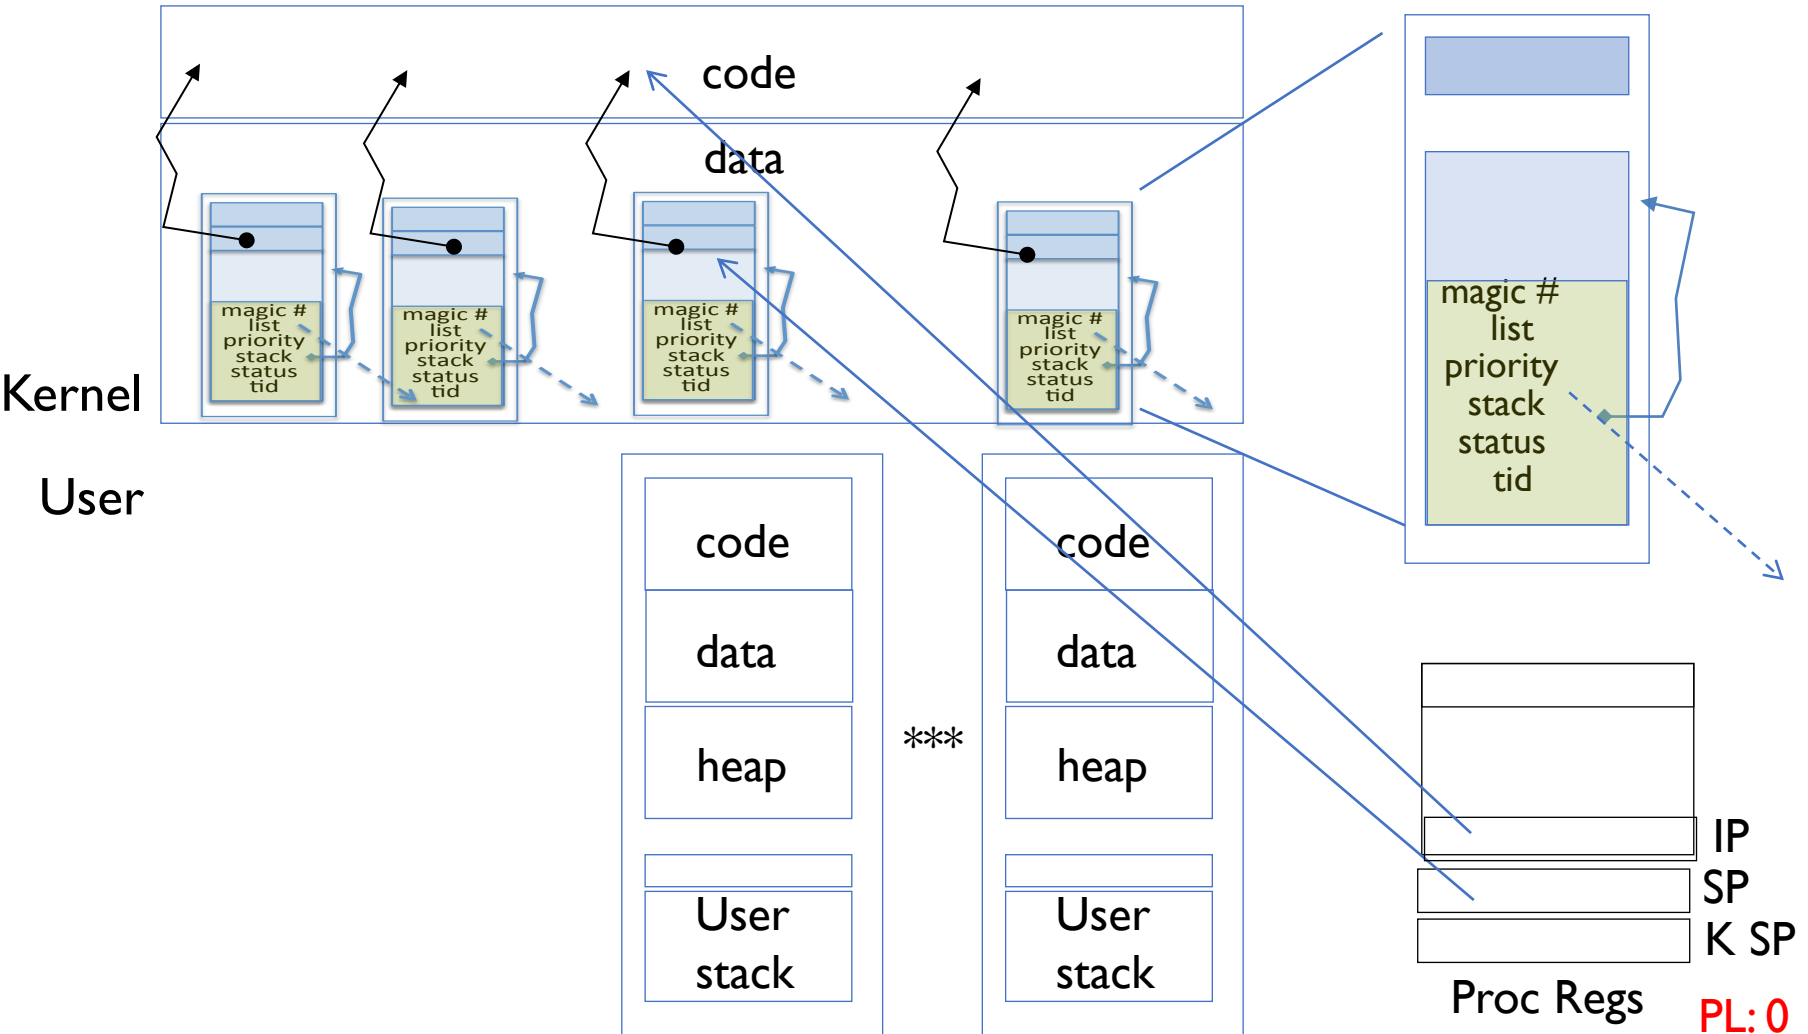
- Each user process/thread associated with a kernel thread, described by a 4kb Page object containing TCB and kernel stack for the kernel thread

In User thread, w/ k-thread waiting



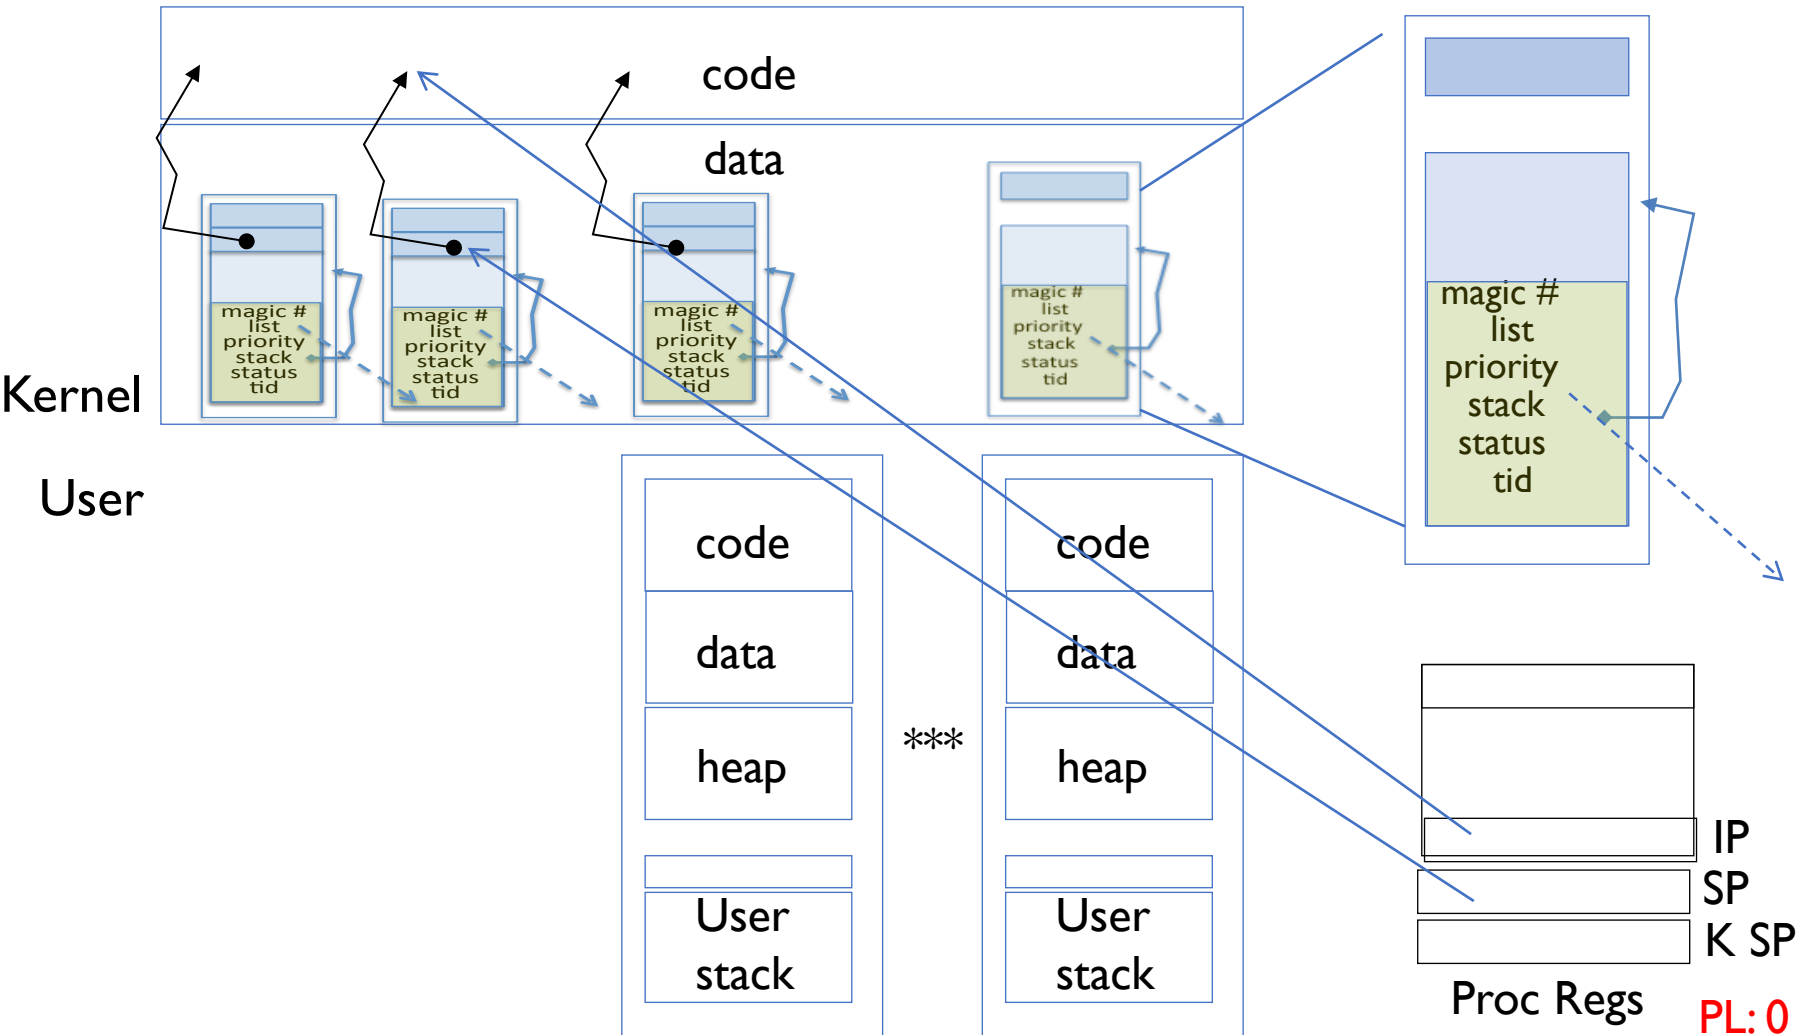
- x86 proc holds interrupt SP high system level
- During user thread exec, associated kernel thread is “standing by”

In Kernel thread



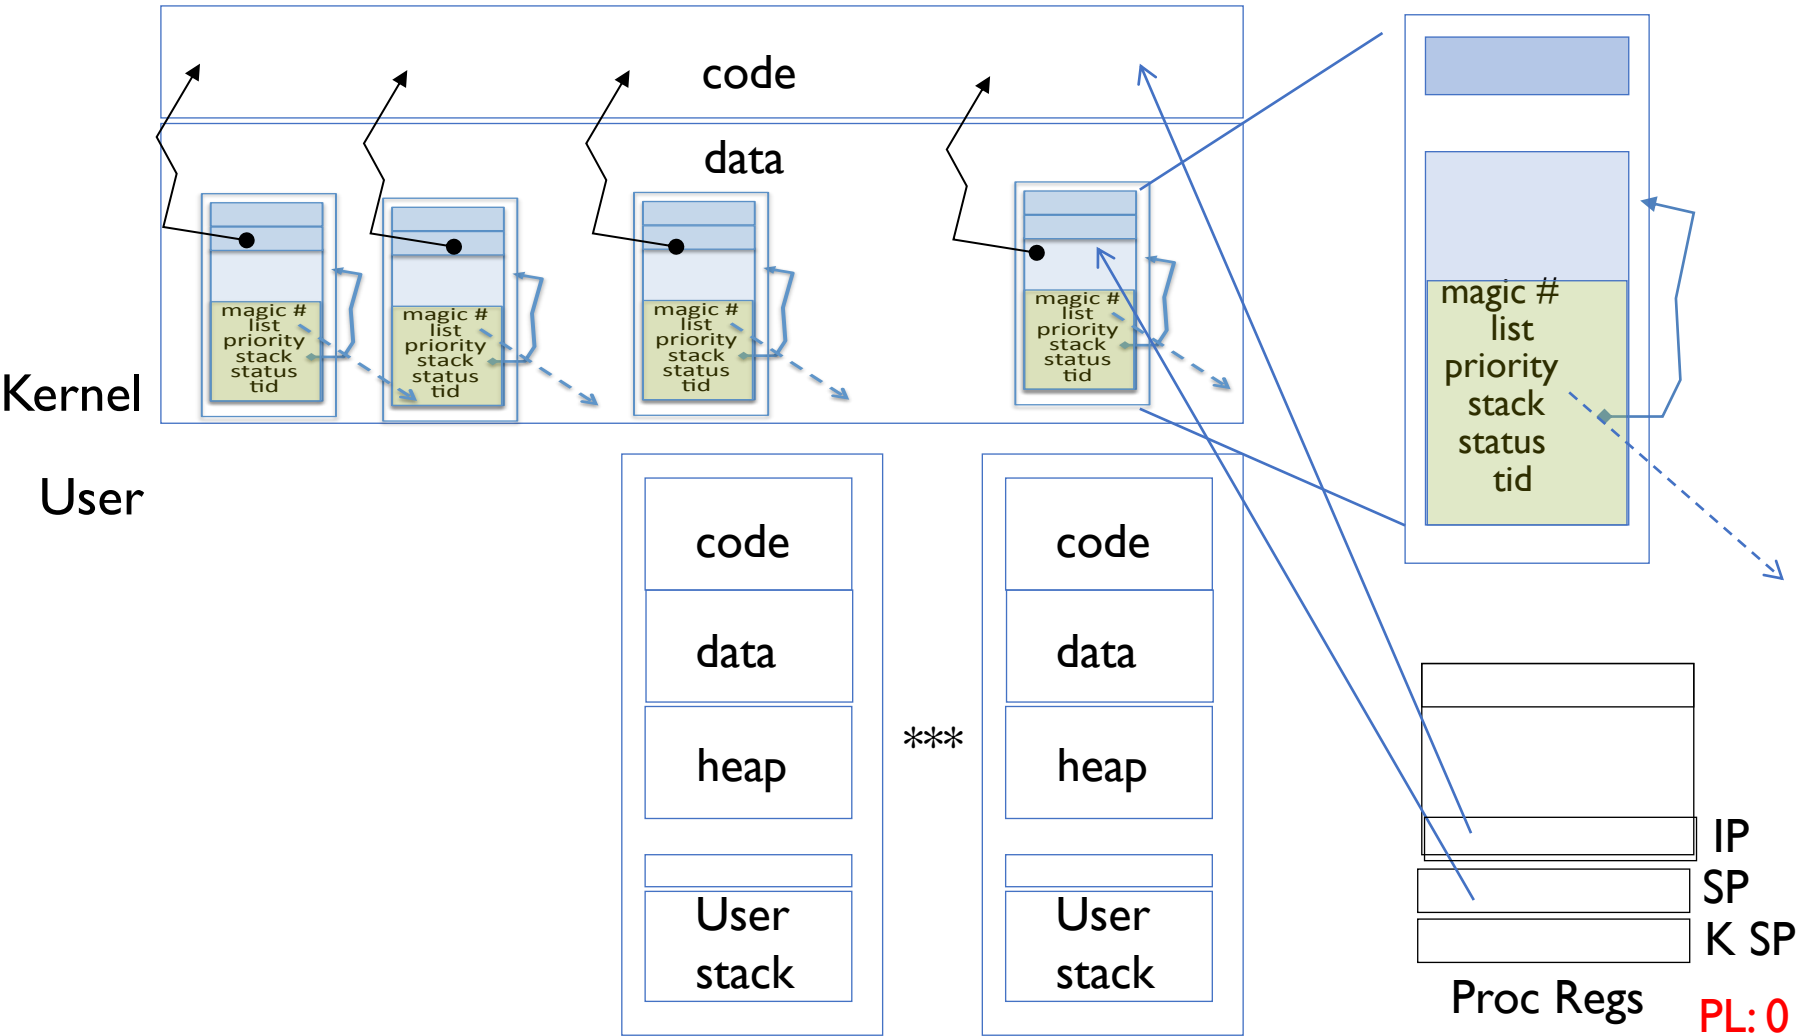
- Kernel threads execute with small stack in thread struct
- Scheduler selects among ready kernel and user threads

Thread Switch (switch.S)

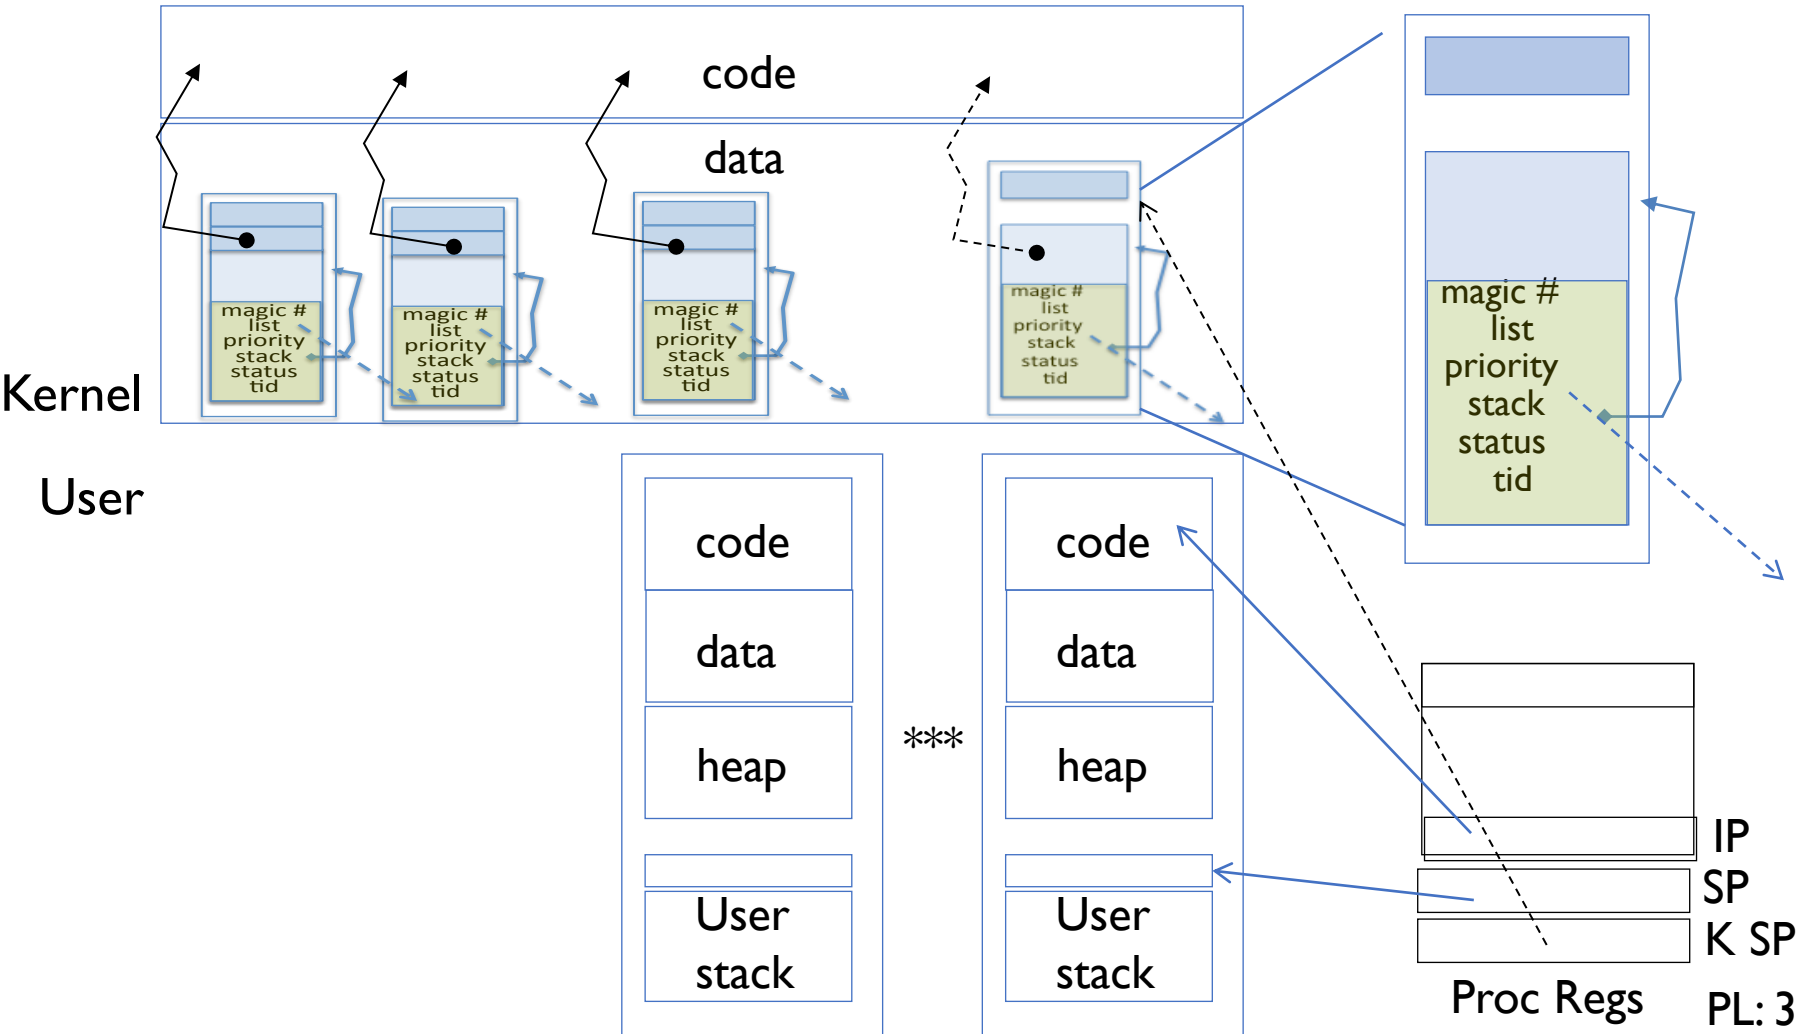


- `switch_threads`: save regs on current small stack, change SP, return from destination threads call to `switch_threads`

Switch to Kernel Thread for Process

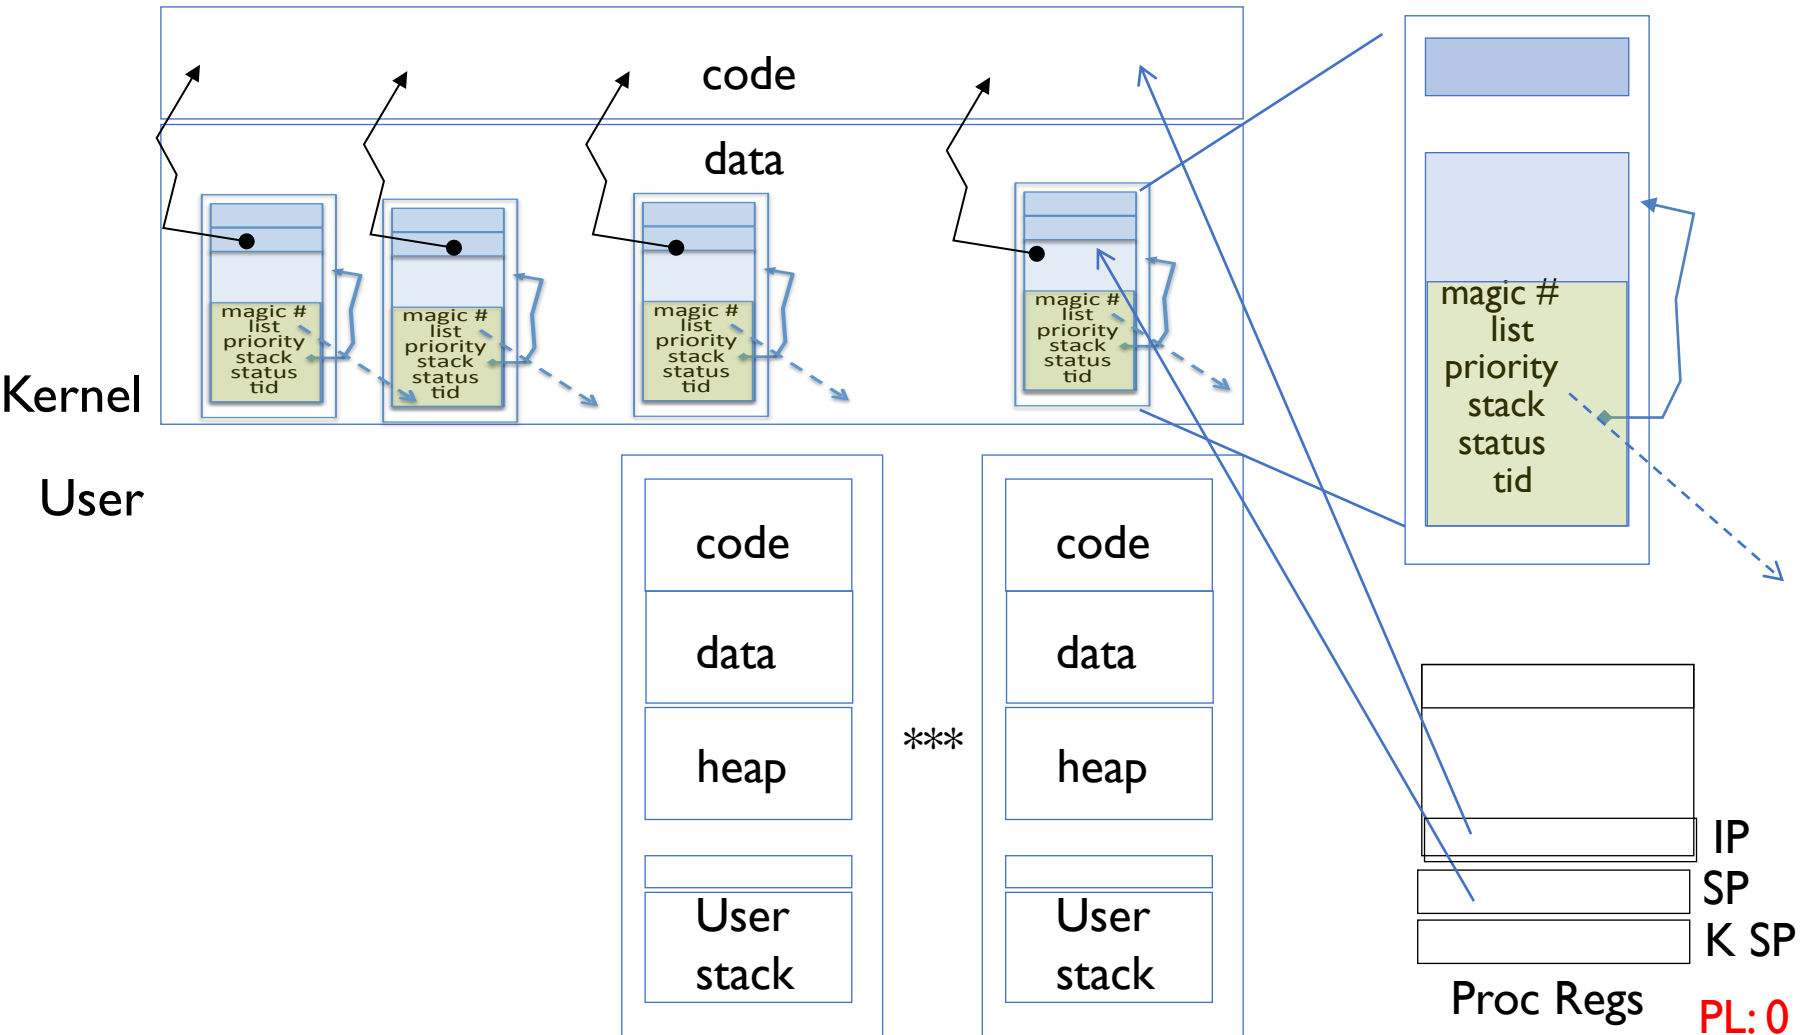


Kernel->User



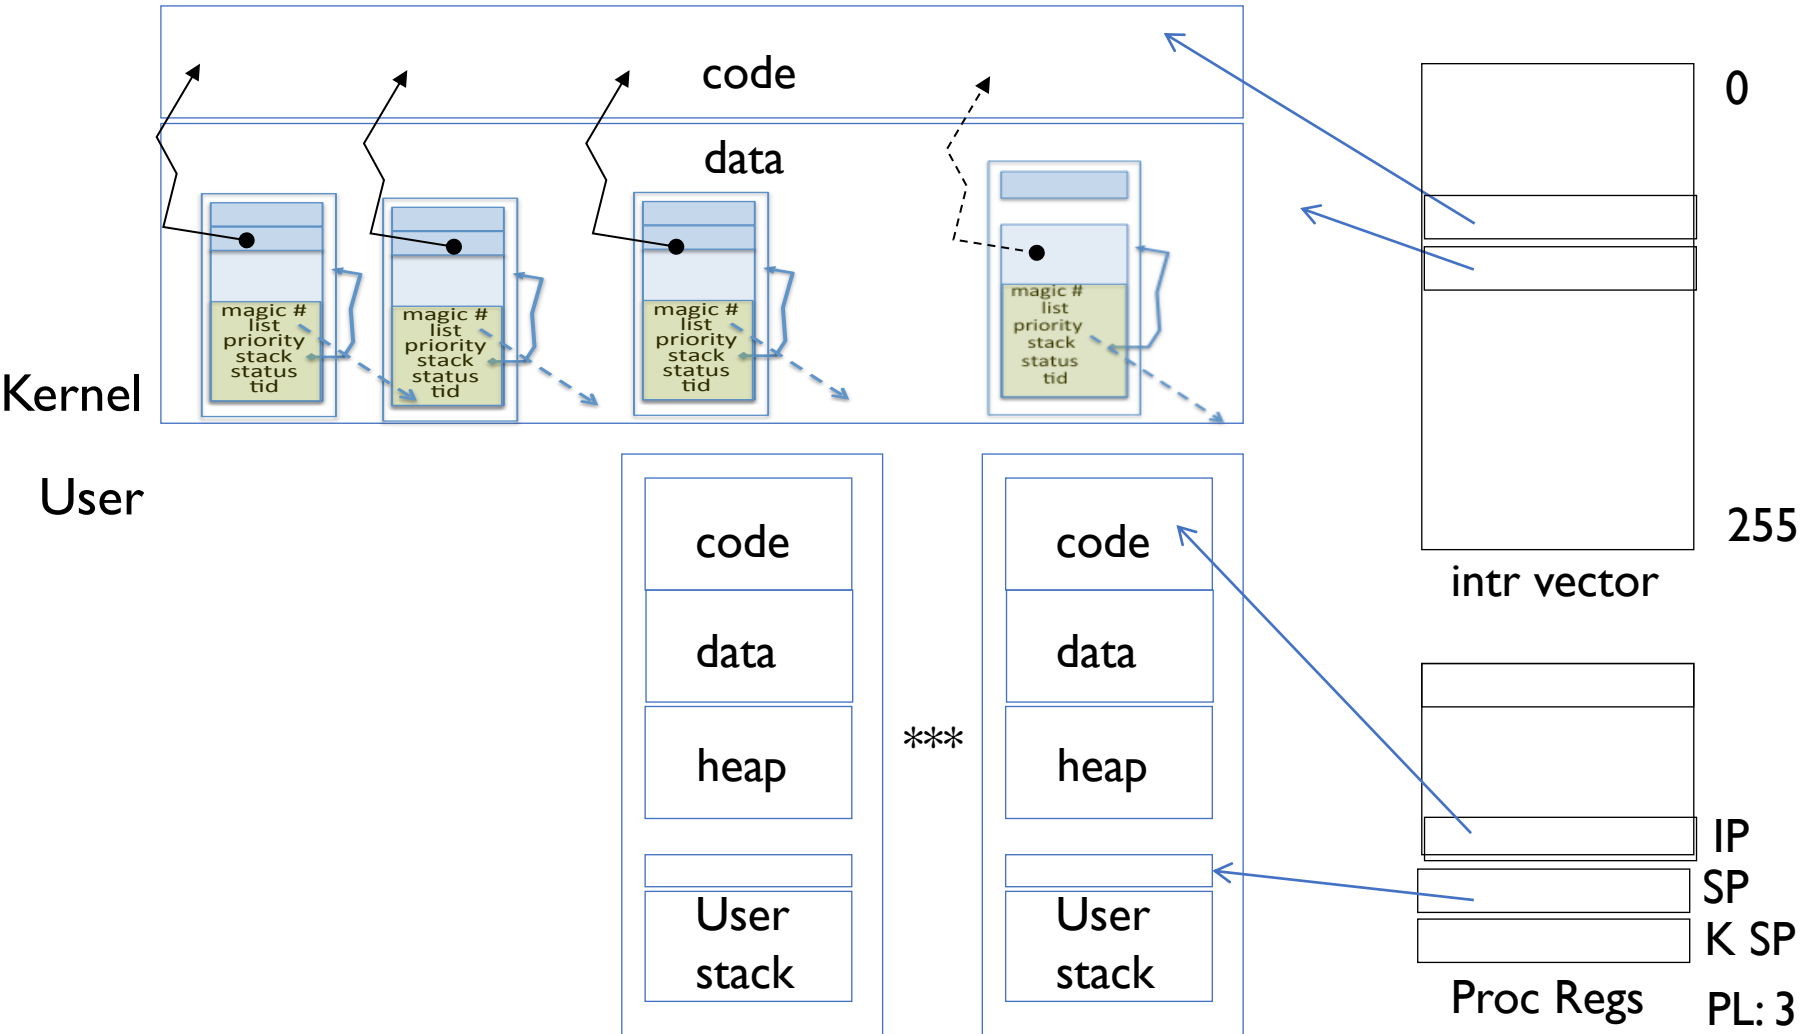
- **iret** restores user stack and PL

User->Kernel



- Mechanism to resume k-thread goes through interrupt vector

User->Kernel via interrupt vector



- Interrupt transfers control through the IV (IDT in x86)
- `iret` restores user stack and PL

Too Much Milk: Correctness

1. At most one person buys milk
2. At least one person buys milk if needed

Solution Attempt #1

- Leave a note
 - Place on fridge before buying
 - Remove after buying
 - Don't go to store if there's already a note
- Leaving/checking a note is atomic (word load/store)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

Attempt #1 in Action

Alice

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

Bob

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

Solution Attempt #2

```
Leave Note;
```

```
if (noMilk) {
```

```
  if (noNote) {
```

```
    leave Note;
```

```
    buy milk;
```

```
  }
```

```
}
```

```
remove Note;
```

← But there's always a note
– you just left one!

At least you don't
buy milk twice...

Solution Attempt #3

- Leave a named note – each person ignores their own

Alice

```
leave note Alice
```

```
if (noMilk) {  
    if (noNote Bob) {  
        buy milk  
    }  
}
```

```
remove note Alice;
```

Bob

```
leave note Bob
```

```
if (noMilk) {  
    if (noNote Alice) {  
        buy milk  
    }  
}
```

```
remove note Bob;
```

Attempt #3 in Action

Alice

leave note Alice

if (noMilk) {

 if (noNote Bob) {

~~buy milk~~

 }

}

remove note Alice

Bob

leave note Bob

if (noMilk) {

 if (noNote Alice) {

~~buy milk~~

 }

remove note Bob

Solution Attempt #4

Alice

```
leave note Alice
while (note Bob) {
    do nothing
}
if (noMilk) {
    buy milk
}
remove note Alice;
```

Bob

```
leave note Bob
if (noNote Alice) {
    if (noMilk) {
        buy milk
    }
}
remove note Bob;
```

- This is a correct solution, but ...

Issues with Solution 4

- Complexity
 - Proving that it works is hard
 - How do you add another thread?
- Busy-waiting
 - Alice **consumes CPU time to wait**
- Fairness
 - Who is more likely to buy milk?

OS Archaeology

- Because of the cost of developing an OS from scratch, most modern OSes have a long lineage:
- Multics → AT&T Unix → BSD Unix → Ultrix, SunOS, NetBSD,...
- Mach (micro-kernel) + BSD → NextStep → XNU → Apple OSX, iPhone iOS
- Linux → Android OS
- CP/M → QDOS → MS-DOS → Windows 3.1 → NT → 95 → 98 → 2000 → XP → Vista → 7 → 8 → phone → ...
- Linux → RedHat, Ubuntu, Fedora, Debian, Suse,...