

CS 162: Operating Systems and Systems Programming

Lecture 7: Synchronization Operations

September 19, 2019

Instructor: David Culler

<https://cs162.eecs.berkeley.edu>

Read: A&D Ch 5

What's a Process?

- The execution instance of a program
- Comprised of
 - One or more threads of execution
 - Processor registers, stack, OS support
 - A virtual address space
 - Page table maps resident pages to memory
 - Set of open file descriptors & buffers
 - User gets handle, OS holds the real thing
 - Whatever else the OS needs to load, run, manage, and terminate it

Grounding Demo

```
(base) CullerMac19:fa19 culler$ ps -al
  UID  PID  PPID      F CPU PRI NI      SZ    RSS WCHAN      S        ADDR  TTY        TIME CMD
    0 55813 55812    4106  0 31  0 4336432    12  -    Ss        0  ttys000    0:00.34 login -pf culler
  501 55814 55813    4006  0 31  0 4325240    948  -    S         0  ttys000    0:00.07 -bash
    0 61830 55814    4106  0 31  0 4270584   1120  -    R+        0  ttys000    0:00.00 ps -al
(base) CullerMac19:fa19 culler$ jobs
(base) CullerMac19:fa19 culler$ top
```

```
Processes: 395 total, 2 running, 393 sleeping, 1868 threads
Load Avg: 1.55, 1.32, 1.23  CPU usage: 0.96% user, 1.20% sys, 97.82% idle  SharedLibs: 280M resident, 48M data, 28M linkedit.
MemRegions: 116759 total, 3602M resident, 119M private, 1370M shared. PhysMem: 10G used (2601M wired), 6316M unused.
VM: 1763G vszize, 1372M framework vszize, 28191649(0) swapins, 29789745(0) swapouts. Networks: packets: 24399684/12G in, 10589805/3743M out.
Disks: 5514555/198G read, 6440971/229G written.
```

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	MEM	PURG	CMPRS	PGRP	PPID	STATE	BOOSTS	%CPU_ME	%CPU_OTHRS	UID	FAULTS	COW	MSGSENT
61839	screencaptur	0.1	00:00.19	5	3	192	13M	0B	0B	61839	1	sleeping	*0[127+]	0.00000	0.03601	501	8500	352	2717+
61838	screencaptur	1.4	00:00.32	2	1	54	3236K+	620K	0B	383	383	sleeping	*0[1]	0.03601	0.00000	501	15815+	207	3234+
61837	top	3.5	00:00.95	1/1	0	27	6024K	0B	0B	61837	55814	running	*0[1]	0.00000	0.00000	0	8675+	107	435408+
61836	QuickLookSat	0.0	00:00.04	2	1	46	3892K	0B	0B	61836	1	sleeping	*0[2]	0.00000	0.00000	501	4470	196	201
61835	quicklookd	0.0	00:00.06	4	1	95	3544K	72K	0B	61835	1	sleeping	*0[2]	0.00000	0.00000	501	4562	222	427
61833	mdworker_sha	0.0	00:00.05	3	1	63	3384K	0B	0B	61833	1	sleeping	*0[1]	0.00000	0.00000	501	4034	207	662
61828	Google Chrom	0.0	00:00.08	13	1	109	13M	4096B	0B	54320	54320	sleeping	*0[4]	0.00000	0.00000	501	9102	1771	401
61821	CoreServices	0.0	00:00.08	3	1	169	4580K	0B	0B	61821	1	sleeping	*0[1]	0.00000	0.00000	501	5023	264	1131+
61818	mdworker_sha	0.0	00:00.03	3	1	63	3316K	0B	0B	61818	1	sleeping	*0[1]	0.00000	0.00000	501	4108	205	631
61806	Google Chrom	0.0	00:00.86	14	1	150	30M	4096B	0B	54320	54320	sleeping	*0[5]	0.00000	0.00000	501	20218	1841	9138
61801	Google Chrom	0.0	00:06.06	14	1	141	37M	4096B	0B	54320	54320	sleeping	*0[6]	0.00000	0.00000	501	40084	1863	36873+
61769	mdworker_sha	0.0	00:00.12	3	1	59	4164K	0B	0B	61769	1	sleeping	*0[1]	0.00000	0.00000	501	7899	221	1303
61768	mdworker_sha	0.0	00:00.13	3	1	59	4292K	0B	0B	61768	1	sleeping	*0[1]	0.00000	0.00000	501	7944	221	1371
61767	mdworker_sha	0.0	00:00.13	3	1	59	4204K	0B	0B	61767	1	sleeping	*0[1]	0.00000	0.00000	501	7943	221	1383
61766	mdworker_sha	0.0	00:00.12	3	1	59	4092K	0B	0B	61766	1	sleeping	*0[1]	0.00000	0.00000	501	7891	221	1315
61762	mdworker_sha	0.0	00:00.05	4	1	52	4452K	0B	0B	61762	1	sleeping	*0[1]	0.00000	0.00000	89	6455	189	613
61726	Google Chrom	0.0	00:00.24	15	2	146	16M	4096B	0B	54320	54320	sleeping	*0[6]	0.00000	0.00000	501	12886	1827	3841
61623	netbiosd	0.0	00:01.07	7	7	29	2912K	0B	0B	61623	1	sleeping	*0[1]	0.00000	0.00000	222	3904	158	159
61621	eapolclient	0.0	00:00.18	3	1	52	3108K	0B	0B	61621	56	sleeping	*0[1]	0.00000	0.00000	501	3902	352	445
61571	mdworker_sha	0.0	00:00.72	3	1	63	14M	0B	0B	61571	1	sleeping	*0[1]	0.00000	0.00000	501	7201	209	938
61569	BoxEditFinde	0.0	00:00.40	3	1	159	5208K	0B	0B	61569	1	sleeping	*0[27]	0.00000	0.00000	501	5410	303	4025+
61565	com.apple.ap	0.0	00:02.83	3	1	349	37M	220K	0B	61565	1	sleeping	*0[144]	0.00000	0.00000	501	33441	406	21061
61478	sandboxd	0.0	00:04.03	4	3	66	5776K	0B	0B	61478	1	sleeping	*0[1]	0.00000	0.00000	0	241756	169	311257
61423	com.microsof	0.0	00:00.44	2	1	46	3640K	0B	2588K	61423	1	sleeping	*0[1]	0.00000	0.00000	0	6620	283	1303
56470	Microsoft Up	0.0	00:03.51	6	2	200-	12M-	128K	4504K	56470	1	sleeping	*0[120]	0.00000	0.00000	501	24520	392	9437+
56468	Microsoft Po	0.0	24:10.79	33	12	2078	644M	189M	107M	56468	1	sleeping	*0[2207]	0.00000	0.00000	501	18023272	677881	7539121
56389	usbmuxd	0.0	00:00.06	3	1	38	1404K	0B	712K	56389	1	sleeping	*0[1]	0.00000	0.00000	213	2287	146	430
56387	zoom.us	0.0	09:26.93	16	1	843	149M	3260K	115M	56387	1	sleeping	*0[345]	0.00000	0.00000	501	14397377	1225	2446416
56332	mdworker_sha	0.0	00:00.16	3	1	58	4272K	0B	2816K	56332	1	sleeping	*0[1]	0.00000	0.00000	501	6110	188	1425
56331	LookupViewSe	0.0	00:00.73	3	1	173	10M	0B	6152K	56331	1	sleeping	*0[246]	0.00000	0.00000	501	16374	365	6213
55956	Google Chrom	0.0	00:44.99	18	2	223	116M	0B	17M	54320	54320	sleeping	*0[5]	0.00000	0.00000	501	229433	7289	201102+
55814	bash	0.0	00:00.07	1	0	21	1020K	0B	452K	55814	55813	sleeping	*0[1]	0.00000	0.00000	501	3189	1525	137
55813	login	0.0	00:00.34	2	1	31	1252K	0B	1236K	55813	55812	sleeping	*0[9]	0.00000	0.00000	0	1810	196	121
55812	Terminal	1.1	00:19.52	8	3	366	53M-	8692K	17M	55812	1	sleeping	*0[544]	0.00100	0.00000	501	399265	451	95590+
55605	PerfPowerSer	0.0	00:11.35	2	1	121	9728K	256K	2424K	55605	1	sleeping	*0[165]	0.00000	0.00000	0	22444	207	55139
55597	BoxEditFinde	0.0	00:01.75	3	1	159	6128K	0B	2020K	55597	1	sleeping	*0[132]	0.00000	0.00000	501	9193	304	16116+
55576	SimulatorTra	0.0	00:01.32	3	1	137	5672K	0B	2132K	55576	1	sleeping	*0[112]	0.00000	0.00000	501	7642	255	14687+
55575	com.apple.Co	0.0	00:01.60	3	1	142	7016K	0B	2960K	55575	1	sleeping	*0[115]	0.00000	0.00000	501	8765	289	14991+
55520	Google Chrom	0.0	00:01.36	6	1	72	46M	0B	9984K	54320	54320	sleeping	*0[4]	0.00000	0.00000	501	31158	1782	6806
55519	Google Chrom	0.0	00:18.44	13	1	151	91M	0B	23M	54320	54320	sleeping	*0[7]	0.00000	0.00000	501	150819	1996	107587

Recall: Scheduling

- **First-Come First-Served:** Simple, vulnerable to convoy effect
- **Round-Robin:** Fixed CPU time quantum, cycle between ready threads
- **Priority:** Respect differences in importance
- **Shortest Job/Remaining Time First:** Optimal for average response time, but unrealistic
- **Multi-Level Feedback Queue:** Use past behavior to approximate SRTF and mitigate overhead

Impacts of Scheduling on ...

- Lot's of attention to algorithmic complexity of operations on the scheduling data structure
 - These queues don't get that long. Otw, buy more hardware
- Interactions of scheduling with memory hierarchy
 - Locality is fundamentally at odds with fairness
 - “Cache / VM / File buffer *affinity*”
- Interactions of scheduling with multiple processors
 - Processor / Core affinity is really about caches
- Memory performance (locality) is critical

System Design ...

- Sophisticated policies (often with deep theoretical basis) boil down into simple manipulation of data structures.
- And understanding multi-dimensional interactions
- We'll return to advanced scheduling (with randomness) later in the term



Going back – to a subtle connection between scheduling and synchronization

- Why do we need to disable interrupts at all?
 - Avoid interruption between checking and setting lock value
 - Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

**Critical
Section**

Recall: Basic Lock Implementation



```
Acquire(*lock) {  
    disable interrupts;  
    if (lock->value == BUSY) {  
        put thread on lock's wait_Q  
        "i.e, Go to sleep"  
        allow a ready thread to run  
    } else {  
        lock->value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release(*lock) {  
    disable interrupts;  
    if (any TCB on lock wait_Q) {  
        "i.e., lock busy";  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        lock->value = FREE;  
    }  
    enable interrupts;  
}
```


Reenabling Interrupts When Waiting

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        run_new_thread()  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

enable interrupts →

enable interrupts →

- Before on the queue?
 - Release might not wake up this thread!
- After putting the thread on the queue?
 - Gets woken up, but immediately switches away

Reenabling Interrupts When Waiting


```
Acquire() {
    disable_interrupts();
    while (true) {
        if (condition) {
            return;
        }
        wait();
    }
    enable_interrupts();
}
```

```
Acquire() {
    disable interrupts;
    if (value == BUSY) {
        put thread on wait queue;
enable interrupts → run_new_thread()
    } else {
        value = BUSY;
    }
    enable interrupts;
}
```

- Best solution: after the current thread suspends
- How?
 - `run_new_thread()` should do it!
 - Part of returning from `switch()`

How to Re-enable After Sleep()

- In scheduler, since interrupts are disabled when you call sleep:
 - Responsibility of the next thread to re-enable ints
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts



User-level threads?


- Can multiple threads be implemented entirely at user level?
- Most other aspects of system virtualize.

Kernel-Supported Threads

- Threads run and block (e.g., on I/O) independently
- One process may have multiple threads waiting on different things
- Two mode switches for every context switch (expensive)
- Create threads with syscalls
- Alternative: multiplex several streams of execution (at user level) on top of a single OS thread
 - E.g., Java, Go, ... (and many many user-level threads libraries before it)

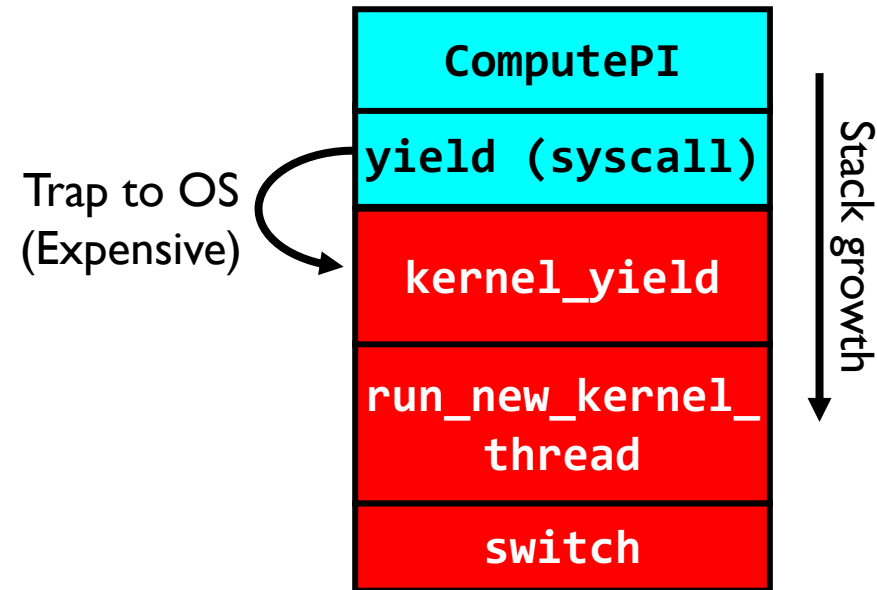
User-Mode Threads

- User program contains its own scheduler
- Several user threads per kernel thd.
- User threads may be scheduled **non-preemptively**
 - Only switch on yield
- Context switches cheaper
 - Copy registers and jump (switch in userspace)




Thread Yield

Kernel-Supported Threads




User-Mode Threads




Thread I/O

Kernel-Supported Threads



User-Mode Threads




- Selects a new *kernel thread* to run
- Bypassing user-level scheduler

User-Mode Threads: Problems

- One user-level thread blocks on I/O: they all do
 - Kernel cannot adjust scheduling among threads it doesn't know about
- Multiple Cores?
- Can't completely avoid blocking (syscalls, page fault)
- One Solution: *Scheduler Activations*
 - Have kernel inform user-level scheduler when a thread blocks
- Evolving the contract between OS and application.

Recall: Multithreaded Server


- **Bounded** pool of worker threads
 - Allocated in **advance**: no thread creation overhead
 - **Queue** of pending requests



Simple Performance Model

- Given that the overhead of a critical section is X
 - User->Kernel Context Switch
 - Acquire Lock
 - Kernel->User Context Switch
 - <perform exclusive work>
 - User->Kernel Context Switch
 - Release Lock
 - Kernel->User Context Switch
- Even if everything else is infinitely fast, with any number of threads and cores
- What is the maximum rate of operations that involve this overhead?

Highly Contended Case – in a picture



Back to system performance

More Practical Motivation

Back to Jeff Dean's "Numbers everyone should know"

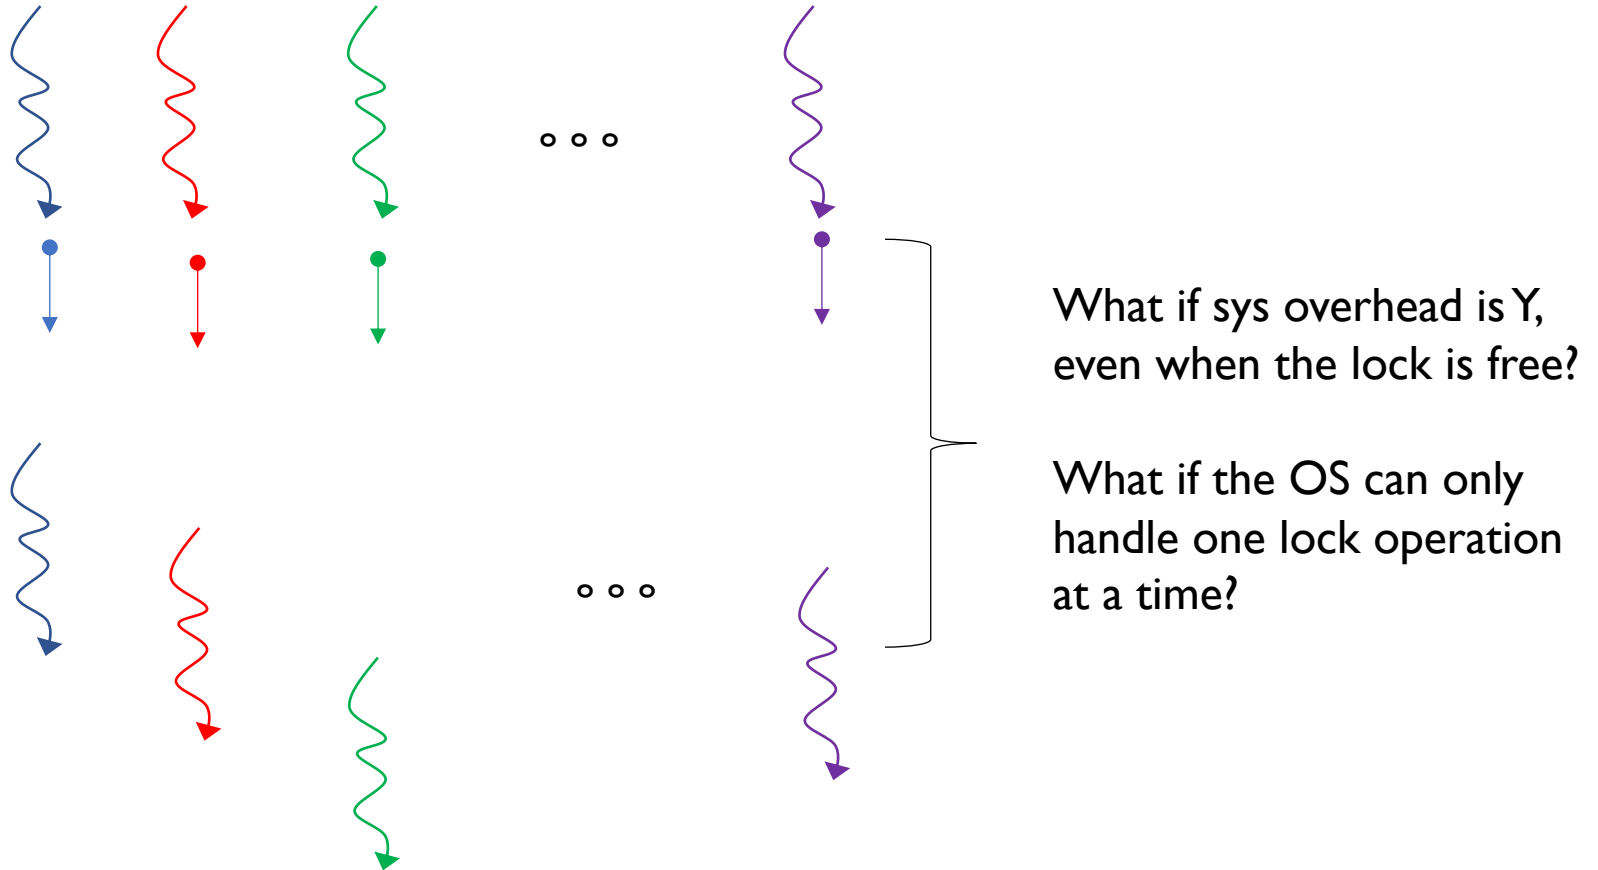
Handle I/O in
separate thread,
avoid blocking
other progress

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns




- $X = 1\text{ms} \Rightarrow 1,000 \text{ ops/sec}$

Uncontended Many-Lock Case



Basic cost of a system call



- Min System call ~ 25x cost of function call
- Scheduling could be many times more
- Streamline system processing as much as possible
- Other optimizations seek to process as much of the call in user space as possible (eg, Linux vDSO)

A Better Lock Implementation

- Interrupt-based solution works for single core, but costly
- Doesn't work well on multi-core machines
 - Disable intr on all cores?
- Solution: Utilize hardware support for **atomic operations**

Recall: Atomic Operations

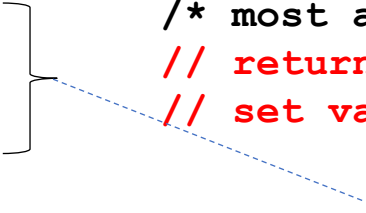
- Definition: **An operation runs to completion or not at all**
- Foundation for synchronization primitives
- Example: Loading or storing a word

Atomic Read-Modify-Write Instructions

- Problems with previous solution:
 - Works only in system Privilege level, not User level
 - Doesn't work well on multiprocessor
 - Disabling interrupts on all processors time consuming and undermines HW parallelism
- Alternative: **atomic instruction sequences**
 - These instructions read a value and write a new value atomically
 - Hardware is responsible for implementing this correctly
 - on both uniprocessors (not too hard)
 - and multiprocessors (requires help from cache coherence protocol)
 - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors
 - Natural extensions to user-level locking

Examples of Read-Modify-Write

- ```
test&set (&address) {
 result = M[address];
 M[address] = 1;
 return result;
}
```



```
/* most architectures */
// return result from "address" and
// set value at "address" to 1
```

as if it happened all at once
- ```
swap (&address, register) {  
    temp = M[address];  
    M[address] = register;  
    register = temp;  
}
```

```
/* x86 */  
// swap register's value to  
// value at "address"
```
- ```
compare&swap (&address, reg1, reg2) { /* 68000 */
 if (reg1 == M[address]) {
 M[address] = reg2;
 return success;
 } else {
 return failure;
 }
}
```

```
// If memory still == reg1,
// then put reg2 => memory

// Otherwise do not change memory
```
- ```
load-linked&store-conditional(&address) { /* R4000, alpha */  
    loop:  
        ll r1, M[address];  
        movi r2, 1;  
        sc r2, M[address];  
        beqz r2, loop;  
}
```

```
// Can do arbitrary computation
```

Implementing Locks with test&set

- Simple, but flawed, solution:

```
int value = 0; // Free
Acquire() {
    while (test&set(value)) {}; // spin while busy
}
Release() {
    value = 0; // atomic store
}
```

- Simple explanation:

- If lock is free, test&set reads 0 and sets value=1, so lock is now busy. It returns 0 so while exits.
- If lock is busy, test&set reads 1 and sets value=1 (no change) It returns 1, so while loop continues.
- When we set value = 0, someone else can get lock.

- **Busy-Waiting**: thread consumes cycles while waiting

- For multiprocessors: every test&set() is a write, which makes value ping-pong around in cache (using lots of memory BW)

Problem: Busy-Waiting for Lock

- Positives for this solution
 - Machine can receive interrupts
 - User code can use this lock (poorly)
 - Works on a multiprocessor
- Negatives
 - This is very inefficient as thread will consume cycles waiting
 - Waiting thread may take cycles away from thread holding lock (no one wins!)
 - **Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock \Rightarrow no progress!
- For semaphores (and monitors), waiting thread may wait for an arbitrary long time!
 - Thus even if busy-waiting was OK for locks, definitely not ok for other primitives
 - Homework/exam solutions should avoid busy-waiting!



Multiprocessor Spin Locks: test&test&set

- A better solution for multiprocessors:

```
int mylock = 0; // Free

Acquire() {
    do {
        while(mylock); // Wait until might be free
    } while(test&set(&mylock)); // exit if get lock
}

Release() {
    mylock = 0;
}
```

- Simple explanation:
 - Wait until lock might be free (only reading – stays in cache)
 - Then, try to grab lock with test&set
 - Repeat if fail to actually get lock
- Issues with this solution:
 - **Busy-Waiting**: thread still consumes cycles while waiting
 - However, it does not impact other processors!

Better Locks using test&set

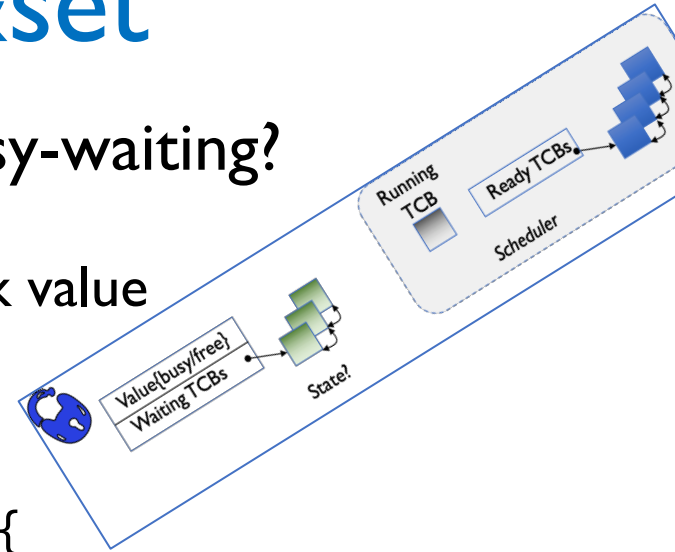
- Can we build test&set locks without busy-waiting?
 - Can't entirely, but can minimize!
 - Idea: only busy-wait to atomically *check* lock value

```
int guard = 0;  
int value = FREE;
```



```
Acquire() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if (value == BUSY) {  
        put thread on wait queue;  
        go to sleep() & guard = 0;  
    } else {  
        value = BUSY;  
        guard = 0;  
    }  
}
```

```
Release() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if anyone on wait queue {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    guard = 0;  
}
```




- Note: sleep has to be sure to reset the guard variable
 - Why can't we do it just before or just after the sleep?

Locks using Interrupts vs. test&set

Recall “disable interrupt” solution

```
int value = FREE;
```



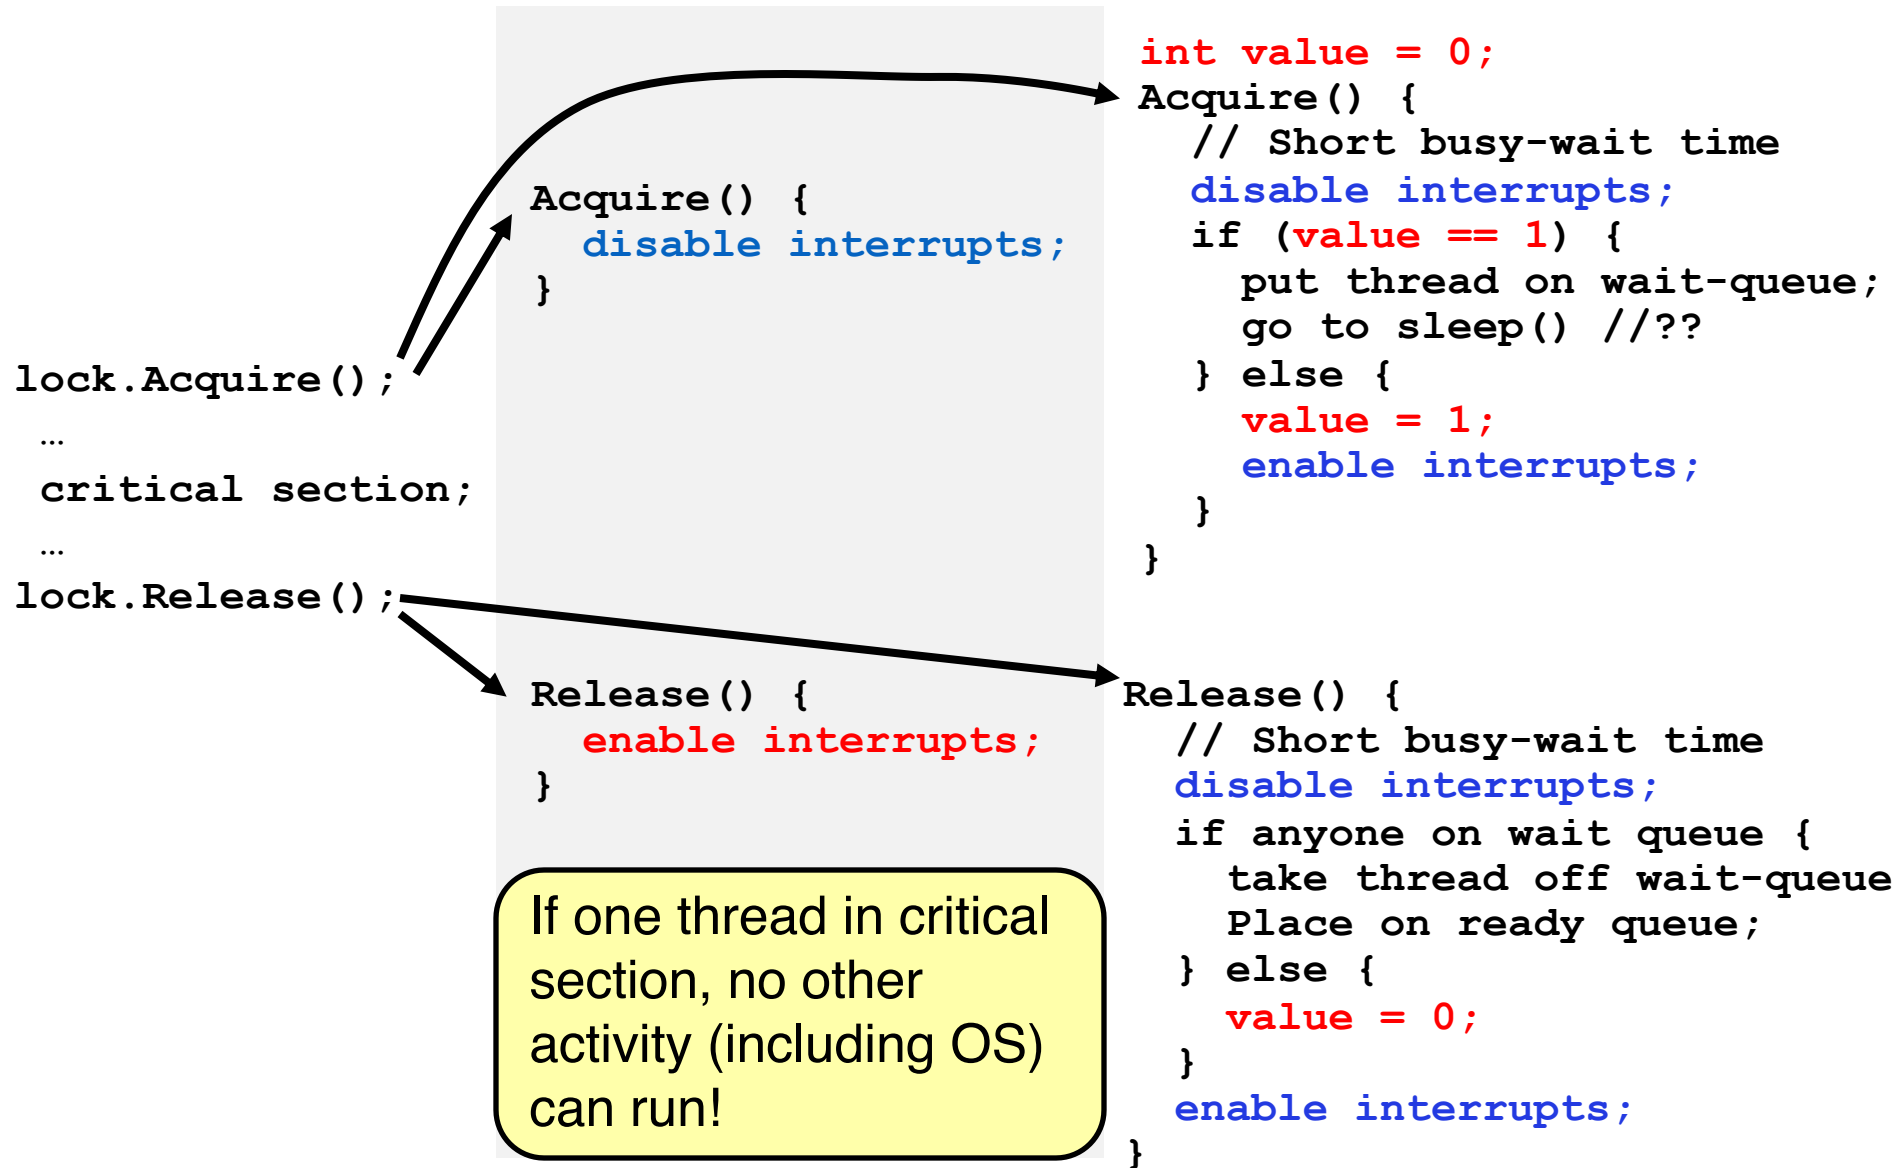
```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

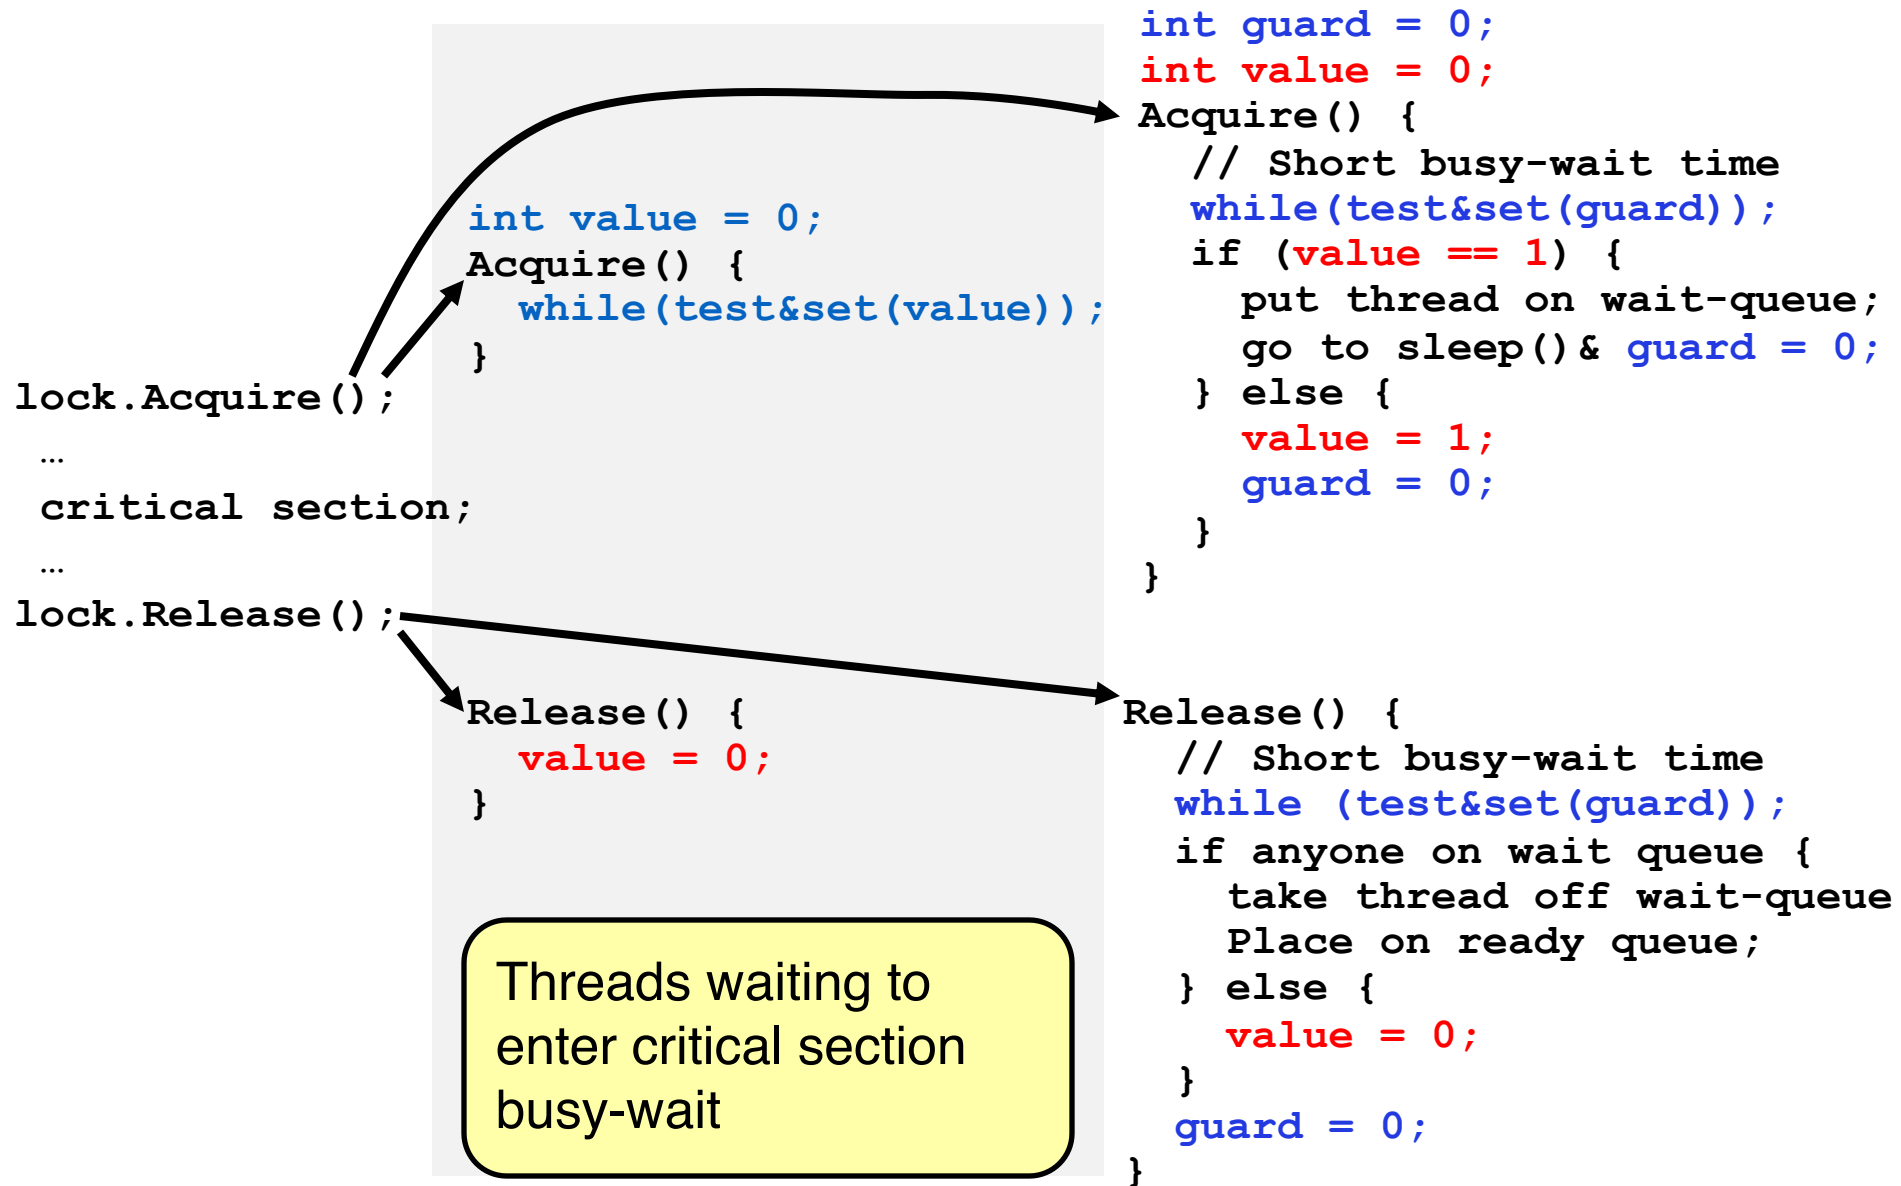
Basically we replaced:

- `disable interrupts` → `while (test&set(guard));`
- `enable interrupts` → `guard = 0;`

Recap: Locks using interrupts




Recap: Locks using test & set



Higher-level Primitives

- Goal of last couple of lectures:
 - What is right abstraction for synchronizing threads that share memory?
 - Want as high a level primitive as possible
- Good primitives and practices important!
 - Since execution is not entirely sequential, really hard to find bugs, since they happen rarely
 - UNIX is pretty stable now, but up until about mid-80s (10 years after started), systems running UNIX would crash every week or so – concurrency bugs
- Synchronization is a way of coordinating multiple concurrent activities that are using shared state
 - Requires both mutual exclusion and cooperation (or orchestration)


Producer-Consumer with a Bounded Buffer



- Problem Definition
 - Producers puts things into a shared buffer
 - Consumers takes them out
- Don't want producers and consumers to have to work in lockstep, so put a buffer (bounded) between them
 - Need synchronization to maintain integrity of the data structure and coordinate producers/consumers
 - Producer needs to wait if buffer is full
 - Consumer needs to wait if buffer is empty
- GCC compiler – simple 1–1
 - `cpp | cc1 | cc2 | as | ld`
- Web servers, Routers, ...

Circular Buffer Data Structure (sequential case)

```
typedef struct buf {  
    int write_index;  
    int read_index;  
    <type> *entries[BUFSIZE];  
} buf_t;
```



- Insert: write & bump write ptr (enqueue)
- Remove: read & bump read ptr (dequeue)
- *How to tell if Full (on insert) Empty (on remove)?*
- *And what do you do if it is?*
- *What needs to be atomic?*


Producer/Consumer Correctness

- With multiple threads, each waits for the other to make process
- Scheduling constraints:
 - Consumer waits for producer if buffer is empty
 - Producer waits for consumer if buffer is full
- Mutual Exclusion: Only one thread manipulates the buffer data structure at a time

Lock Solution – first cut

mutex buf_lock = <initially unlocked>

```
Producer(item) {  
    lock buffer  
    while (buffer full) {}; // Wait for a free slot  
    Enqueue(item);  
    unlock buffer  
}
```



Will we ever come out of the wait loop?


```
Consumer() {  
    lock buffer  
    while (buffer empty) {}; // Wait for arrival  
    item = queue();  
    unlock buffer  
    return item  
}
```

Lock Solution – 2nd cut

mutex buf_lock = <initially unlocked>

```
Producer(item) {  
    lock buffer  
    while (buffer full) {unlock; lock;};  
    Enqueue(item);  
    unlock buffer  
}
```

```
Consumer() {  
    lock buffer  
    while (buffer empty) {unlock; lock;};  
    item = queue();  
    unlock buffer  
    return item  
}
```



What happens when one is waiting for the other?

- Multiple cores ?
- Single core ?

Explore semaphore solution

- **One semaphore per constraint**
 1. Mutex (mutual exclusion)
 2. Filled Slots (consumer waits if necessary)
 3. Empty Slots (producer waits if necessary)

Producer/Consumer Code

```
Semaphore fullSlots = 0; // Buffer empty to start
Semaphore emptySlots = bufSize; // All slots empty
Semaphore mutex = 1; // No one in critical sect.
```

```
Producer(item) {
    emptySlots.P(); // Wait for a free slot
    mutex.P();      // down
    Enqueue(item);
    mutex.V();      // up
    fullSlots.V(); // Tell consumers about new data
}
```

Producer/Consumer Code

```
Semaphore fullSlots = 0; // Queue empty to start  
Semaphore emptySlots = bufSize; // All slots empty  
Semaphore mutex = 1; // No one in critical sect.
```

```
Consumer() {  
    fullSlots.P(); // Wait for an item to be present  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V(); // Tell producers about new slot  
    return item;  
}
```

Producer/Consumer Code

```
Semaphore fullSlots = 0; // Queue empty to start
Semaphore emptySlots = bufSize; // All slots empty
Semaphore mutex = 1; // No one in critical sect.
```

```
Producer(item) {
    emptySlots.P();
    mutex.P();
    Enqueue(item);
    mutex.V();
    fullSlots.V();
}
```

```
Consumer() {
    fullSlots.P();
    mutex.P();
    item = Dequeue();
    mutex.V();
    emptySlots.V();
    return item;
}
```

Discussion

- **What if we wrote the following?**

```
Producer(item) {  
    mutex.P();  
    emptySlots.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}
```

```
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

Deadlock ... More on this later

Discussion

- **What about this?**

```
Producer(item) {  
    emptySlots.P();  
    mutex.P();  
    Enqueue(item);  
    fullSlots.V();  
    mutex.V();  
}
```

```
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

Correct, possibly less efficient

Problems with Semaphores

- More powerful (and primitive) than locks
- Argument: Clearer to have separate constructs for
 - Mutual Exclusion: One thread can do something at a time
 - Waiting for a condition to become true
- Need to make sure a thread calls $P()$ for every $V()$
 - Other tools are more flexible than this

Break

Condition Variables

- **Collection of threads waiting *inside* a critical section**
- **Operations:**
 - **wait(&lock):** Atomically release lock and go to sleep. **Re-acquire** the lock before returning.
 - **signal():** Wake up on waiting thread (if there is one)
 - **broadcast():** Wake up all waiting threads
- **Rule:** Hold lock when using a condition variable


Lock Solution – 2nd cut

mutex buf_lock = <initially unlocked>

Condvar buf_signal = <initially nobody>

```
Producer(item) {  
    lock buffer  
    while (buffer full) {cond_wait(buf_signal, buf_lock) };  
    Enqueue(item);  
    unlock buffer  
}
```

```
Consumer() {  
    lock buffer  
    while (buffer empty) {cond_wait(buf_signal, buf_lock) };  
    item = queue();  
    unlock buffer  
    return item  
}
```




Release lock; signal others to run; reacquire on resume

n.b. OS must do the reacquire
Why User must recheck?

Why the while Loop?

- When a thread is woken up by `signal()`, it is simply put on the ready queue
- It may or may not reacquire the lock immediately!
 - Another thread could be scheduled first and "sneak in" to empty the queue
 - Need a loop to re-check condition on wakeup

Monitors




- **Lock:** protects access to shared data
- Always acquire lock when accessing
- Queue of threads waiting to enter the monitor

Monitors in practice

- Locks for mutual exclusion
- Condition variables for waiting
- A **monitor** is a lock and zero or more condition variables with some associated data and operations
 - Java provides this natively
 - POSIX threads: Provides **locks** and **condvars**, have to build your own

Monitors



- **Condition Variables:** queue of threads waiting for something to become true inside critical sect.
- Atomically release lock and start waiting
 - Another thread using the monitor will signal them
- The condition: Some function of monitor's data

Why the `while` Loop?

- Can we "hand off" the lock directly to the signaled thread so no other thread "sneaks in?"
 - Yes. Called **Hoare-Style Monitors**
 - Many textbooks describe this scheme
- Most OSs implement **Mesa-Style Monitors**
 - Allows other threads to sneak in
 - Much easier to implement
 - Even easier if you allow "spurious wakeups"
 - `wait()` can return when no signal occurred, in rare cases
 - POSIX allows spurious wakeups

Interlude: Concurrency Is Hard


- Even for practicing engineers trying to write mission-critical, bulletproof code!
- Therac-25: Radiation Therapy Machine with Unintended Overdoses (reading on course site)
- Mars Pathfinder Priority Inversion ([JPL Account](#))
- Toyota Uncontrolled Acceleration ([CMU Talk](#))
 - 256.6K Lines of C Code, ~9-11K global variables
 - Inconsistent mutual exclusion on reads/writes

Comparing Synchronization

- Semaphores can implement locks
 - `Acquire()` { `semaphore.P()`; }
 - `Release()` { `semaphore.V()`; }
- and Condition Variables
- Monitors combine locks and CVs in a structured fashion
- Modern view: concurrent objects (e.g., Java)
- Can monitors implement semaphores?
- Are there other important common patterns?

Time Permitting ...

Reader/Writer Problem



- Shared Database
 - Many readers – never modify the database
 - Few writers – read and modify database
- Single lock sufficient?

Reader/Writer Correctness

- Readers can access when no writers
- Writers can access when no readers **and no other writers**
- A lock will satisfy these requirements
 - But we want to allow **multiple readers**
 - Better efficiency

Reader/Writer with Monitors

```
Reader() {  
    Wait until no active writers  
    Access database  
    Maybe wake up a writer  
}  
  
Writer() {  
    Wait until no active readers or writers  
    Access database  
    Maybe wakeup reader or writer  
}
```

Lock (for mutual exclusion)

int activeReaders, condVar okToRead

int activeWriters, condVar okToWrite

Reader Version I

```
Reader() {  
    // First check self into system  
    lock.Acquire();  
    while (AW > 0) { // Is it safe to read?  
        okToRead.wait(&lock); // Sleep on cond var  
    }  
    AR++; // Now we are active!  
    lock.release();  
    // Perform actual read-only access  
    AccessDatabase(ReadOnly);  
    // Now, check out of system  
    lock.Acquire();  
    AR--; // No longer active  
    if (AR == 0) // No other active readers  
        okToWrite.signal(); // Wake up one writer  
    lock.Release();  
}
```

Writer Version I

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while (AR > 0 || AW > 0) { // Is it safe to write?  
        okToWrite.wait(&lock); // Sleep on cond var  
    }  
    AW++; // Now we are active!  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--; // No longer active  
    okToWrite.signal(); // Wake up one writer  
    okToRead.broadcast(); // Wake up all readers  
    lock.Release();  
}
```

Writer Version I: Starvation

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while (AR > 0 || AW > 0) {  
        okToWrite.wait(&lock);  
    }  
    AW++;  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--;  
    okToWrite.signal();  
    okToRead.broadcast();  
    lock.Release();  
}
```

If there are always readers, this is always true! Writer starves

Writer Version I: Conflict

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while (AR > 0 || AW > 0) { // Is it safe to write?  
        okToWrite.wait(&lock); // Sleep on cond var  
    }  
    AW++; // Now we are active!  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--; // No longer active  
    okToWrite.signal(); // Wake  
    okToRead.broadcast(); // Wake  
    lock.Release();  
}
```

**Relies on waiting threads
double-checking condition**

Writer Version 1: Conflict

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while (AR > 0 || AW > 0) { // Is it safe to write?  
        okToWrite.wait(&lock); // Sleep on cond var  
    }  
    AW++; // Now we are active!  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--; // No longer active  
    okToWrite.signal(); // Wake  
    okToRead.broadcast(); // Wake  
    lock.Release();  
}
```

Everyone races, all but 1 thread just goes back to sleep

Reader/Writer with Monitors v2

```
Reader() {  
    Wait until no active or waiting writers  
    Access database  
    Maybe wake up a writer  
}  
Writer() {  
    Wait until no active readers or writers  
    Access database  
    If waiting writer, wake it up;  
    Otherwise, wakeup readers;  
}  
  
int waitingWriters
```

Reader Version 2

```
Reader() {  
    // First check self into system  
    lock.Acquire();  
    while (AW > 0 || WW > 0) { // Is it safe to read?  
        okToRead.wait(&lock); // Sleep on cond var  
    }  
    AR++; // Now we are active!  
    lock.release();  
    // Perform actual read-only access  
    AccessDatabase(ReadOnly);  
    // Now, check out of system  
    lock.Acquire();  
    AR--; // No longer active  
    if (AR == 0) // No other active readers  
        okToWrite.signal(); // Wake up one writer  
    lock.Release();  
}
```

Writer Version 2

```
Writer() {  
    // First check self into system  
    lock.Acquire();  
    while (AR > 0 || AW > 0) { // Is it safe to write?  
        WW++;  
        okToWrite.wait(&lock); // Sleep on cond var  
        WW--;  
    }  
    AW++; // Now we are active!  
    lock.release();  
    // Perform actual read/write access  
    AccessDatabase(ReadWrite);  
    // Now, check out of system  
    lock.Acquire();  
    AW--; // No longer active  
    if (WW > 0)  
        okToWrite.signal(); // Wake up one writer  
    else  
        okToRead.broadcast(); // Wake up all readers  
    lock.Release();  
}
```

Simulation of Reader/Writer

- Sequence of arrivals: R1, R2, W1, R3

- On entry each reader checks

```
while (AW > 0 || WW > 0) { // Is it safe to read?  
    okToRead.wait(&lock); // Sleep on cond var  
}
```

- First R1 enters (no waiting)
 - **AR = 1**, AW = 0, WW = 0
- Then R2 enters (no waiting)
 - **AR = 2**, AW = 0, WW = 0

Simulation of Reader/Writer

- Sequence of arrivals: R1, R2, *W1, R3
- **R1, R2 still running (AR = 2)**
- **W1 does a check: $AR > 0$, waits on okToWrite**

```
while (AR > 0 || AW > 0) { // Is it safe to write?  
    WW++;  
    okToWrite.wait(&lock); // Sleep on cond var  
    WW--;  
}
```

- Now $AR = 2, AW = 0, \mathbf{WW} = 1$
- R3: $\mathbf{WW} > 0$, waits on okToRead

Simulation of Reader/Writer

- R1 finishes, does not wake anyone up
 - $AR = 1, AW = 0, WW = 1$
- R2 finishes
 - $AR = 0, AW = 0, WW = 1$
 - Wakes up W1 (signals **okToWrite**)
- W1 runs and finishes
 - $AR = 1, AW = 1$ then $0, WW = 0$
 - Wakes up R3 (**okToRead.Broadcast()**)

Reader/Writer Design Choices

- Reader starvation:

```
while (AW > 0 || WW > 0) { // Safe to read?  
    okToRead.wait(&lock); // Sleep on cond var  
}
```

- "Writer-biased" Lock

- Can favor readers by changing conditions on wait loops
- Other possibilities, e.g. track readers waiting since before current writer started

Summary

- Scheduling and Synchronization are Deeply Interrelated
- Synchronization overhead is a critical performance factor
- User-level Threads can remove OS-switch cost of synchronization, but lose the connection with scheduler
 - With lots of cores, this matters less
- Disabling interrupts is brute-force way to implement synchronization operations.
 - Does not play well with multiple cores. Cannot be used at User Level
- Hardware atomic read-modify-write provides a better solution
- Must be constructed carefully – spin on simple read (test & test-and-set)
- Synchronization involves both Mutual Exclusion and Signaling
 - Locks for Mutex, Condition Variables for signaling (cooperation)
- Semaphores: More primitive & general than locks, but used in both ways
- Alternative: Monitors
 - One lock, zero or more condition variables
- Reader/Writer Synchronization
 - Treat readers differently from writers for efficiency