

# Introduction to Communication - Sockets

David E. Culler  
CS162 – Operating Systems and Systems  
Programming  
Lecture 8  
Sept 24, 2019

**HW 2 out, Due 10/12**  
**Proj 1**

# Producer/Consumer Problem

---

- With multiple threads, each waits for the other to make process
- Scheduling constraints:
  - Consumer waits for producer if buffer is empty
  - Producer waits for consumer if buffer is full
- Mutual Exclusion: Only one thread manipulates the buffer data structure at a time

# Condition Variables

---

- Collection of threads waiting *inside* a critical section
- Operations:
  - **wait(&lock)**: Atomically release lock and go to sleep. **Re-acquire** the lock before returning.
  - **signal()**: Wake up one waiting thread (if there is one)
  - **broadcast()**: Wake up all waiting threads
- **Rule**: Hold lock when using a condition variable

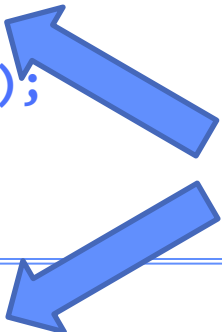
# Producer/Consumer Buffer – 2<sup>nd</sup> cut

mutex buf\_lock = <initially unlocked>

Condvar buf\_signal = <initially nobody>

```
Producer(item) {  
    lock buffer  
    while (buffer full) { cond_wait(buf_signal, buf_lock) };  
    Enqueue(item);  
    cond_broadcast(buf_signal);  
    unlock buffer  
}
```

```
Consumer() {  
    lock buffer  
    while (buffer empty) { cond_wait(b  
    item = queue();  
    cond_broadcast(buf_signal, buf_lock);  
    unlock buffer  
    return item  
}
```



**Release lock; signal others  
to run; reacquire on  
resume**

**n.b. OS must do the  
reacquire**

**Why User must recheck?**

# Why the while Loop?

---

- When a thread is woken up by **signal()** or **broadcast()**, it is simply put on the ready queue
- It may or may not reacquire the lock immediately!
  - Another thread could be scheduled first and "sneak in" to empty the queue
  - Need a loop to re-check condition on wakeup

# Recall: Semaphore Solution

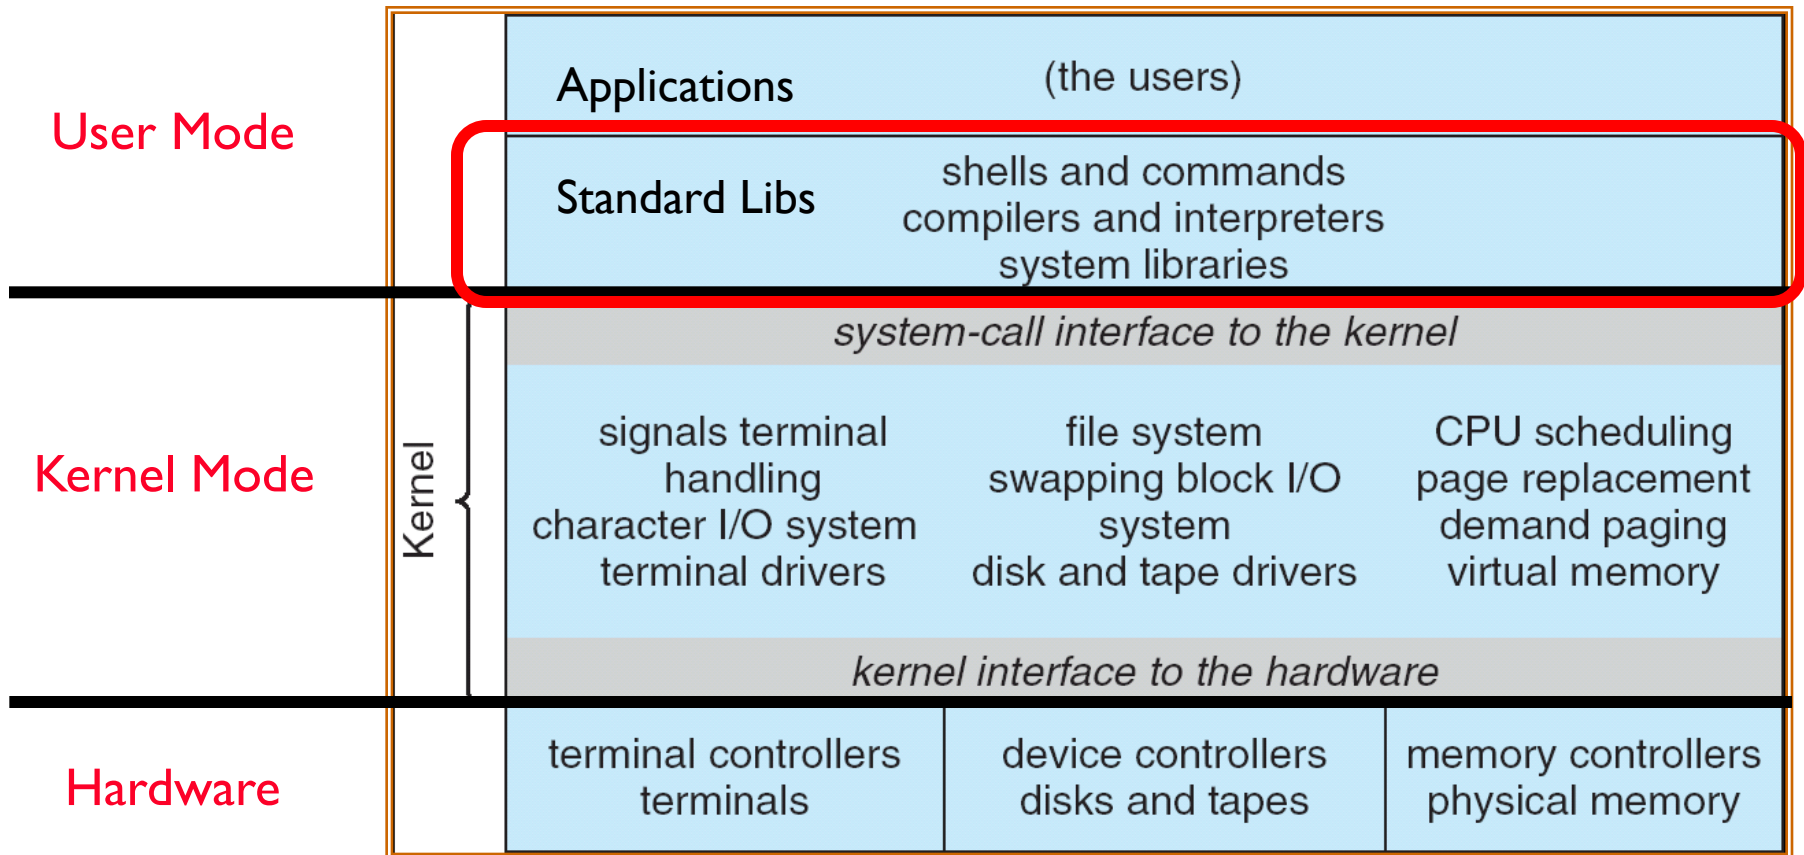
---

```
Semaphore fullSlots = 0; // Queue empty to start
Semaphore emptySlots = bufSize; // All slots empty
Semaphore mutex = 1; // No one in critical sect.
```

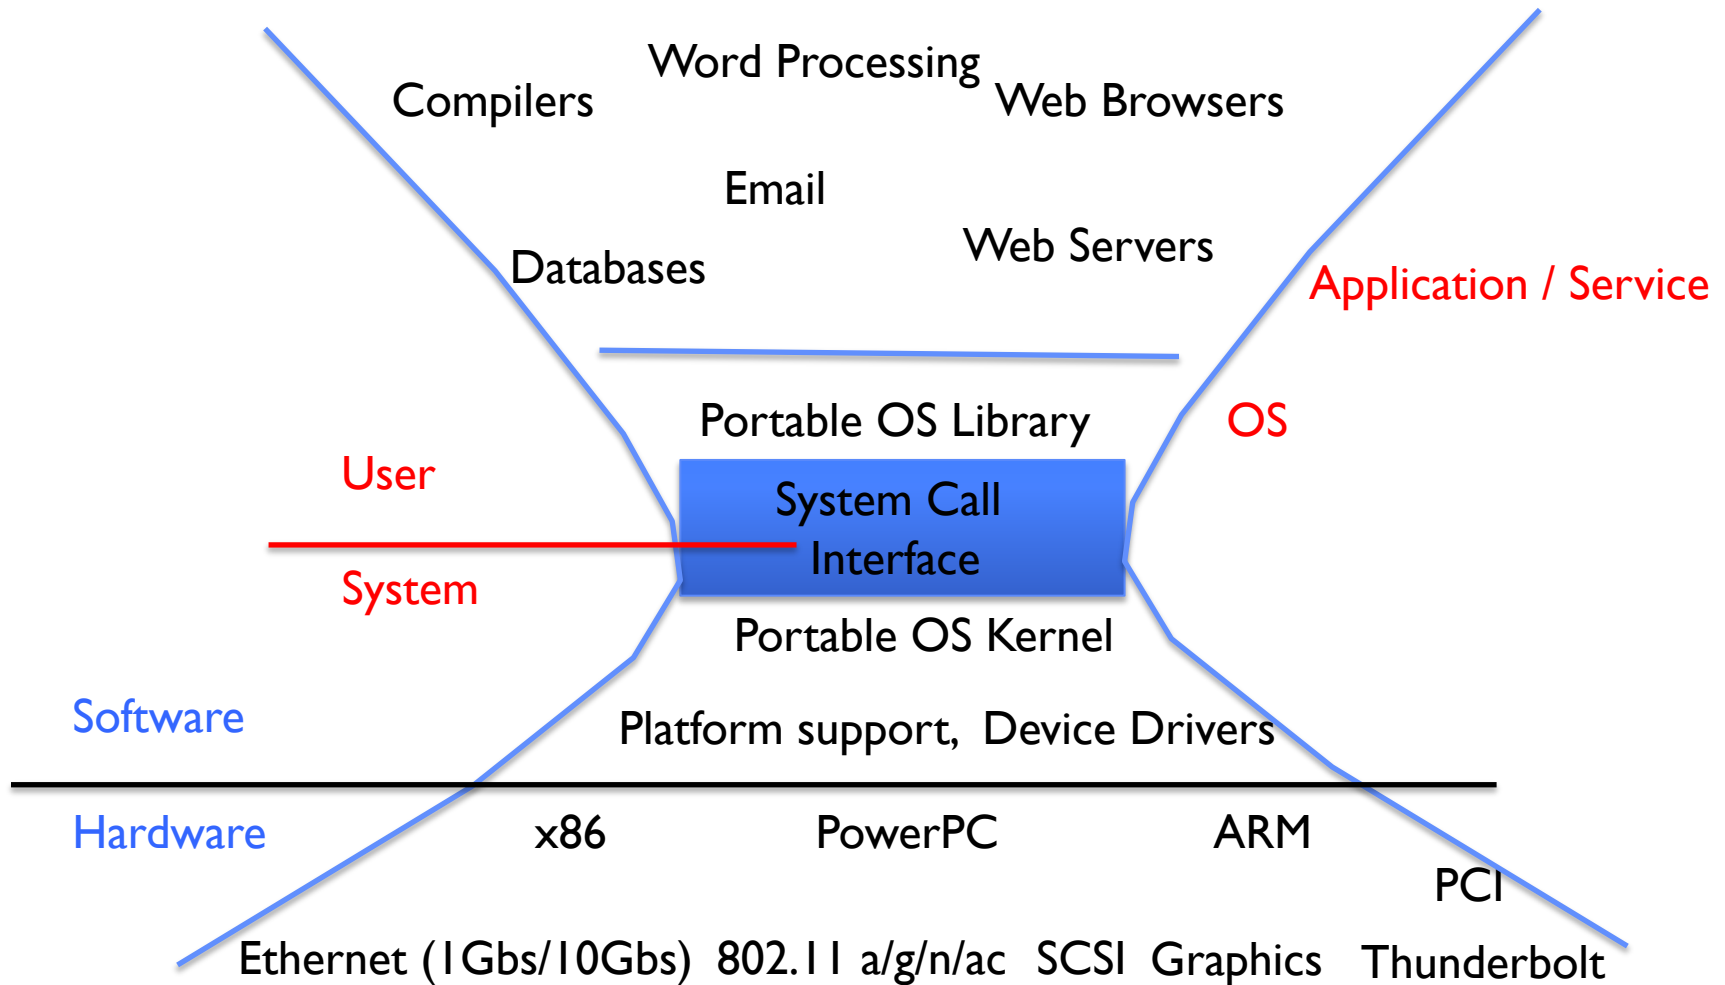
```
Producer(item) {
    emptySlots.P();
    mutex.P();
    Enqueue(item);
    mutex.V();
    fullSlots.V();
}
```

```
Consumer() {
    fullSlots.P();
    mutex.P();
    item = Dequeue();
    mutex.V();
    emptySlots.V();
    return item;
}
```

# Recall: UNIX System Structure




# A Kind of Narrow Waist





# Putting it together: web server




# Recall: Key Unix I/O Design Concepts

---

- Uniformity
  - file operations, device I/O, and interprocess communication through `open`, `read/write`, `close`
  - Allows simple composition of programs
    - » `find | grep | wc ...`
- Open before use
  - Provides opportunity for access control and arbitration
  - Sets up the underlying machinery, i.e., data structures
- Byte-oriented
  - Even if blocks are transferred, addressing is in bytes
- Kernel buffered reads
  - Streaming and block devices looks the same
  - read blocks process, yielding processor to other task
- Kernel buffered writes
  - Completion of out-going transfer decoupled from the application, allowing it to continue
- Explicit close

# How can a process communicate with another?

```
write(wfd, wbuf, wlen);
```



```
n = read(rfd, rbuf, rmax);
```


- Producer and Consumer of a file may be distinct processes
  - Also separated in time (one writes and then one later reads)
- However, what if data written once and consumed once?
  - Don't we want something more like a queue?
  - Can still look like File I/O!

# Communication between processes

---

- Can we view files as communication channels?

```
write(wfd, wbuf, wlen);
```




```
n = read(rfd, rbuf, rmax);
```

- We have seen one example – pipes
- Routinely used with the shell

```
>>> grep list src/*/*.c | more
```

# Communication Across the world looks like file IO

```
write(wfd, wbuf, wlen);
```



```
n = read(rfd, rbuf, rmax);
```

- Connected queues over the Internet
  - But what's the analog of open?
  - What is the namespace?
  - How are they connected in time?







# What Is A Protocol?

---

- A protocol is an **agreement on how to communicate**
- Includes
  - **Syntax**: how a communication is specified & structured
    - » Format, order messages are sent and received
  - **Semantics**: what a communication means
    - » Actions taken when transmitting, receiving, or when a timer expires
- Described formally by a state machine
  - Often represented as a message transaction diagram

# Examples of Protocols in Human Interactions

- Telephone

1. (Pick up / open up the phone)
2. Listen for a dial tone / see that you have service
3. Dial
4. Should hear ringing ...
5.  Callee: "Hello?"
6. Caller: "Hi, it's John...."  
Or: "Hi, it's me" (← what's *that* about?) 
7. Caller: "Hey, do you think ... blah blah blah ..." **pause**
1.  Callee: "Yeah, blah blah blah ..." **pause**
2. Caller: Bye 
3.  Callee: Bye
4. Hang up 

# Examples of Protocols in Human Interactions

## Asking a question

1. Raise your hand
2. Wait to be called on
3. Or: wait for speaker to **pause** and vocalize




# Clients and Servers

---

- Client program
  - Running on end host
  - Requests service
  - E.g., Web browser

**GET /index.html**




# Clients and Servers

---

- Client program
  - Running on end host
  - Requests service
  - E.g., Web browser
- Server program
  - Running on end host
  - Provides service
  - E.g., Web server

`GET /index.html`



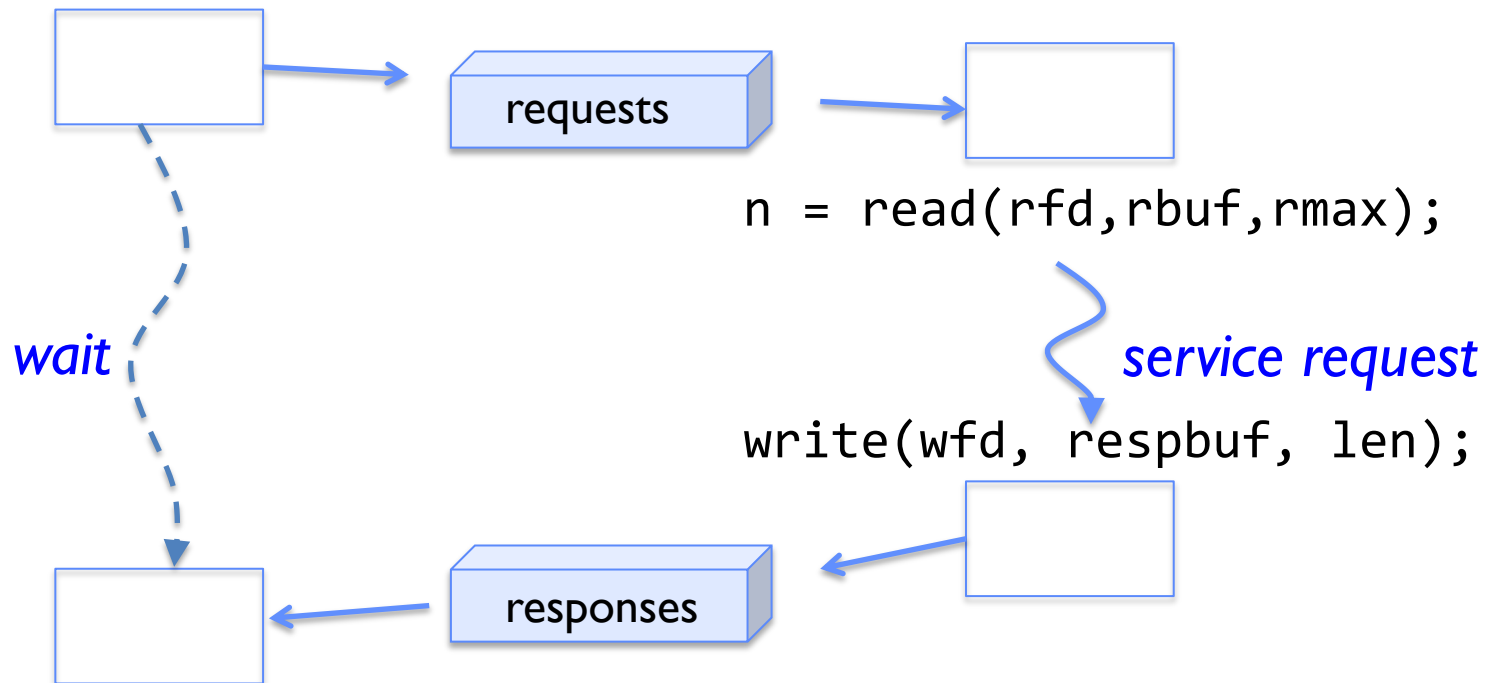
`"Site under construction"`

# Request Response Protocol

Client (issues requests)

Server (performs operations)

```
write(rqfd, rqbuf, buflen);
```



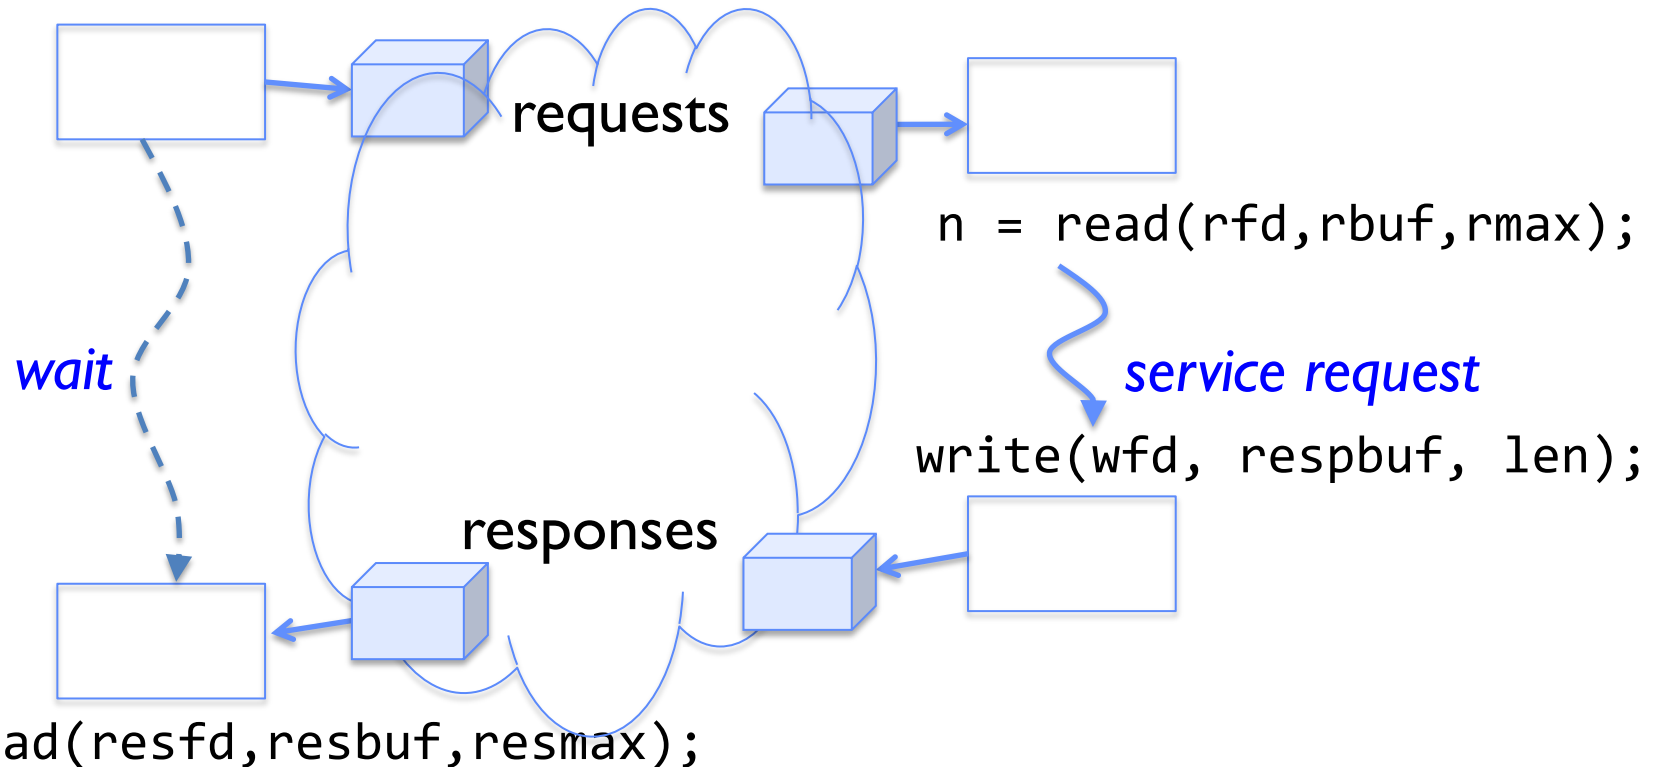
```
n = read(resfd, resbuf, resmax);
```

# Request Response Protocol

Client (issues requests)


Server (performs operations)

```
write(rqfd, rqbuf, buflen);
```



# Client-Server Models

---




- File servers, web, FTP, Databases, ...
- Many clients accessing a common server

# Client-Server Communication

---

- Client “sometimes on”
  - Initiates a request to the server when interested
  - E.g., Web browser on your laptop or cell phone
  - Doesn’t communicate directly with other clients
  - Needs to know the server’s address
- Server is “always on”
  - Services requests from many client hosts
  - E.g., Web server for the *www.cnn.com* Web site
  - Doesn’t initiate contact with the clients
  - Needs a fixed, well-known address



# Sockets

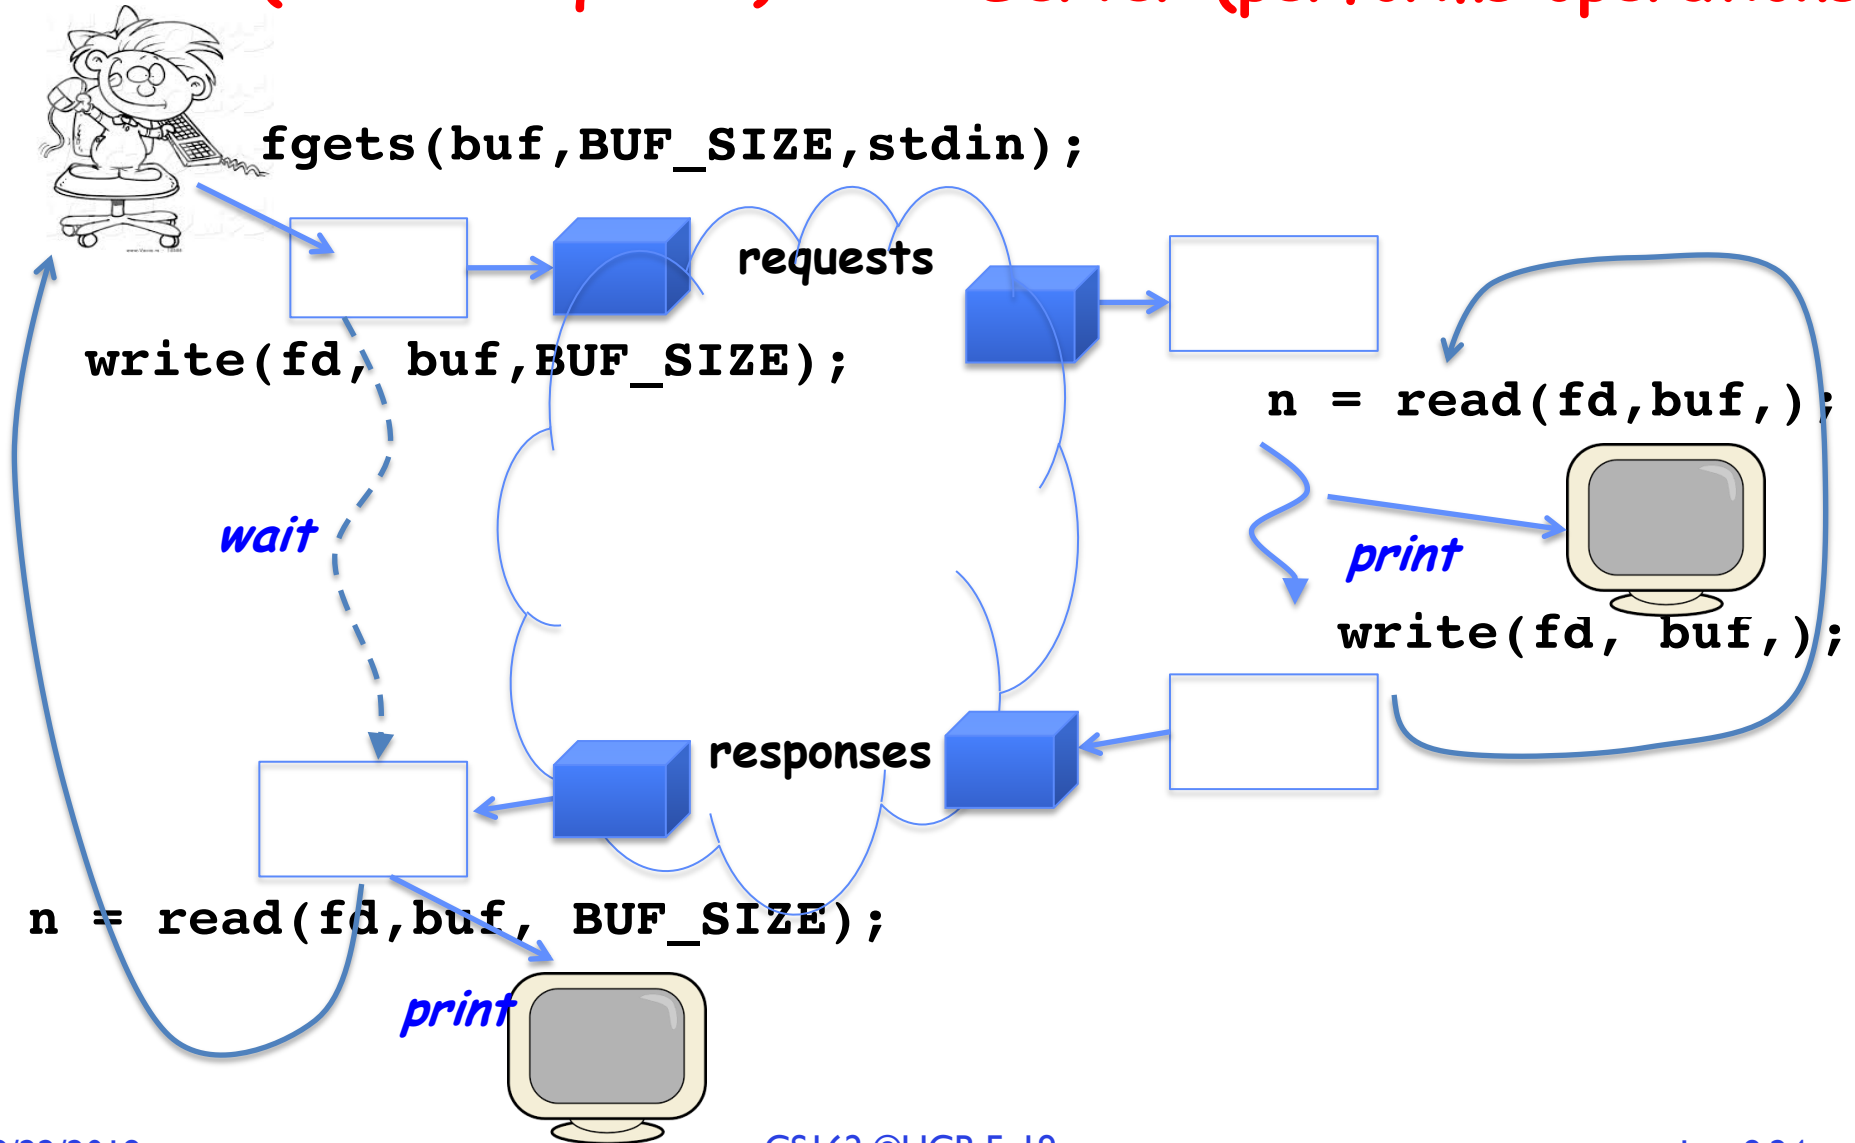
---

- **Socket:** an abstraction of a network I/O queue
  - Mechanism for inter-process communication
  - Embodies one side of a communication channel
    - » Same interface regardless of location of other end
    - » Could be local machine (called “UNIX socket”) or remote machine (called “network socket”)
  - First introduced in 4.2 BSD UNIX: big innovation at time
    - » Now most operating systems provide some notion of socket
- Data transfer like files
  - Read / Write against a descriptor
- Over ANY kind of network
  - Local to a machine
  - Over the internet (TCP/IP, UDP/IP)
  - OSI, Appletalk, SNA, IPX, SIP, NS, ...

# Silly Echo Server – running example

Client (issues requests)

Server (performs operations)

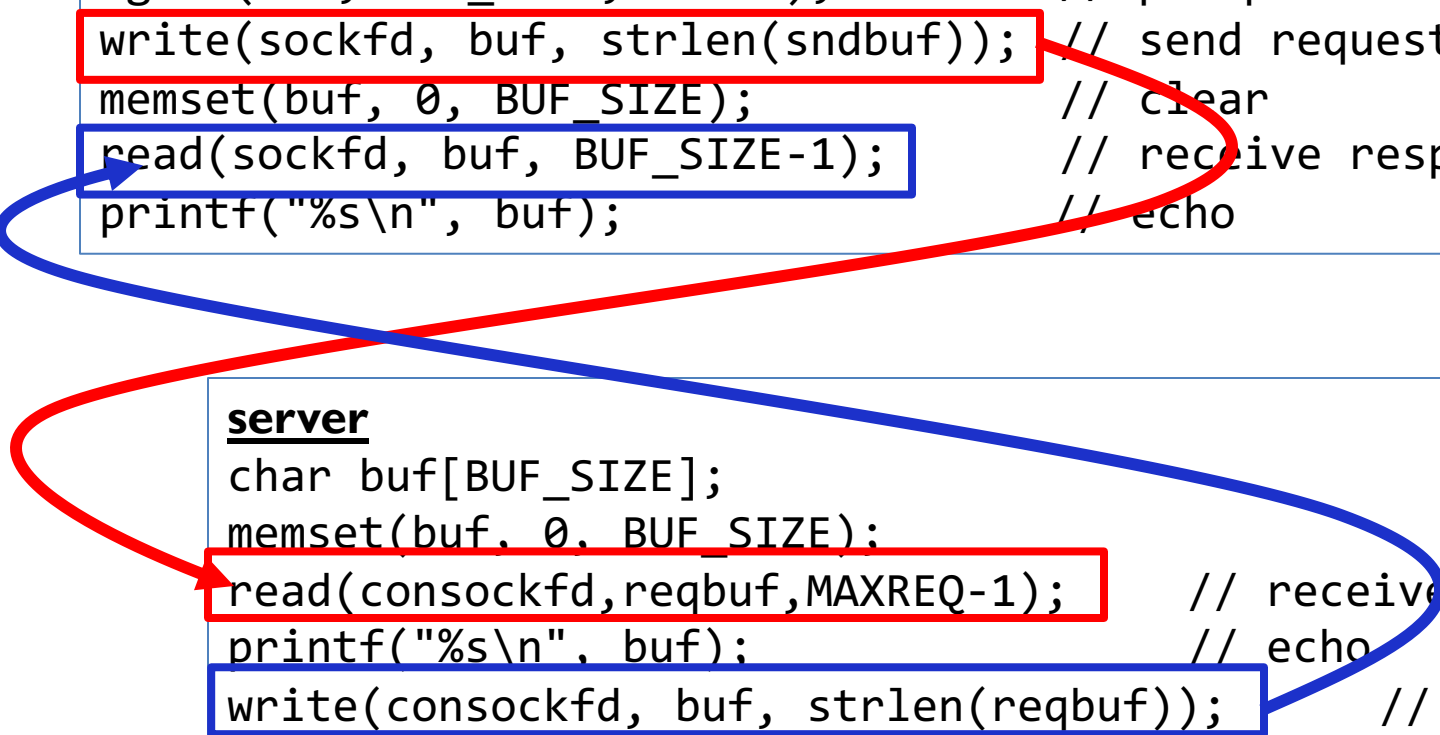




# Echo client-server example

## client

```
char buf[BUF_SIZE];  
fgets(buf, BUF_SIZE, stdin);           // prompt  
write(sockfd, buf, strlen(sndbuf));     // send request  
memset(buf, 0, BUF_SIZE);              // clear  
read(sockfd, buf, BUF_SIZE-1);          // receive response  
printf("%s\n", buf);                   // echo
```



## server

```
char buf[BUF_SIZE];  
memset(buf, 0, BUF_SIZE);  
read(consockfd, reqbuf, MAXREQ-1);     // receive  
printf("%s\n", buf);                   // echo  
write(consockfd, buf, strlen(reqbuf)); // send response
```

# What assumptions are we making?

---

- Reliable
  - Write to a file => Read it back. Nothing is lost.
  - Write to a (TCP) socket => Read from the other side, same.
  - Like pipes
- In order (sequential stream)
  - Write X then write Y => read gets X then read gets Y
- When ready?
  - File read gets whatever is there at the time. Assumes writing already took place.
  - Like pipes!

# Socket creation and connection

---

- File systems provide a collection of permanent objects in structured name space
  - Processes open, read/write/close them
  - Files exist independent of the processes
- Sockets provide a means for processes to communicate (transfer data) to other processes.
- Creation and connection is more complex
- Form 2-way pipes between processes
  - Possibly worlds away
- How do we name them?
- How do these completely independent programs know that the other wants to “talk” to them?

# Namespaces for communication over IP

---


- Hostname
  - www.eecs.berkeley.edu
- IP address
  - 128.32.244.172 (IPv4, 32-bit Integer)
  - 2607:f140:0:81::f (IPv6, 128-bit Integer)
- Port Number
  - 0-1023 are “well known” or “system” ports
    - » Superuser privileges to bind to one
  - 1024 – 49151 are “registered” ports (registry)
    - » Assigned by IANA for specific services
  - 49152–65535 ( $2^{15}+2^{14}$  to  $2^{16}-1$ ) are “dynamic” or “private”
    - » Automatically allocated as “ephemeral Ports”

# How do they “shake hands”?

---


- How does the server know that a client wants to make a request of them?
- How does a client know that the server is accepting requests?

# Socket Setup over TCP/IP



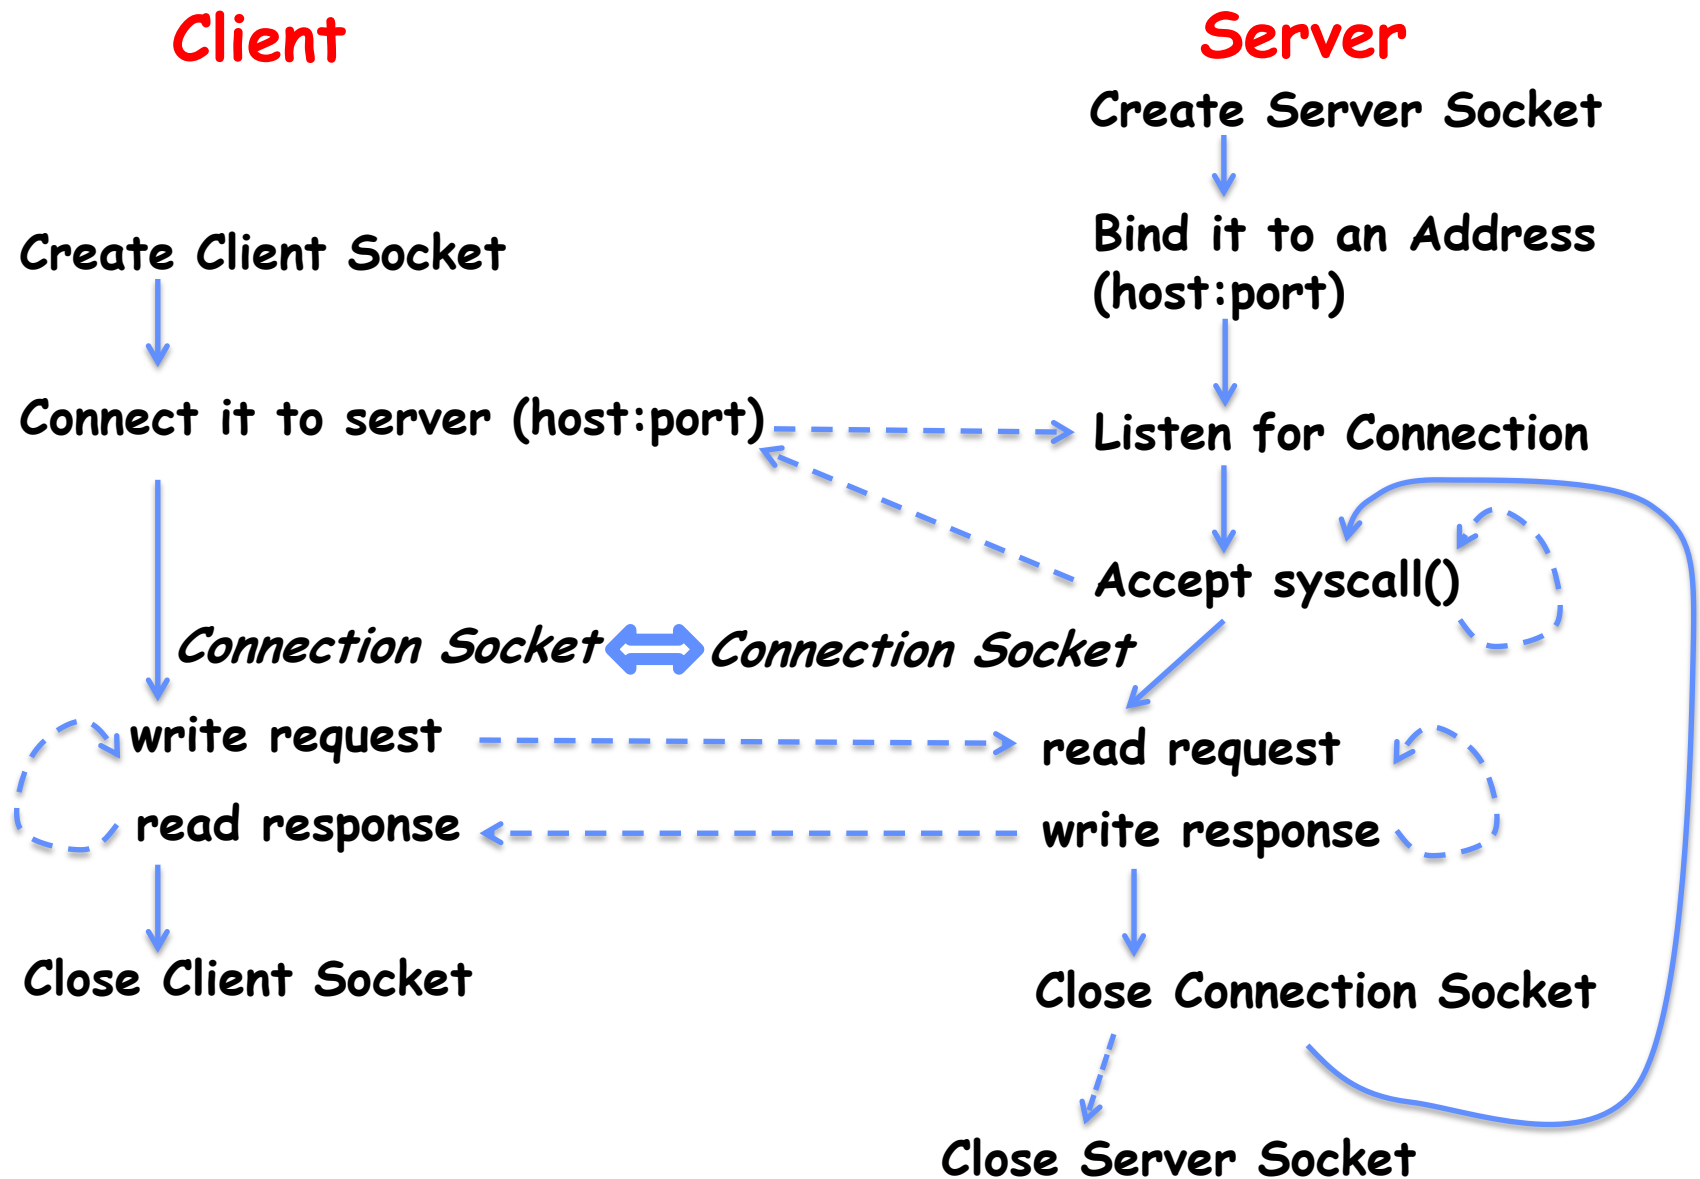
- Special kind of socket: **server socket**
  - Has file descriptor
  - Can't read or write
- Two operations:
  1. **listen()**: Start allowing clients to connect
  2. **accept()**: Create a *new socket* for a *particular* client connection

# Socket Setup over TCP/IP



- 5-Tuple identifies each connection:
  1. Source IP Address
  2. Destination IP Address
  3. Source Port Number
  4. Destination Port Number
  5. Protocol (always TCP here)
- Often, Client Port “randomly” assigned
  - Done by OS during client socket setup
- Server Port often “well known”
  - 80 (web), 443 (secure web), 25 (sendmail), etc
  - Well-known ports from 0—1023

# Sockets in Schematic





# Client Protocol

---

```
char *host_name, port_name;
```

```
// Create a socket
```

```
struct addrinfo *server = lookup_host(host_name, port_name);  
int sock_fd = socket(server->ai_family, server->ai_socktype,  
                    server->ai_protocol);
```

```
// Connect to specified host and port
```

```
connect(sock_fd, server->ai_addr, server->ai_addrlen);
```

```
// Carry out Client-Server protocol
```

```
run_client(sock_fd);
```

```
/* Clean up on termination */
```

```
close(sock_fd);
```

# Server Protocol (v1)

---

```
// Create socket to listen for client connections
char *port_name;
struct addrinfo *server = setup_address(port_name);
int server_socket = socket(server->ai_family,
                           server->ai_socktype, server->ai_protocol);

// Bind socket to specific port
bind(server_socket, server->ai_addr, server->ai_addrlen);
// Start listening for new client connections
listen(server_socket, MAX_QUEUE);

while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    serve_client(conn_socket);
    close(conn_socket);
}

close(server_socket);
```

# How does the server protect itself?

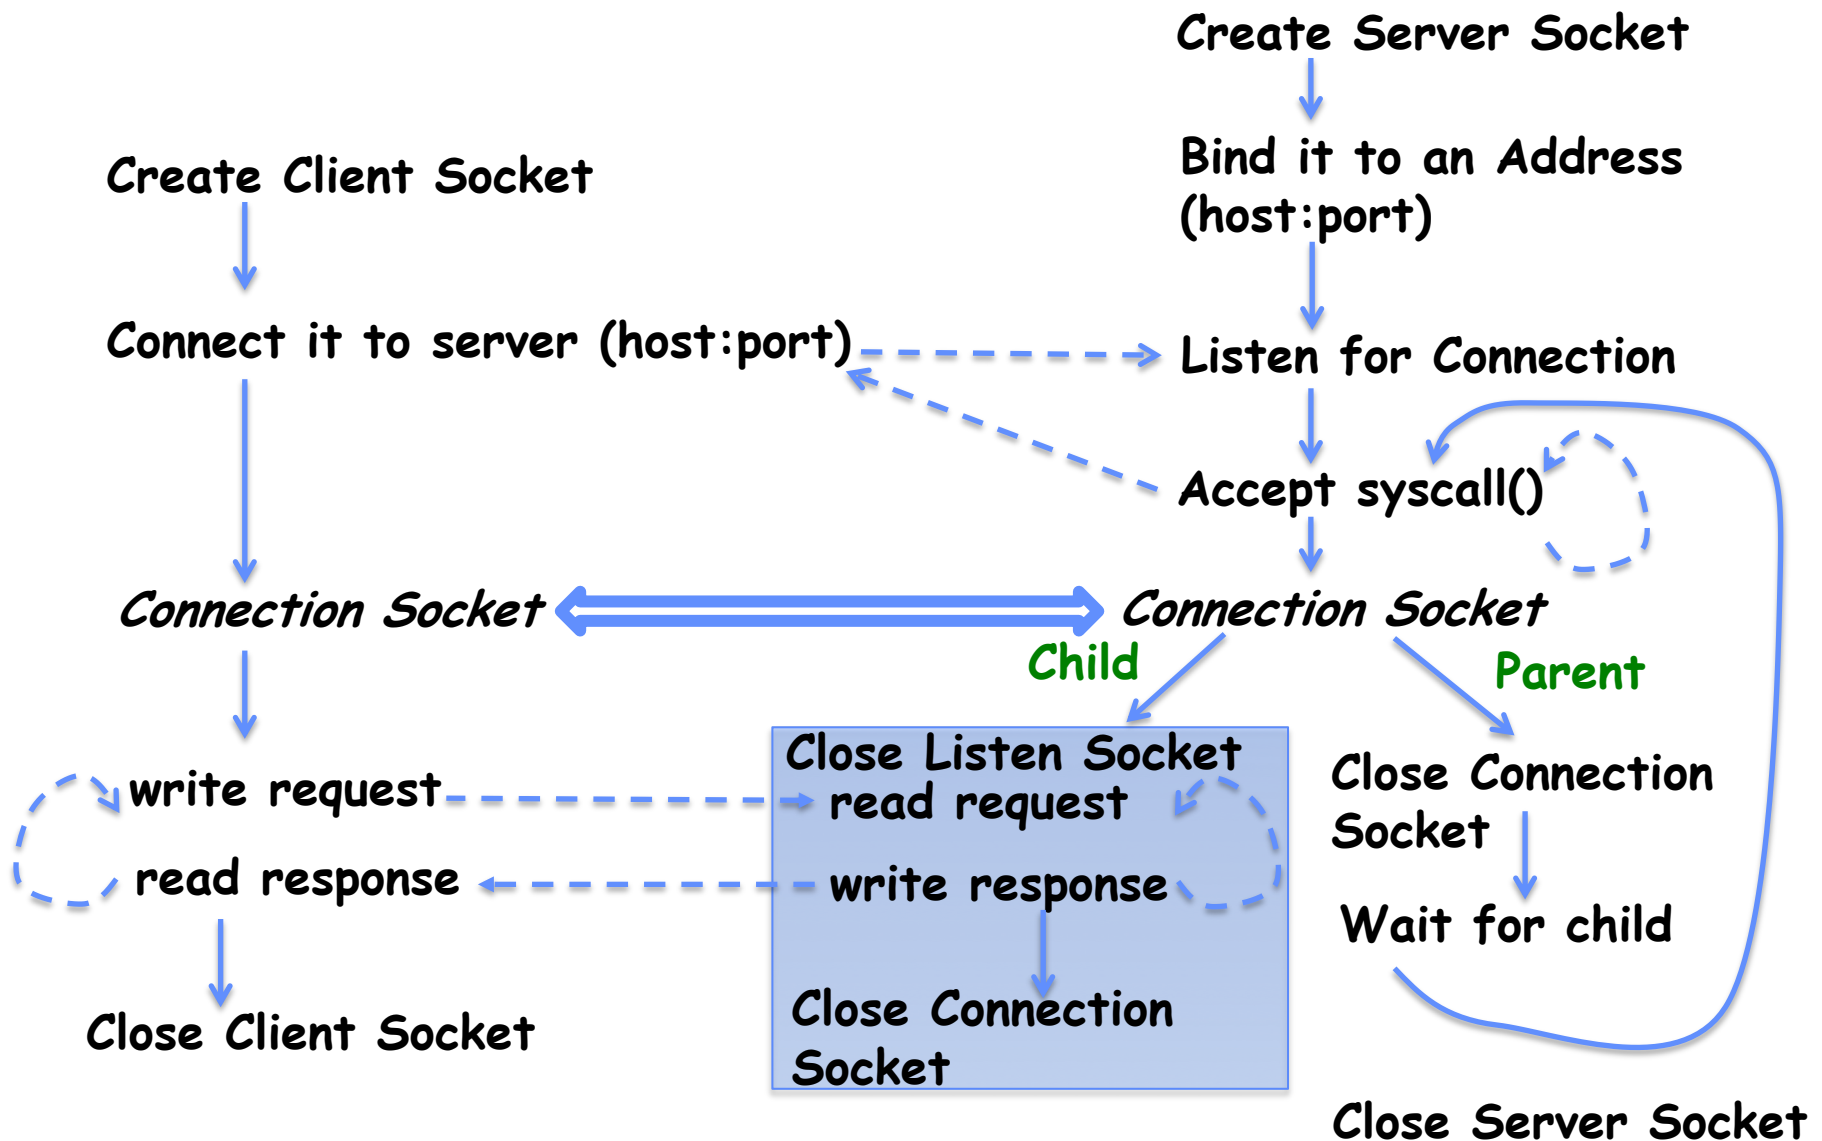
---

- Isolate the handling of each connection
- By forking it off as another process

# Sockets With Protection

## Client

## Server



# Server Protocol (v2)

---

```
// Socket setup code elided...
```

```
while (1) {  
    // Accept a new client connection, obtaining a new socket  
    int conn_socket = accept(server_socket, NULL, NULL);  
    pid_t pid = fork();  
    if (pid == 0) {  
        close(server_socket);  
        serve_client(conn_socket);  
        close(conn_socket);  
        exit(0);  
    } else {  
        close(conn_socket);  
        wait(NULL);  
    }  
}  
  
close(server_socket);
```

# Concurrent Server

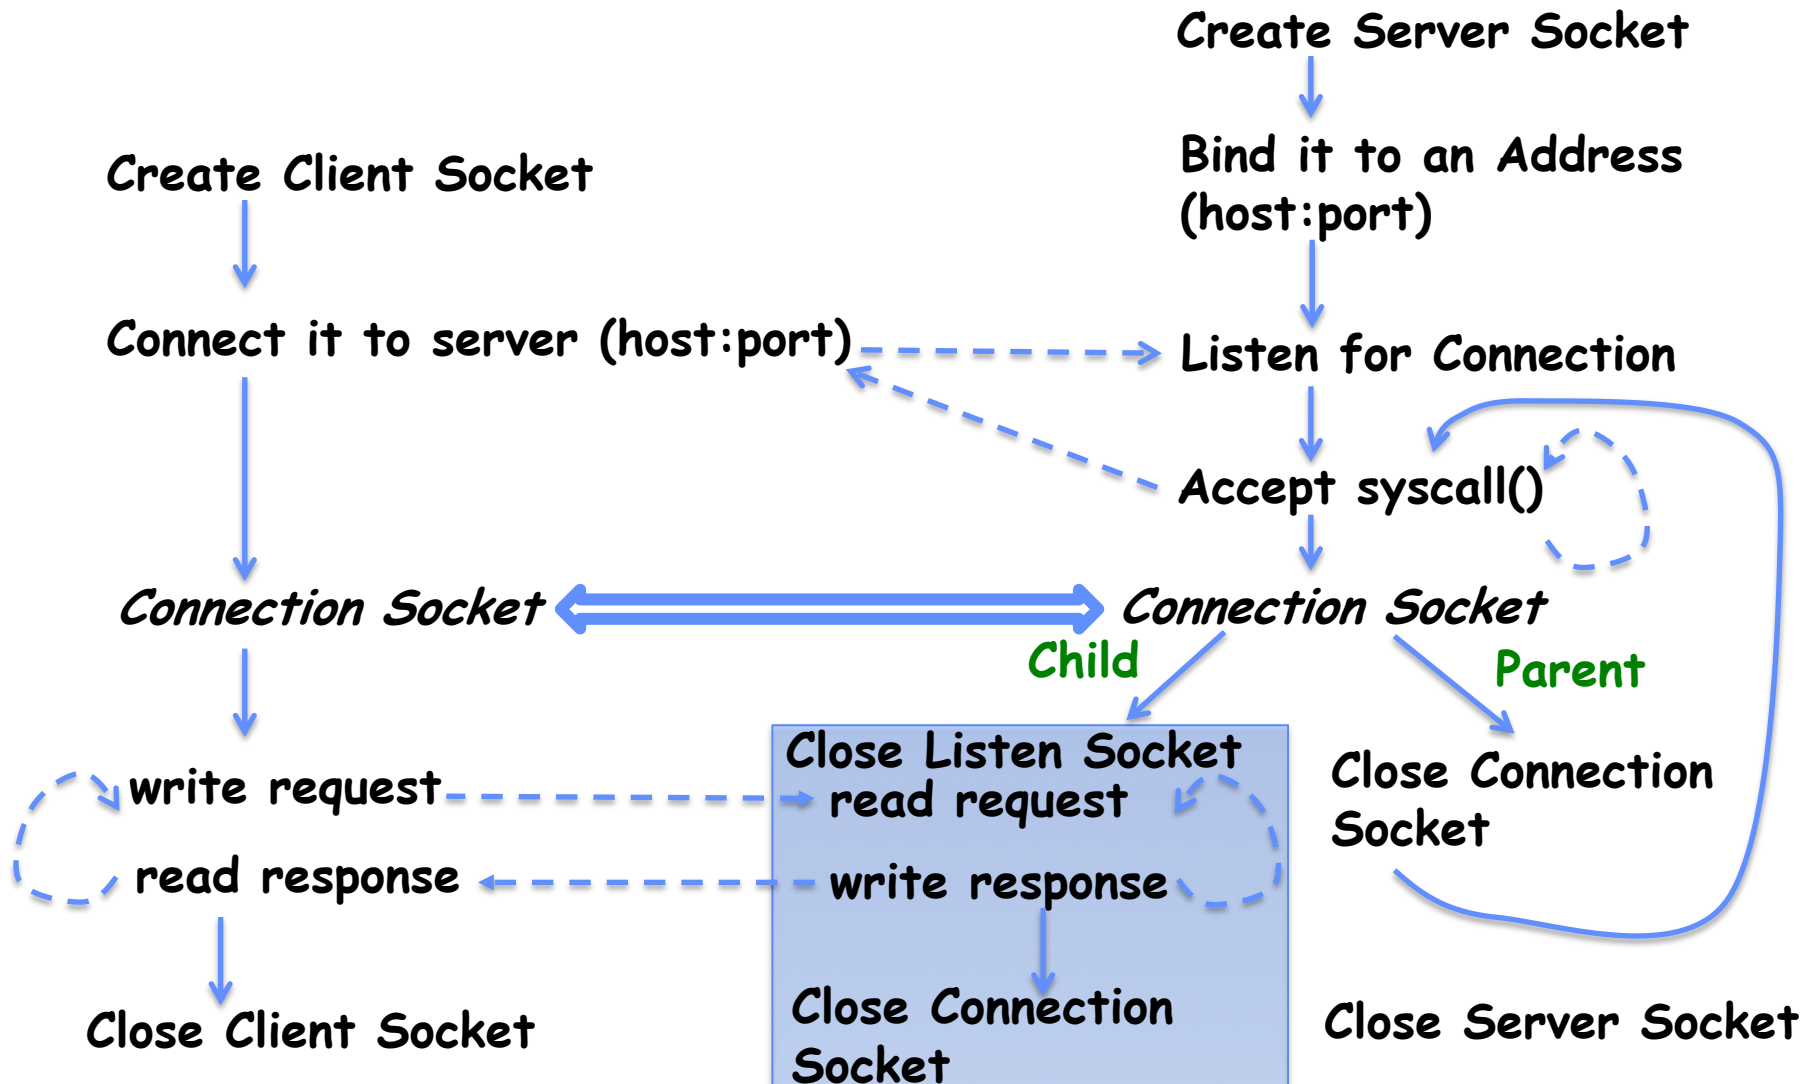
---

- Listen will queue requests
- Buffering present elsewhere
- But server waits for each connection to terminate before initiating the next

# Sockets With Protection and Parallelism

## Client

## Server



# Server Protocol (v3)

---

```
// Socket setup code elided...
```

```
while (1) {  
    // Accept a new client connection, obtaining a new socket  
    int conn_socket = accept(server_socket, NULL, NULL);  
    pid_t pid = fork();  
    if (pid == 0) {  
        close(server_socket);  
        serve_client(conn_socket);  
        close(conn_socket);  
        exit(0);  
    } else {  
        close(conn_socket);  
        //wait(NULL);  
    }  
}  
  
close(server_socket);
```

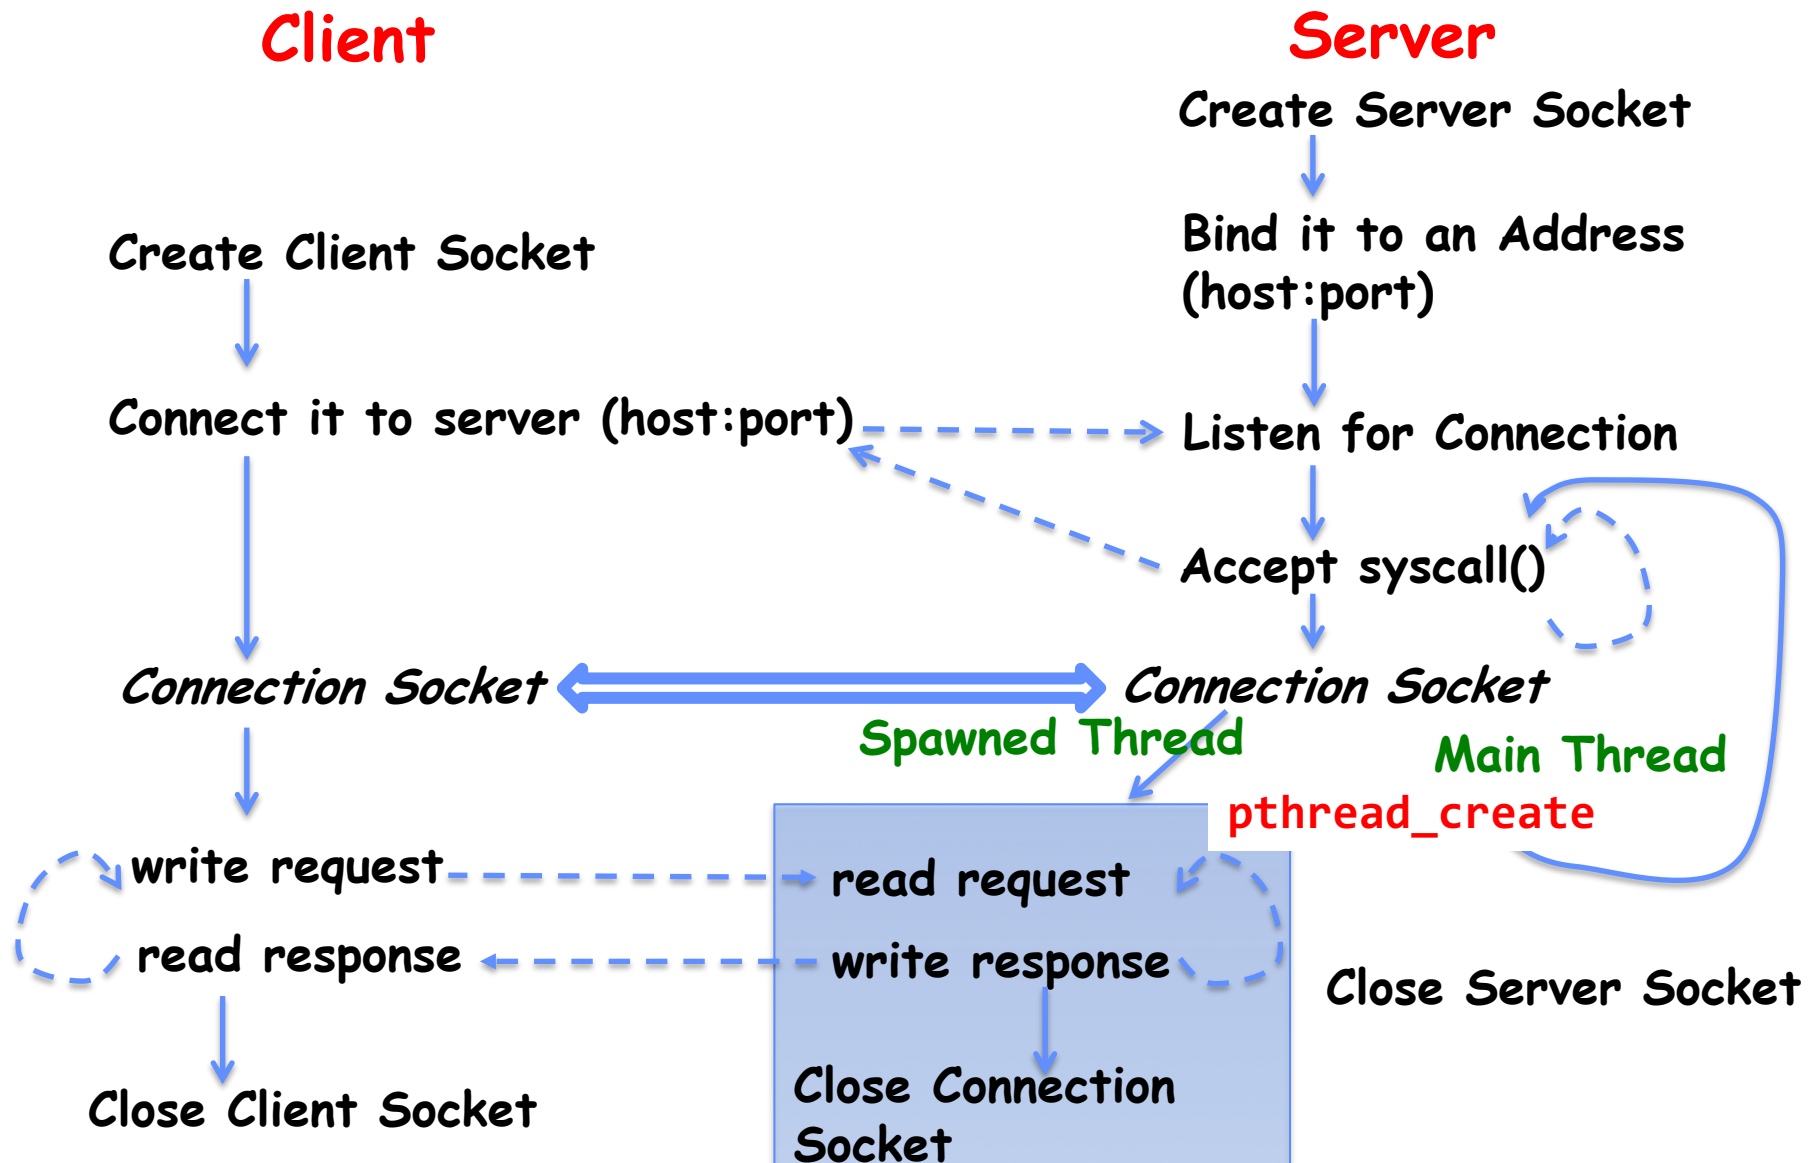


# Concurrent Server without Protection

---

- Spawn a new thread to handle each connection
- Main thread initiates new client connections without waiting for previously spawned threads
- Why give up the protection of separate processes?
  - More efficient to create new threads
  - More efficient to switch between threads

# Sockets With Parallelism, Without Protection



# Server Address - itself

---

```
struct addrinfo *setup_address(char *port) {  
    struct addrinfo *server;  
    struct addrinfo hints;  
    memset(&hints, 0, sizeof(hints));  
    hints.ai_family = AF_UNSPEC;  
    hints.ai_socktype = SOCK_STREAM;  
    hints.ai_flags = AI_PASSIVE;  
    getaddrinfo(NULL, port, &hints, &server);  
    return server;  
}
```

- Simple form
- Internet Protocol, TCP
- Accepting any connections on the specified port

# Client: getting the server address

---

```
struct addrinfo *lookup_host(char *host_name, char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;
    hints.ai_socktype = SOCK_STREAM;

    int rv = getaddrinfo(host_name, port_name,
                        &hints, &server);

    if (rv != 0) {
        printf("getaddrinfo failed: %s\n", gai_strerror(rv));
        return NULL;
    }
    return server;
}
```

# Conclusion (I)

---

- System Call Interface is “narrow waist” between user programs and kernel
- Streaming I/O: modeled as a stream of bytes
  - Most streaming I/O functions start with “f” (like “**fread**”)
  - Data buffered automatically by C-library functions
- Low-level I/O:
  - File descriptors are integers
  - Low-level I/O supported directly at system call level
- **STDIN / STDOUT** enable composition in Unix
  - Use of pipe symbols connects **STDOUT** and **STDIN**
    - » `find | grep | wc ...`

# Conclusion (II)

---

- Device Driver: Device-specific code in the kernel that interacts directly with the device hardware
  - Supports a standard, internal interface
  - Same kernel I/O system can interact easily with different device drivers
- File abstraction works for inter-processes communication (local or Internet)
- Socket: an abstraction of a network I/O queue
  - Mechanism for inter-process communication