# CS 162: Operating Systems and Systems Programming

Lecture 10: Understanding System Performance & Fair Scheduling

October 1, 2019

Instructor: David E. Culler

https://cs162.eecs.berkeley.edu

Read: 3easy ch 9

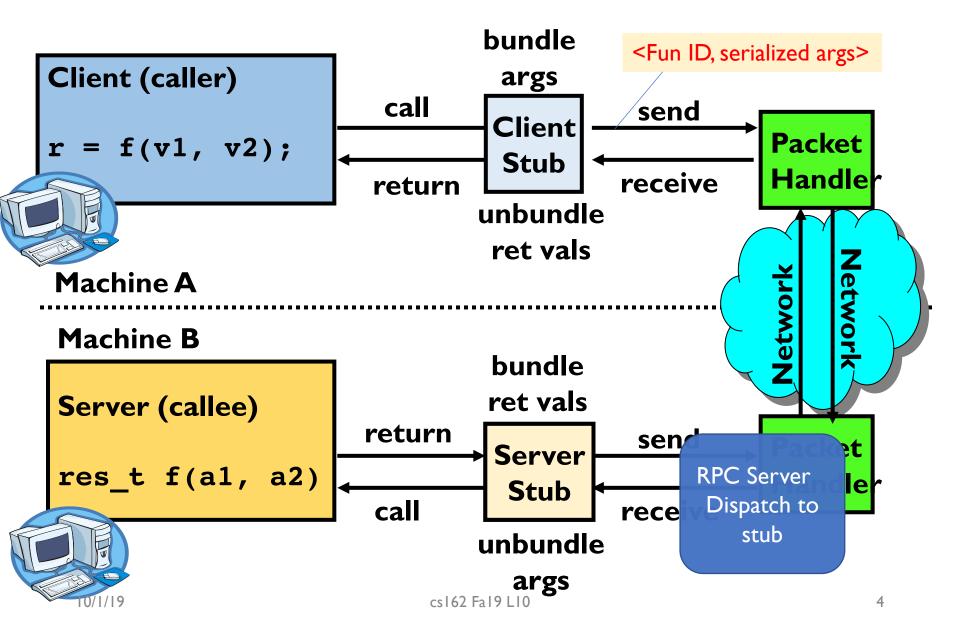
#### Recap

- Process: instance of an executing program
  - one or more threads, address space, systems resources (file descr., locks, ...)
  - Isolated from other processes (its own address space)
  - Threads communicate primarily through shared variables, including locks, but could send messages
- Inter-Process Communication (IPC)
  - External means of communication information between processes, i.e., across distinct address spaces
    - File write/read, Socket write/read, message channels
    - Enables "inter-machine" communication among processes
- Serialization
  - Conversion of an "in-memory" data structure into an external canonical form that can be recovered by another process
    - Independent of binary representation in the machine
- Remote Procedure Call (RPC)
  - Methodology used to make interactions across processes familiar and convenient
  - Involves serialization and IPC
  - Also dispatch to the proper remote method by the RPC server
  - Stubs provide the glue on both side

#### Recall: Request/Response Protocol

```
Client (issues requests)
                                      Server (performs operations)
write(rqfd, rqbuf, buflen);
                           re que sts
                                    n = read(rfd,rbuf,rmax);
                 Serialized Objects
                                                   service request
       wait
                                    write(wfd, respbuf, len);
                           responses
n = read(resfd, resbuf, resmax);
```

#### **RPC Information Flow**



#### Recall: RPC Six steps

- I. The client calls the client stub. The call is a local procedure call, with parameters pushed on to the stack in the normal way.
- 2. The client stub packs the parameters into a message and makes a system call to send the message. Packing the parameters is called marshaling.
- 3. The client's local operating system sends the message from the client machine to the server machine.
- 4. The local operating system on the server machine passes the incoming packets to the RPC server which dispatches the call to the server stub.
- 5. The server stub unpacks the parameters from the message. Unpacking the parameters is called unmarshaling.
- 6. Finally, the server stub <u>calls</u> the server procedure. The reply traces the same steps in the reverse direction

### Systems Design:

 Translate important high-level goals down to simple, robust mechanisms – and implement them

• Example: maintain "performance" in the presence of diverse, perhaps increasing demand

- Schedulers are a critical algorithmic component
  - Boil down to specific operations on certain (fast) data structures

### Recall: Evaluating Schedulers

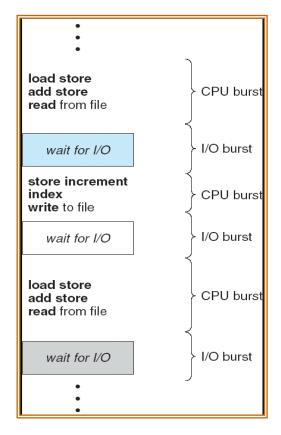
- Response Time (ideally low)
  - What user sees: from keypress to character on screen

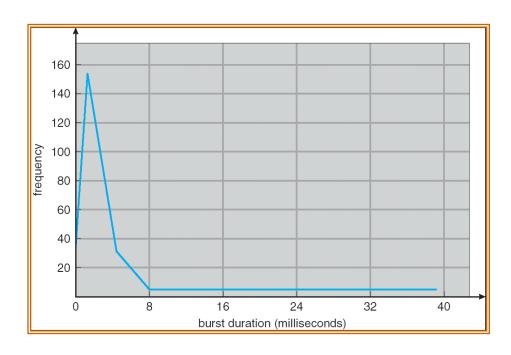
- Throughput (ideally high)
  - Total operations per second
  - Problem: Overhead (e.g. from context switching)

#### Fairness

• Conflicts with best avg. throughput/resp. time

#### CPU & I/O Bursts





#### Support interactive programs: prefer I/O-bound tasks

### Basics of System Performance

- "Back of the Envelope" calculation and modeling
- "Mean value analysis"
- Get the rough picture first ... and don't lose sight of it

## Times (s) and Rates (ops/s)

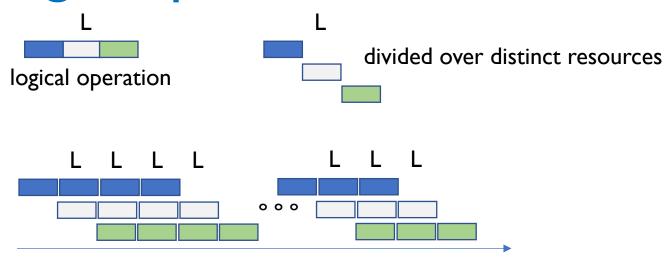
- Latency time to complete a task
  - Measured in units of time (s, ms, us, ..., hours, years)
- **Response Time** time to initiate and operation and get its response
  - Able to issue one that depends on the result
  - Know that it is done (anti-dependence, resource usage)
- Throughput or Bandwidth rate at which tasks are performed
  - Measured in units of things per unit time (ops/s, GLOP/s)
- Performance ???
  - Operation time (4 mins to run a mile...)
  - Rate (mph, mpg, ...)

#### Sequential Server Performance



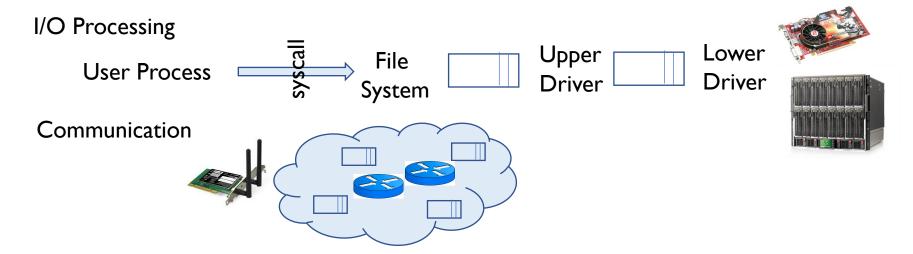
- Single sequential "server" that can deliver a task in time L operates at rate  $\leq \frac{1}{L}$  (on average, in steady state, ...).
  - L = 10 ms => B = 100 ops/sec
  - L = 2 years  $\Rightarrow B = 0.5$  ops/year
- a processor, a disk drive, a Person, a TA, ...

#### Single Pipelined Server



- Single pipelined server of k stages for tasks of length L (i.e., time L/k per stage) delivers at rate  $\leq \frac{k}{L}$ . (asymptote)
  - L = 10 ms, k = 4 => B = 400 ops/sec
  - L = 2 years,  $k = 2 \Rightarrow B = I$  ops/year
- In 61C you saw "synchronous" pipelines in processors
  - Systems present lots of asynchronous pipeline-like settings

## Example Systems "Pipelines"



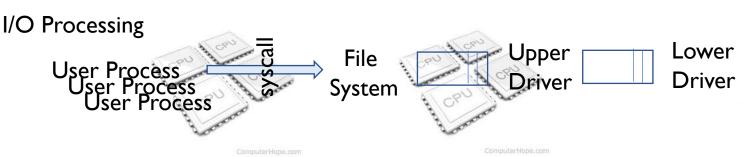
- Anything with queues between operational process behaves roughly "pipeline like"
- Important difference is that "initiations" are decoupled from processing
  - May have to queue up a burst of operations
  - Not synchronous and deterministic like in 61C

#### Multiple Servers



- k servers handling tasks of length L (i.e., time L/k per stage) delivers at rate  $\leq \frac{k}{L}$ .
  - L = 10 ms, k = 4 => B = 400 ops/sec
  - L = 2 years, k = 2 => B = I ops/year
- In 61C you saw multiple processors (cores)
  - Systems present lots of multiple parallel servers
  - Often with lots of queues

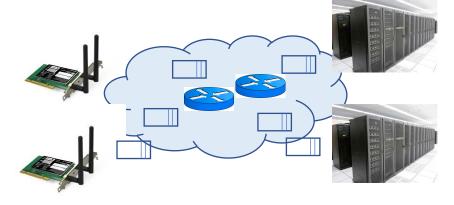
## Example Systems "Parallelism"





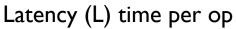


Communication



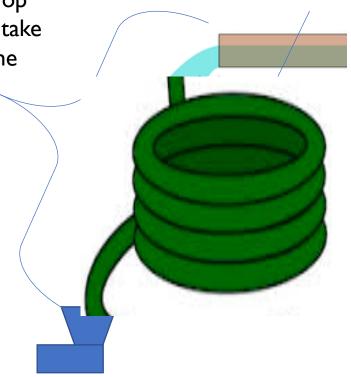
Parallel Computation, Databases, ...

#### A Simple Systems Performance Model



 How long does it take to flow through the system

"Service Time"



Bandwidth (B): Rate, Op/s

e.g., flow: gal per min

If 
$$B = 2\frac{gal}{s}$$
 and  $L = 3s$   
How much water is "in the system?"

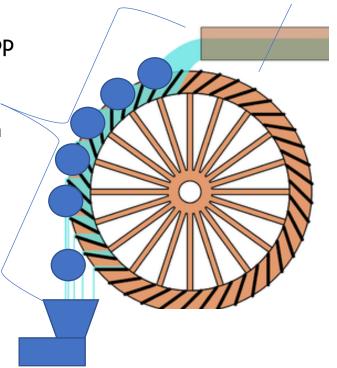
#### A Simple Systems Performance Model

Bandwidth (B): Rate, Op/s

Latency (L) time per op

How long does an operation to flow through the system

Service Time



If  $B = 2\frac{ops}{s}$  and L = 3sHow many ops are "in the system?"

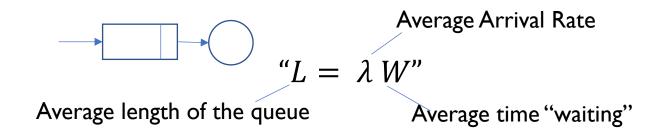
#### Little's Law

• The number of "things" in the systems is equal to the bandwidth times the latency (on average)

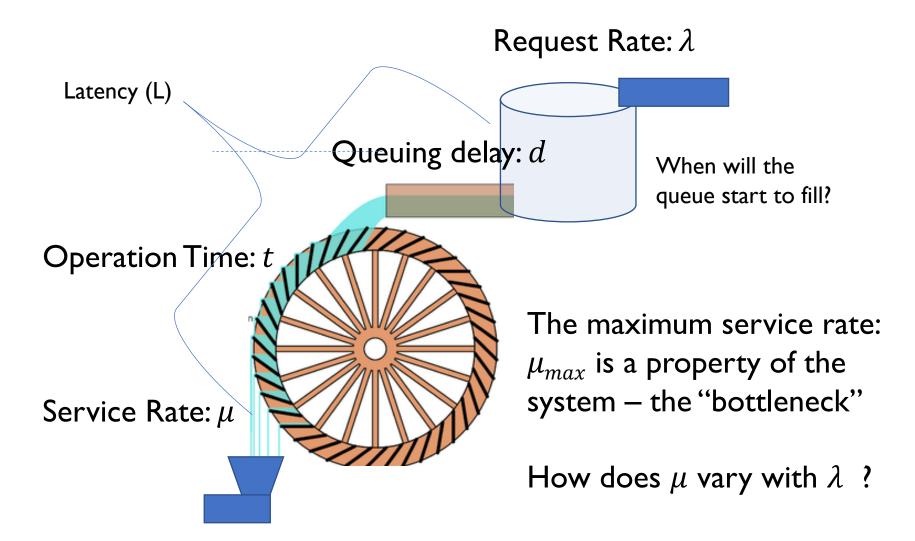
$$n = L B$$

- In networks, the bandwidth-delay product
- Include the queues, the processing stages, parallelism, whatever


"in other words"



#### A Simple Systems Performance Model



### Idealized System Performance

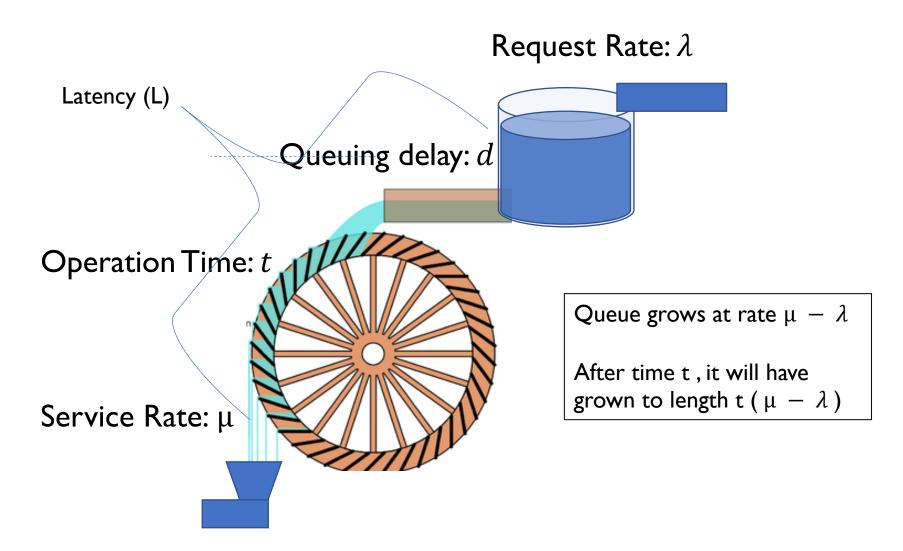


Request Rate ( $\lambda$ ) - "offered load"

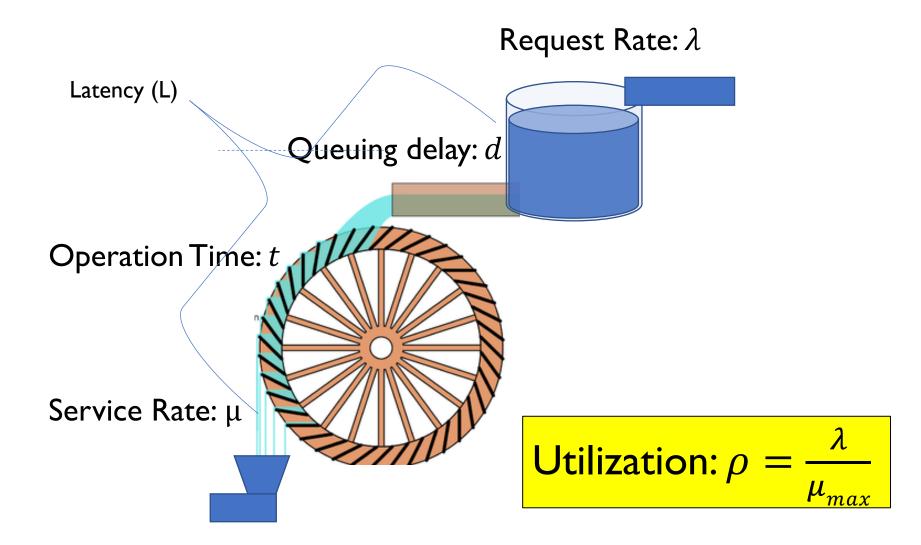
## Queuing

- What happens when request rate exceeds max service rate?
- Short bursts can be absorbed by the queue
  - If on average  $\lambda < \mu$ , it will drain eventually
- Prolonged  $\lambda > \mu$  queue will grow arbitrarily

#### A Simple Systems Performance Model



#### A Simple Systems Performance Model

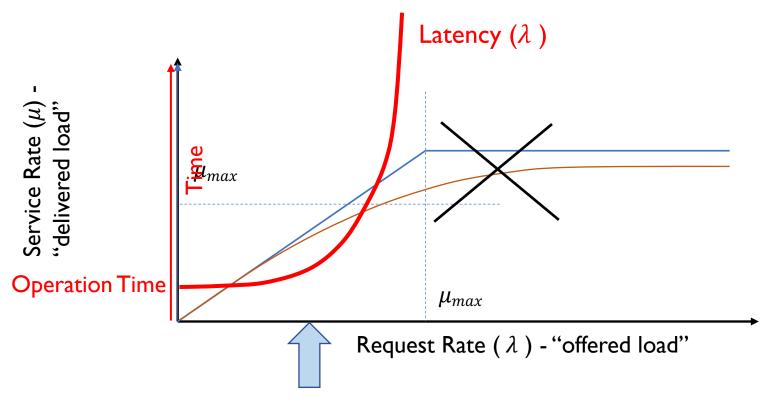


### Length of the queue ...

 Queuing Theory describes steady-state (stationary, not transient) behavior in the presence of stochastic variations – primarily the request rate distribution, i.e., bursts and lulls.

• Ave Number of Operations in the system =  $\frac{\rho}{1-\rho}$ 

#### Idealized System Performance

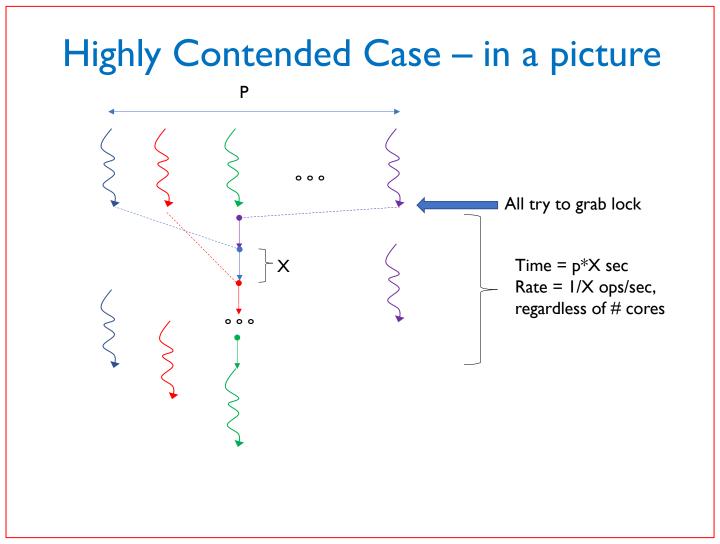


"Half-Power Point": load at which system delivers half of peak performance

- Design and provision systems to operate roughly in this regime
- Latency low and predictable, utilization good: ~50%25

# What's the bottleneck that determines $\mu_{max}$ ?

## Time to service a highly contended lock?



#### Potential Bottlenecks ...

- Actual operation time servicing request apply parallelism
- System calls if the operating system is only able to process one or few calls at a time
- System call overhead to grab even uncontended locks
- Socket accept/bind/fork (thread create)
- Time for scheduler to select the next thread to run

- Time to write to a log file
- ... <your idea / concern goes here >

## Do real systems really hit a wall as utilization approaches 100%?

And how do we measure and test it?

## server $\mu_{max}$ **Closed System** server $\mu_{max}$

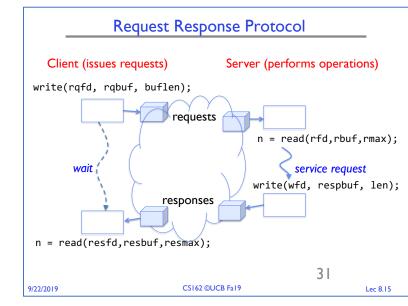
Open System

### **Closed System**

- Clients generating the load depend on completion of previous requests
  - Request-response protocols
  - Humans in-the-loop waiting for results
- Model of client: {request, wait}+ repeat
  - Request rate determined by length of wait
- In closed system, wait time depends on response time (latency = operation time + queuing delay)
- As system saturates (utilization  $\rightarrow 100\%$ ) delay increases, request rate is limited by service rate
  - Queueing smooths bursts, but does not grow unbounded due to rate mismatch

#### What causes systems to "close"?

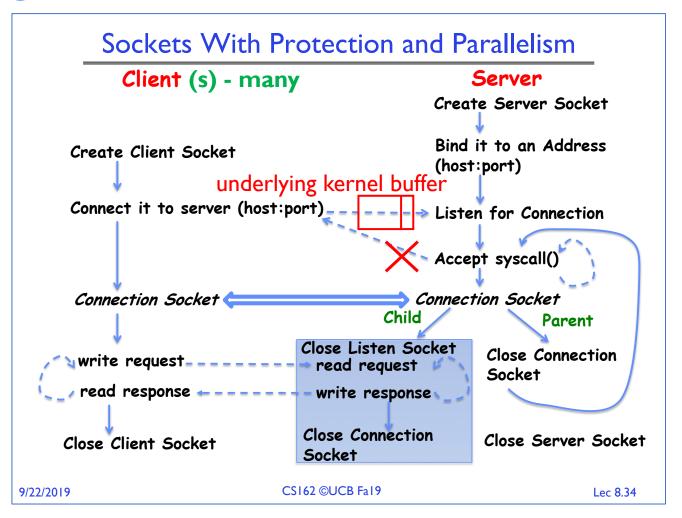
- Protocols are designed to have self-limited behavior
  - Request-response, bounded number of outstanding requests per client
- Underlying system induces "back pressure" even if higher level services and applications don't
  - Bounded size queues (not just because of memory size)
  - What happens when it fills up?



## What do you do when a queue fills up?

- Within a machine?
  - Block the thread (or process) slowing it down
  - Lot's of cheap ways to set things aside
  - Scheduling take on only what you can make progress on
  - Trying to do too much causes thrashing
  - Design as network of bounded queues with finite servicing
- Across machines
  - Drop the request
  - Don't respond, keep the client hanging
  - But, all those millions are clients are independent !!! Doesn't that make it behave "open"?

### Limiting excessive load



- By not accepting connections, server slows ALL clients
- 10/1/19\* But some still make progression 10

#### Break

#### Recall: First-Come First-Served



• Just run processes in order of arrival

 Convoy Effect: Short processes stuck behind long processes

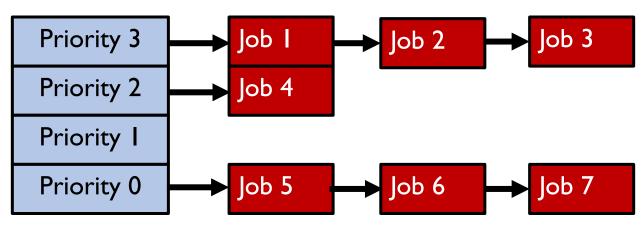
#### Recall: Round Robin

• Give out *small* units of CPU time ("time quantum")

Preempt a thread when its quantum expires

Cycle through ready threads

#### Recall: Priority Scheduling



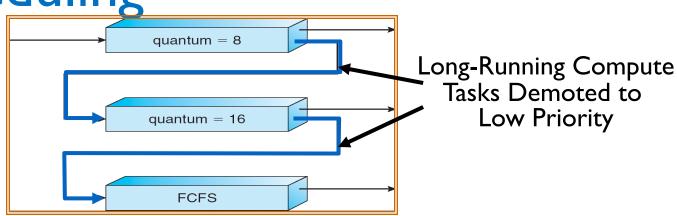
- Something gives jobs (processes) priority
  - User sets it explicitly
  - System manipulates priorities in pursuit of some policy
- Always run the ready thread with highest priority

# Recall: Shortest Job First & Shortest Remaining Time First

• **Provably Optimal** with respect to Response Time

- Key Idea: remove convoy effect
  - Short jobs always stay ahead of long ones

Recall: Multi-Level Feedback Scheduling



- Observe process behavior to approximate SRTF
- Starvation still possible without a workaround
- Basis for real OS schedulers, e.g. Linux O(I)

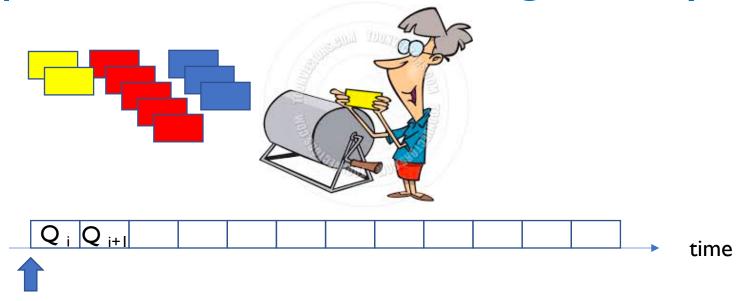
## Goals, Ends and Means

- Priorities were a means, not an end
- Our end goal was to serve a mix of CPU-bound, I/O bound, and Interactive jobs effectively on common hardware
  - Give the I/O bound ones enough CPU to issue their next file operation and wait (on those slow discs)
  - Give the interactive ones enough CPU to respond to an input and wait (on those slow humans)
  - Let the CPU bound ones grind away without too much disturbance

# Changing landscape of scheduling

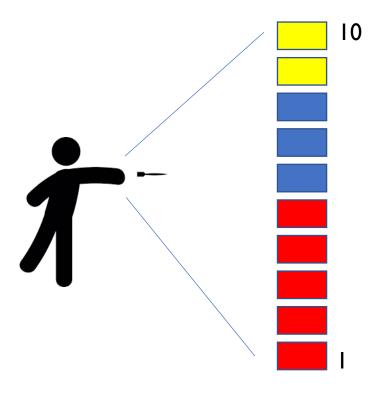
- Priority-based scheduling rooted in "time-sharing"
  - Allocating precious, limited resources across a diverse workload
    - CPU bound, vs interactive, vs I/O bound
- 80's brought about personal computers, workstations, and servers on networks
  - Different machines of different types for different purposes
  - Shift to fairness and avoiding extremes (starvation)
- 90's emergence of the web, rise of internet-based services, the data-center-is-the-computer
  - Server consolidation, massive clustered services, huge flashcrowds
  - It's about predictability, 95<sup>th</sup> percentile performance guarantees

## Proportion Share Scheduling: lottery



- Given a set of jobs (the mix), provide each with a proportional share of a resource
  - e.g., 50% of the CPU for Job A, 30% for Job B, and 20% for Job C
- Idea: Give out tickets according to the portion each should receive,
- Every quanta (tick) draw one at random, schedule that job (thread) to run

## Simpler Mechanism



- $N_{ticket} = \sum N_i$
- Pick a number d in 1..  $N_{ticket}$  as the random "dart"
- Jobs record their N<sub>i</sub> of allocated tickets
- Order them by N<sub>i</sub>
- Select the first j such that  $\sum N_i$  up to j exceeds d.

Various additional measures to allow a job to subdivide its tickets to divide up its share.

#### **UnFairness**

• E.g., Given two jobs A and B of same run time (# Qs) that are each supposed to receive 50%,

U = finish time of first / finish time of last

As a function of run time

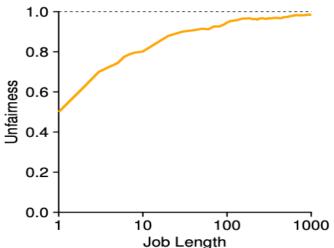


Figure 9.2: Lottery Fairness Study

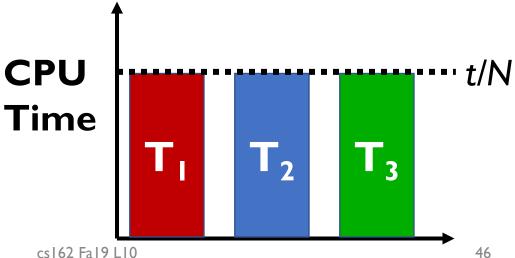
## Stride Scheduling

- Achieve proportional share scheduling without resorting to randomness, and overcome the "law of small numbers" problem.
- "Stride" of each job is  $\frac{big\#W}{N_i}$ 
  - The larger your share of tickets, the smaller your stride
  - Ex:W = 10,000, A=100 tickets, B=50, C=250
  - A stride: I00, B: 200, C: 40
- Each job as a "pass" counter
- Scheduler: pick job with lowest pass, runs it, add its stride to its pass
- Low-stride jobs (lots of tickets) run more often
  - Job with twice the tickets gets to run twice as often
- Some messiness of counter wrap-around, new jobs, ...

## Linux Completely Fair Scheduler

- Goal: Each process gets an equal share of CPU
- N threads "simultaneously" execute on I/Nth of CPU

At any time t we would observe:



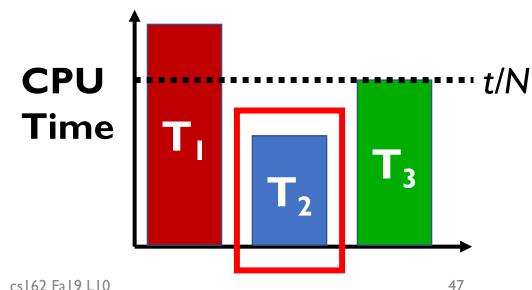
10/1/19

## Linux Completely Fair Scheduler

- Can't do this with real hardware
  - Still need to give out full CPU in time slices
- Instead: track CPU time given to a thread so far

#### **Scheduling Decision:**

- "Repair" illusion of complete fairness
- Choose thread with minimum CPU time



10/1/19

## Linux CFS: Responsiveness

- Goal: Preserve response time
- Constraint 1: Target Latency
  - Period of time over which every process gets service
  - Quanta = Target\_Latency / n
- Target Latency: 20ms, 4 Processes
  - Each process gets 5ms time slice
- Target Latency: 20 ms, 200 Processes
  - Each process gets 0.1 ms time slice (!!!)
  - Recall Round-Robin: large context switching overhead if slice gets to small

#### Linux CFS: Overhead

- Goal: throughput
  - avoid excessive overhead,
- Constraint 2: Minimum Granularity
  - Minimum length of any time slice
- Target Latency 20ms, Minimum Granularity Ims, 200 processes
  - Each process gets Ims time slice

## Aside: Priority in Unix – Being Nice

- The industrial operating systems of the 60s and 70's provided priority to enforced desired usage policies.
- When it was being developed at Berkeley, instead it provided ways to "be nice".
- nice values range from -20 to 19
  - Negative values are "not nice"
  - If you wanted to let you friends get more time, you would nice up you job
- Schedule puts higher nice (lower priority) to sleep more ...

## Linux CFS: beyond equal share

- What if we want to give more to some and less to others (proportional share)?
- Reuse nice value to reflect share, rather than priority
- Key Idea: Assign a weight  $w_i$  to each process i
- Basic equal share:  $Q = \text{Target Latency } * \frac{1}{N}$
- Weighted Share:

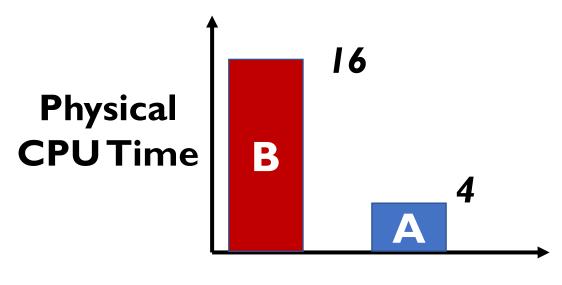
$$Q_i = (w_i / \sum_p w_p) * \text{Target Latency}$$

## Linux CFS: weighted share

- Target Latency = 20ms,
- Minimum Granularity = Ims
- Two CPU-Bound Threads
  - Thread A has weight I
  - Thread B has weight 4
- Time slice for A? 4 ms
- Time slice for B? 16 ms

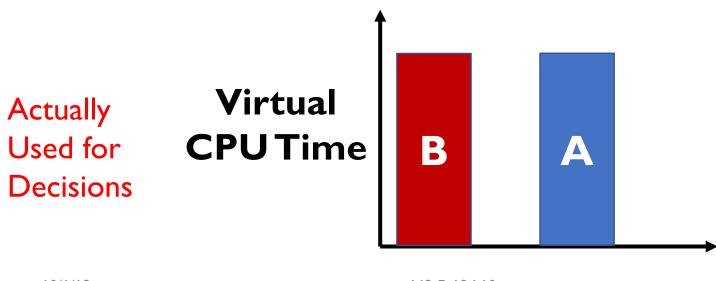
# Linux CFS: equal weighted shares

- Track a thread's virtual runtime rather than its true physical runtime
- Higher weight: Virtual runtime increases more slowly
- Lower weight: Virtual runtime increases more quickly



## Linux CFS: equal weighted shares

- Track a thread's *virtual* runtime rather than its true physical runtime
- Higher weight: Virtual runtime increases more slowly
- Lower weight: Virtual runtime increases more quickly



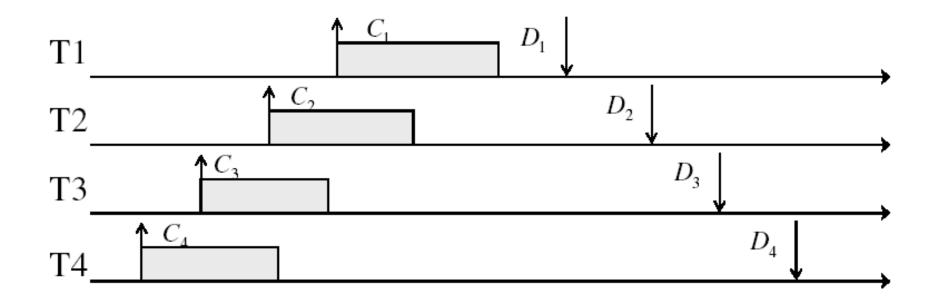
## Real-Time Scheduling

- Goal: Guaranteed Performance
  - Meet deadlines even if it means being unfair or slow
  - Limit how bad the worst case is

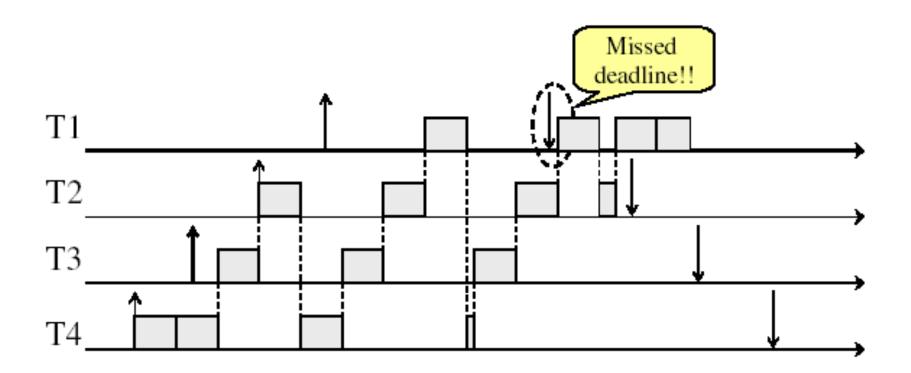
- Hard real-time:
  - Meet all deadlines (if possible)
  - Ideally: determine in advance if this is possible

## Real-Time Example

Preemptible tasks with known deadlines (D) and known burst times (C)

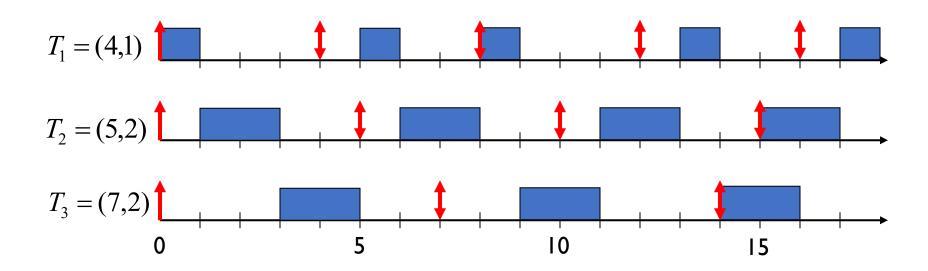


# What if we try Round-Robin?



#### Earliest Deadline First

- Priority scheduling with preemption
- Priority proportional to time until deadline
- Example with periodic tasks:



# **EDF: Feasibility Testing**

- Even EDF won't work if you have too many tasks
- For *n* tasks with computation time *C* and deadline *D*, a feasible schedule exists if:

$$\sum_{i=1}^{n} \left( \frac{C_i}{D_i} \right) \le 1$$

# Choosing the Right Scheduler

I Care About:	Then Choose:
CPU Throughput	FCFS
Avg. Response Time	SRTF Approximation
I/O Throughput	SRTF Approximation
Fairness (CPU Time)	Linux CFS
Fairness – Wait Time to Get CPU	Round Robin
Meeting Deadlines	EDF
Favoring Important Tasks 10/1/19 cs 162 Fa	Priority 60

## Further reading

 Lottery scheduling: flexible proportionalshare resource management, Waldspurger and Weihl, OSDI 1994