



CS 162: Operating Systems and Systems Programming


Lecture 21: Distributed Key-Value Store & 2-Phase Commit

November 12, 2019

Instructor: David E. Culler

<https://cs162.eecs.berkeley.edu>

Distributed File Systems



- VFS (Virtual File System) interface allows client “front end” to carry out protocol with server “back end”
 - E.g., NFS RPC, stateless, open, read/write close
- Unit of access: NFS blocks, AFS whole file
- Client caching and consistency management
 - NFS eventual consistency, AFS last writer wins, notification on close
 - Enforce single-write, multiple reader discipline
- Handling failures



Recall: Distributed

- **Transparent access to files located on remote disks**
 - Caching for performance
 - Blocks or whole files
 - Introduces consistency issues
 - » File save vs write
 - » Remote vs actively shared
 - NFS: Check periodically for changes to server copy
 - AFS: Server notifies client of changes



The Promise of Dist. Systems

- **Availability:** One machine goes down, overall system stays up
- **Durability:** One machine loses data, but *system* does not lose anything
- **Security:** Easier to secure each component of the system individually?



Dist. System – the darker side

- **Availability:** Failure in one machine causes others to hang waiting for it
 - Two sides of Fate sharing
- **Durability:** Lots of machines that might lose your data
- **Security:** More components means more points of attack
- **Engineering of distributed systems – both cloud and end hosts – are fundamentally more reliable than in the 80's and 90's when the approach emerged**



Sharing Data, rather than Files ?

- **Key:Value stores are used everywhere**
- **Native in many programming languages**
 - Associative Arrays in Perl
 - Dictionaries in Python
 - Maps in Go
 - ...
- **What about a collaborative key-value store rather than message passing or file sharing?**
 - An alternative basis for building distributed systems
 - Especially distributed within the cloud
- **Can we make it scalable and reliable?**



Key Value Storage

Simple interface

- `put(key, value);` // Insert/write "value" associated with key
- `get(key);` // Retrieve/read value associated with key
- Remember wordcount? Word:Count
 - `AddCount(char *word, int count)` over KV?



Why Key Value Storage?

- **Easy to Scale**
 - Handle huge volumes of data (e.g., petabytes)
 - Uniform items: distribute easily and roughly equally across many machines
- **Relatively simple consistency properties**
 - vs general database transactions or file system ops over multiple blocks
- **Used as a simpler but more scalable "database"**
 - Or as a building block for a more capable DB

Key Values: Examples



- **Amazon:**
 - Key: customerID
 - Value: customer profile (e.g., buying history, credit card, ..)



- **Facebook, Twitter:**
 - Key: UserID
 - Value: user profile (e.g., posting history, photos, friends, ...)

- **iCloud/iTunes:**
 - Key: Movie/song name
 - Value: Movie, Song






KV Storage Systems in the Wild

- **Amazon**
 - **DynamoDB**: internal key value store used to power Amazon.com (shopping cart)
 - **Simple Storage System (S3)**
- **BigTable/HBase/Hypertable: distributed, scalable data storage**
 - All the different services share distributed systems infrastructure
- **Cassandra: “distributed data management system” (developed by Facebook)**
- **Memcached: in-memory key-value store for small chunks of arbitrary data (strings, objects)**



Key Value Store

- Also called Distributed Hash Tables (DHT)
- Main idea: partition set of key-value pairs across many machines
 - Easy to be a local key:value store





Important Questions

- **put(key, value):**
 - **where** do you store a new (key, value) tuple?
- **get(key):**
 - **where** is the value associated with a given “key” stored?
- **And, do the above while providing**
 - Fault Tolerance
 - Scalability
 - Consistency



How to solve the “where?”

- **Hashing**

- Given M nodes, all share a hash function: $\text{Key} \rightarrow \{0, \dots, M-1\}$
- Send Put/Get to node $\text{id} == \text{Hash}(\text{key})$
- Process it with local KV store

- **Challenges**

- But what if you don’t know “who” all the nodes are that are participating?
- Perhaps they come and go ...
- What if some keys are *really* popular? (load balance)
- Replicate K:V in multiple storage servers?

- **Have something keep track - directory**

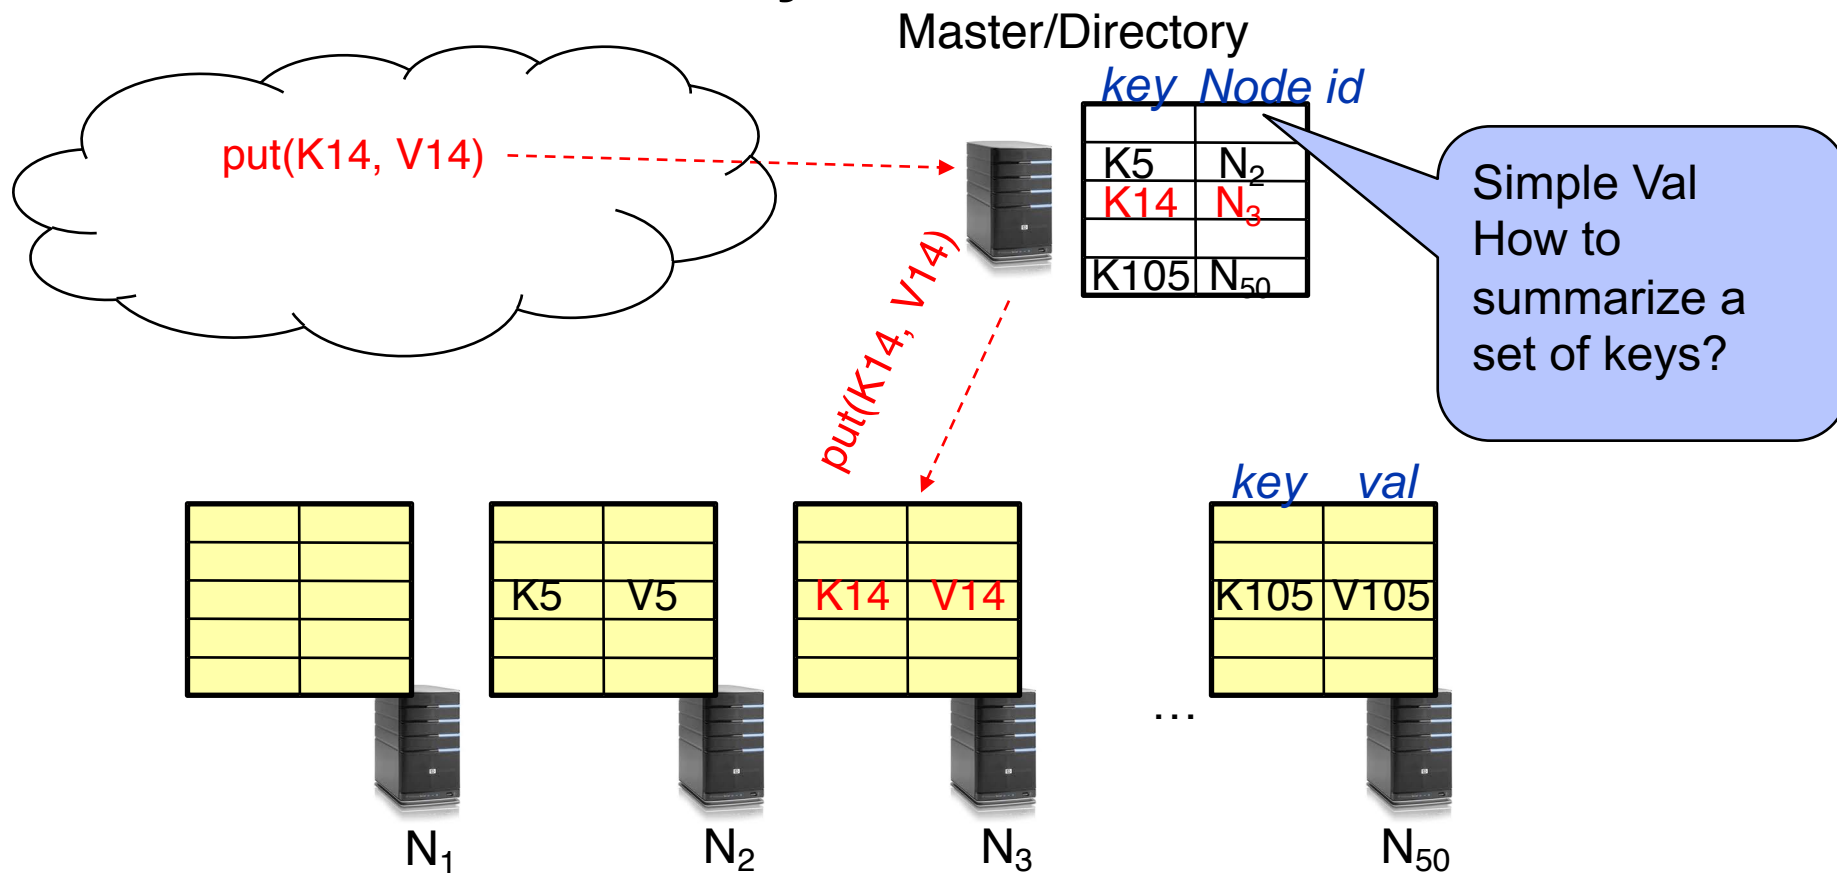
- **Lookup**

- Hmm, won’t this be a bottleneck and single point of failure?

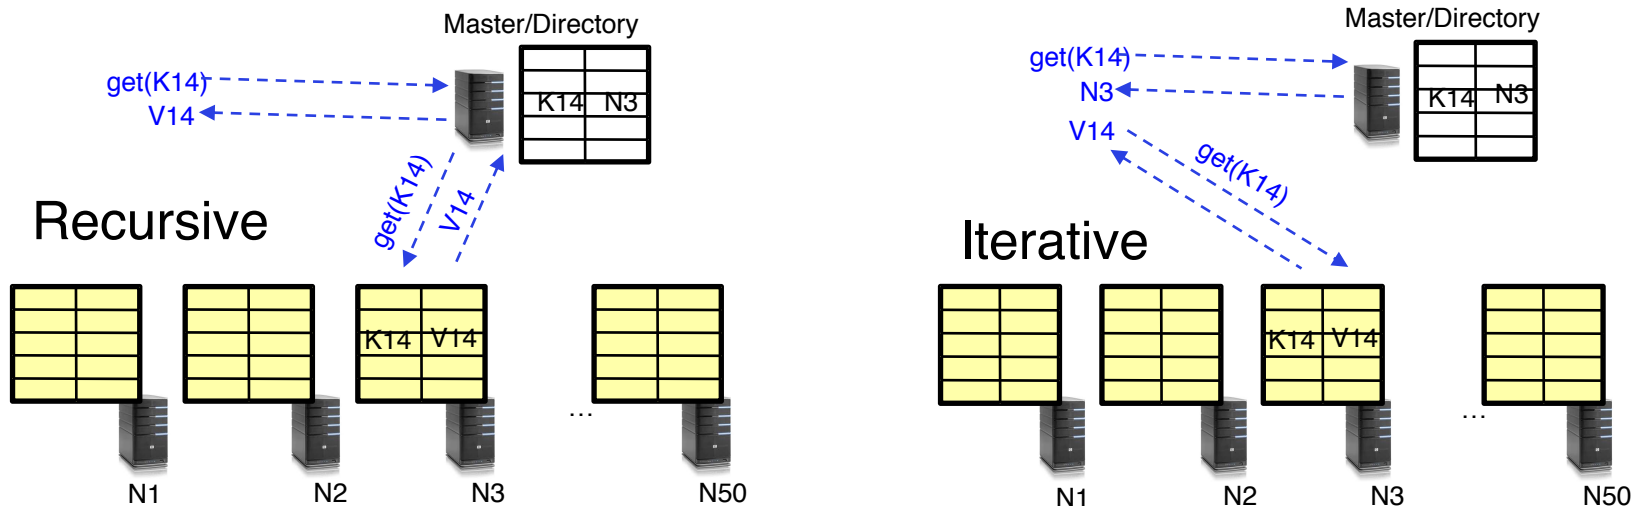


Directory-Based Architecture

Have a node maintain the mapping between keys and the *machines* (nodes) that store the values associated with the keys



Iterative vs. Recursive Query



- **Recursive Query: Directory Server Delegates**
- **Iterative Query: Client Delegates**



Iterative vs Recursive Query


Recursive

- + Faster, as directory server is typically close to storage nodes
- + Easier for consistency: directory can enforce *an order* for all puts and gets
- Directory is a performance bottleneck

Iterative


- + More scalable, clients do more work
- Harder to enforce consistency

Challenges



- **Fault Tolerance:** handle machine failures without losing data and without degradation in performance
- **Scalability:**
 - Need to scale to thousands of machines
 - Need to allow easy addition of new machines

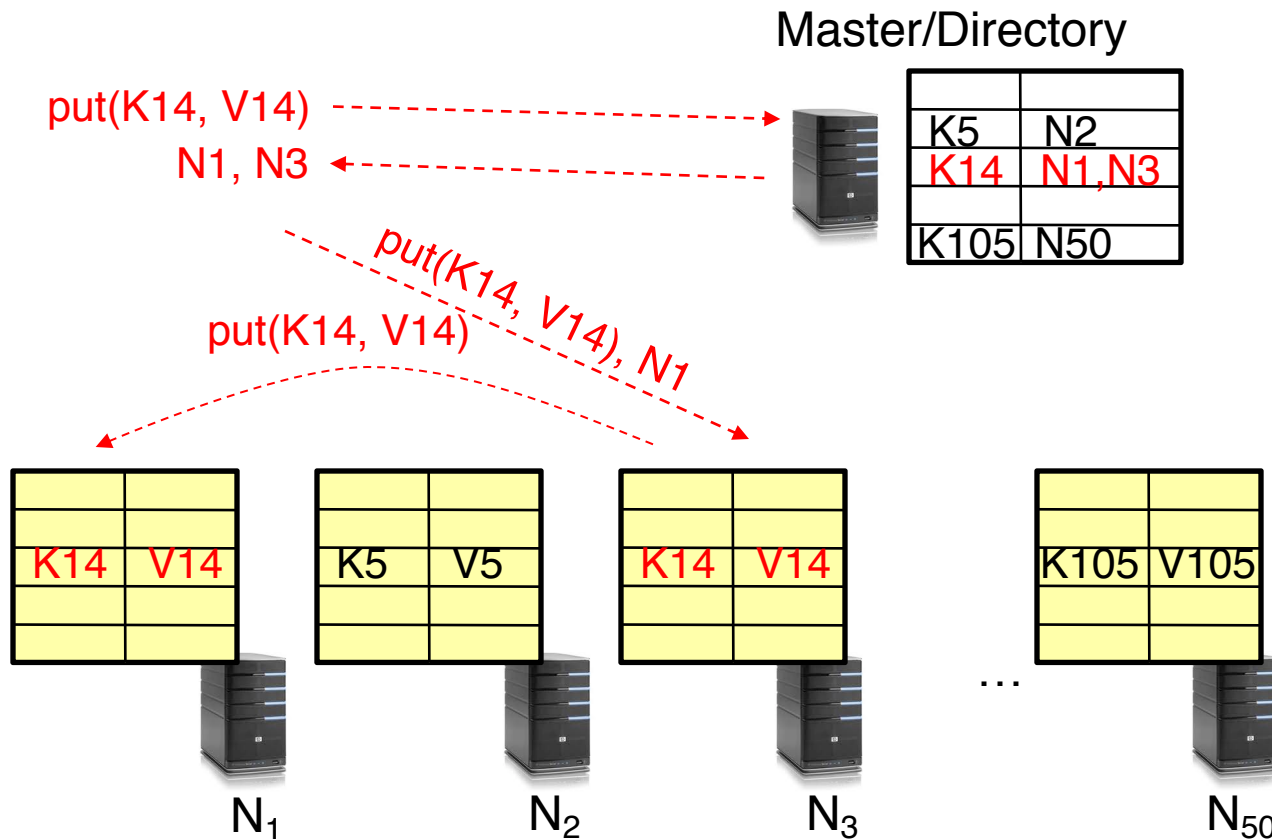
Challenges



- **Consistency:** maintain data consistency in face of node failures and message losses
- **Heterogeneity** (esp. if deployed as peer-to-peer systems):
 - Latency: 1ms to 1000ms
 - Bandwidth: 32 Kb/s to 1 Gb/s

Fault Tolerance

- Replicate value on several nodes
- Usually, place replicas on different racks in a datacenter to guard against rack failures



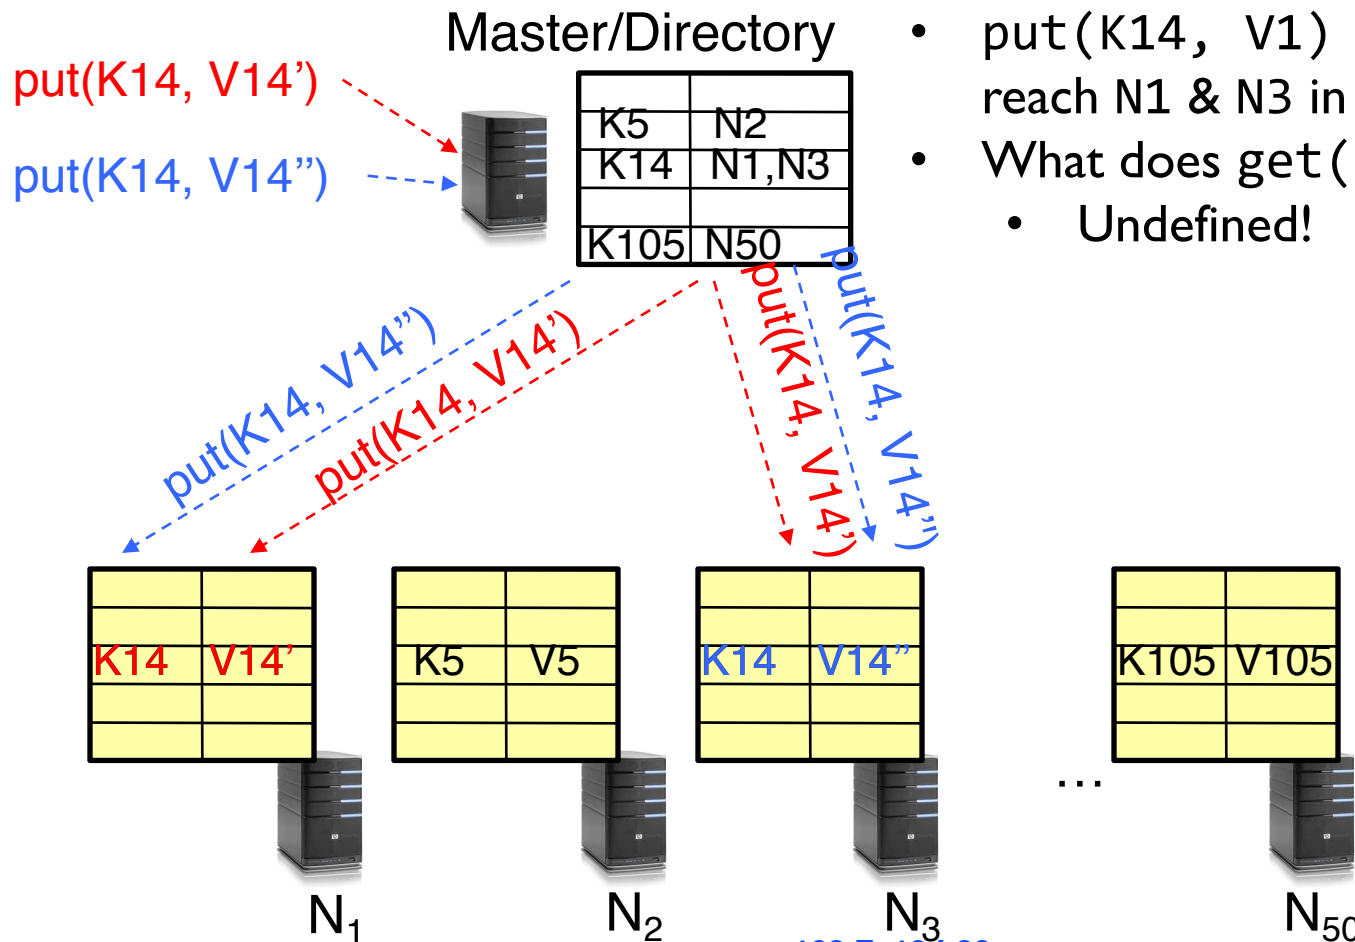


Consistency

- Replication is essential for fault tolerance (and performance)
 - But introduces inherent challenges
- Need to make sure a value is replicated correctly
- How do you know a value is replicated on every expected node?
- **Wait** for acknowledgements from all expected nodes ???

Consistency

- If concurrent updates (i.e., puts to same key) may need to make sure that updates happen in the “same order”



- put(K14, V1) and put(K14, V2) reach N1 & N3 in reverse order
- What does get(K14) return?
 - Undefined!



How to ensure order?

- Wait – for explicit acknowledgement
- Hmm..



Consistency & Fault Tolerance

- **What happens if a node fails during replication?**
 - Pick another node and try again
- **What happens if a node is slow?**
 - Slow down entire put? Pick another node?
- **In general with multiple replicas: slow put and fast get operations**




Quorum Consensus

- Improve **put** and **get** operation performance by reducing the # of replicas to hear back from
- Define a replica set of size N
 - **put** waits for acknowledgements from at least W replicas
 - **get** waits for responses from at least R replicas
 - $W + R > N$
- Why does it work?
 - There is at least one node that contains the update
- Why might you use $W+R > N+1$?

Quorum Consensus Example


- $N=3$, $W=2$, $R=2$
- Replica set for K14: {N1, N2, N4}





Quorum Consensus Example

- Now, issuing get to any two nodes out of three will return the answer





Scalability

How easy is it to make the system bigger?

- **Storage: Use more nodes**
- **Number of Requests**
 - Can serve requests from all nodes on which a value is stored in parallel
 - Master can replicate a popular item on more nodes
- **Master/Directory Scalability**
 - Replicate It (multiple identical copies)
 - » Maintain consistency across them
 - Partition it, so different keys are served by different directories



Scalability: Load Balancing

- **Directory tracks available storage at each node**
 - Prefer to insert at nodes with more storage available
- **What happens when a new node is added?**
 - Cannot insert only new values at new node
 - Move values from heavily loaded nodes to new node
- **What happens when a node fails?**
 - Replicate values from failed node to other nodes



Scaling Up Directory

- Directory contains number of entries equal to number of key/value pairs in entire system
 - Could be tens or hundreds of billions of pairs
- Solution: **Consistent Hashing**
 - The set of storage nodes may change dynamically
 - » fail, enter, leave
 - Assign each node a unique ID in large namespace $[0..2^m-1]$
 - » m bit namespace, s.r., $M \ll 2^m$
 - » Each node can pick its ID at random !
 - hash keys in a manner that everyone assigns same range of IDs to a node
 - Each (key,value) stored at node with *smallest ID larger than hash(key)*
- Important property: Adding a new bucket doesn't require moving lots of existing values to new buckets



Key to Node Mapping Example

Partitioning example with
 $m = 6 \rightarrow$ ID space: 0..63

0: Node 4 maps keys [59, 4]

1: Node 8 maps keys [5,8]

2: Node 15 maps keys [9,15]

3: Node 20 maps keys [16, 20]


...

M-1: Node 58 maps [45, 58]

$n := \text{Hash}(\text{key})$

Find first i in [4, 8, 15, 20, ...]

s.t., $N_i > n \pmod{M}$






Performing a Lookup

- **Any one with a list of Node IDs can hash a key and pick the node responsible for that range**
 - But the mapping may have changed due to nodes join/leave
- **Each node knows about its successor and predecessor in the "circle"**
 - If it is no longer responsible for the hash(key), it can find node that is
 - All that is strictly needed for correctness
- **Fully decentralized**
 - Any node can act as a directory for clients
 - Still works if a node leaves the network
- **Faster lookups: Each node maintains a routing table, allows client to get closer to destination in one hop**



Example: Chord





Logistics Break


- **Key:Value, DHT questions**
- **12/14 Eric Brewer guest lecture – Google VP Infra**
- **Work for remainder of term released**
 - Provide as much flexibility as possible, given other demands
 - Do both HW5a and HW6a and one of HW{5 or 6}b
 - Extra credit applied to midterm for the other part b
 - » HW6b still in beta
- **Project 3**
 - Simplified concurrency aspects
 - Buffer Cache, inode extension, dirent
- **... on to 2-phase commit**



Consistency Review

- **Problem: shared state replicated across multiple clients, do they see a consistent view?**
 - **Propagation: Writes become visible to reads**
 - **Serializability: The order of writes seen by each client's series of reads and writes is *consistent* with a total order**
 - » **As if all writes and reads had been serviced at a single point**
 - » **The total order is not actually generated, but it could be**
- **Many distributed systems provide weaker semantics**
 - **Eventual consistency**

Unfinished Business: Multiple Servers



Finish Previous Lecture

- What happens if cannot update all the replicas?
- Availability => Inconsistency



In Everyday Life

Where do we meet?

Where do we meet?

Where do we meet?
At Nefeli's
At Top Dog

~~Where do we meet?
At Nefeli's~~

Where do we meet?

At Nefeli's

Where do we meet?
At Nefeli's
At Top Dog

~~Where do we meet?
At Top Dog
At Nefeli's~~

Where do we meet?

At Top Dog

Where do we meet?
At Nefeli's
At Top Dog

~~Where do we meet?
At Nefeli's
At Top Dog~~



Durability and Atomicity

- **How do you make sure transaction results persist in the face of failures (e.g., server node failures)?**
- **Replicate store / database**
 - Commit transaction to each replica
- **What happens if you have failures during a transaction commit?**
 - Need to ensure atomicity: either transaction is committed on all replicas or none at all

Distributed Consensus Making



- **Consensus problem**
 - All nodes propose a value
 - Some nodes might crash and stop responding
 - Eventually, all remaining nodes decide on the same value from set of proposed values
- **Distributed Decision Making**
 - Choose between “true” and “false”
 - Or Choose between “commit” and “abort”
- **Equally important (but often forgotten!): make it durable!**
 - How do we make sure that decisions cannot be forgotten?
 - » This is the “D” of “ACID” in a regular database
 - In a global-scale system?
 - » What about erasure coding or massive replication?
 - » Like **BlockChain** applications!



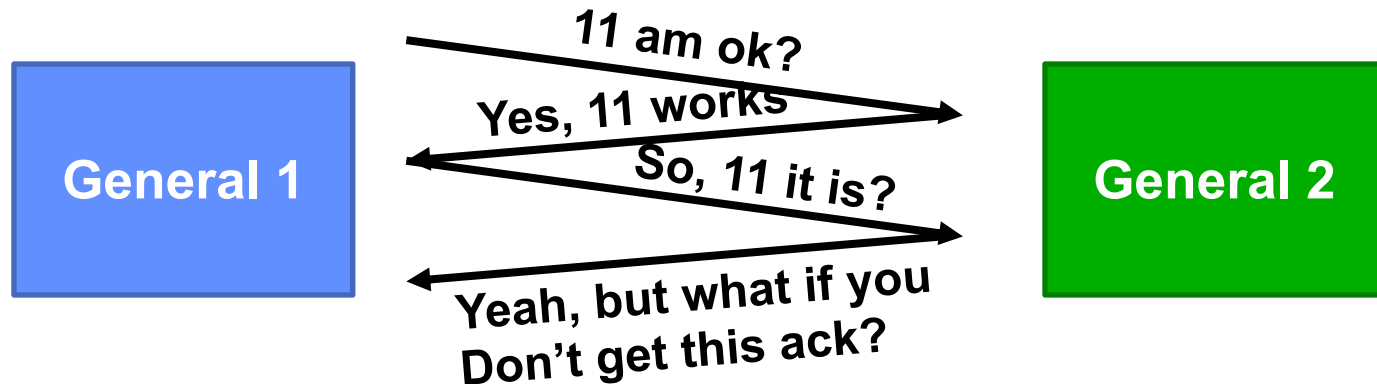
“Byzantine General’s Paradox”

- **Two generals located on opposite sides of their enemy’s position**
- **Can only communicate via messengers**
- **Messengers go through enemy territory: might be captured**
- **Problem: Need to coordinate time of attack**
 - Both generals lose unless they attack at same time
 - If they attack at same time, they win
- **How do you ever know your message got through?**
 - Requires a message... that might not get through



General's Paradox

- Can messages over an unreliable network be used to guarantee two entities do something simultaneously?
- No, even if all messages go through





Two-Phase Commit

- **We can't solve the General's Paradox**
 - No simultaneous action
 - But we can solve a related problem
- **Distributed Transaction: Two (or more) machines agree to do something *or not do it* atomically**
- **Extra tool: Persistent Log**
 - If machine fails, it will remember what happened
 - Assume log itself can't be corrupted



Two Phase (2PC) Commit

- 2PC is a distributed protocol
- High-level problem statement
 - If no node fails and all nodes are ready to commit, then all nodes **COMMIT**
 - Otherwise **ABORT** at all nodes
- Developed by Turing award winner Jim Gray (first Berkeley CS PhD, 1969)



2PC Algorithm

- One coordinator
- N workers (replicas)
- High level algorithm description
 - Coordinator asks all workers if they can commit
 - If all workers reply “**VOTE-COMMIT**”, then coordinator broadcasts “**GLOBAL-COMMIT**”,
Otherwise coordinator broadcasts “**GLOBAL-ABORT**”
 - Workers obey the **GLOBAL** messages



Two-Phase Commit: Setup

- One machine (*coordinator*) initiates the protocol
- It asks *every* machine to vote on transaction
- Two possible votes:
 - Commit
 - Abort
- Commit transaction only if unanimous approval



Two-Phase Commit: Preparing

Agree to Commit

- Machine has guaranteed that it will accept transaction
- Must be recorded in log so machine will remember this decision if it fails and restarts

Agree to Abort

- Machine has guaranteed that it will never accept this transaction
- Must be recorded in log so machine will remember this decision if it fails and restarts



Two-Phase Commit: Finishing

Commit Transaction

- Coordinator learns *all machines have agreed to commit*
- Apply transaction, inform voters
- Record decision in local log

Abort Transaction

- Coordinator learns *at least one machine has voted to abort*
- Do not apply transaction, inform voters
- Record decision in local log



Two-Phase Commit: Finishing

Commit Transaction

- Coordinator learns *all machines have agreed to commit*
- Apply transaction, inform voters
- Record decision in local log

Abort Transaction

- Coordinator learns *at least one machine has voted to abort*
- Do not apply transaction, inform voters
- Record decision in local log

Because no machine can take back its decision, exactly one of these will happen



Detailed Algorithm

Coordinator Algorithm

Coordinator sends **VOTE-REQ** to all workers

- If receive **VOTE-COMMIT** from all N workers, send **GLOBAL-COMMIT** to all workers
- If doesn't receive **VOTE-COMMIT** from all N workers, send **GLOBAL-ABORT** to all workers


Worker Algorithm

- Wait for **VOTE-REQ** from coordinator
- If ready, send **VOTE-COMMIT** to coordinator
- If not ready, send **VOTE-ABORT** to coordinator
 - And immediately abort

- If receive **GLOBAL-COMMIT** then commit
- If receive **GLOBAL-ABORT** then abort

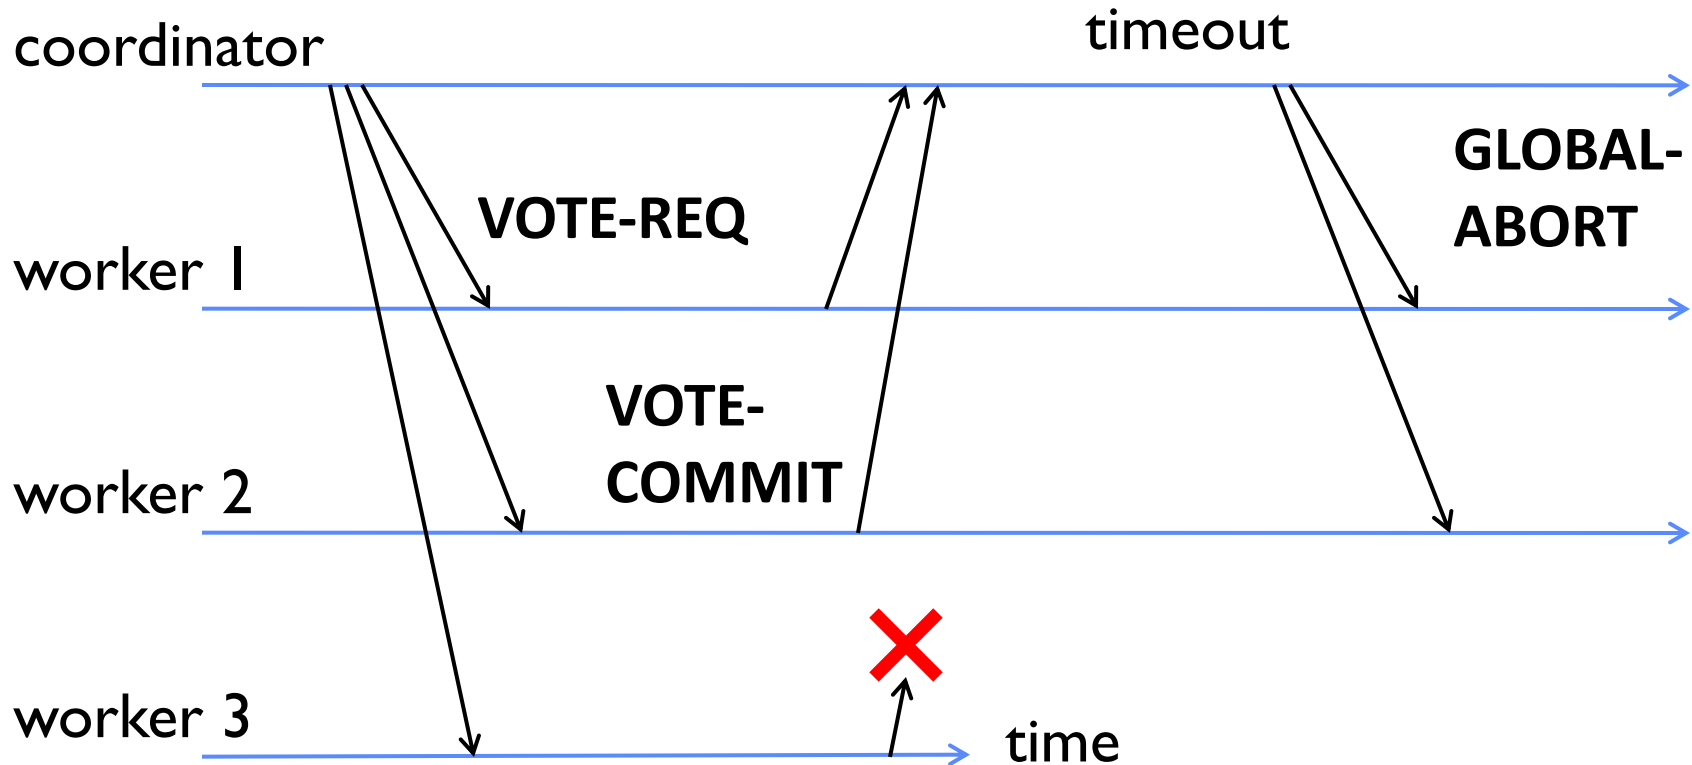


Example: Failure-Free 2PC



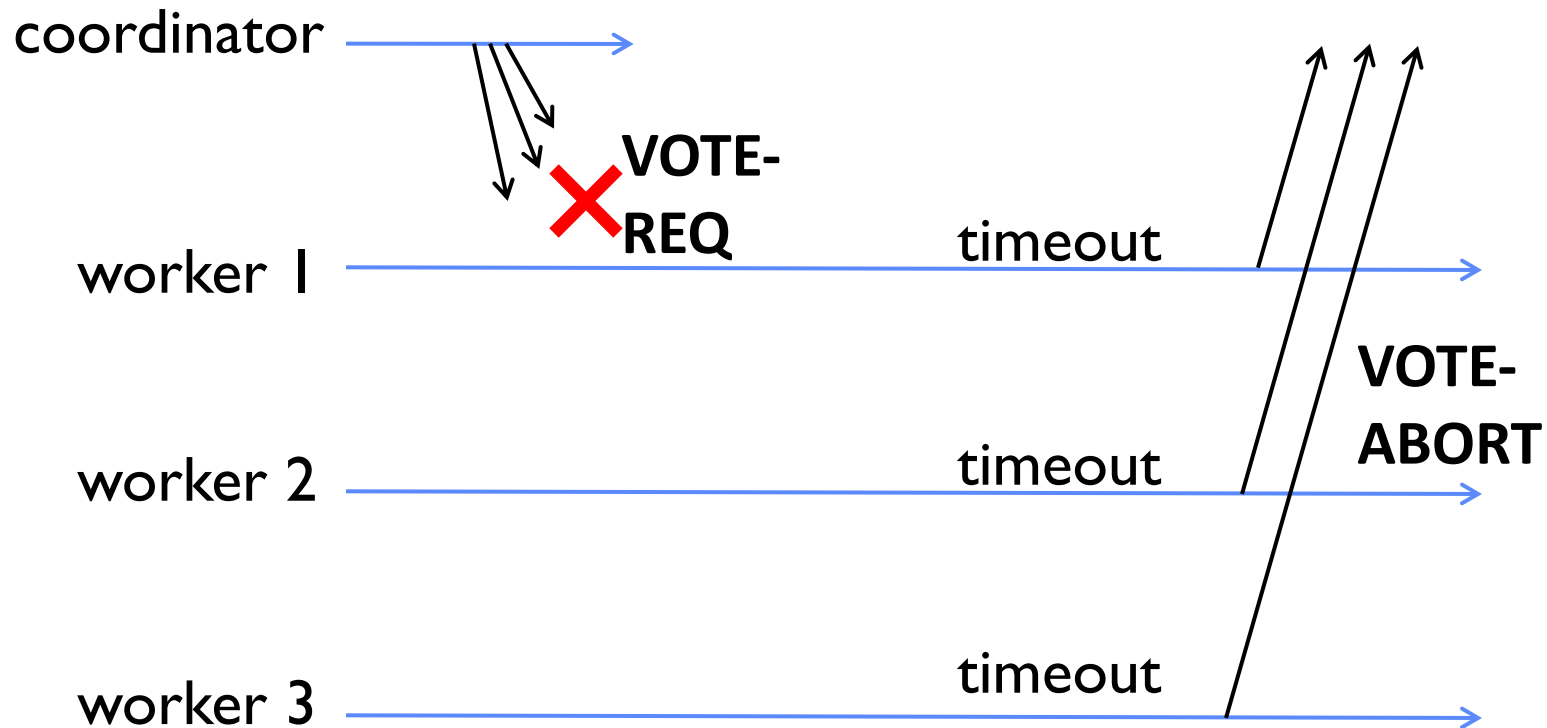


Example of Worker Failure





Example of Coordinator Failure





Formalizing Two-Phase Commit

- ***N* workers (replicas): actually perform *transactions***
- **One coordinator (may also serve a worker)**
 - Asks each worker to vote on transaction
 - Tells every machine result of the vote (workers don't need to ask each other)



Messages in Two-Phase Commit

Coordinator → Worker

- **VOTE-REQ**

Worker → Coordinator

- **VOTE-COMMIT**
- **VOTE-ABORT**

Coordinator → Worker

- **GLOBAL-COMMIT**
- **GLOBAL-ABORT**

No taking back: always
logged before sending

Actual result of transaction
attempt

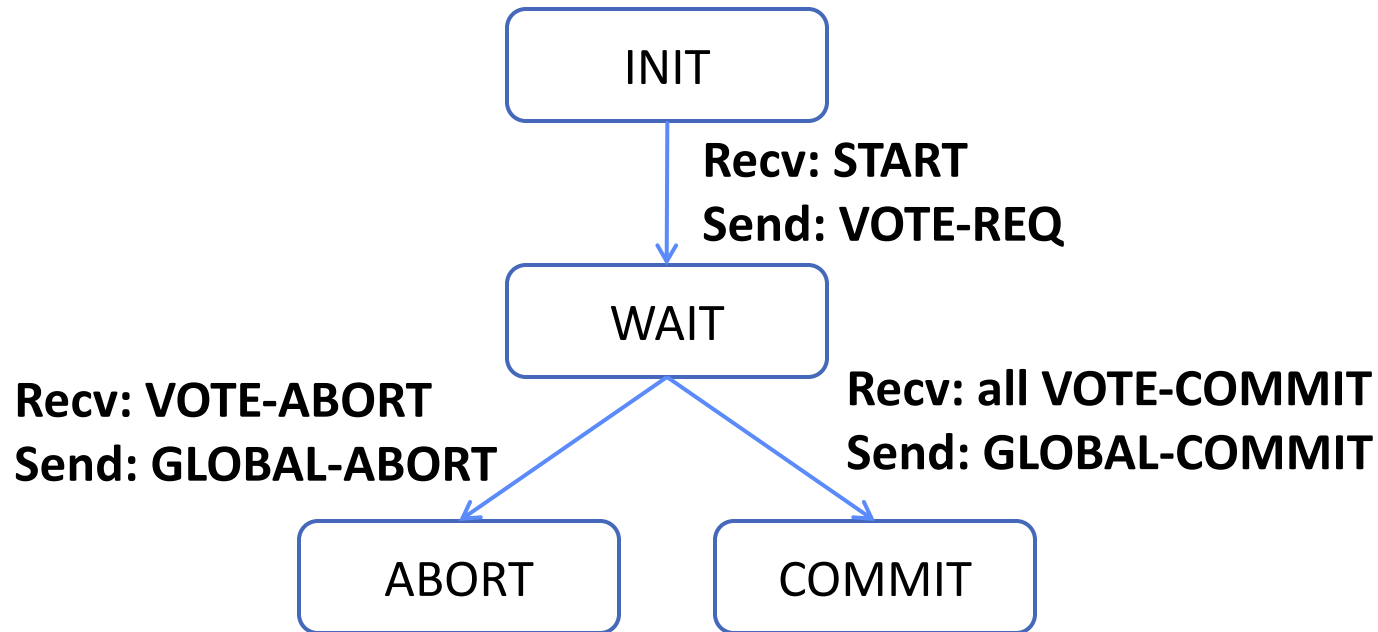


State Machines

- Distributed systems are hard to reason about
- Want a *precise* way to express each node's behavior that is also *easy to reason about*
- One approach: State Machine
 - Every node is in a *state*
 - When the node receives a message (or timeout),
 - it *transitions* to another state and
 - Sends zero or more messages

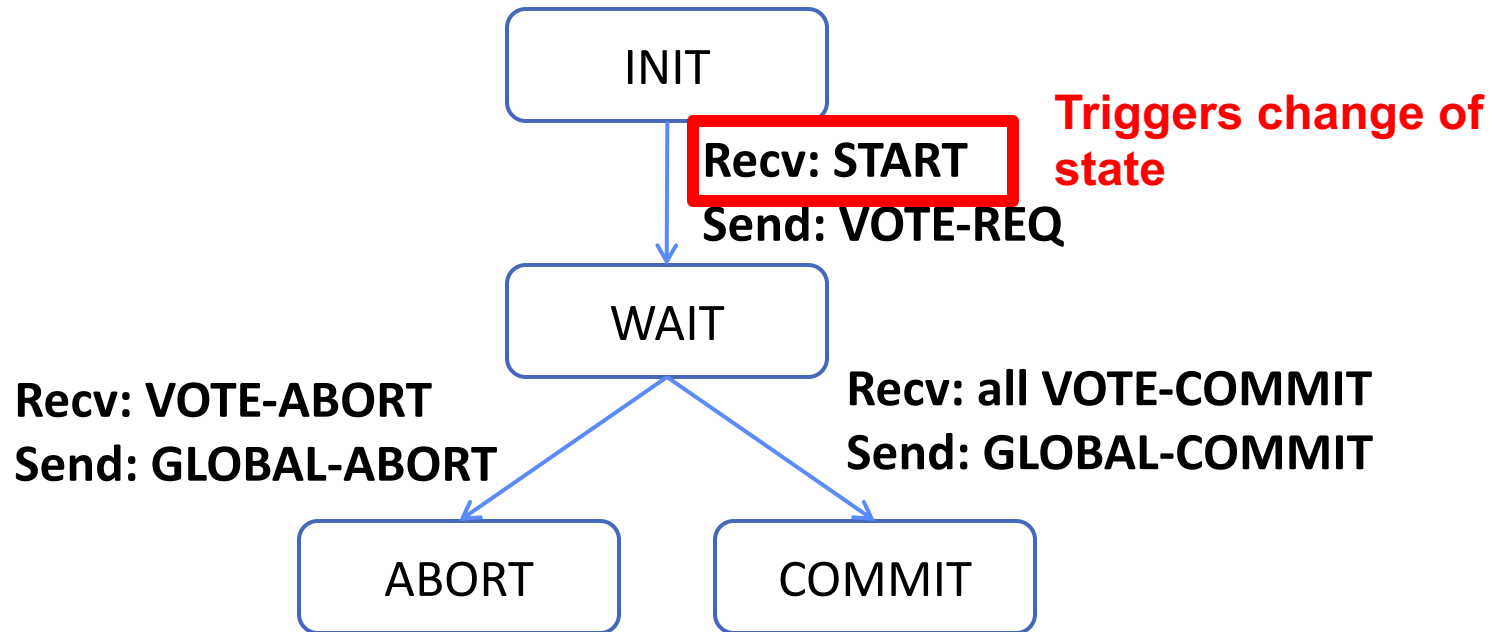


Coordinator's State Machine



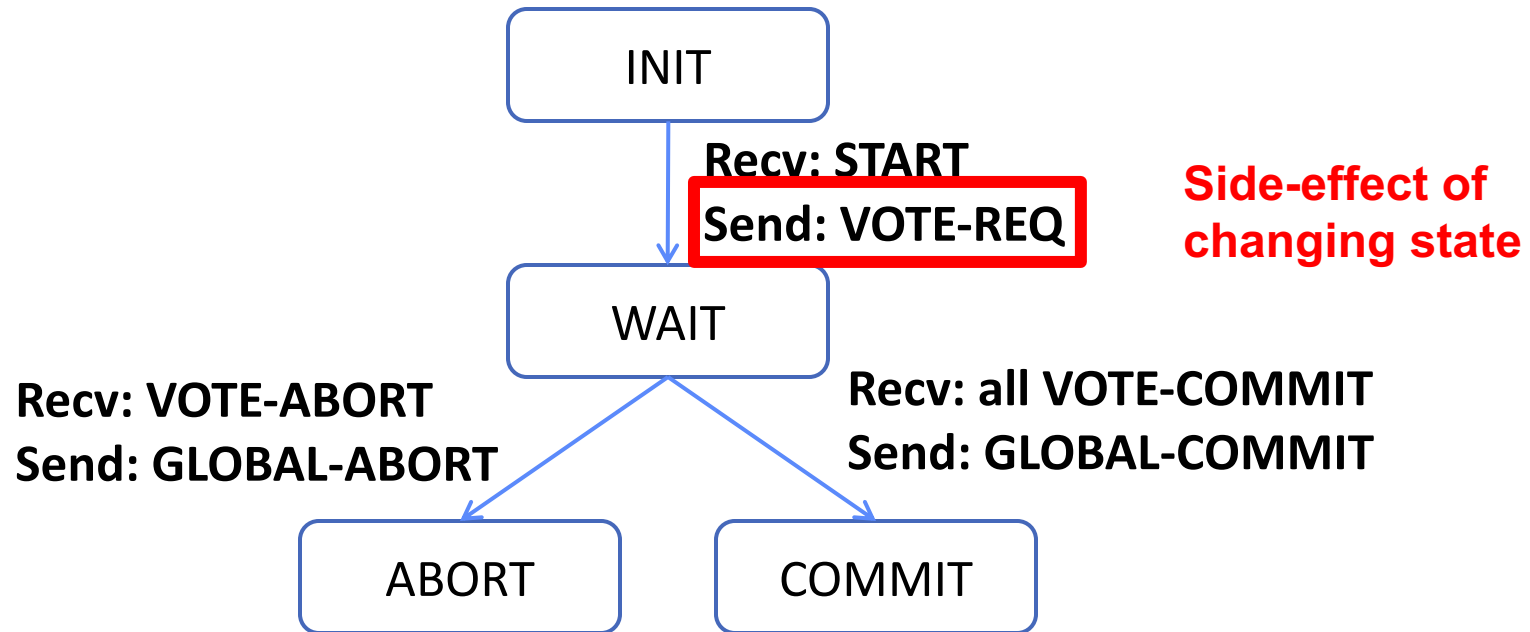


Coordinator's State Machine



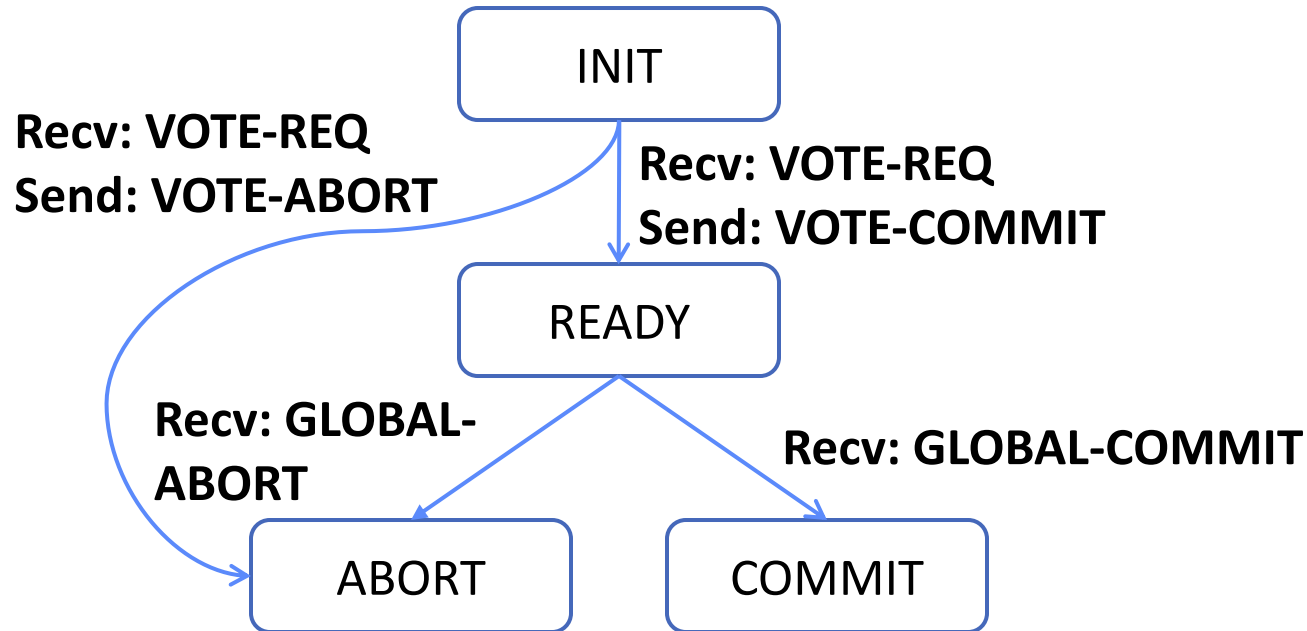


Coordinator's State Machine



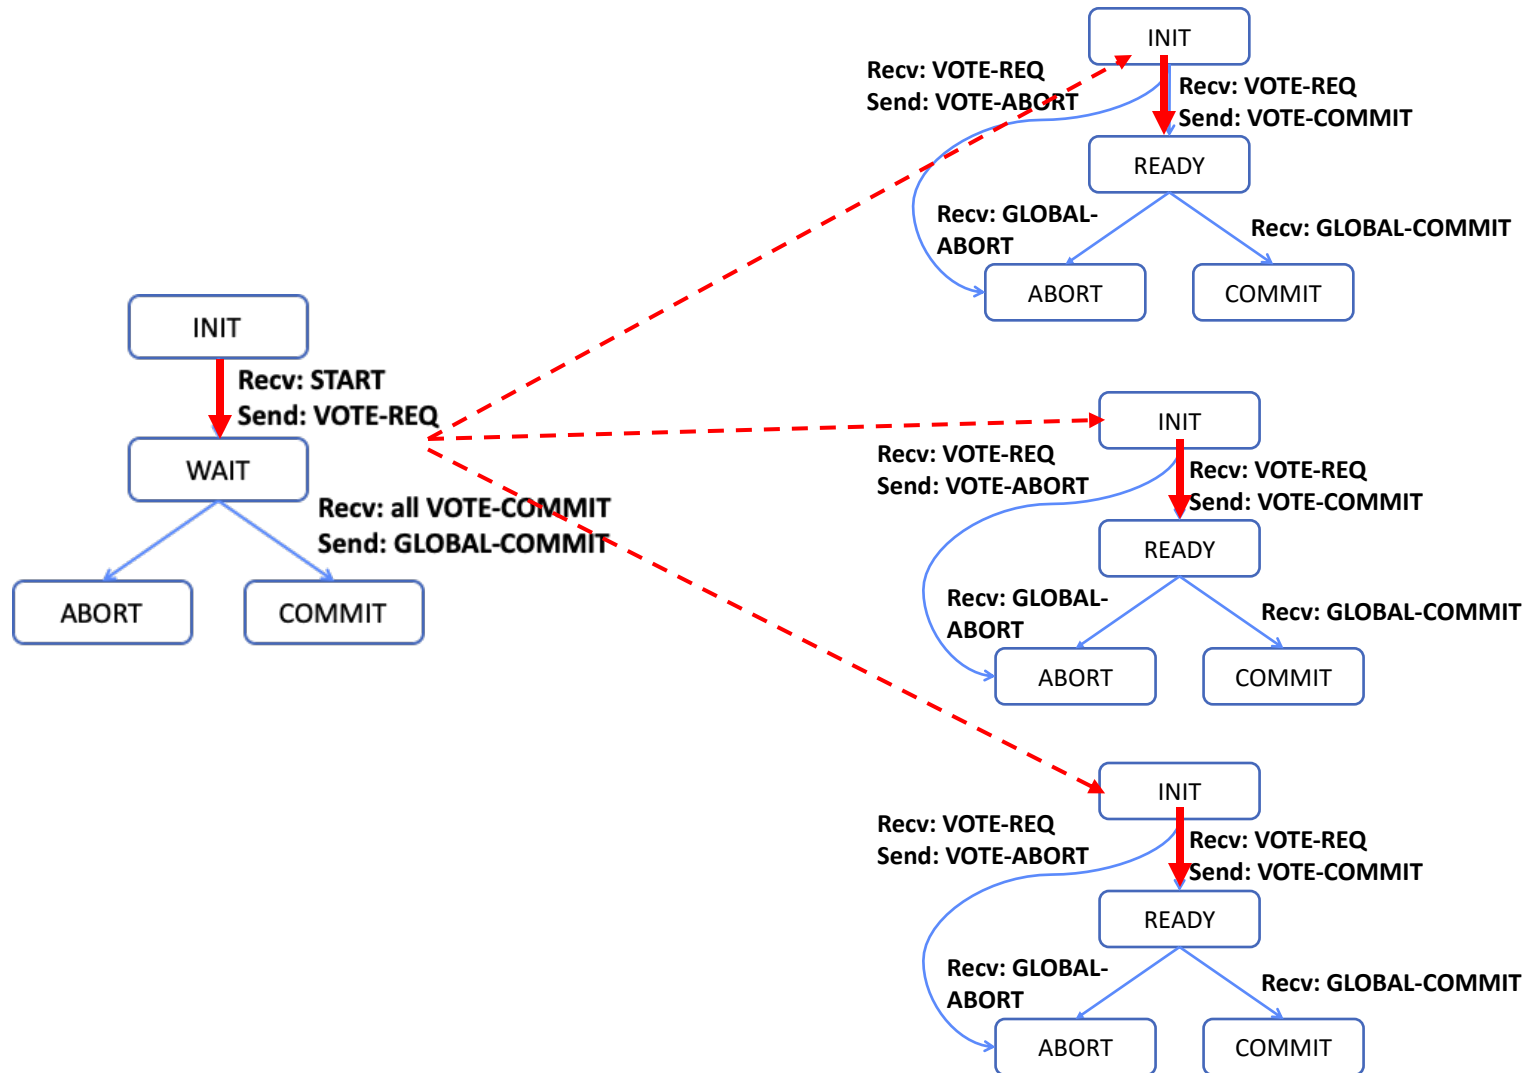


Worker's State Machine



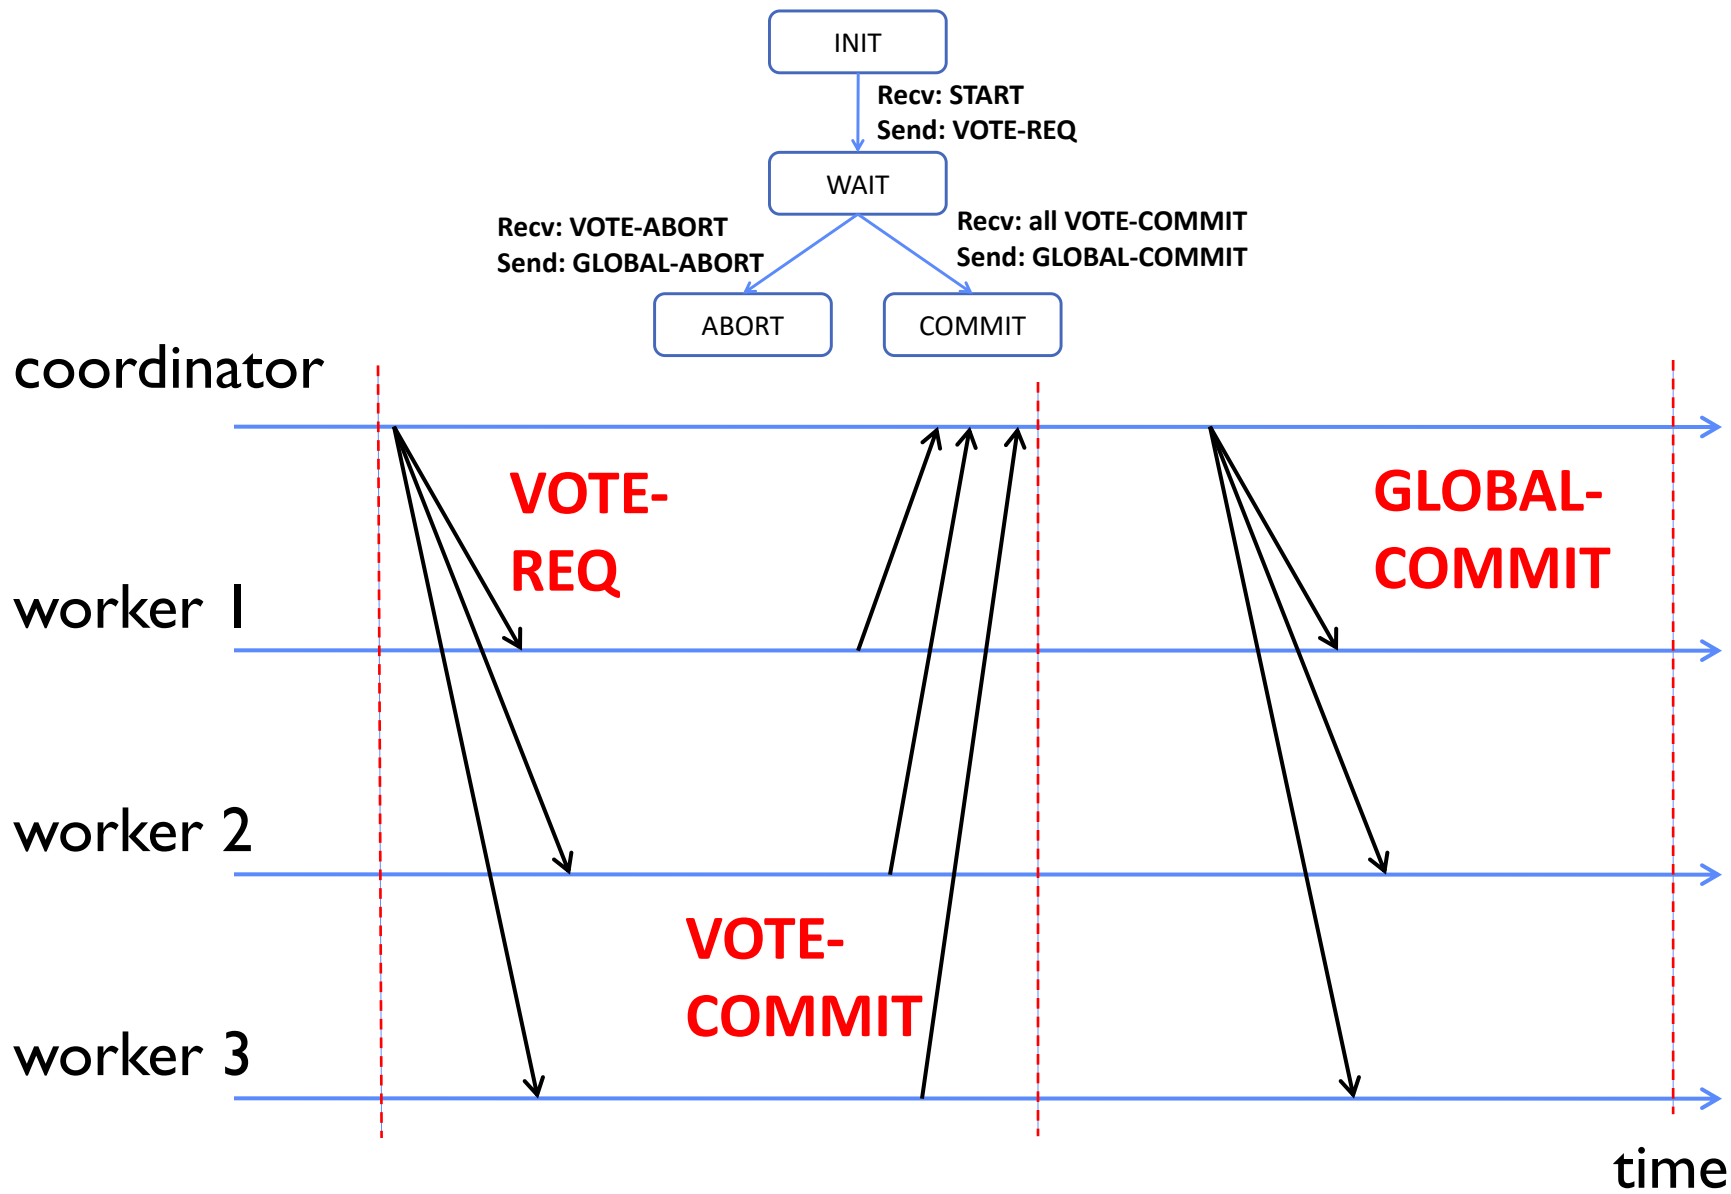


Protocol: cooperating state machines





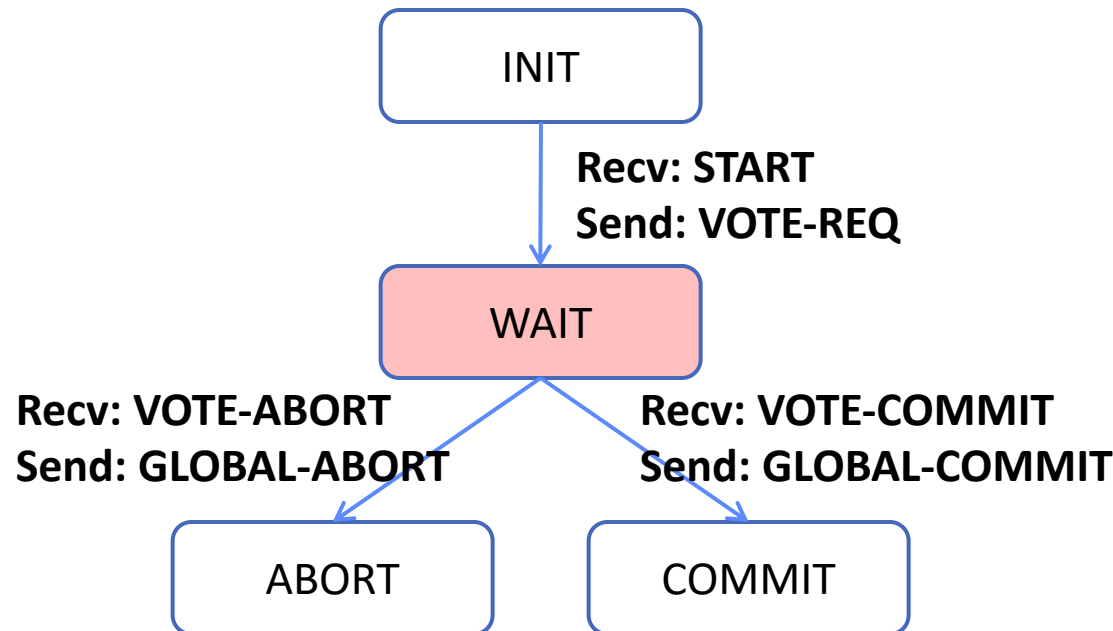
Example: Failure-Free 2PC (w/ state)





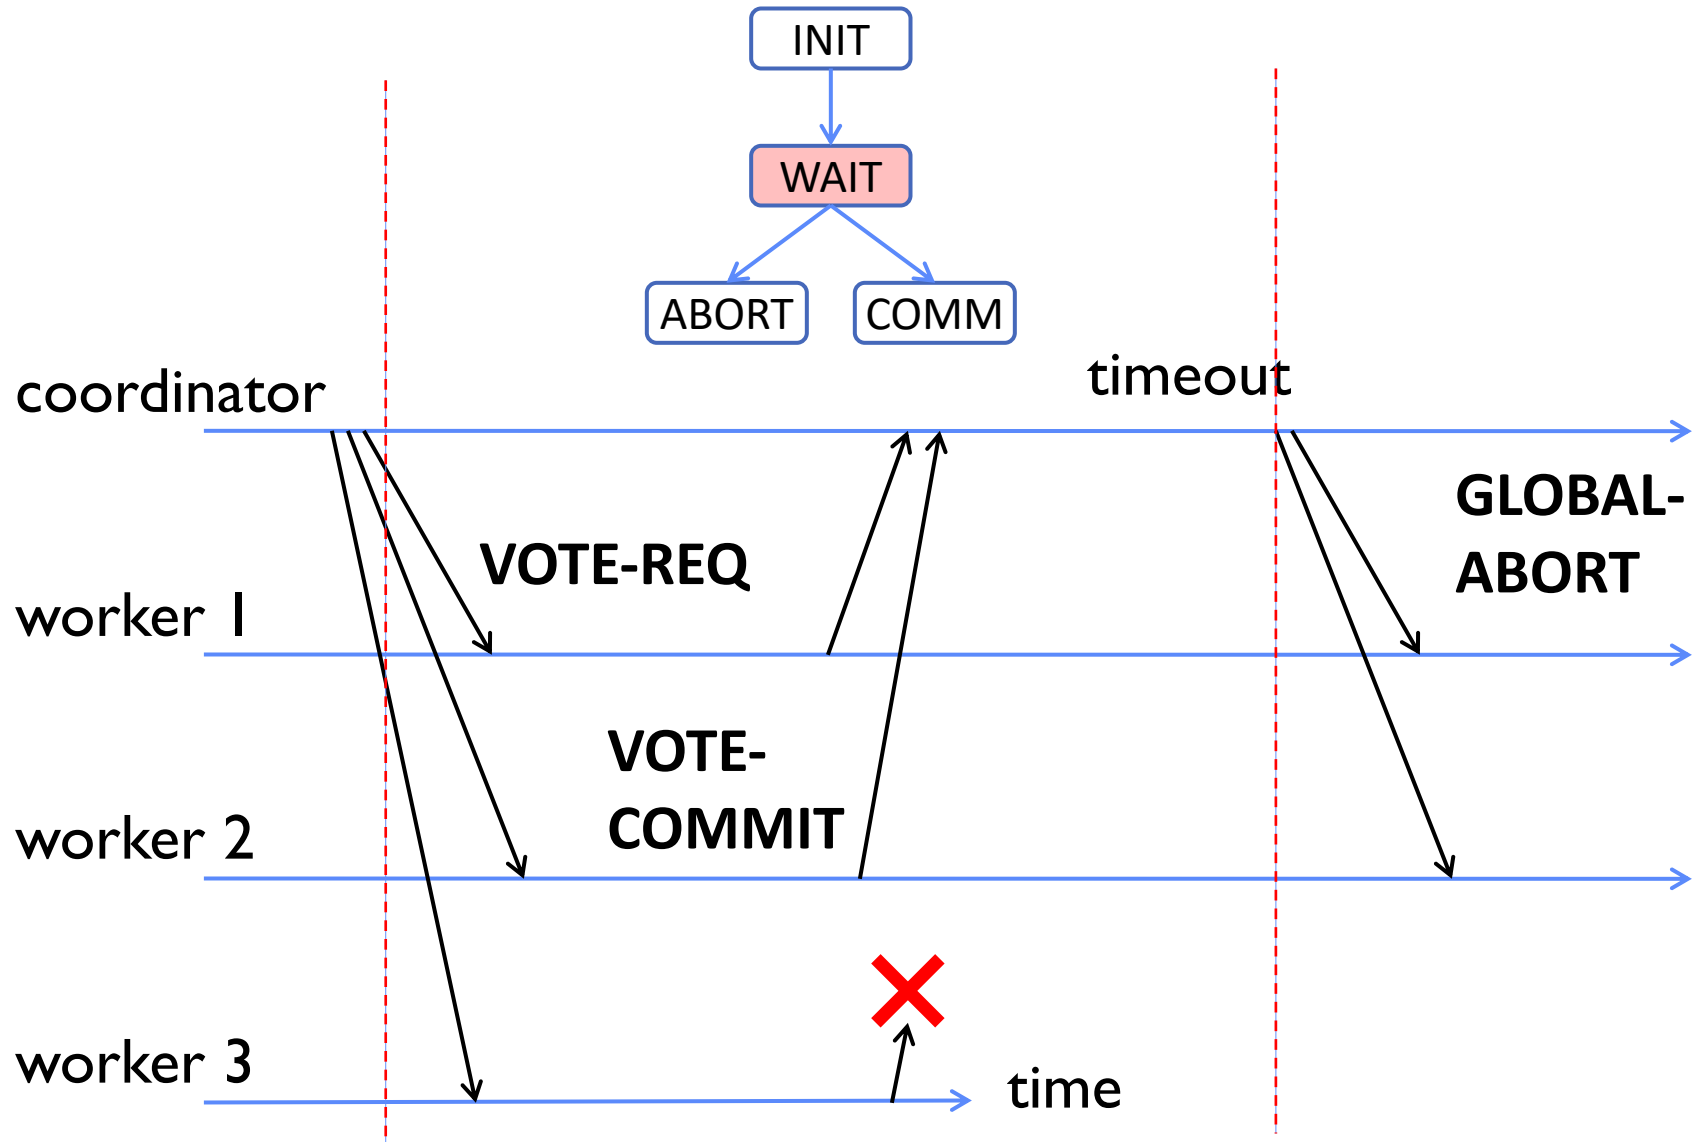
Dealing with Worker Failures

- Failure only affects states in which the coordinator is waiting for messages
- In WAIT, if coordinator doesn't receive N votes, it times out and sends GLOBAL-ABORT





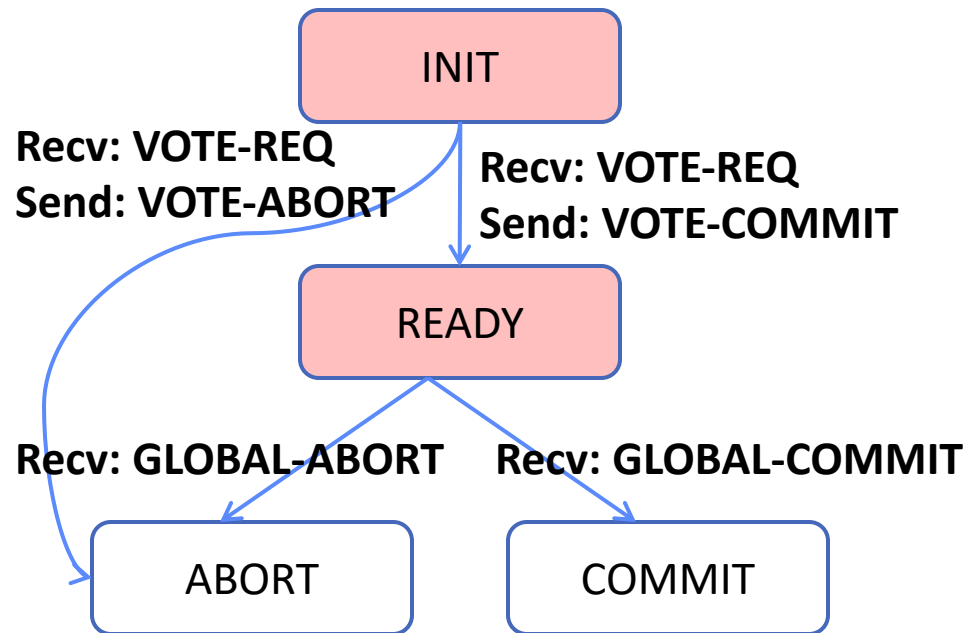
Example of Worker Failure (w/ state)





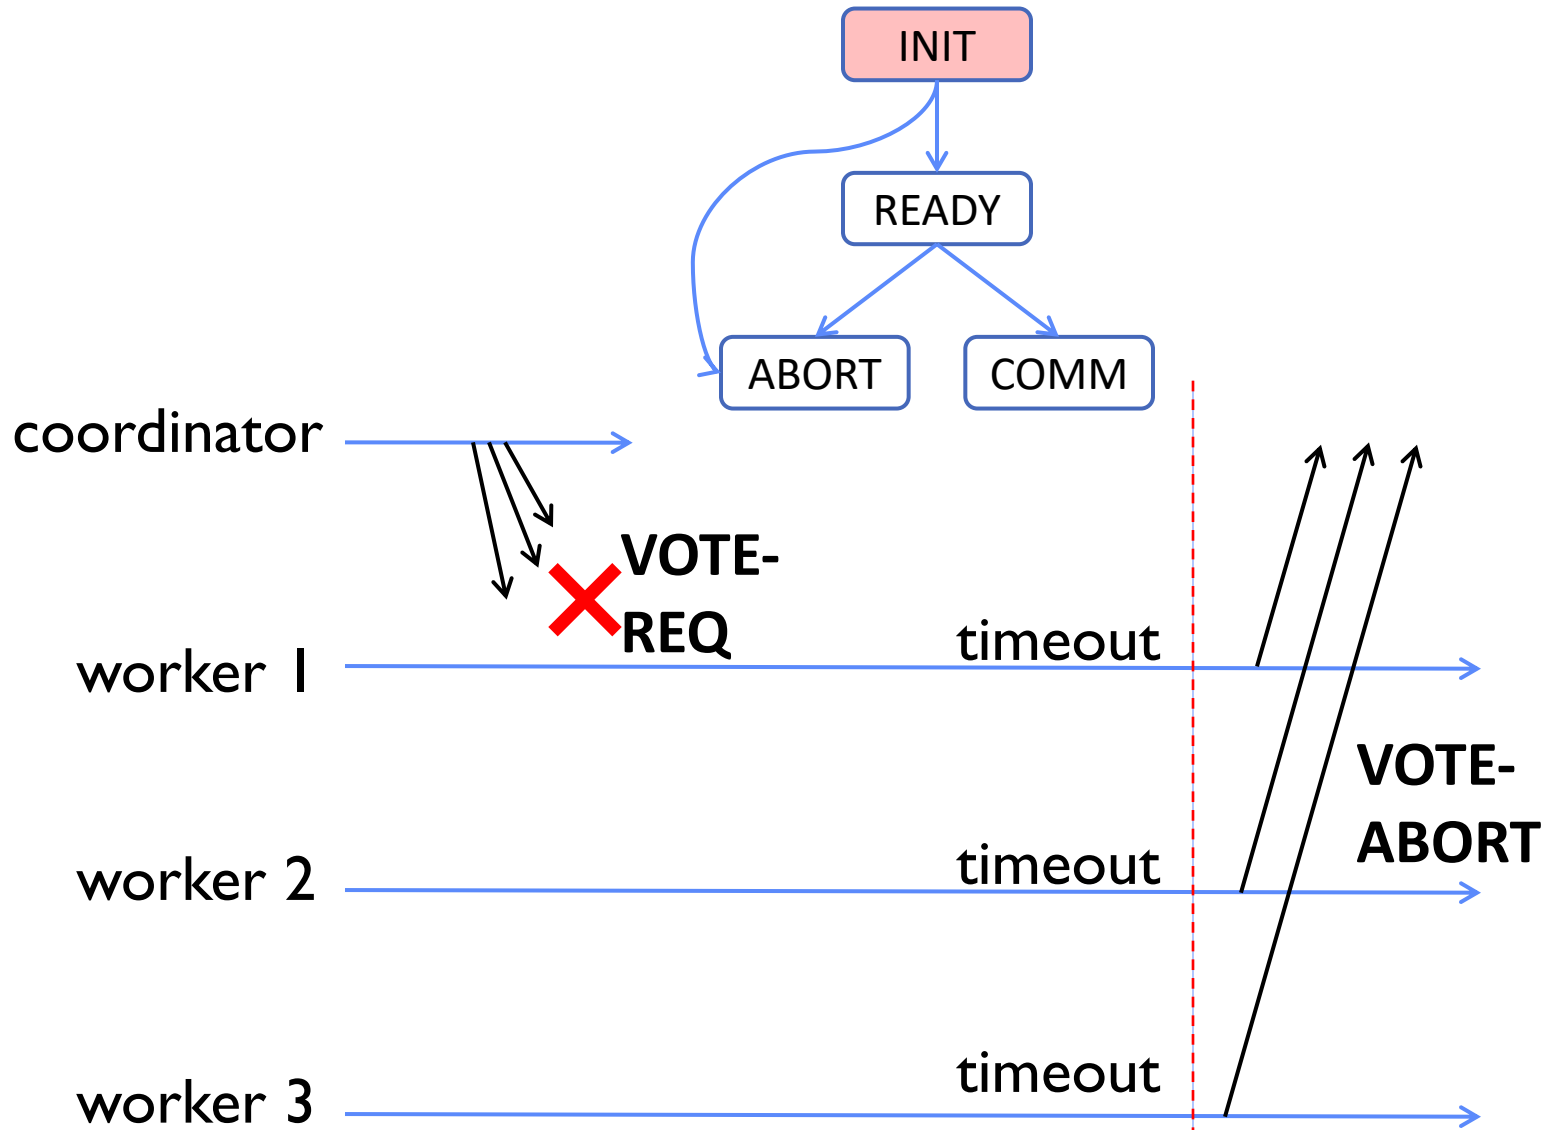
Dealing with Coordinator Failure

- Worker waits for VOTE-REQ in INIT
 - Worker can time out and abort (coordinator handles it)
- Worker waits for GLOBAL-* message in READY
 - If coordinator fails, workers must **BLOCK** waiting for coordinator to recover and send GLOBAL_* message



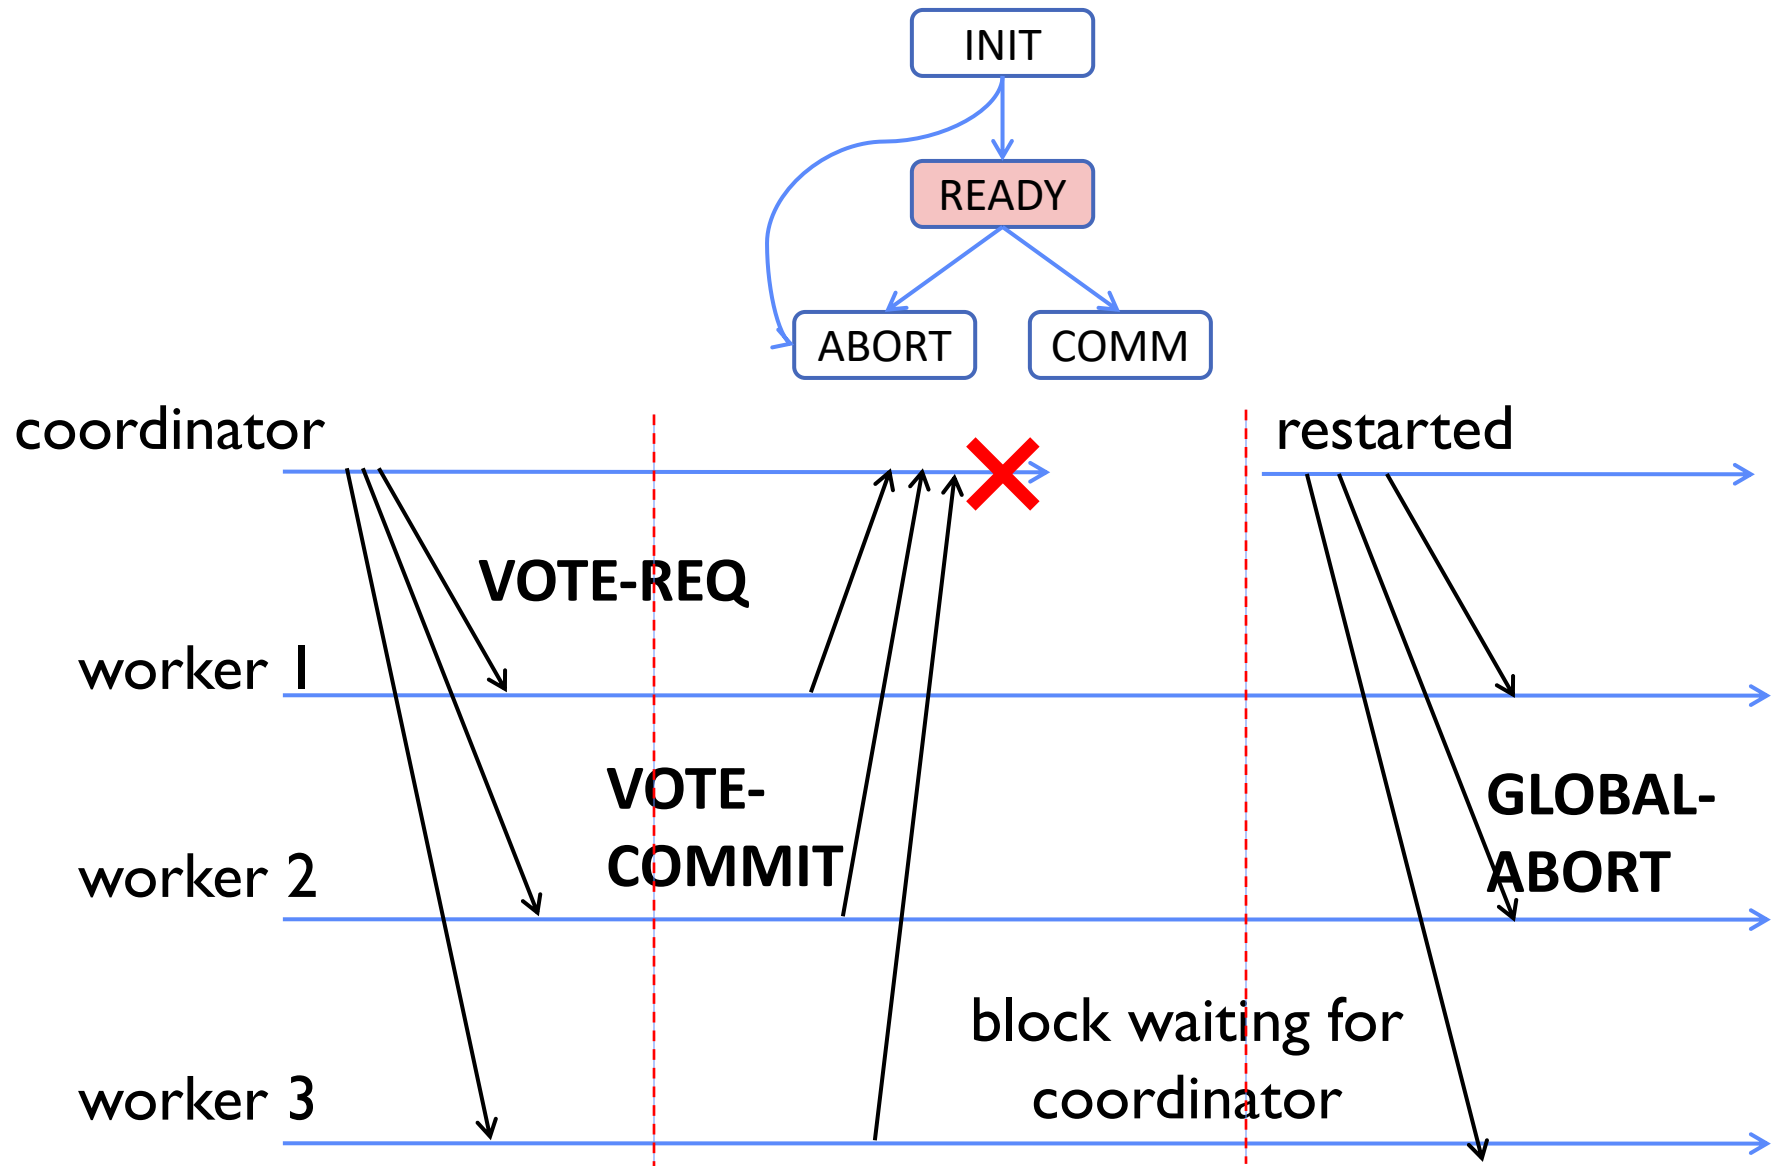


Example of Coordinator Failure (1)





Example of Coordinator Failure (2)





Failure Recovery

- Nodes need to know what state they are in when they come back from a failure
- How? Log events on local disk, SSD, NVRAM
- Then we have the following recovery rules:
 - Coordinator **aborts** transaction if it was in the INIT, WAIT, or ABORT states
 - Coordinator **commits** transaction if it was in COMMIT
 - Worker **aborts** if in INIT or ABORT states
 - Worker **commits** if it was in COMMIT state
 - Worker “**asks**” coordinator what to do if in READY state

Blocking for Coordinator to Recover

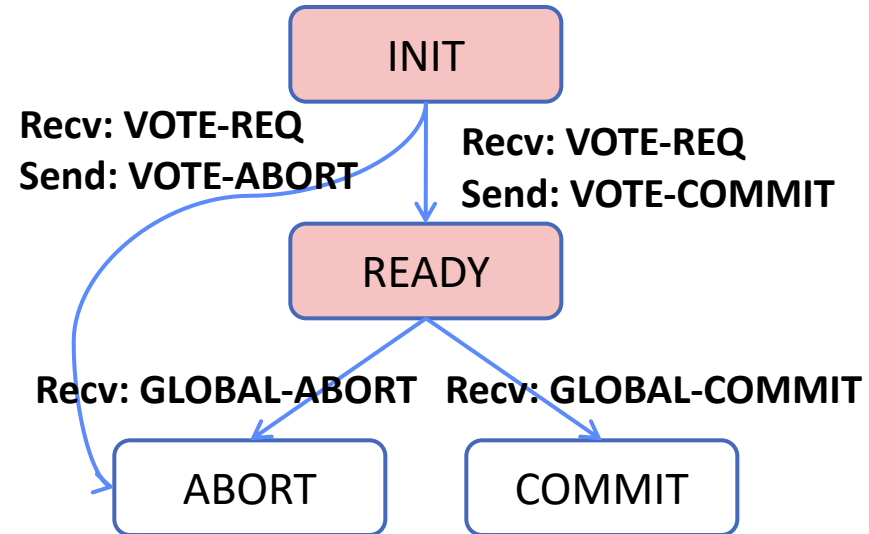


- A worker waiting for global decision can ask fellow workers about their state

- If another worker is in ABORT or COMMIT state then coordinator must have sent GLOBAL-*
 - » Thus, worker can safely abort or commit, respectively

- If another worker is still in INIT state then both workers can decide to abort

- If all workers are in ready, need to **BLOCK** (don't know if coordinator wanted to abort or commit)





Blocking

- What if *both* coordinator and a worker fail?
- The remaining workers can still consult each other
- But they can't reach a conclusion on what to do!

Why?

- If all workers in INIT, we still don't know state of failed worker w
- w may have been first to be notified of a commit, and then coordinator and w crashed



Blocking for Coordinator

- What if *both* coordinator and a worker fail?
- The remaining workers can still consult each other
- But they can't reach a conclusion on what to do!

This problem motivated *Three Phase Commit*



Distributed Consensus

- **Two- and Three-Phase commit make a decentralized decision**
- **Example: Changing the value of a key among all replicas for the key**
- **But they are hardly the only solutions to this problem**



Parallel vs Distributed

- **Distributed: different machines responsible for different parts of task**
 - Usually no centralized state
 - Usually about different responsibilities or redundancy
- **Parallel: different parts of same task performed on different machines**
 - Usually about performance



Summary

- **Key Value Store: Simple put and get operations**
 - Fault tolerance: replication
 - Scalability: Add nodes, balance load, no central directory
 - Consistency: Quorum consensus for better performance
- **Consensus Goal: Everyone agrees on the state of the distributed system**
 - Doesn't depend who you ask
 - Doesn't matter if nodes go down
- **Distributed Transactions**
 - Atomic, can't revert once agreement is reached



Summary: Two-Phase Commit

- **Voting protocol requires unanimity**
- **Transaction committed if and only if: all workers and coordinator vote to commit**
- **Nodes never take back their vote**
 - Logged for durability
- **Nodes work in lock step (for an item)**
 - Don't perform new transactions until old one is resolved
 - Stall until transaction is resolved