

# Section 3: Files and Basic Scheduling

CS162

September 18-20, 2019

## Contents

<b>1</b>	<b>Vocabulary</b>	<b>2</b>
<b>2</b>	<b>Problems</b>	<b>3</b>
2.1	Signals . . . . .	3
2.2	Signal Handlers . . . . .	3
2.3	Files . . . . .	4
2.3.1	Files vs File Descriptor . . . . .	4
2.3.2	Quick practice with write and seek . . . . .	4
2.4	Reading and Writing with File Pointers vs. Descriptors . . . . .	5
2.5	Storing Ints . . . . .	6
2.6	Dup and Dup2 . . . . .	8
2.6.1	Warmup . . . . .	8
2.6.2	Redirection: executing a process after dup2 . . . . .	8
2.6.3	Redirecting in a new process . . . . .	9
2.7	Round Robin Scheduling . . . . .	10

# 1 Vocabulary

- **system call** - In computing, a system call is how a program requests a service from an operating system's kernel. This may include hardware-related services, creation and execution of new processes, and communication with integral kernel services such as process scheduling.
- **file descriptors** - File descriptors are an index into a file-descriptor table stored by the kernel. The kernel creates a file-descriptor in response to an open call and associates the file-descriptor with some abstraction of an underlying file-like object; be that an actual hardware device, or a file-system or something else entirely. Using file descriptors, a process's read or write calls are routed to the correct place by the kernel. When your program starts you have 3 file descriptors.

File Descriptor	File
0	stdin
1	stdout
2	stderr

- **int open(const char \*path, int flags)** - open is a system call that is used to open a new file and obtain its file descriptor. Initially the offset is 0.
- **size\_t read(int fd, void \*buf, size\_t count)** - read is a system call used to read `count` bytes of data into a buffer starting from the file offset. The file offset is incremented by the number of bytes read.
- **size\_t write(int fd, const void \*buf, size\_t count)** - write is a system call that is used to write up to `count` bytes of data from a buffer to the file offset position. The file offset is incremented by the number of bytes written.
- **size\_t lseek(int fd, off\_t offset, int whence)** - lseek is a system call that allows you to move the offset of a file. There are three options for whence
  - SEEK\_SET - The offset is set to `offset`.
  - SEEK\_CUR - The offset is set to `current_offset + offset`
  - SEEK\_END - The offset is set to the size of the file + `offset`
- **int dup(int oldfd)** - creates an alias for the provided file descriptor and returns the new fd value. dup always uses the smallest available file descriptor. Thus, if we called dup first thing in our program, it would use file descriptor 3 (0, 1, and 2 are already signed to stdin, stdout, stderr). The old and new file descriptors refer to the same open file description and may be used interchangeably.
- **int dup2(int oldfd, int newfd)** - dup2 is a system call similar to dup. It duplicates the `oldfd` file descriptor, this time using `newfd` instead of the lowest available number. If `newfd` was open, it closed before being reused. This becomes very useful when attempting to redirect output, as it automatically takes care of closing the file descriptor, performing the redirection in one elegant command. For example, if you wanted to redirect standard output to a file, then you would
- **signals** - A signal is a software interrupt, a way to communicate information to a process about the state of other processes, the operating system, and the hardware. A signal is an interrupt in the sense that it can change the flow of the program when a signal is delivered to a process, the process will stop what its doing, either handle or ignore the signal, or in some cases terminate, depending on the signal.
- **int signal(int signum, void (\*handler)(int))** - signal() is the primary system call for signal handling, which given a signal and function, will execute the function whenever the signal is delivered. This function is called the signal handler because it handles the signal.

- **SIG\_IGN, SIG\_DFL** Usually the second argument to signal takes a user defined handler for the signal. However, if you'd like your process to drop the signal you can use SIG\_IGN. If you'd like your process to do the default behavior for the signal use SIG\_DFL.
- **scheduler** - The process scheduler is a part of the operating system that decides which process runs at a certain point in time. It usually has the ability to pause a running process, move it to the back of the running queue and start a new process;
- **FIFO Scheduling** - First-In-First-Out (aka First-Come-First-Serve) scheduling runs jobs as they arrive. Turnaround time can degrade if short jobs get stuck behind long ones (convoy effect);
- **round-robin Scheduling** - Round-Robin scheduling runs each job in fixed-length time slices (quanta). The scheduler preempts a job that exceeds its quantum and moves on, cycling through the jobs. It avoids starvation and is good for short jobs, but context switching overhead can become important depending on quanta length;

## 2 Problems

### 2.1 Signals

The following is a list of standard Linux signals:

Signal	Value	Action	Comment
SIGHUP	1	Terminate	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Terminate	Interrupt from keyboard (Ctrl - c)
SIGQUIT	3	Core Dump	Quit from keyboard (Ctrl - \)
SIGILL	4	Core Dump	Illegal Instruction
SIGABRT	6	Core Dump	Abort signal from abort(3)
SIGFPE	8	Core Dump	Floating point exception
SIGKILL	9	Terminate	Kill signal
SIGSEGV	11	Core Dump	Invalid memory reference
SIGPIPE	13	Terminate	Broken pipe: write to pipe with no readers
SIGALRM	14	Terminate	Timer signal from alarm(2)
SIGTERM	15	Terminate	Termination signal
SIGUSR1	30,10,16	Terminate	User-defined signal 1
SIGUSR2	31,12,17	Terminate	User-defined signal 2
SIGCHLD	20,17,18	Ignore	Child stopped or terminated
SIGCONT	19,18,25	Continue	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at tty
SIGTTIN	21,21,26	Stop	tty input for background process
SIGTTOU	22,22,27	Stop	tty output for background process

### 2.2 Signal Handlers

Fill in the blanks for the following function using syscalls such that when we type Ctrl-C, the user is prompted with a message: "Do you really want to quit [y/n]? ", and if "y" is typed, the program quits. Otherwise, it continues along.

```

void sigint_handler(int sig)
{
    char c;
    printf("Ouch, you just hit Ctrl-C?. Do you really want to quit [y/n]?");
    c = getchar();
    if (c == "y" || c == "Y")
        exit(0);
}

int main() {
    signal(SIGINT, sigint_handler);
    ...
}

```

## 2.3 Files

### 2.3.1 Files vs File Descriptor

What's the difference between `fopen` and `open`?

`fopen` is implemented in `libc` whereas `open` is a syscall. `fopen` will use `open` in it's implementation. `fopen` will return a `FILE *` and `open` will return an `int`. The `FILE *` object allows you to call utility methods from `stdio.h` like `fscanf`. Also the `FILE *` object comes with some library level buffering of writes.

```

-----
|  libc      |
-----
| syscall    |
-----

```

### 2.3.2 Quick practice with write and seek

What will the `test.txt` file look like after I run this program? (Hint: if you write at an offset past the end of file, the bytes inbetween the end of the file and the offset will be set to 0.)

```

int main() {
    char buffer[200];
    memset(buffer, 'a', 200);
    int fd = open("test.txt", O_CREAT|O_RDWR);
    write(fd, buffer, 200);
    lseek(fd, 0, SEEK_SET);
    read(fd, buffer, 100);
    lseek(fd, 500, SEEK_CUR);
    write(fd, buffer, 100);
}

```

The first write gives us 200 bytes of a. Then we seek to the offset 0 and read 100 bytes to get to offset 100. Then we seek to offset 100 + 500 to offset 600. Then we write 100 more bytes of a.

At then end we will have a from 0-200, 0 from 200-600, and a from 600-700

## 2.4 Reading and Writing with File Pointers vs. Descriptors

Write a utility function, `void copy(const char *src, const char *dest)`, that simply copies the file contents from src and places it in dest. You can assume both files are already created. Also assume that the src file is at most 100 bytes long. First, use the file pointer library to implement this. Fill in the code given below:

```
void copy(const char *src, const char *dest) {
    char buffer [100];
    FILE* read_file = fopen(_____, ____);
    int buf_size = fread(_____, ____, _____, _____);
    fclose(read_file);

    FILE* write_file = fopen(_____, ____);
    fwrite(_____, ____, _____, _____);
    fclose(write_file);
}
```

```
void copy(const char *src, const char *dest) {
    char buffer [100];
    FILE* read_file = fopen(src, "r");
    int buf_size = fread(buffer, 1, sizeof(buffer), read_file);
    fclose(read_file);

    FILE* write_file = fopen(dest, "w");
    fwrite(buffer, 1, buf_size, write_file);
    fclose(write_file);
}
```

Next, use file descriptors to implement the same thing.

```

void copy(const char *src, const char *dest) {
    char buffer [100];
    int read_fd = open(_____, _____);
    int bytes_read = 0;
    int buf_size = 0;

    while ((bytes_read = read(_____, _____, _____)) > 0) {
        _____
    }
    close(read_fd);

    int bytes_written = 0;
    int write_fd = open(_____, _____);
    while (_____) {
        _____ += write(_____, _____, _____);
    }
    close(write_fd);
}

```

```

void copy(const char *src, const char *dest) {
    char buffer [100];
    int read_fd = open(src, O_RDONLY);
    int bytes_read = 0;
    int buf_size = 0;

    while ((bytes_read = read(read_fd, &buffer[buf_size], sizeof(buffer) - buf_size)) > 0) {
        buf_size += bytes_read;
    }
    close(read_fd);

    int bytes_written = 0;
    int write_fd = open(dest, O_WRONLY);
    while (bytes_written < buf_size) {
        bytes_written += write(write_fd, &buffer[bytes_written], buf_size - bytes_written);
    }
    close(write_fd);
}

```

Compare the file pointer implementation to the file descriptor implementation. In the file descriptor implementation, why does **read** and **write** need to be called in a loop?

Read and write need to be called in a loop because there is no guarantee that both functions will actually process the specified number of bytes (they can return less bytes read / written). However, this functionality is already handled in the file pointer library, so a single call to **fread** and **fwrite** would suffice.

## 2.5 Storing Ints

You are working for BigStore and your boss has tasked you with writing a function that takes an array of ints and writes it to a specified file for later use. He also informs you that a major bug has been found

in the C file pointer library and wants you to use file descriptors. Fill in the following function:

```
void write_to_file(const char *file, int *a, int size) {
    int write_fd = open(_____, _____);

    char *write_buf = _____
    int buf_size = _____
    int bytes_written = 0;

    // Write a to file.
    _____
    _____
    _____
    close(write_fd);
}
```

```
void write_to_file(const char *file, int *a, int size) {
    int write_fd = open(file, O_WRONLY);

    char *write_buf = (char *) &a[0];
    int buf_size = size * sizeof(int);
    int bytes_written = 0;

    while (bytes_written < buf_size) {
        bytes_written += write(write_fd, &write_buf[bytes_written], buf_size - bytes_written);
    }
    close(write_fd);
}
```

Now, write the function that retrieves previously saved integers and places them in a int array.

```
void read_from_file(const char *file, int *a, int size) {
    int read_fd = open(_____, _____);

    char *read_buf = _____
    int buf_size = _____

    // Read a from a file.
    _____
    _____
    _____
    _____
    close(read_fd);
}
```

```
void read_from_file(const char *file, int *a, int size) {
    int read_fd = open(file, O_RDONLY);
```

```
char *read_buf = (char *) &a[0];
int buf_size = size * sizeof(int);

int bytes_read = 0;
int total_read = 0;
while ((bytes_read = read(read_fd, &read_buf[total_read], buf_size - total_read)) > 0) {
    total_read += bytes_read;
}
close(read_fd);
}
```

Your coworker opens up one of the files that you used to store ints on his text editor and complains its full of junk! Explain to him why this might be the case.

Currently, we are reading and writing the contents of memory directly to disk. This is convenient for us, because we do not have to do any parsing of the input. However, the memory representation of an int array is unlikely to be human readable.

## 2.6 Dup and Dup2

### 2.6.1 Warmup

What does C print in the following code?

```
int main(int argc, char **argv)
{
    int pid, status;
    int newfd;

    if ((newfd = open("output_file.txt", O_CREAT|O_TRUNC|O_WRONLY, 0644)) < 0) {
        exit(1);
    }
    printf("The last digit of pi is...");
    dup2(newfd, 1);
    printf("five\n");
    exit(0);
}
```

This prints "The last digit of pi is..." to standard output. Unfortunately, "five" gets written to the output\_file.txt and our joke is left incomplete.

### 2.6.2 Redirection: executing a process after dup2

Describe what happens, and what the output will be.

```
int main(int argc, char **argv)
{
    int pid, status;
    int newfd;
    char *cmd[] = { "/bin/ls", "-al", "/", 0 };
}
```



```

    if (argc != 2) {
        fprintf(stderr, "usage: %s output_file\n", argv[0]);
        exit(1);
    }
    if ((newfd = open(argv[1], O_CREAT|O_TRUNC|O_WRONLY, 0644)) < 0) {
        perror(argv[1]);    /* open failed */
        exit(1);
    }
    printf("writing output of the command %s to \"%s\"\n", cmd[0], argv[1]);
    dup2(newfd, 1);
    execvp(cmd[0], cmd);
    perror(cmd[0]);    /* execvp failed */

    exit(1);
}

```

We get the name of the output file from the command line and set that to be the standard output but now execute a command (`ls -al /` in this example). The command sends its output to the standard output stream, which is now the file that we created.

### 2.6.3 Redirecting in a new process

Modify the above code such that the result of `ls -al` is written to the file specified by the input argument and immediately after "all done" is printed to the terminal. (Hint: you'll need to use `fork` and `wait`.)

```

int main(int argc, char **argv)
{
    int pid, status;
    int newfd;
    char *cmd[] = { "/bin/ls", "-al", "/", 0 };
    if ((pid = fork()) < 0) {
        perror();
        exit(1);
    }
    if (pid == 0) {
        if (argc != 2) {
            fprintf(stderr, "usage: %s output_file\n", argv[0]);
            exit(1);
        }
        if ((newfd = open(argv[1], O_CREAT|O_TRUNC|O_WRONLY, 0644)) < 0) {
            perror(argv[1]);    /* open failed */
            exit(1);
        }
        printf("writing output of the command %s to \"%s\"\n", cmd[0], argv[1]);
        dup2(newfd, 1);
        execvp(cmd[0], cmd);
        perror(cmd[0]);    /* execvp failed */
        exit(1);
    }
}

```

```
wait(&status);  
printf("all done");  
exit(0);  
}
```

## 2.7 Round Robin Scheduling

Which of the following are true about Round Robin Scheduling?

1. The average wait time is less than that of FCFS for the same workload.
2. Is supported by `thread_tick` in Pintos.
3. It requires pre-emption to maintain uniform quanta.
4. If quanta is constantly updated to become the # of cpu ticks since boot, Round Robin becomes FIFO.
5. If all threads in the system have the same priority, Priority Schedulers **must** behave like round robin.
6. Cache performance is likely to improve relative to FCFS.
7. If no new threads are entering the system all threads will get a chance to run in the cpu every `QUANTA*SECONDS_PER_TICK*NUMTHREADS` seconds. (Assuming `QUANTA` is in ticks).
8. This is the default scheduler in Pintos
9. It is the fairest scheduler

2,3,4,8

1. Easy to find counter example. 2. True. 3. True. 4. True. 5. False. Not a requirement. 6. False. More context switches means worse cache performance. 7. Trick question. There is some overhead. 8. True. 9. Trick question. Needs definition of fair.